

# IAGO ATTACKS: WHY THE SYSTEM CALL API IS A BAD UNTRUSTED RPC INTERFACE

Stephen Checkoway and Hovav Shacham

---

March 19, 2013

# A vulnerable program

---



```
#include <stdlib.h>

int main() {
    void *p = malloc(100);
}
```

# Problem setting

---



- ❖ Trusted application:



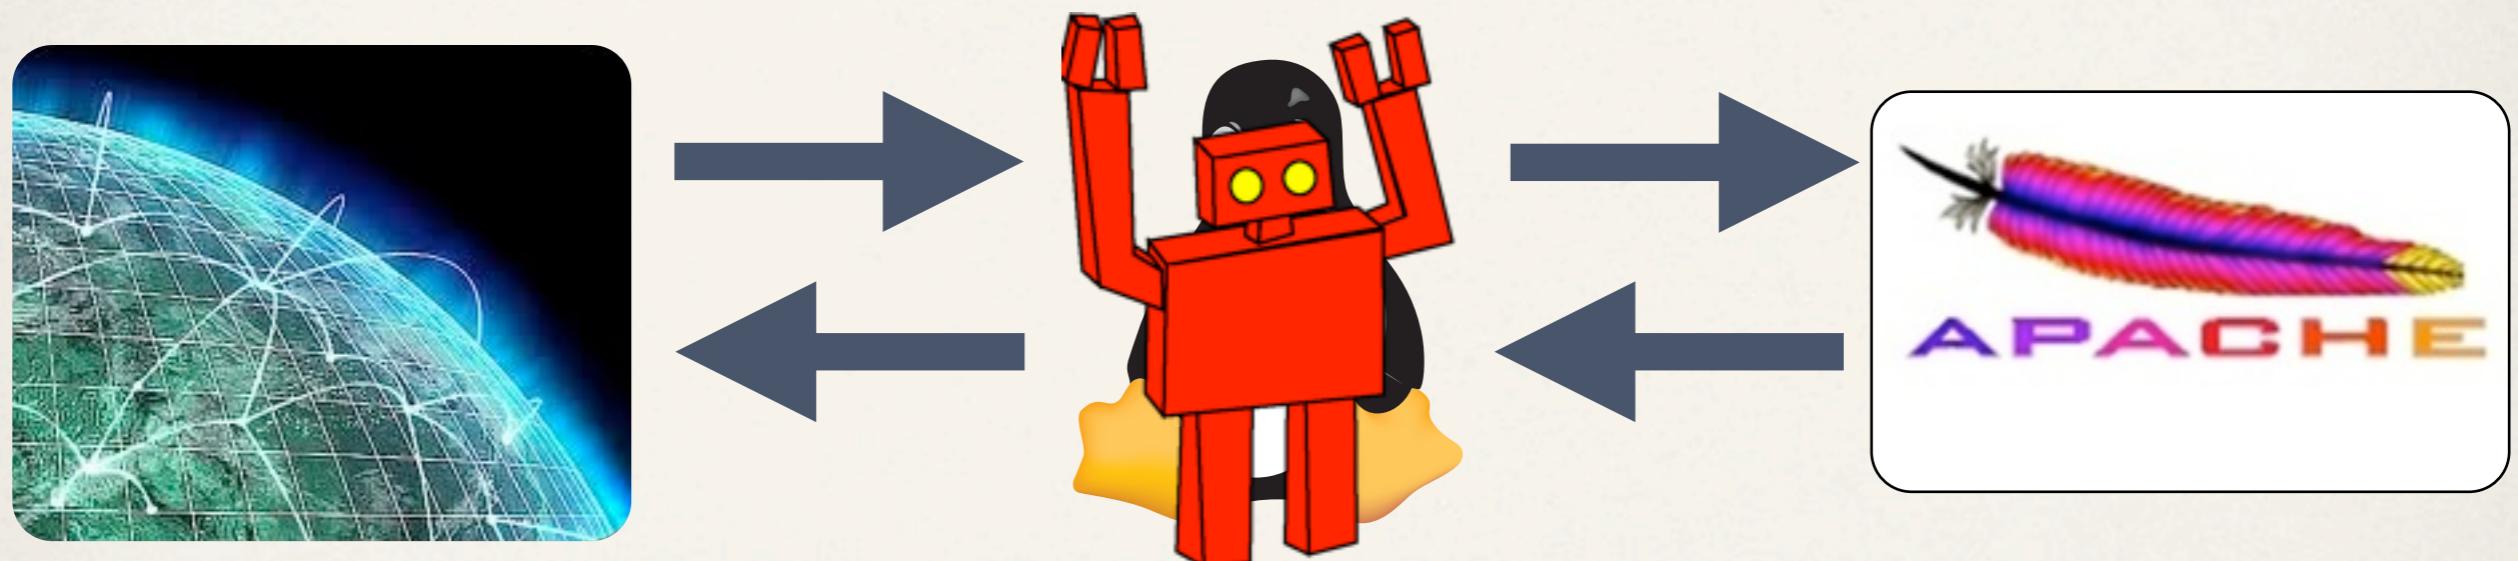
- ❖ Untrusted operating system:



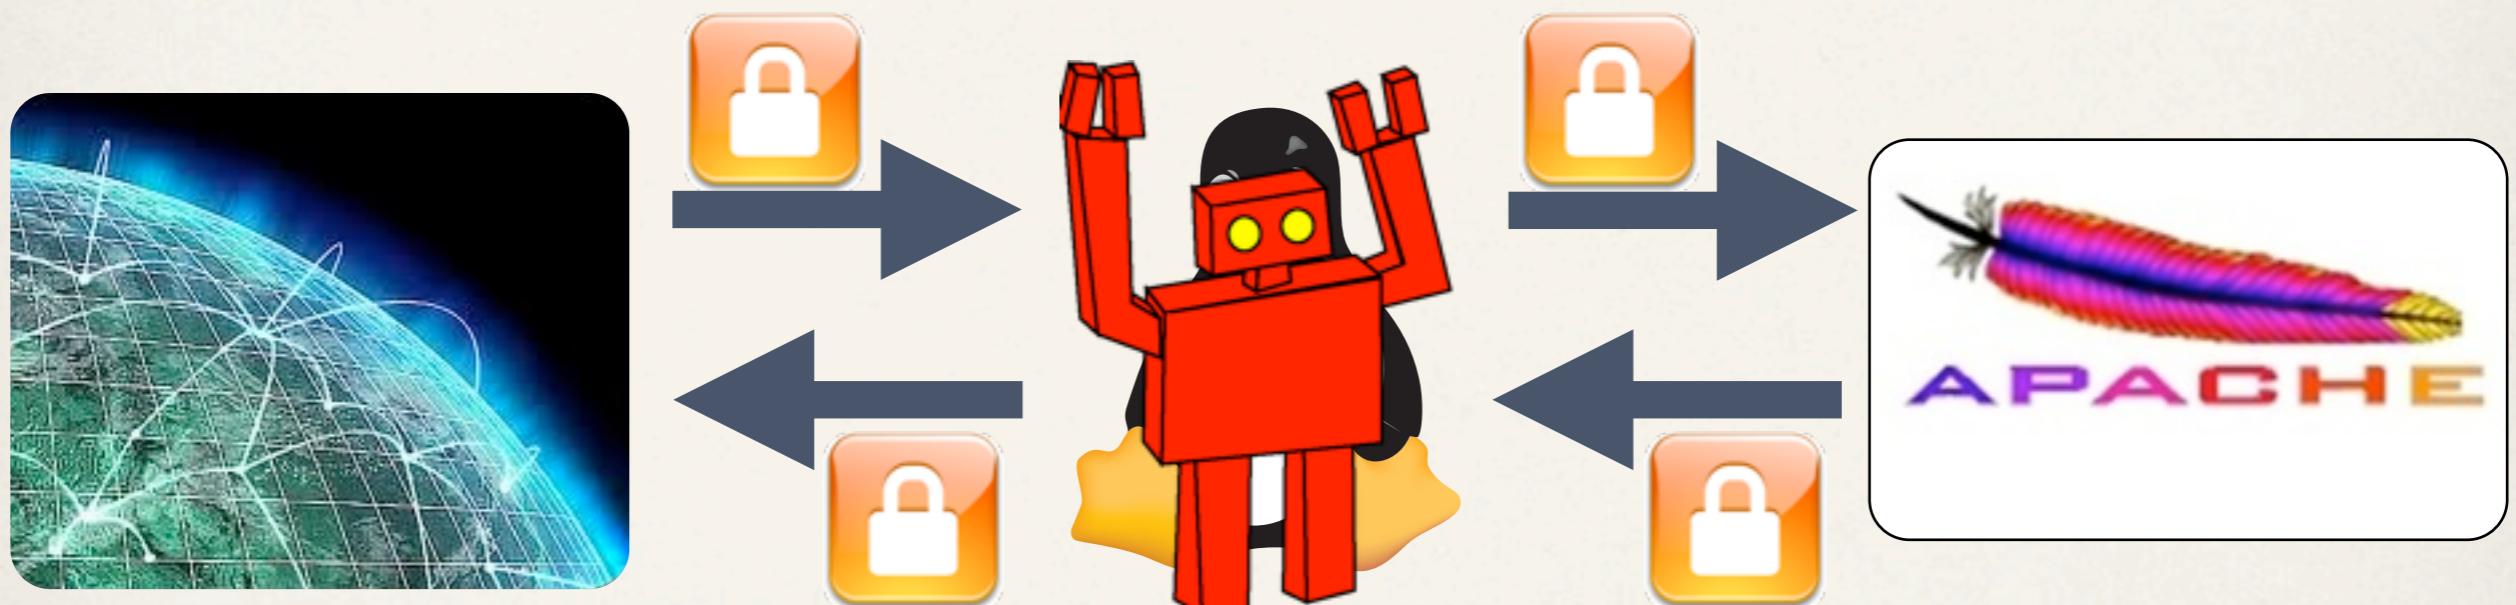
# Problem motivation



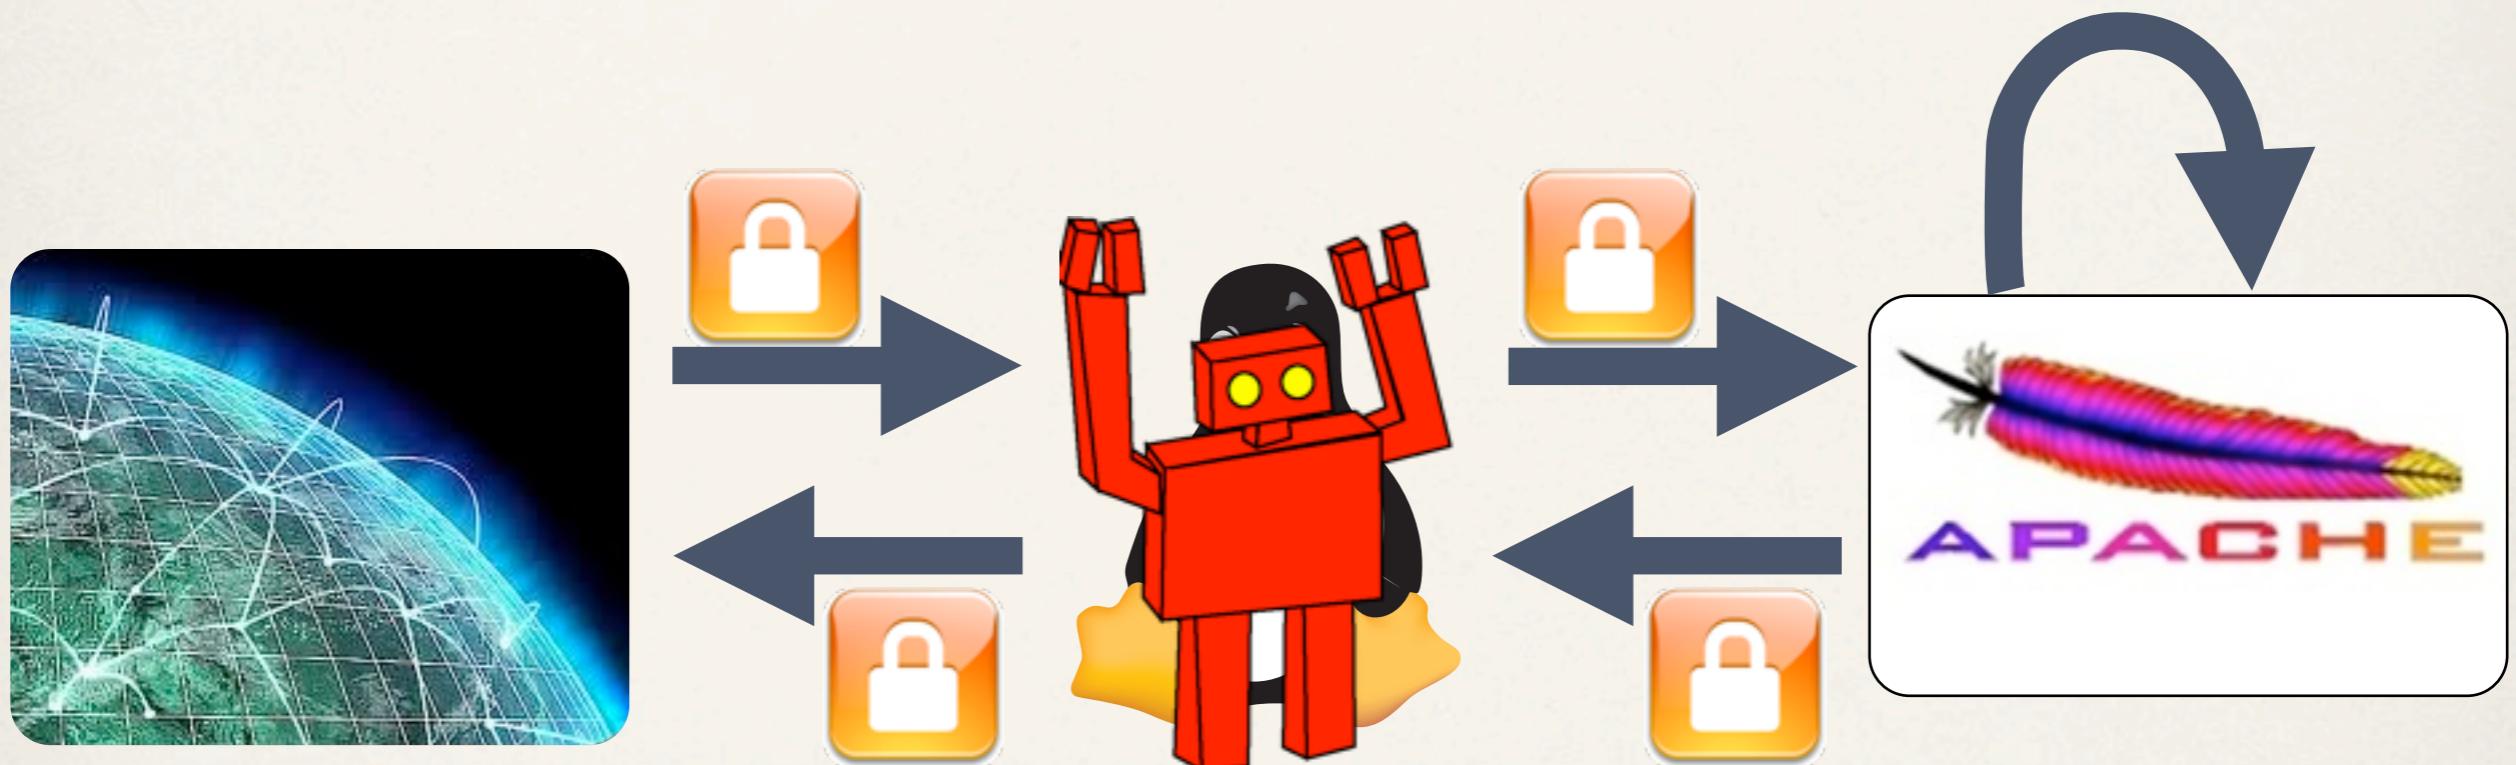
# Problem motivation



# Problem motivation



# Problem motivation



# Possible solutions

---



- ❖ Reimplement in a secure environment (e.g.,  $\mu$ kernel)
- ❖ Hardware-based solutions (e.g., XOM processor)
- ❖ Multiple virtual machines (e.g., Proxos)
- ❖ Hypervisor-assisted (e.g., Overshadow)

# Possible solutions

---



- ❖ Reimplement in a secure environment (e.g.,  $\mu$ kernel)
- ❖ Hardware-based solutions (e.g., XOM processor)
- ❖ Multiple virtual machines (e.g., Proxos)
- ❖ Hypervisor-assisted (e.g., Overshadow)

# The Overshadow approach



Application



Operating system

Chen et al. Overshadow: A Virtualization-Based Approach to Retrofitting Protection in Commodity Operating Systems. ASPLOS'08

# The Overshadow approach



Application

Shim



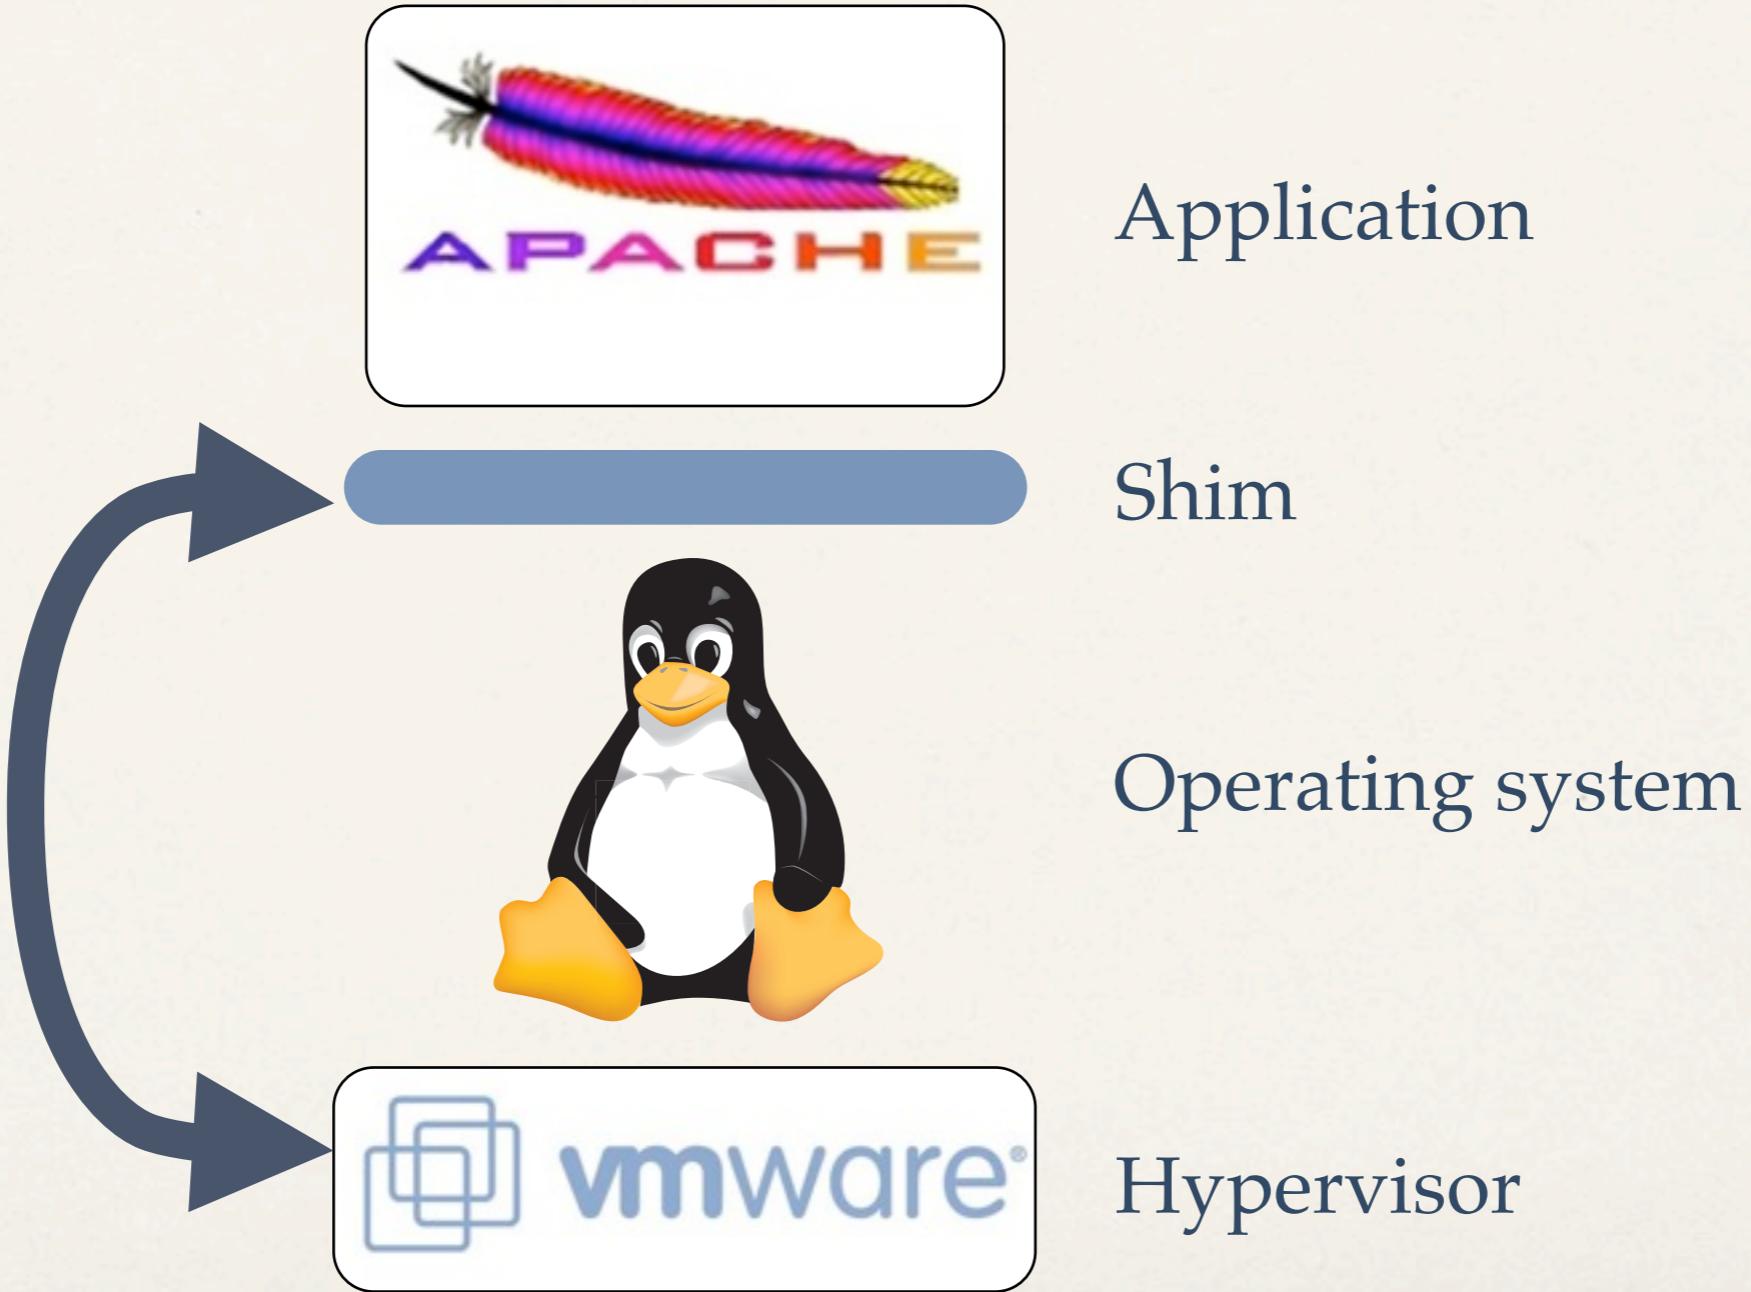
Operating system



Hypervisor

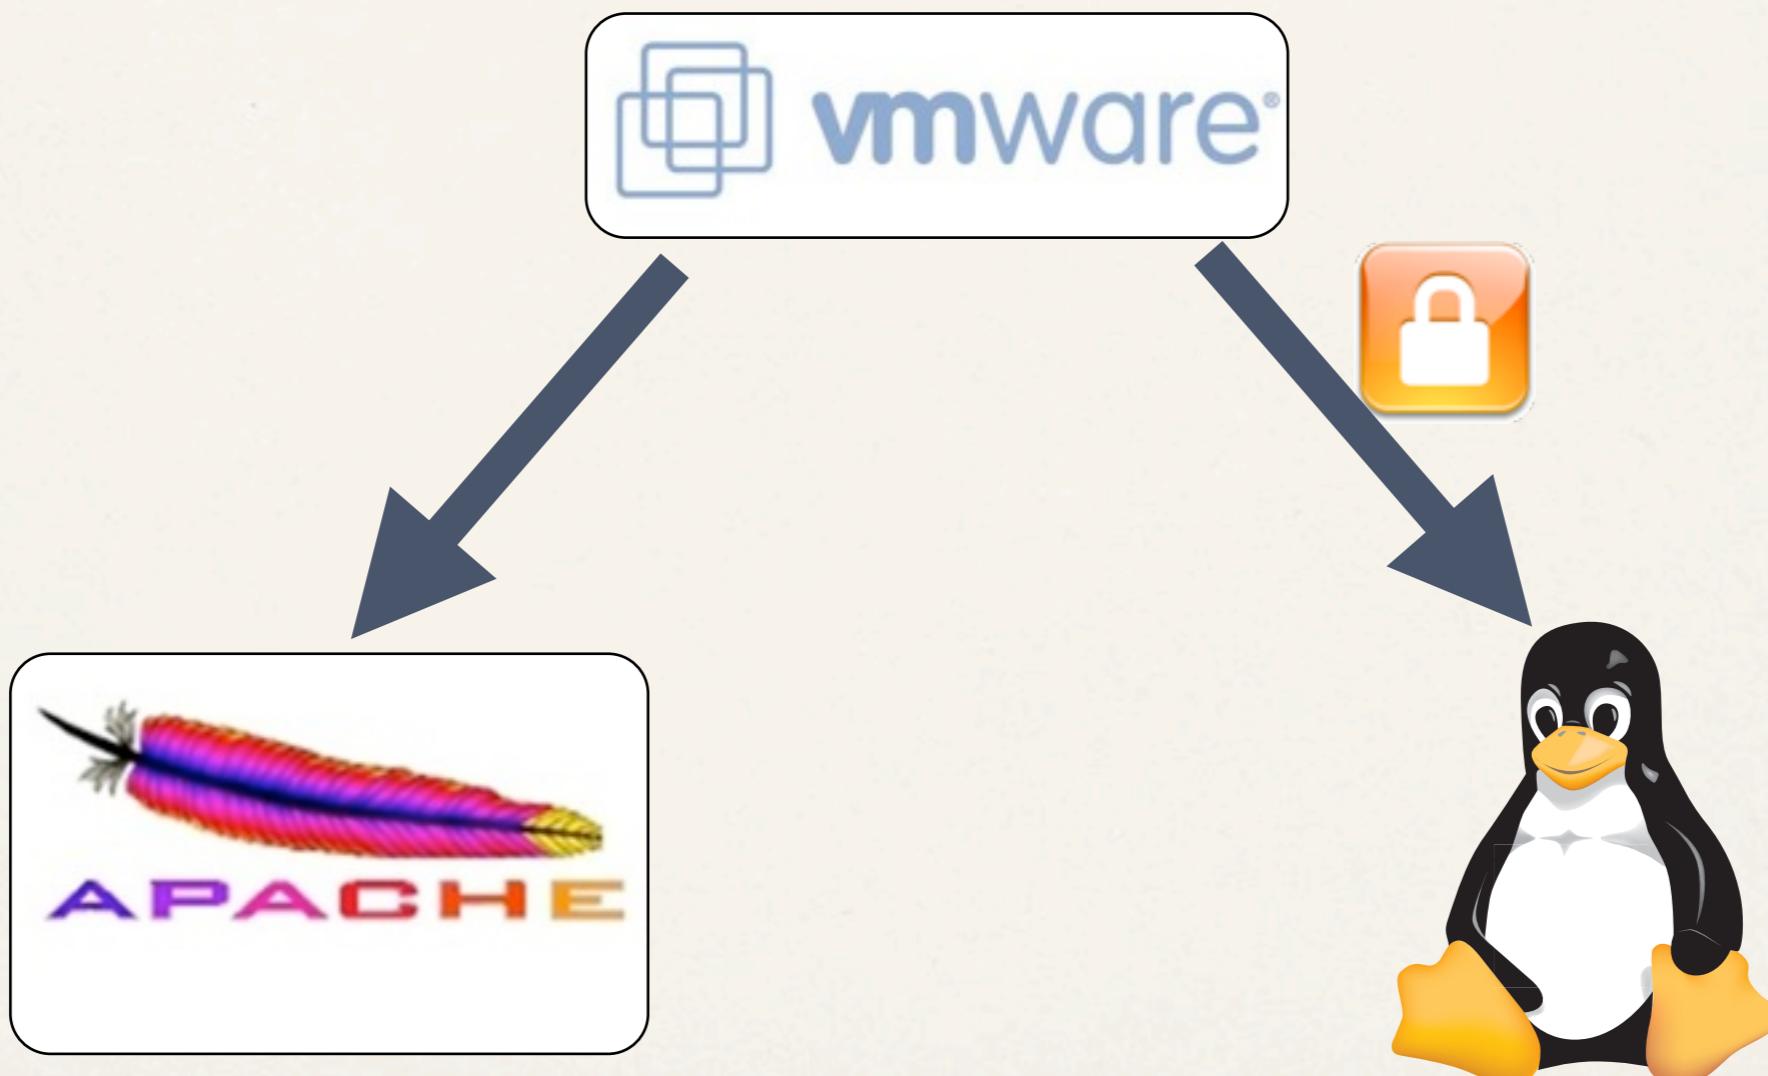
Chen et al. Overshadow: A Virtualization-Based Approach to Retrofitting Protection in Commodity Operating Systems. ASPLOS'08

# The Overshadow approach



Chen et al. Overshadow: A Virtualization-Based Approach to Retrofitting Protection in Commodity Operating Systems. ASPLOS'08

# Cloaking: Two views of application memory



# The shim

---



- Marshals arguments and return values for system calls
- Communicates directly with the hypervisor

*A majority of system calls can be passed through to the OS with no special handling. These include calls with scalar arguments that have no interesting side effects, such as `getpid`, `nice`, and `sync`.*

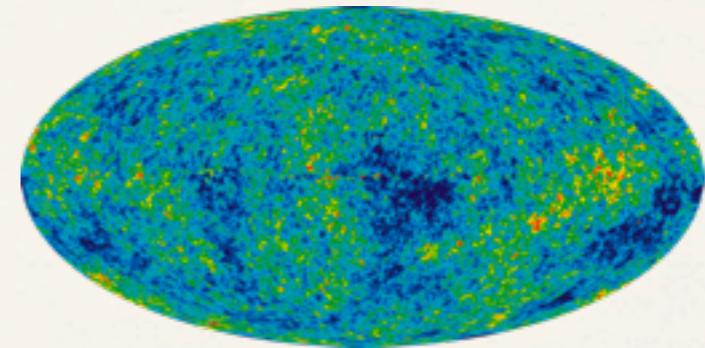
— Chen et al. ASPLOS'08

# Warmup: A thought experiment

---

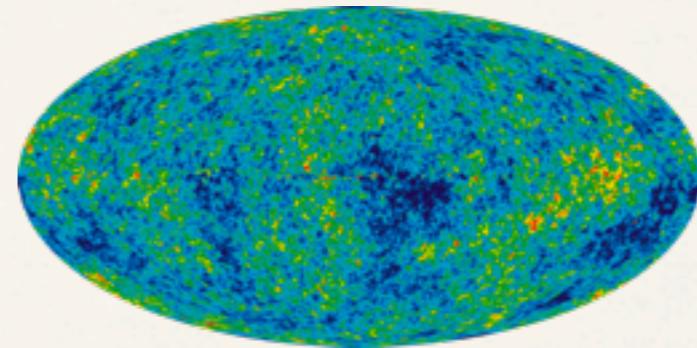


Main Apache process

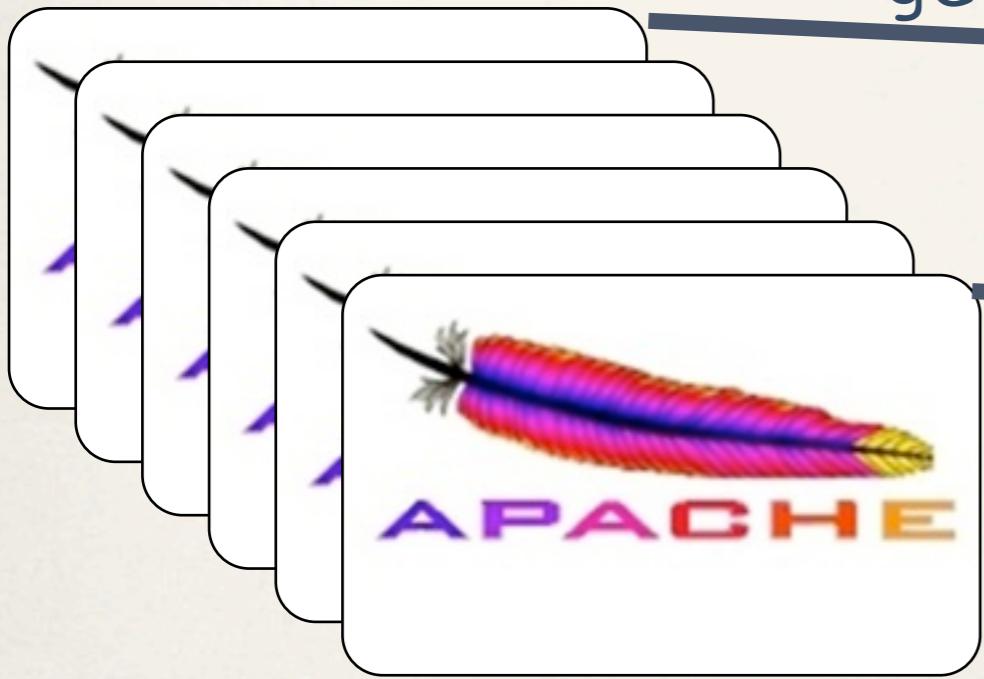


Entropy pool

# Warmup: A thought experiment



Main Apache process

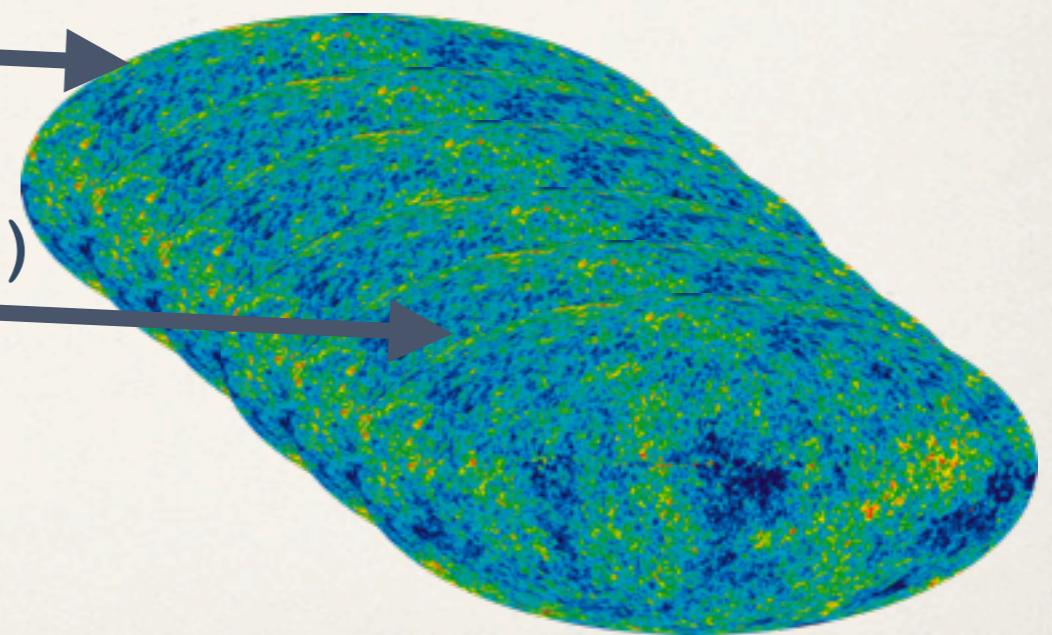


getpid()

⋮

getpid()

Entropy pool



Workers

Workers' entropy pools

10

# Technical goals

---



- ❖ Abstract away details of Overshadow
- ❖ Develop a malicious operating system kernel to attack protected applications
- ❖ Cause the protected application to act against its interests

# Threat model

---

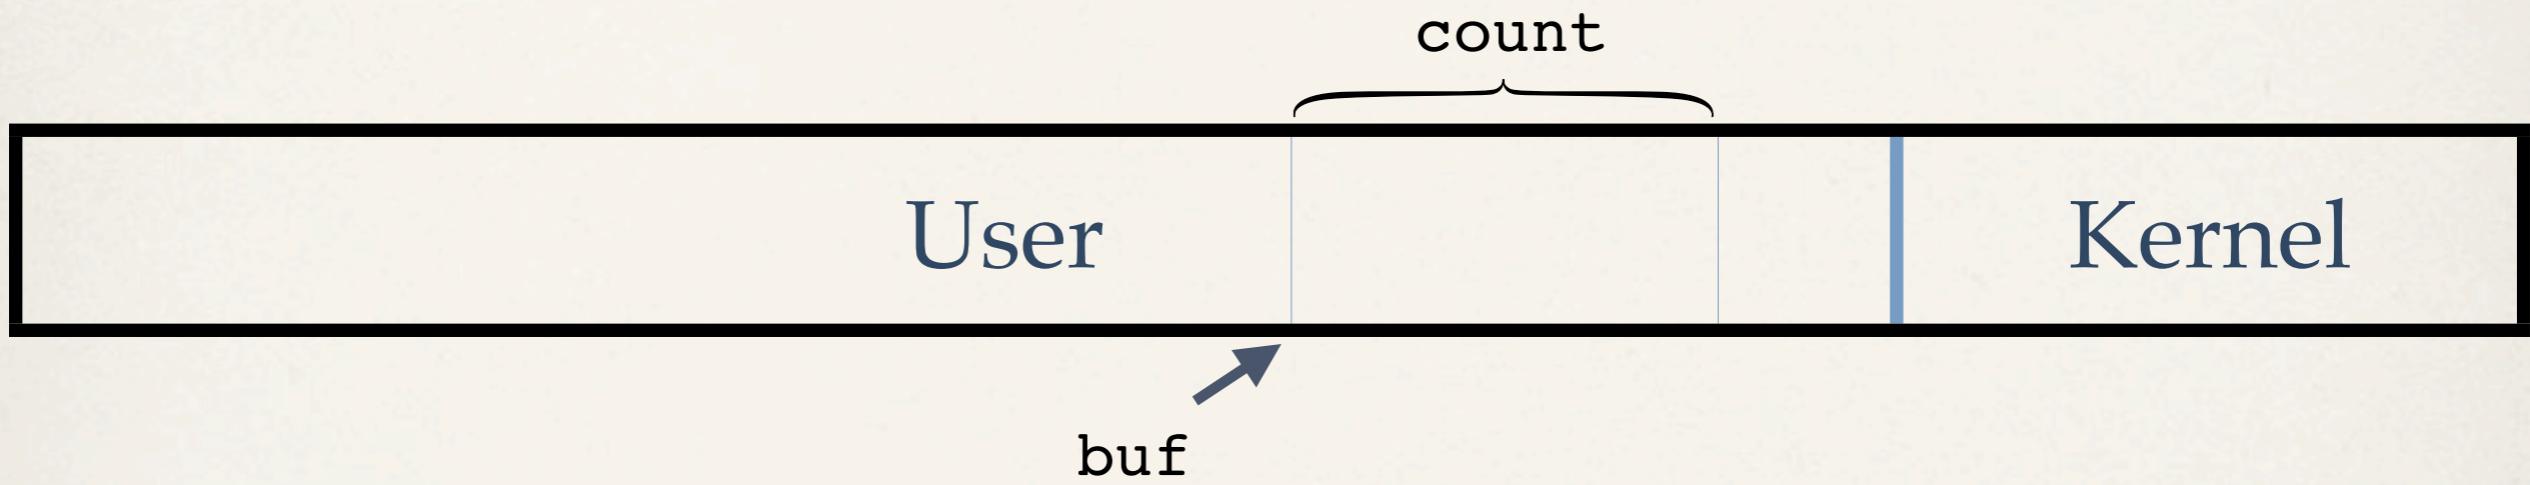


- ❖ Trusted, legacy application
- ❖ Unmodified system libraries
- ❖ Kernel cannot read or modify application state
- ❖ Kernel responds to system calls normally except for return values



# Threat model: example

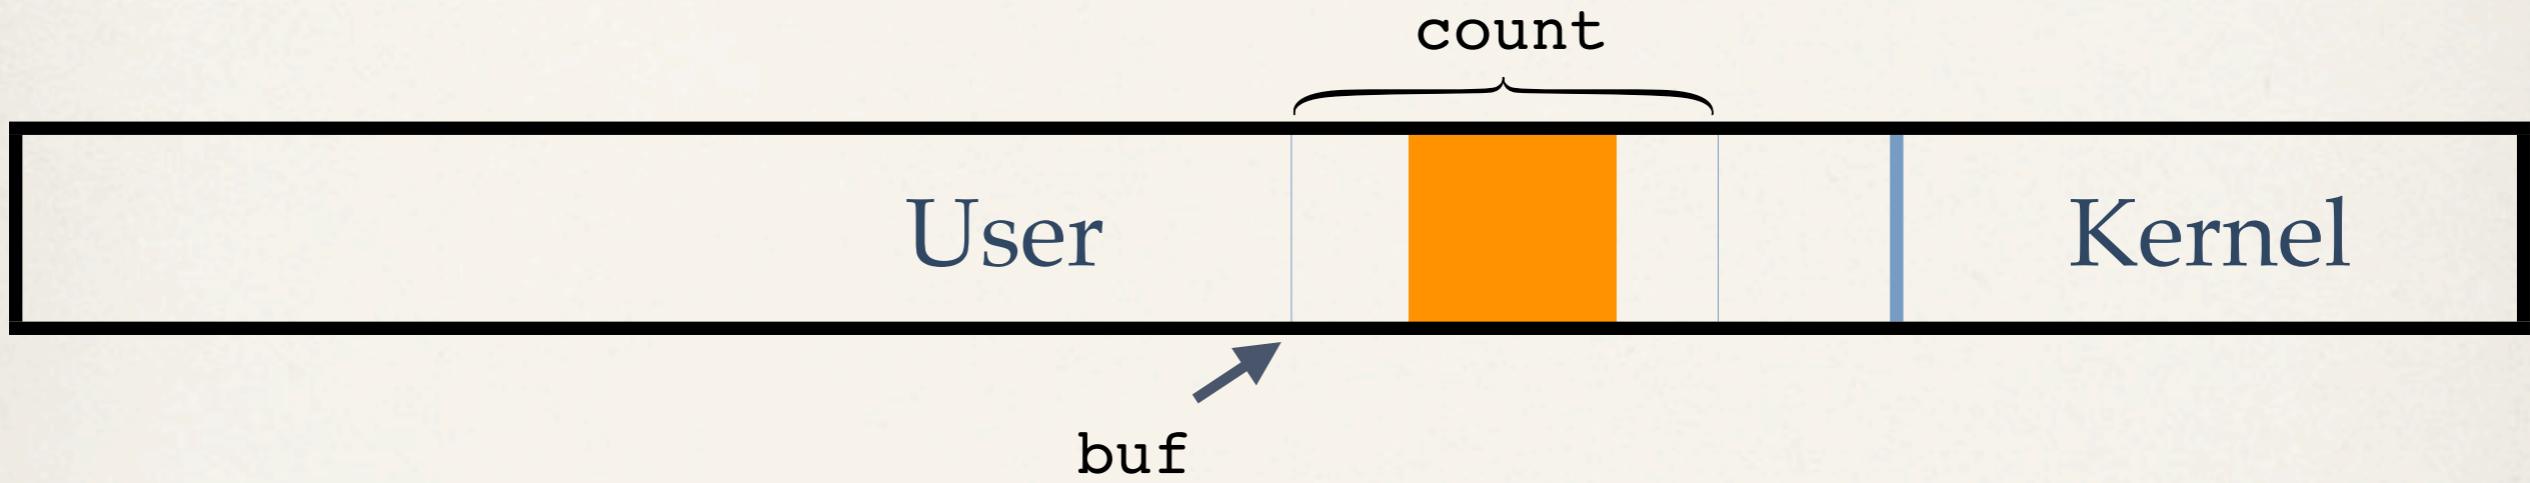
```
asmlinkage long  
sys_read(unsigned int fd, char __user *buf, size_t count);
```





# Threat model: example

```
asmlinkage long  
sys_read(unsigned int fd, char __user *buf, size_t count);
```



- Write arbitrary data, but only inside the supplied buffer
- Arbitrary return value

# Abstraction



- ❖ Malicious kernel (modified Linux)
  - ❖ No reading/writing application memory
  - ❖ Handle all “unsafe” system calls correctly
  - ❖ Can handle “safe” system calls maliciously
- ❖ Unmodified user space



# Recall our vulnerable program

---



```
#include <stdlib.h>

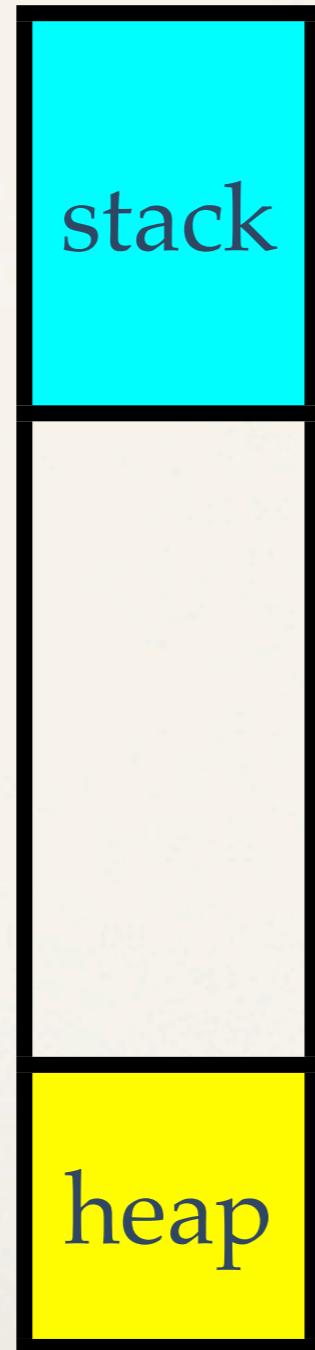
int main() {
    void *p = malloc(100);
}
```

# Step 1: mmap(2)/read(2); normal behavior

---



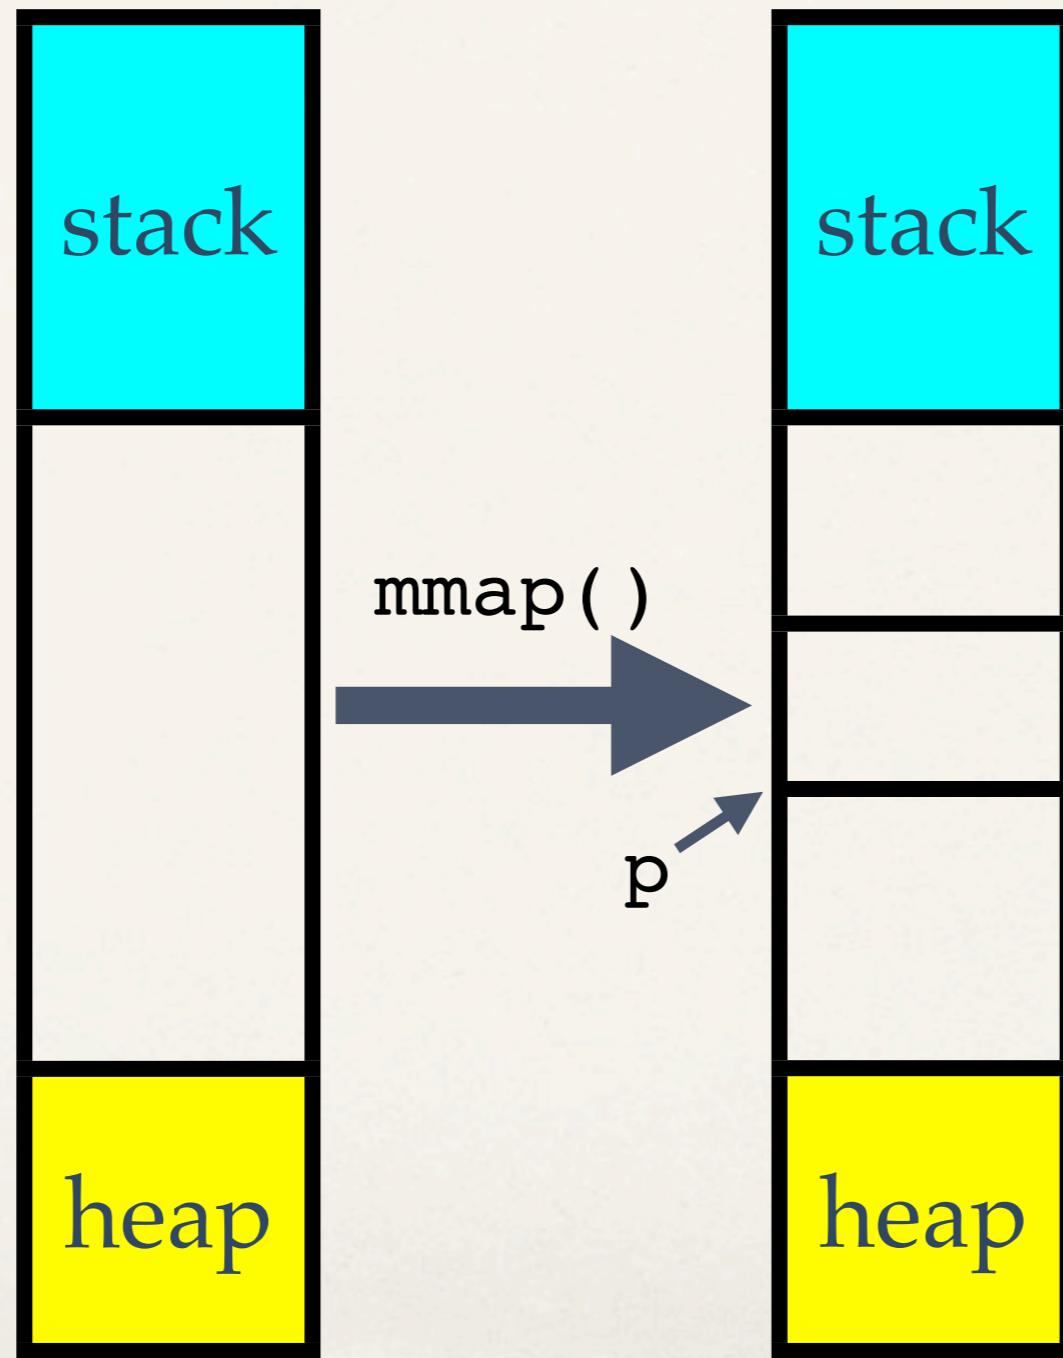
```
void *p;  
p = mmap(4096);  
read(0,p,4096);
```



# Step 1: mmap( 2 )/read( 2 ); normal behavior



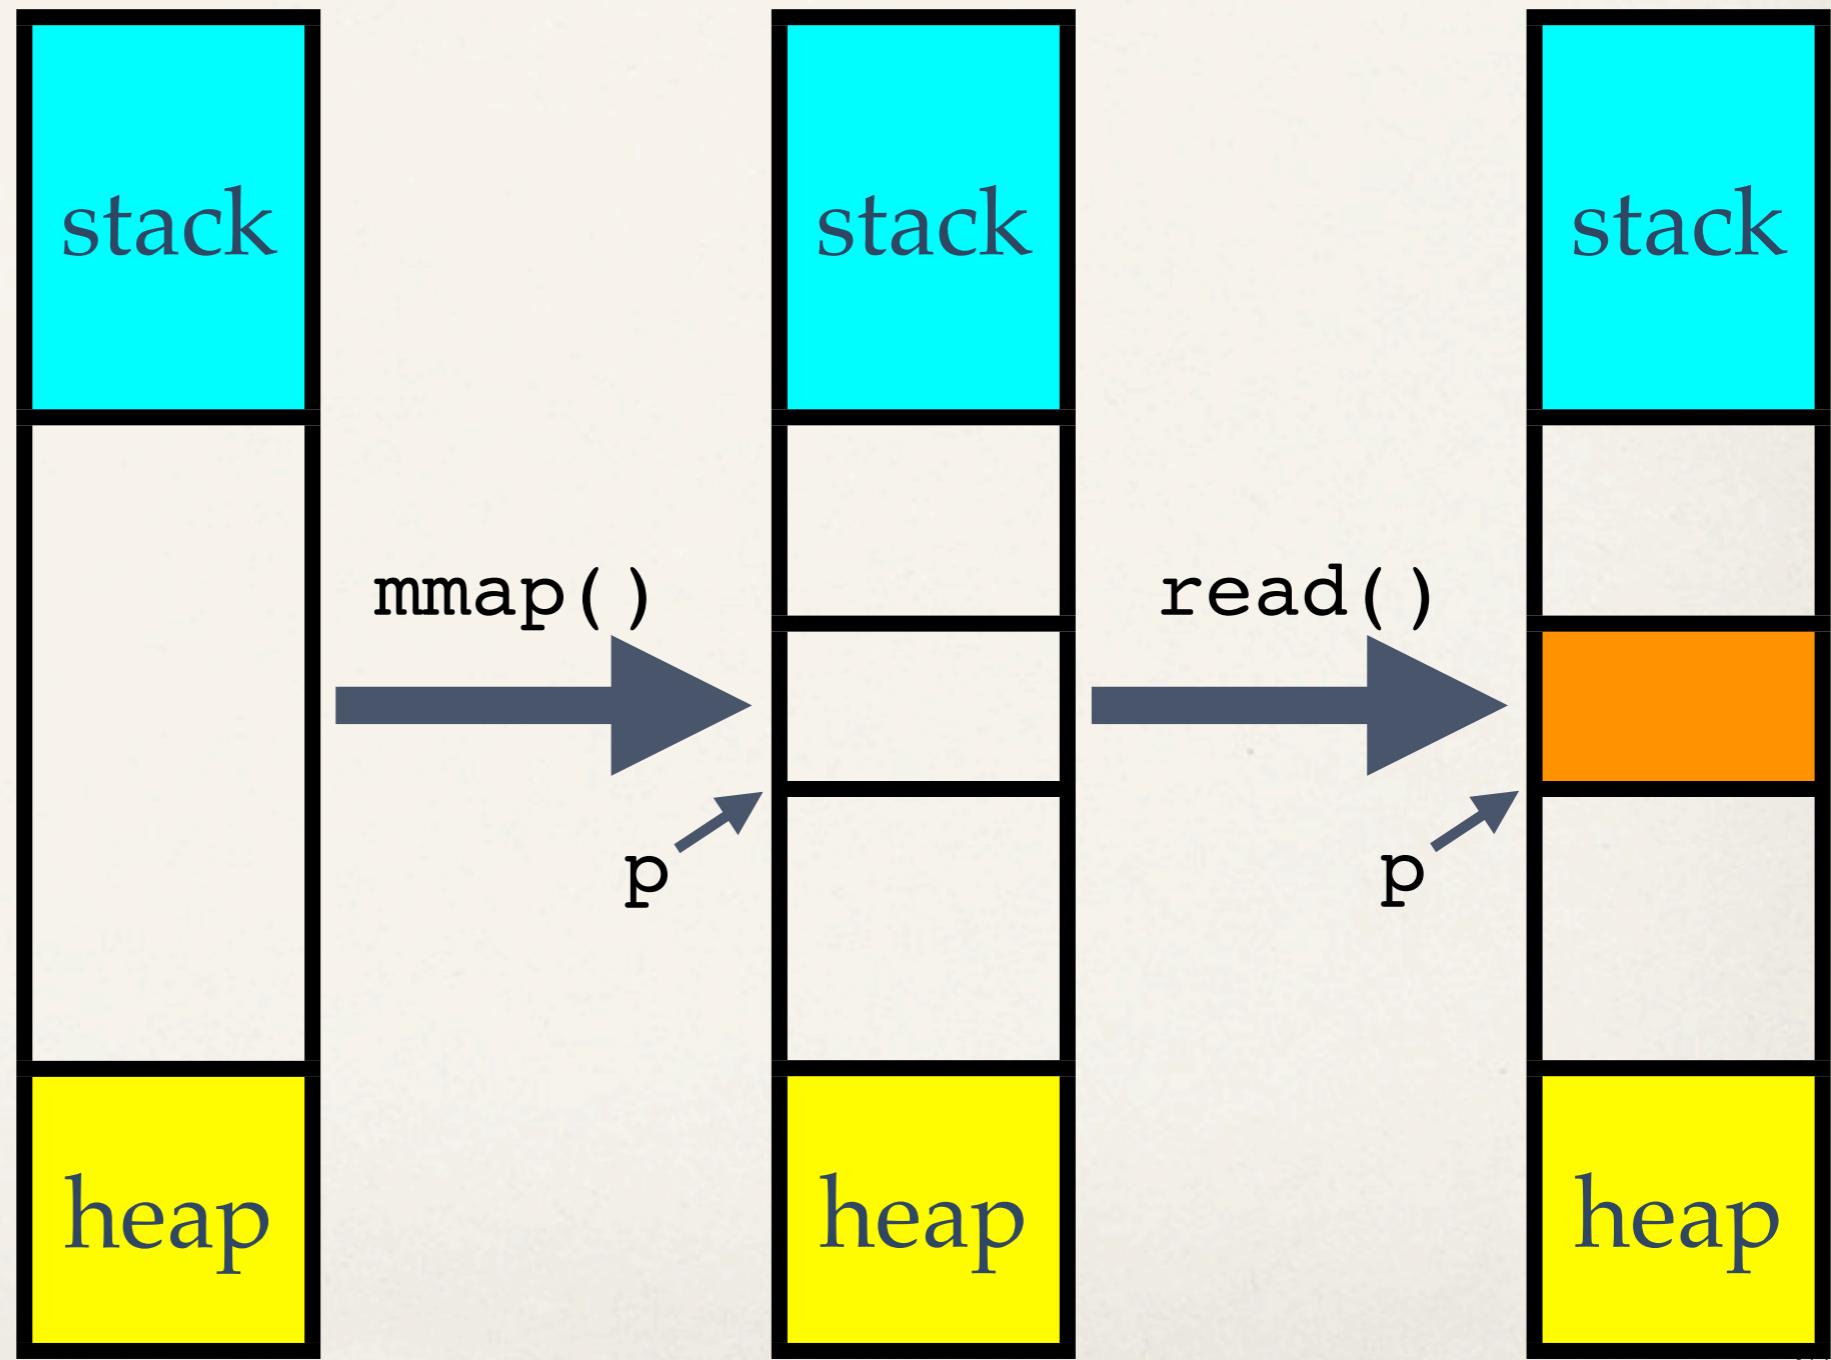
```
void *p;  
p = mmap( 4096 );  
read( 0, p, 4096 );
```



# Step 1: mmap( 2 )/read( 2 ); normal behavior



```
void *p;  
p = mmap( 4096 );  
read( 0, p, 4096 );
```

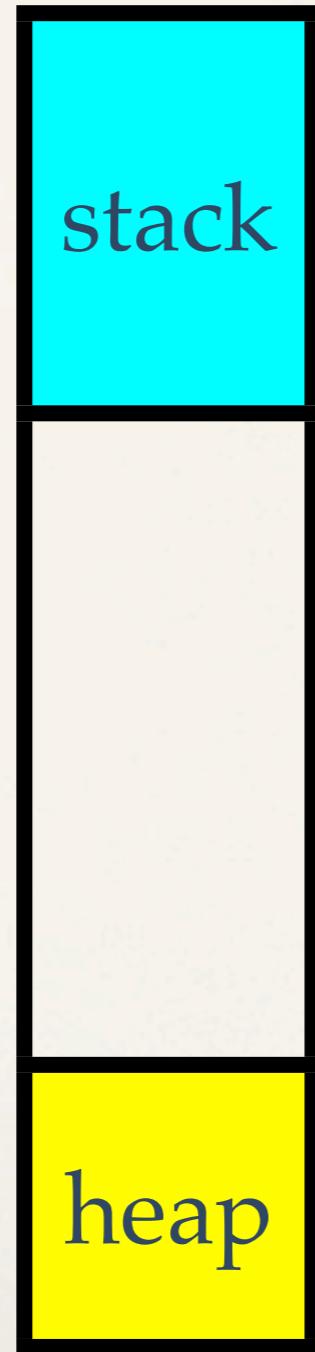


# Step 1: mmap(2)/read(2); malicious behavior

---



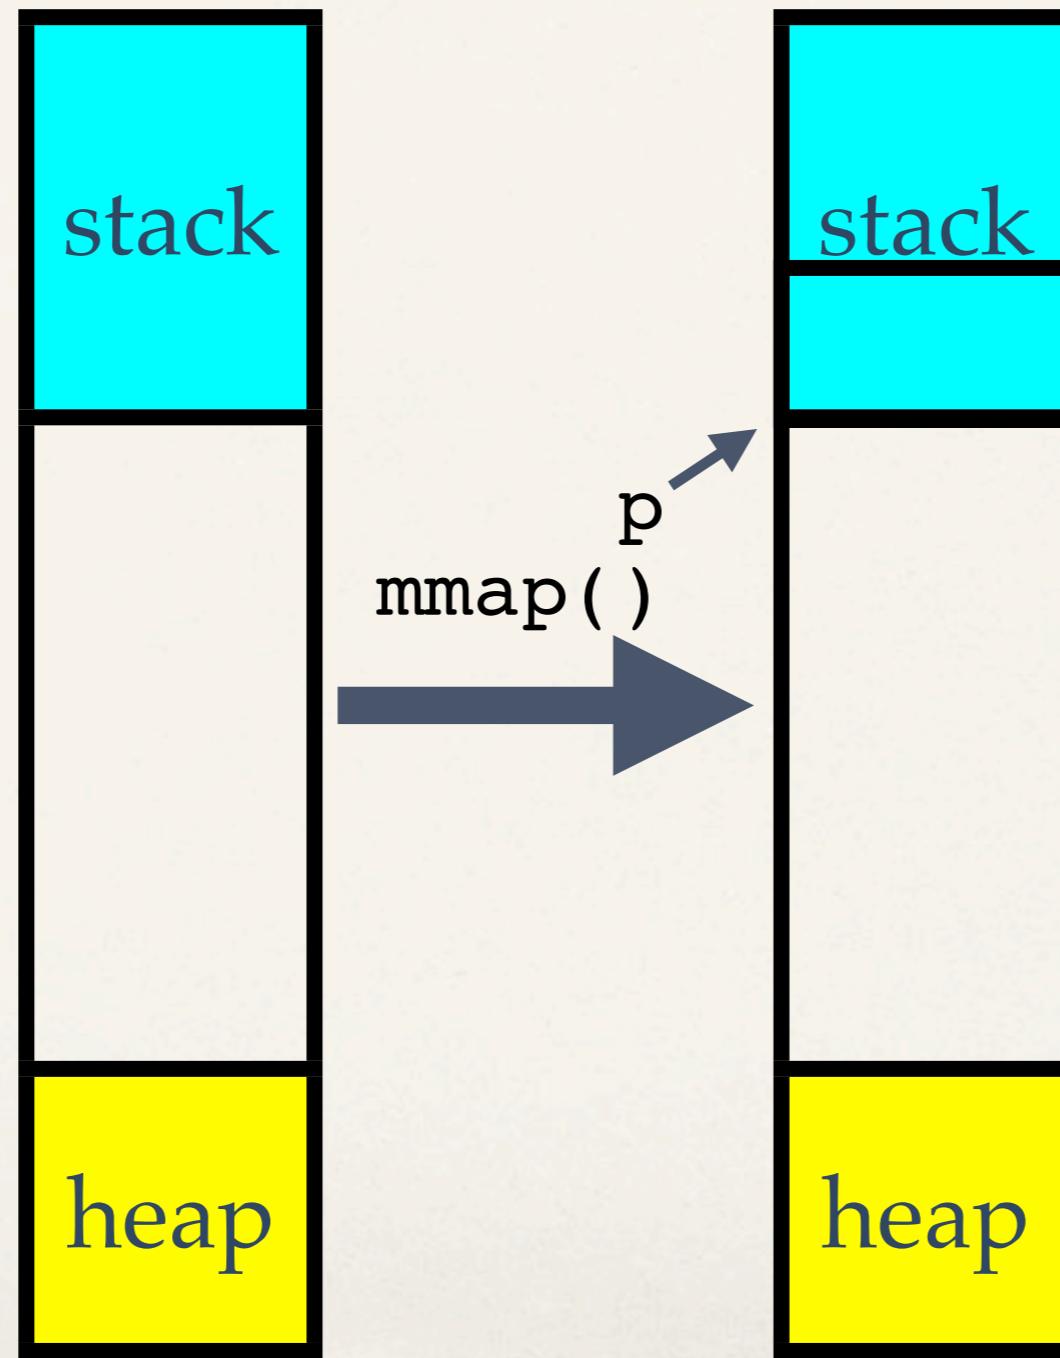
```
void *p;  
p = mmap(4096);  
read(0,p,4096);
```



# Step 1: mmap( 2 )/read( 2 ); malicious behavior



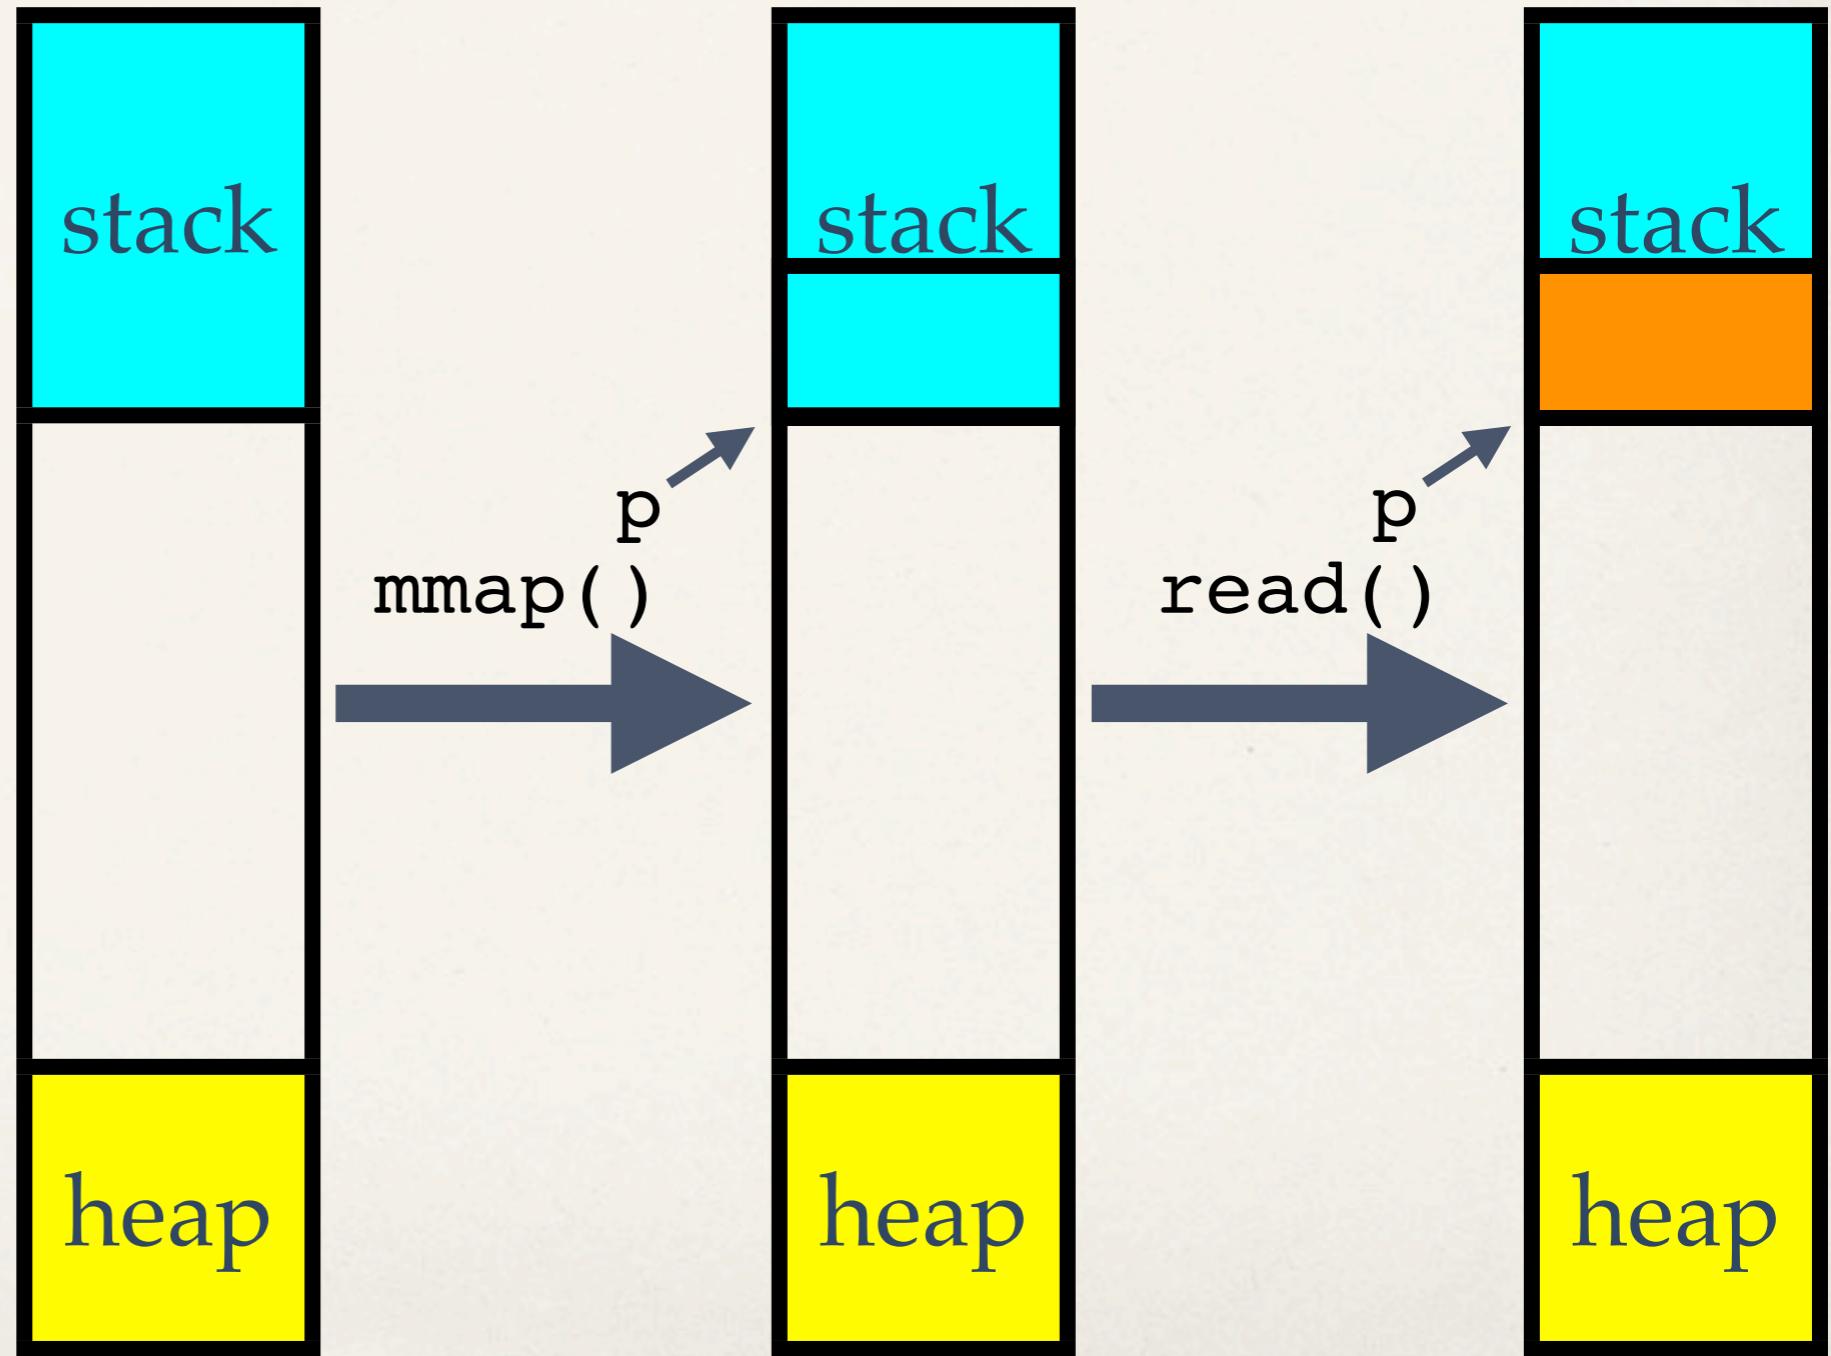
```
void *p;  
p = mmap( 4096 );  
read( 0, p, 4096 );
```



# Step 1: mmap( 2 )/read( 2 ); malicious behavior



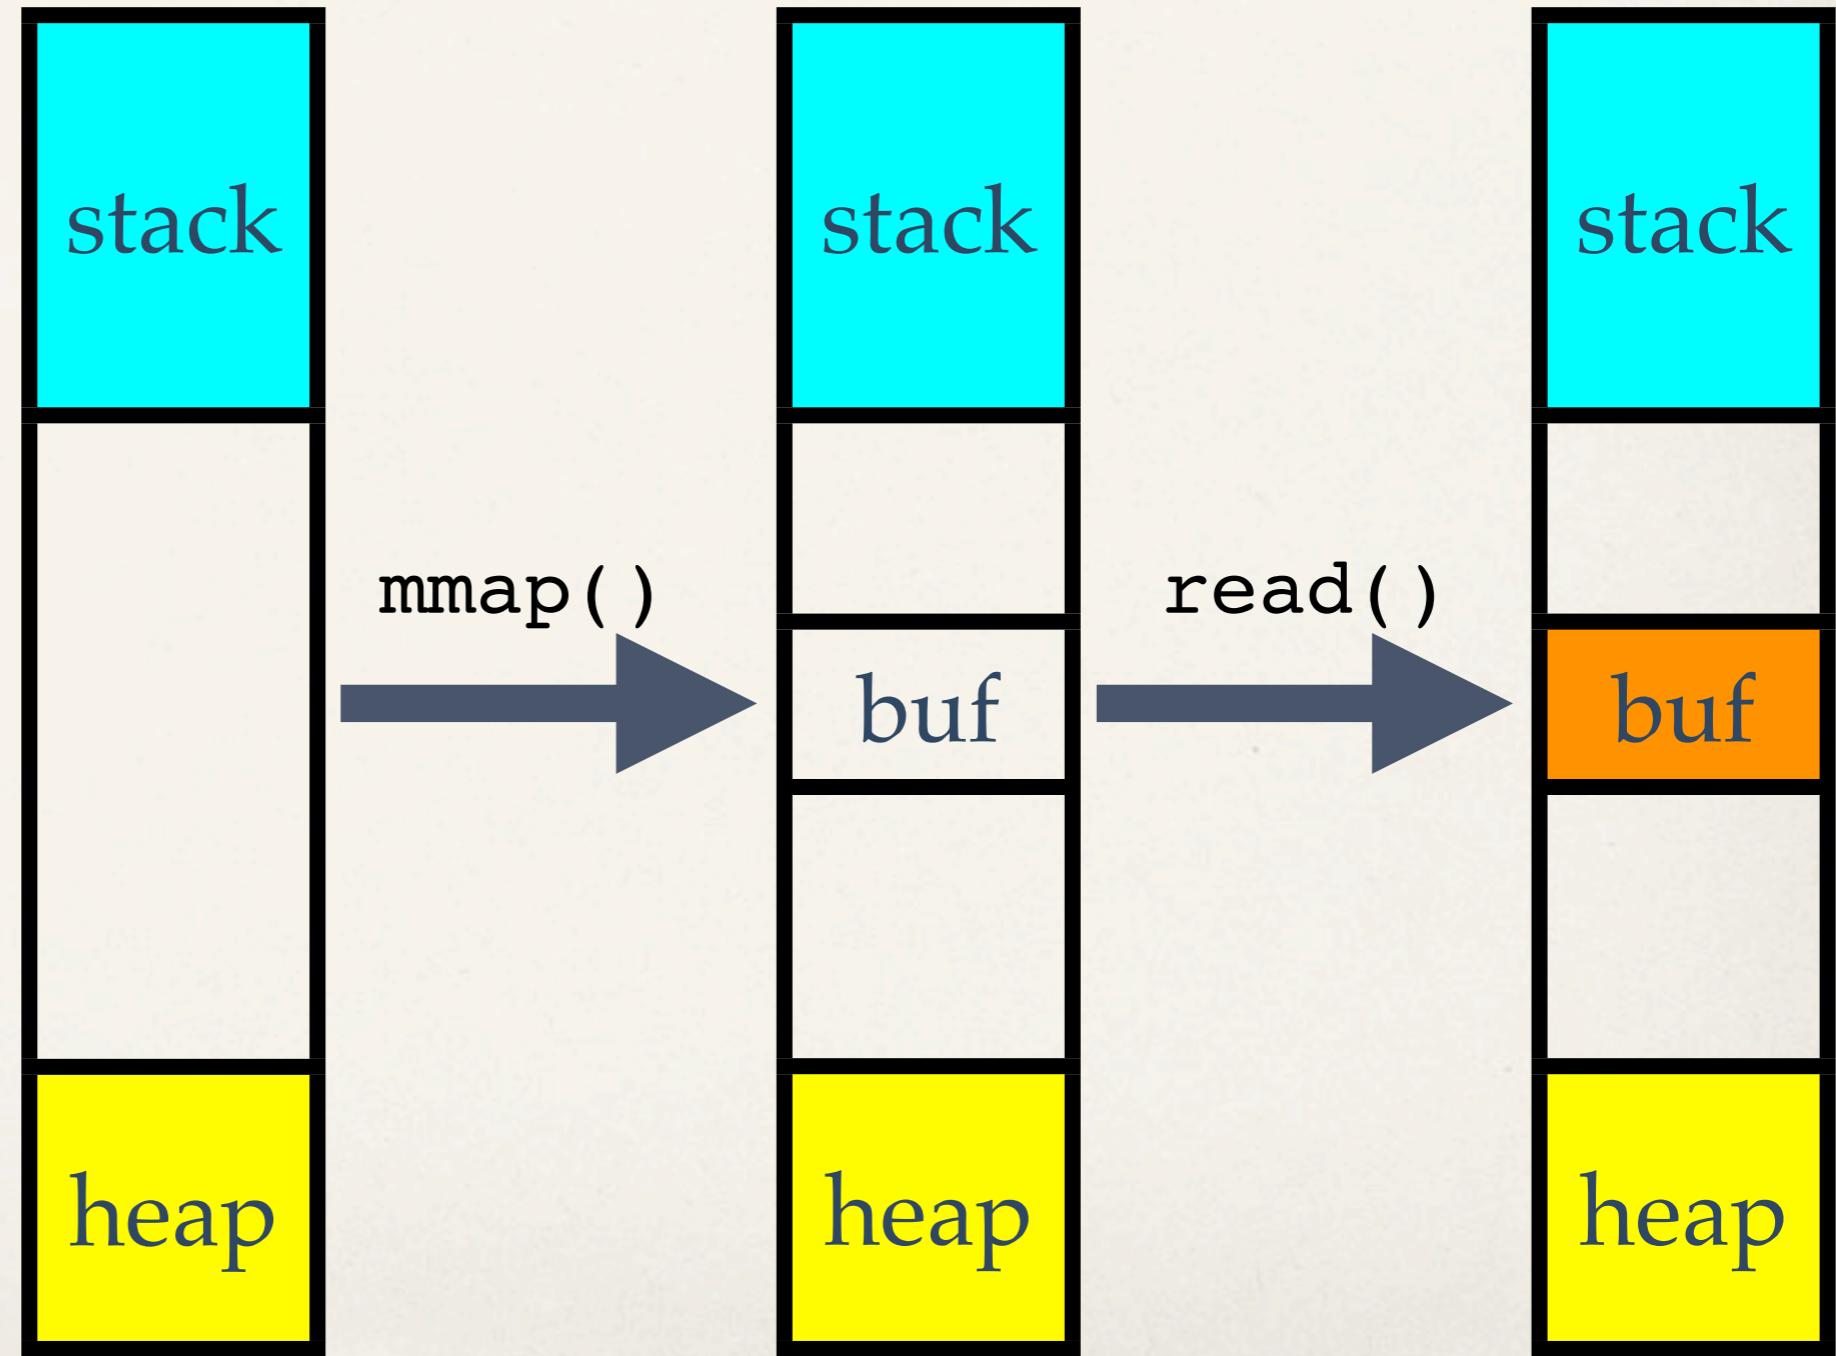
```
void *p;  
p = mmap( 4096 );  
read( 0, p, 4096 );
```



# Step 2: Standard I/O; normal behavior



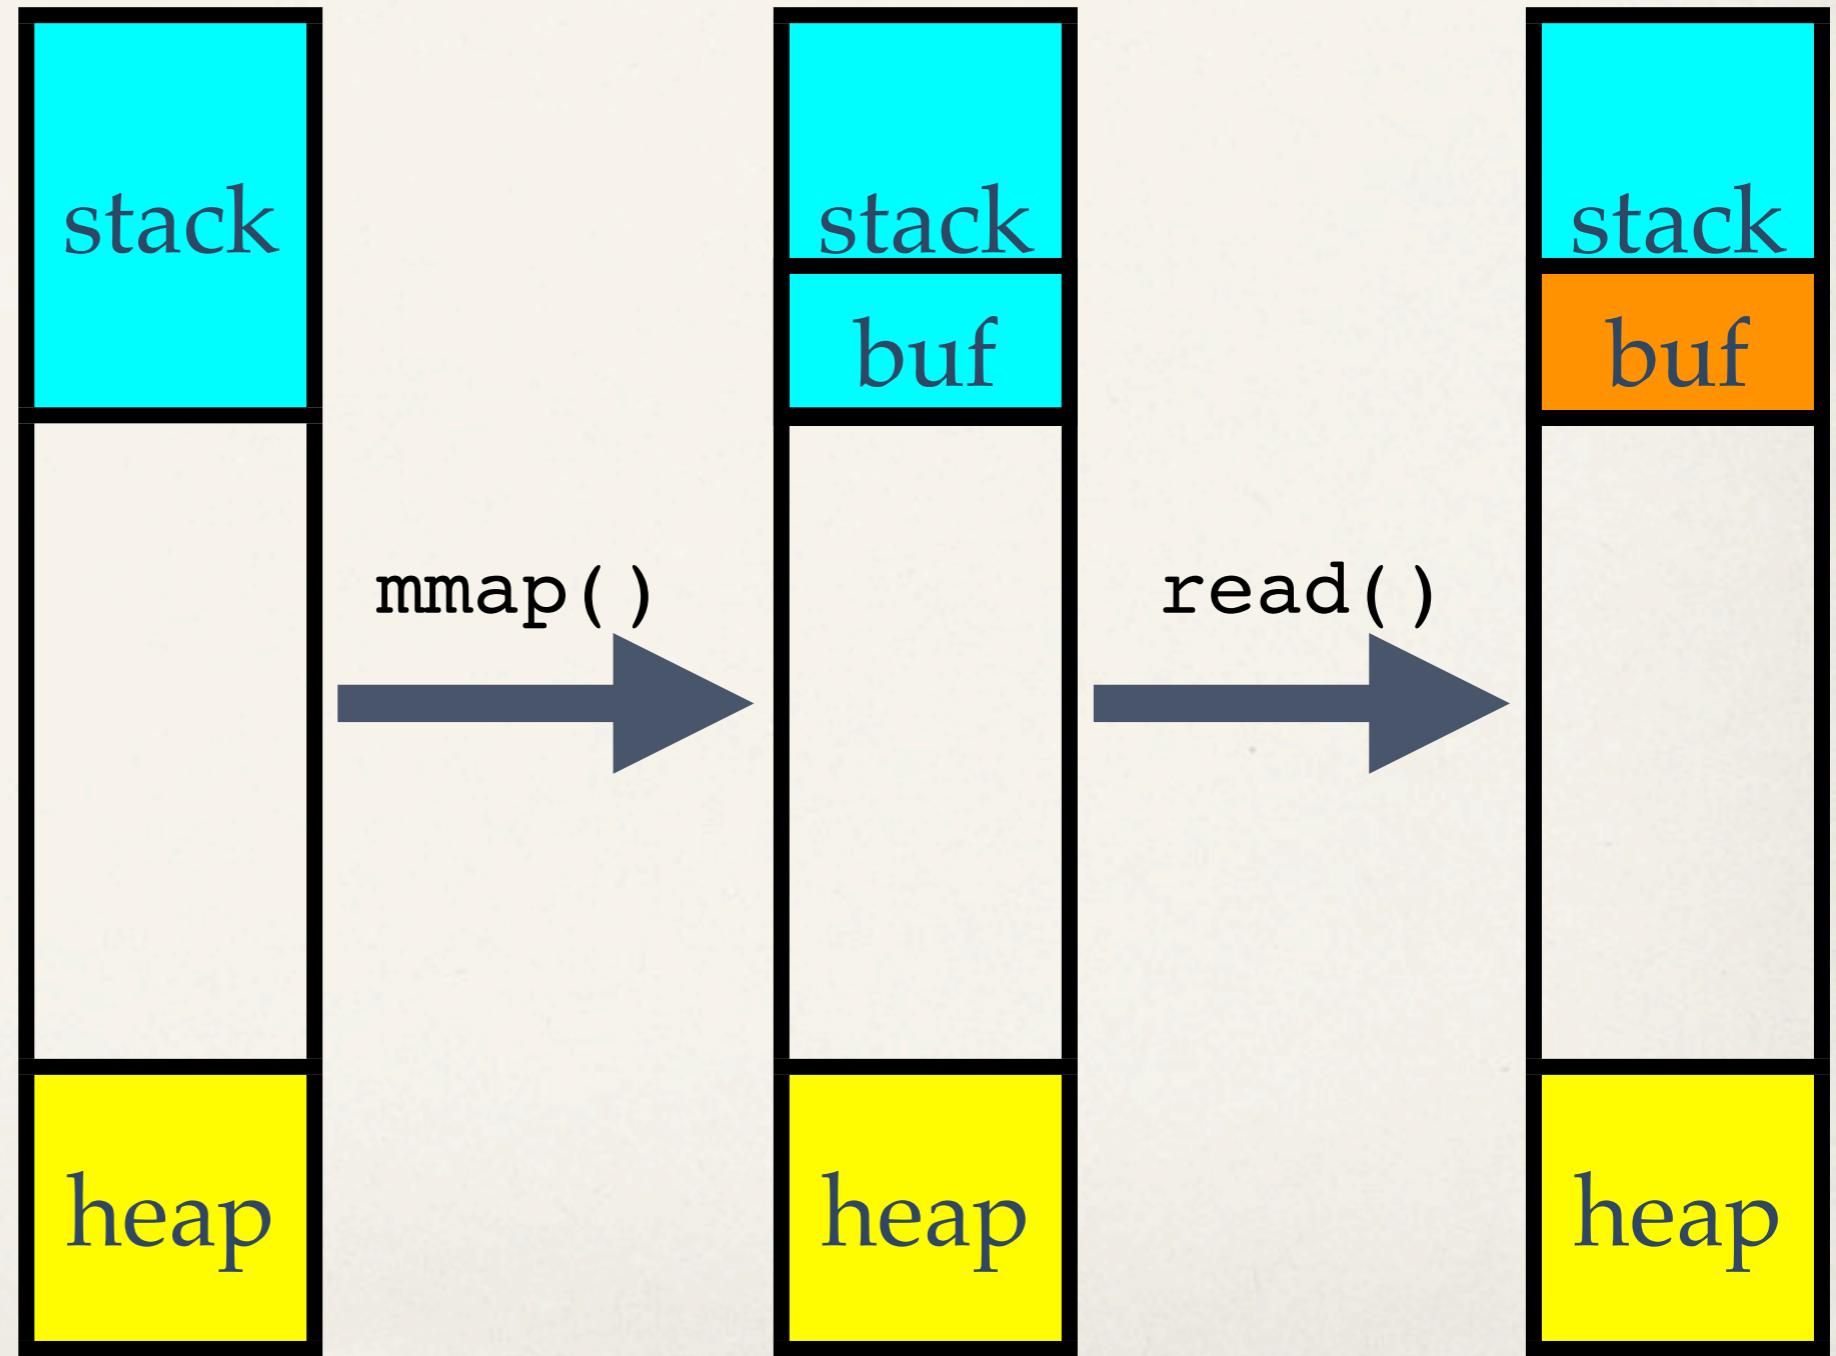
- `fgetc()`
- `fgets()`
- `freopen()`
- `fscanf()`
- `getc()`
- `getchar()`
- `getdelim()`
- `getline()`
- **`gets()`**
- `scanf()`
- `vfscanf()`
- `vscanf()`
- ...



# Step 2: Standard I/O; malicious behavior

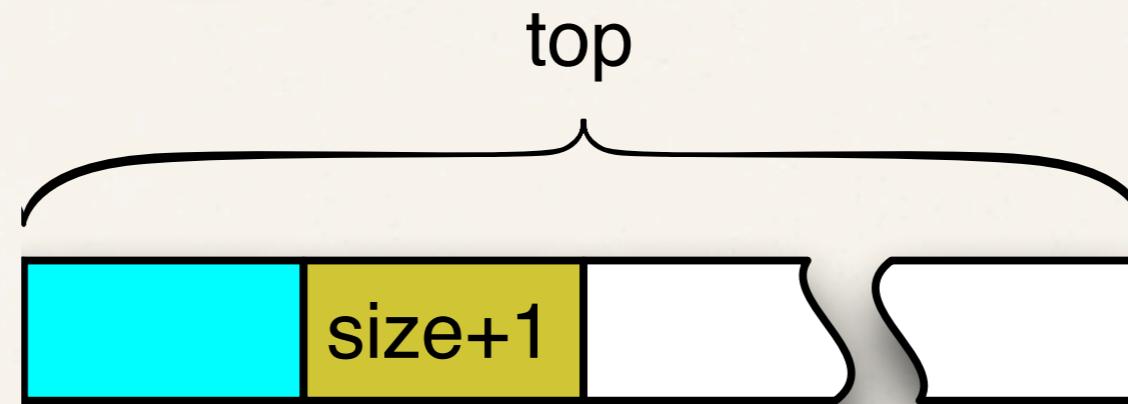


- `fgetc()`
- `fgets()`
- `freopen()`
- `fscanf()`
- `getc()`
- `getchar()`
- `getdelim()`
- `getline()`
- **`gets()`**
- `scanf()`
- `vfscanf()`
- `vscanf()`
- ...



20

# Step 3: LibC's malloc



- ❖ Split into upper and lower halves
  - ❖ Upper half: manages chunks, free lists, handles `malloc()` and `free()`
  - ❖ **Lower half: requests memory from the OS**
- ❖ Maintains a top region of unallocated memory from the OS
  - ❖ Metadata (including size) inline

# The lower half algorithm



---

First call to `malloc(n)` [creating the top chunk]:

1.  $nb \leftarrow n + 4$  rounded up to a multiple of 8 bytes
2. Determine the start of the heap via **brk** system call
3. Increase the size of the heap via **brk**
4. Increase the size again to maintain 8-byte alignment via **brk**  
(updates the start  $S$  of the heap)
5. If step 4 failed, determine the end  $E$  of the heap (last **brk**'s return value)
6. Carve off a chunk of size  $nb$
7. Write the size  $E - S - nb$  of the remaining top chunk at  $S + nb + 4$

# malloc( n ) example



- 
1.  $nb \leftarrow n + 4$  rounded up to a multiple of 8 bytes

# malloc( n ) example



2. Determine the start of the heap via **brk** system call



# malloc( n ) example



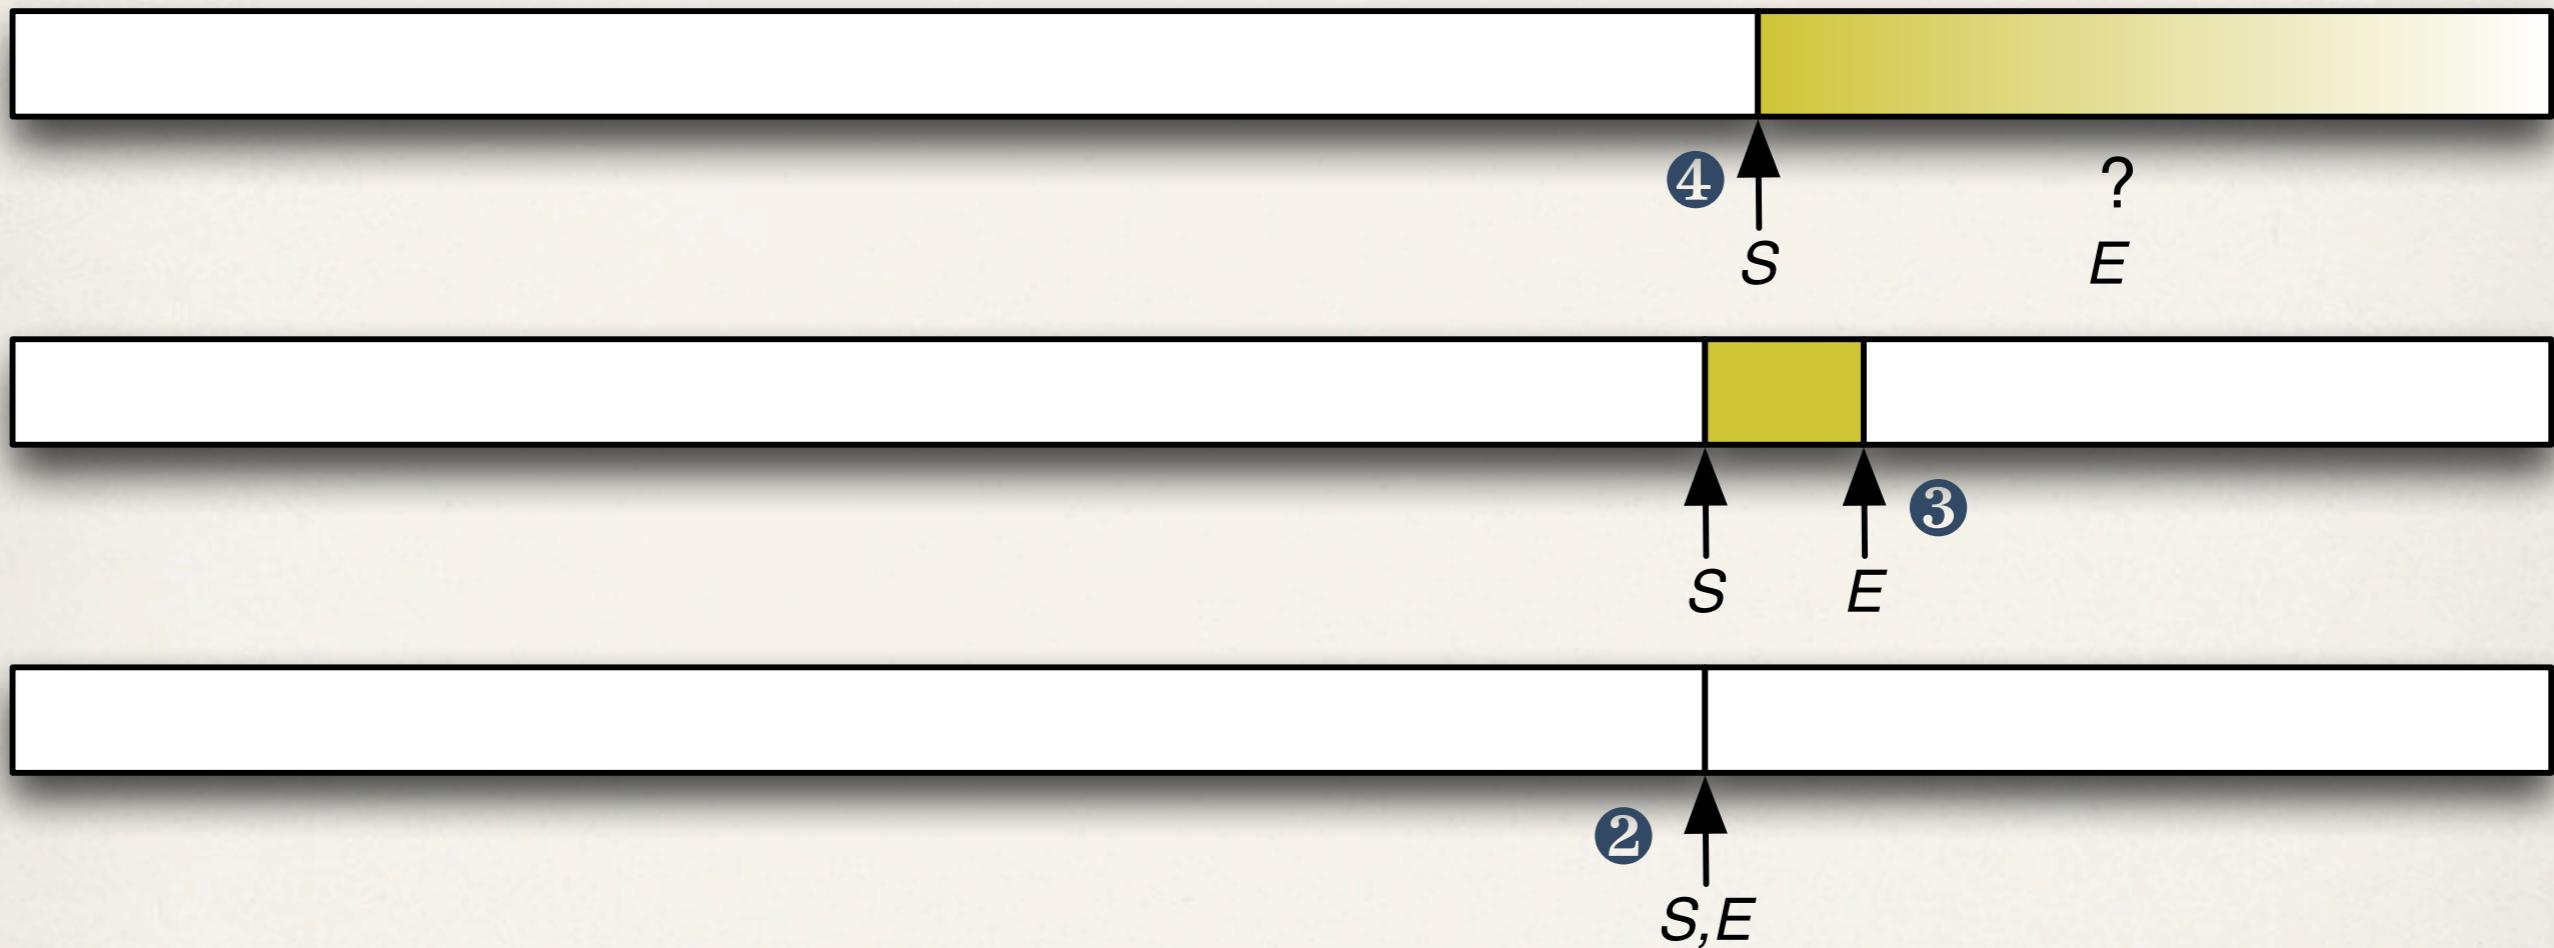
3. Increase the size of the heap via **brk**



# malloc( n ) example



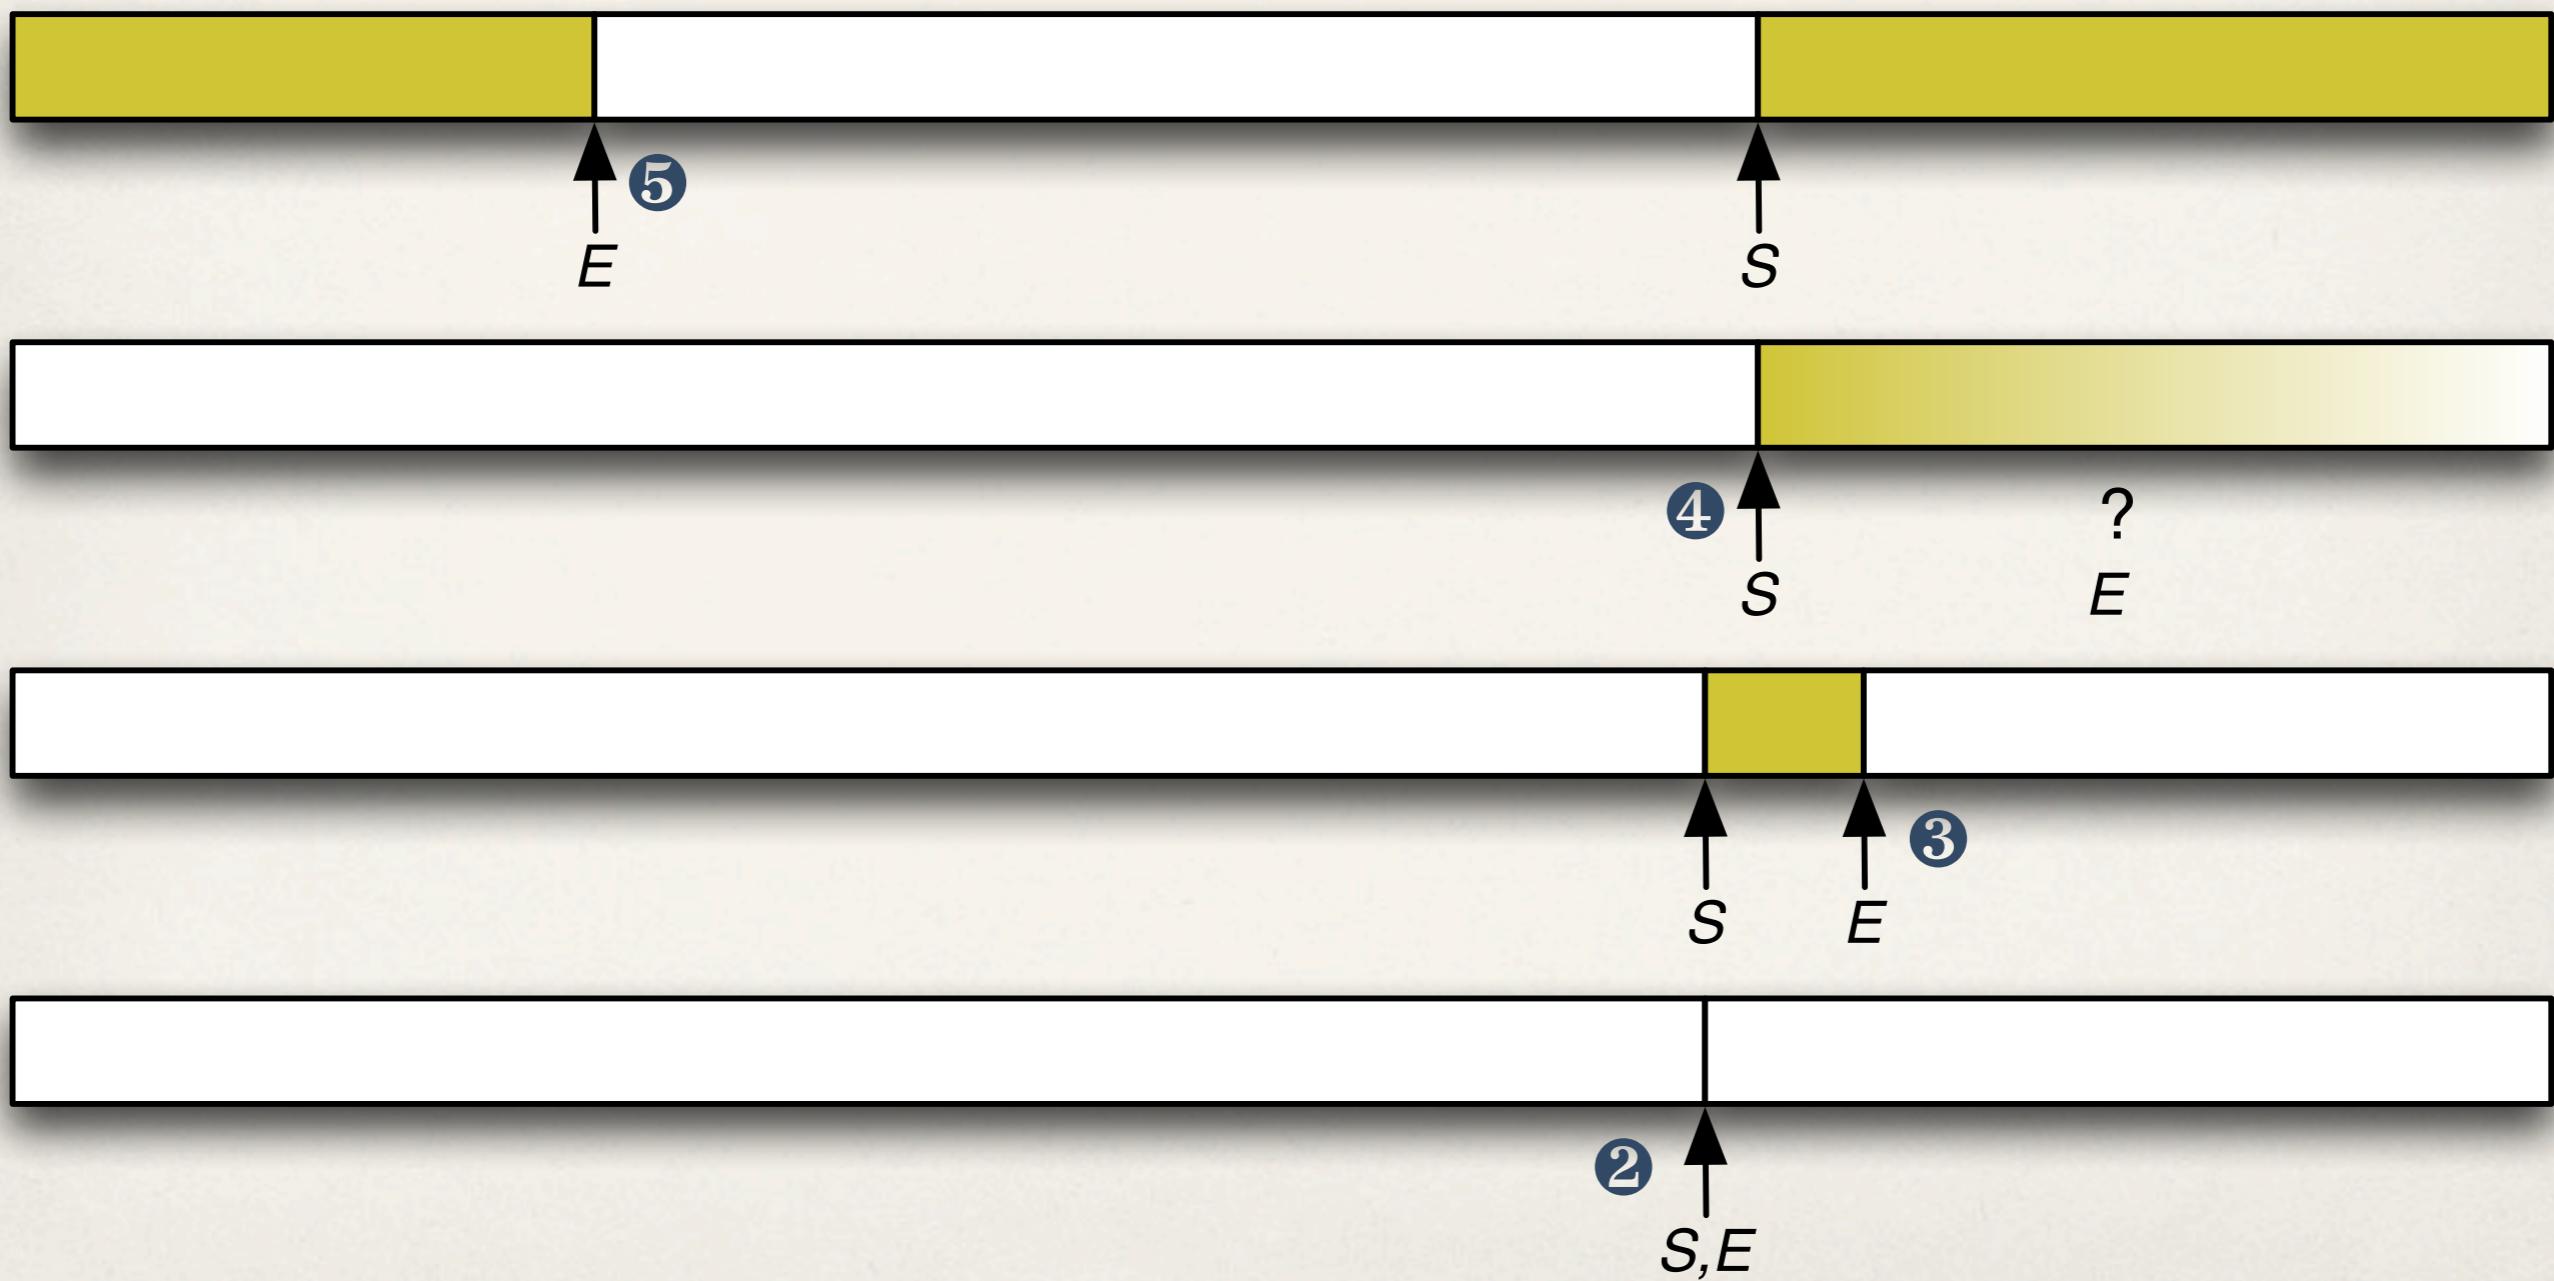
4. Increase the size again to maintain 8-byte alignment via **brk**





# malloc( n ) example

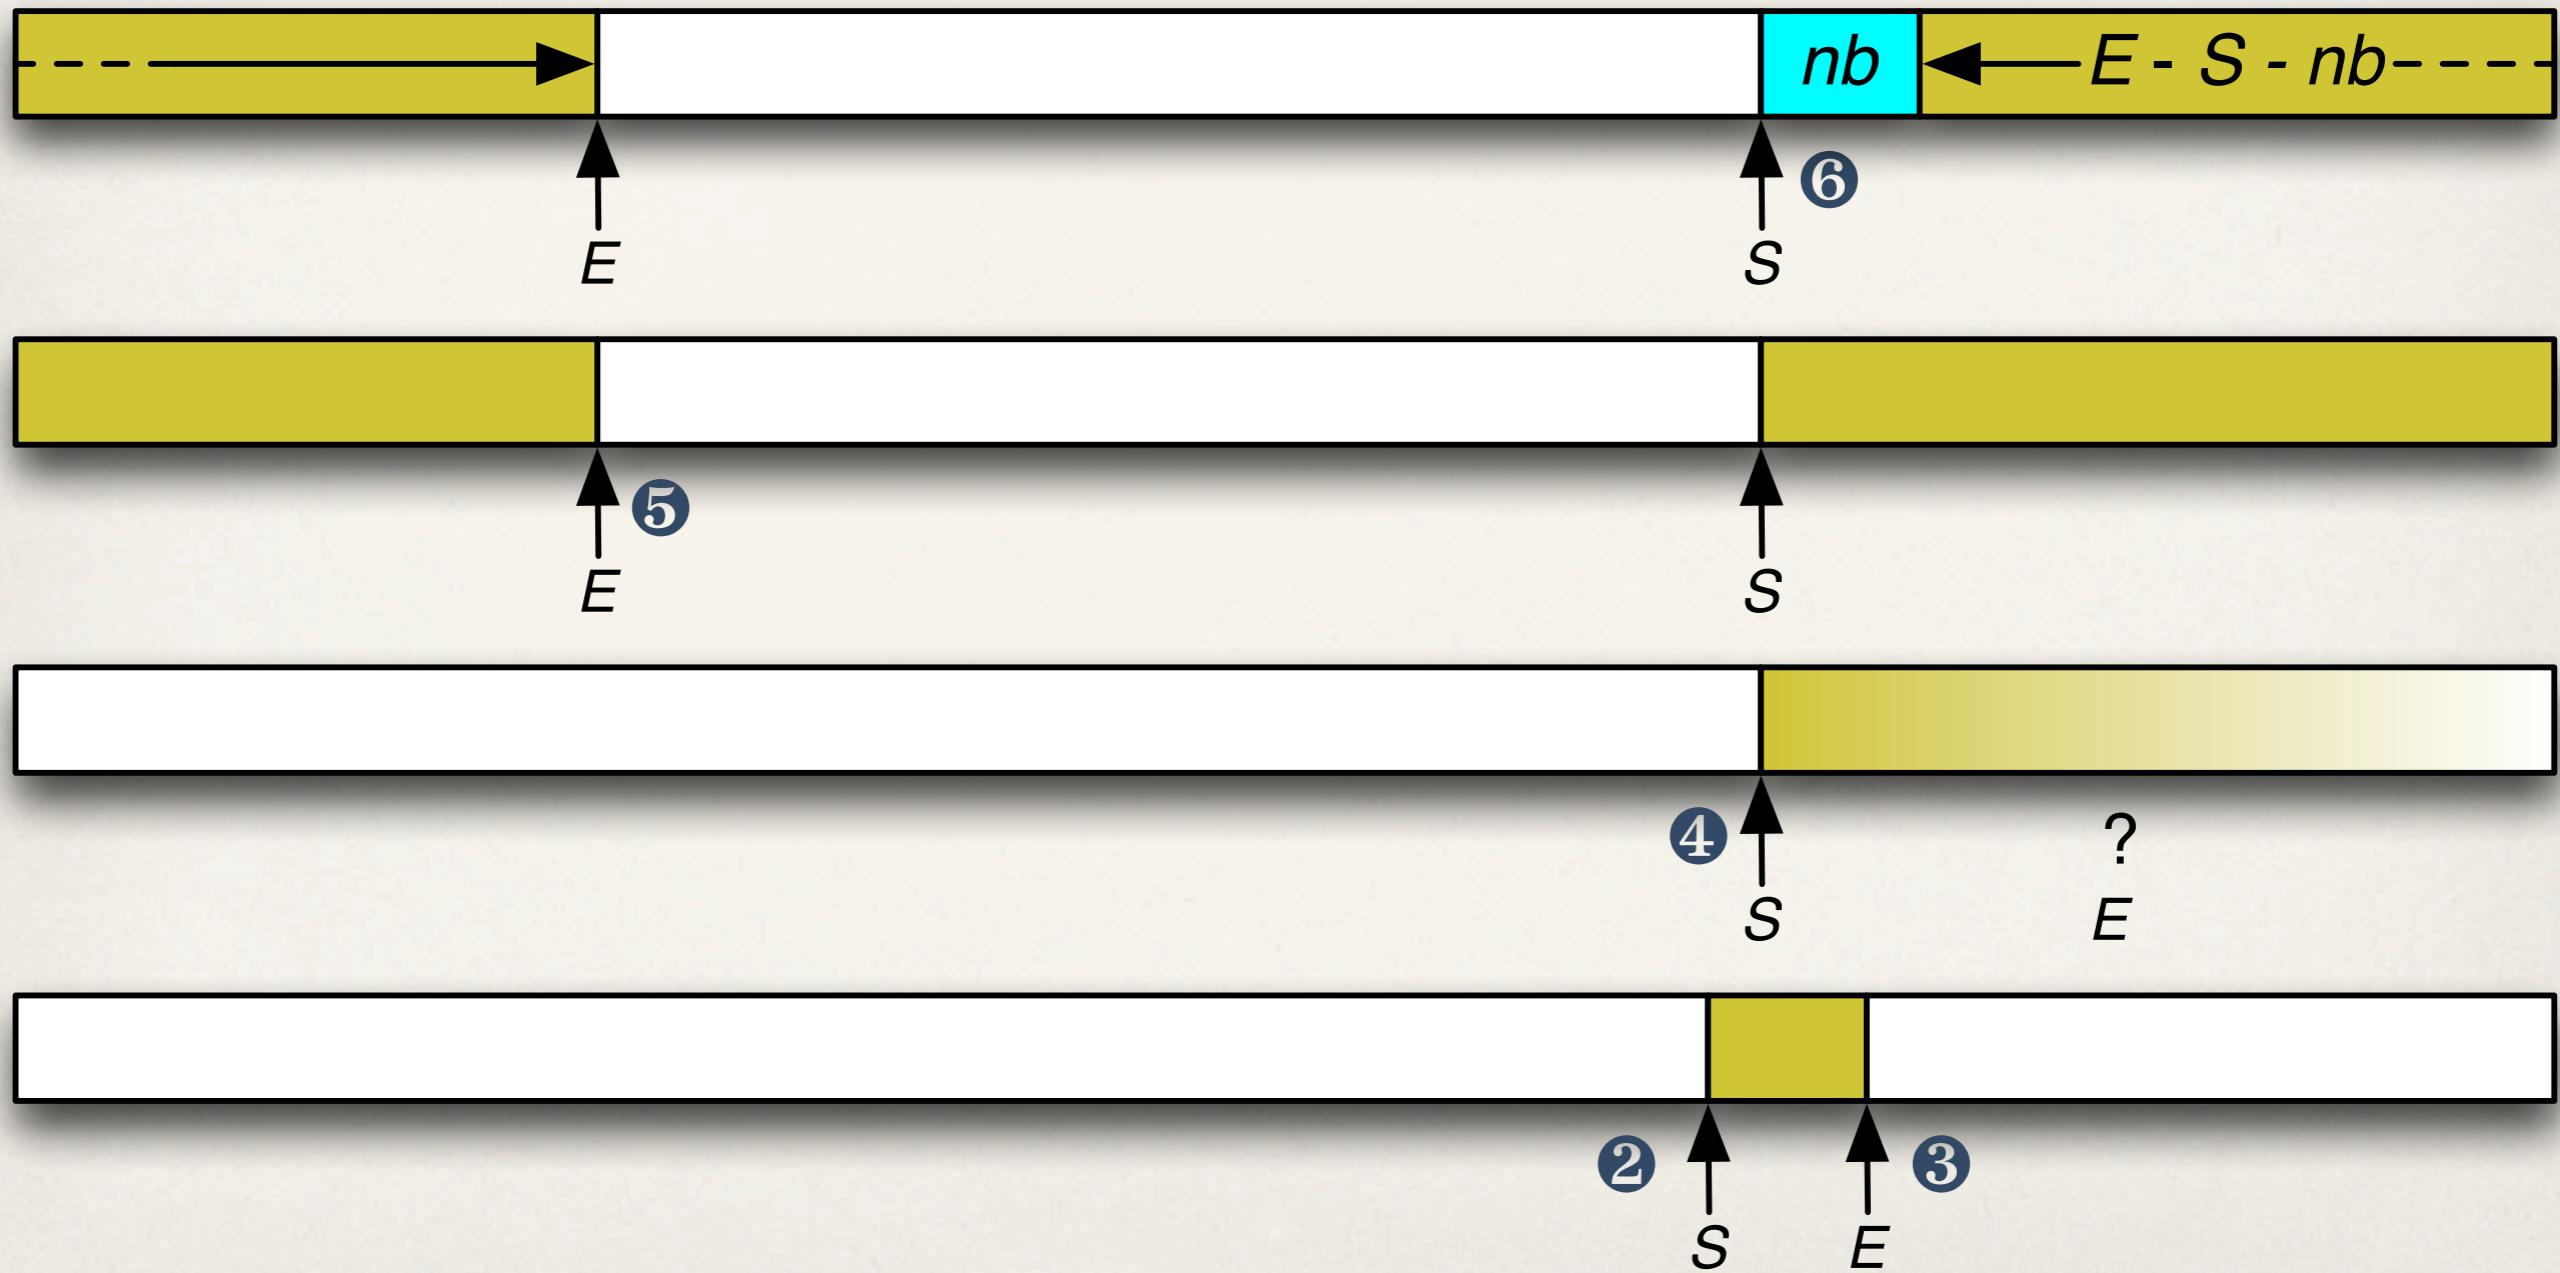
5. If step 4 failed, determine the end  $E$  of the heap (last brk's return value)



# malloc( n ) example



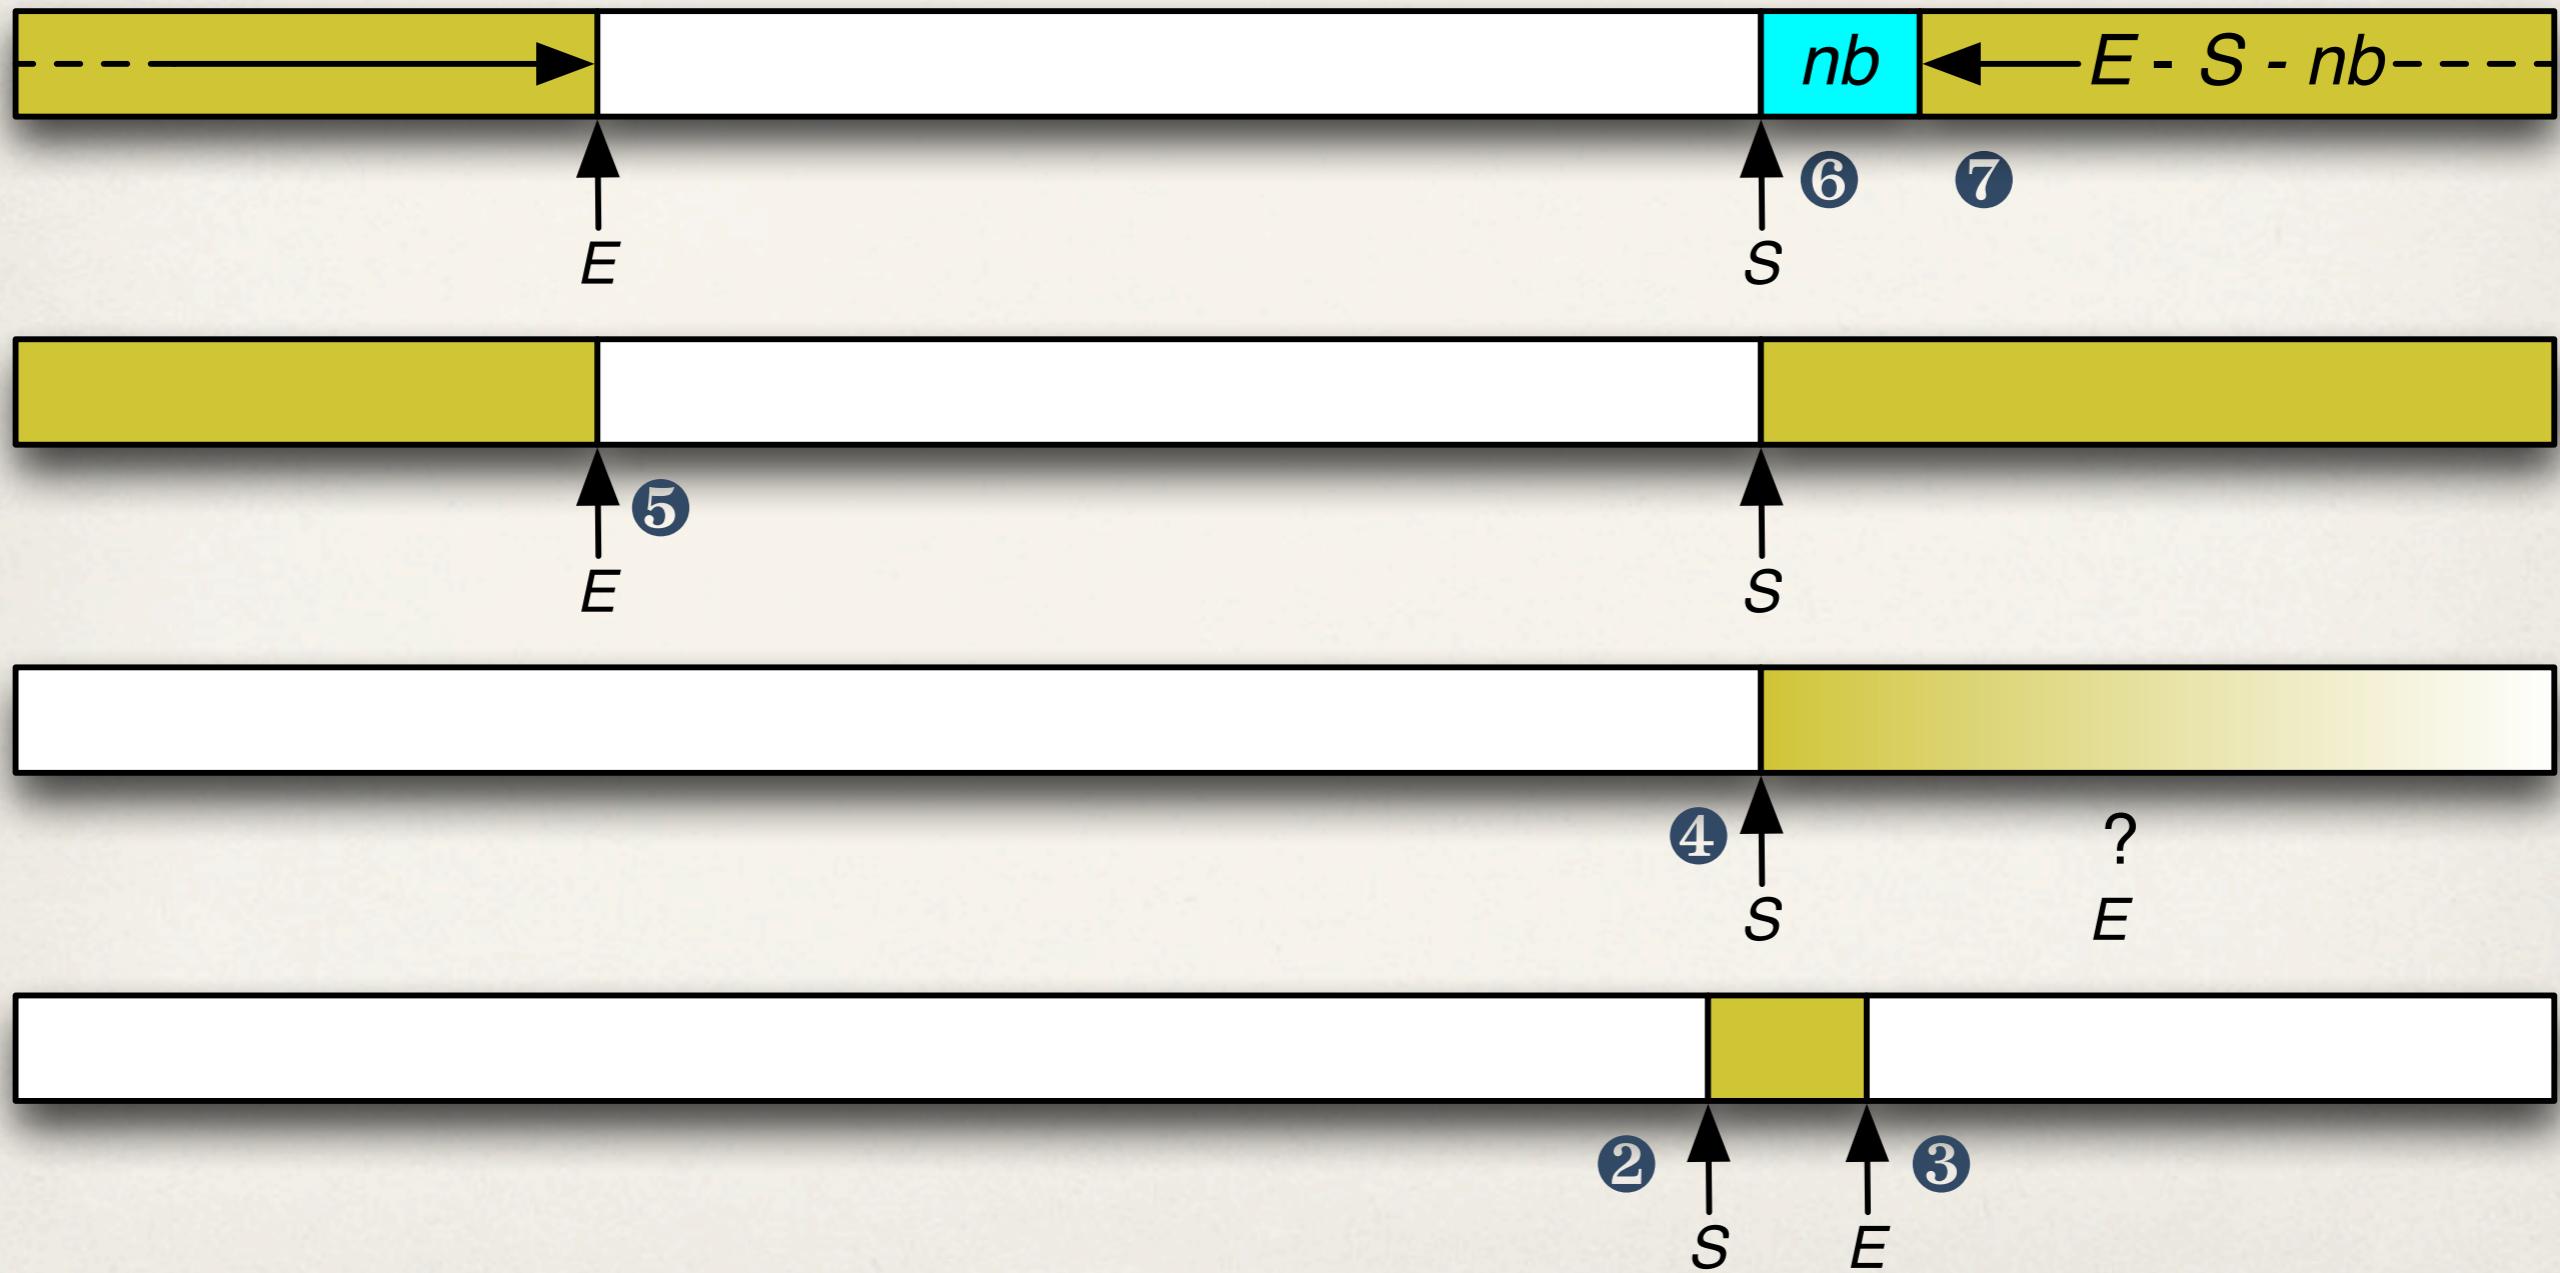
6. Carve off a chunk of size  $nb$



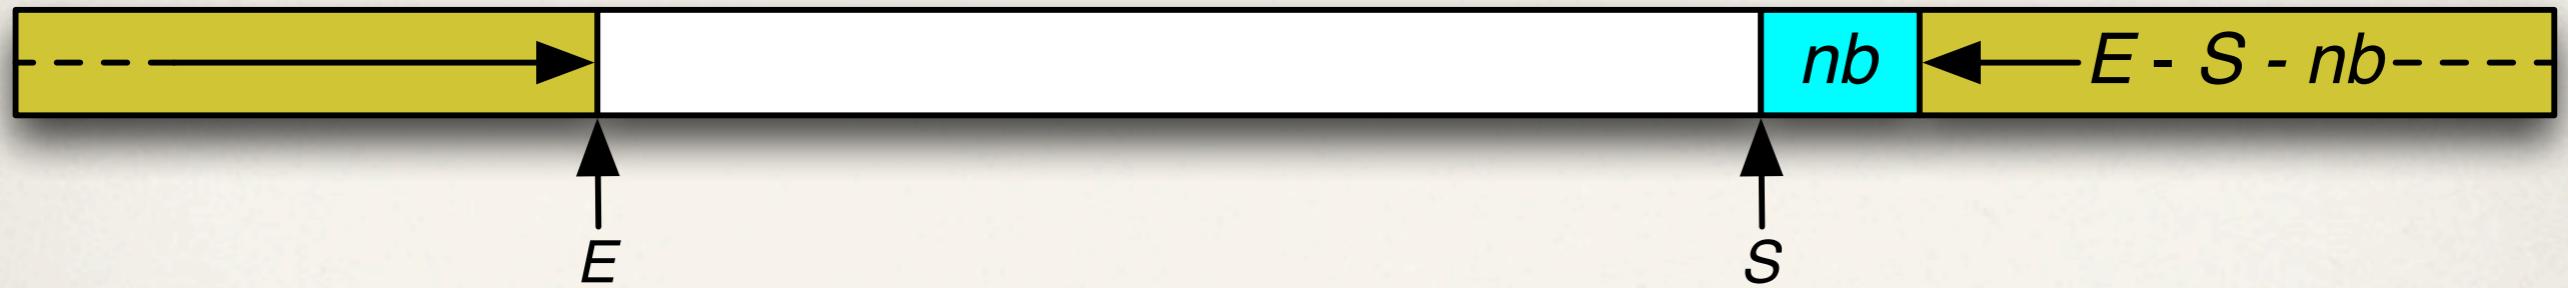
# malloc( n ) example



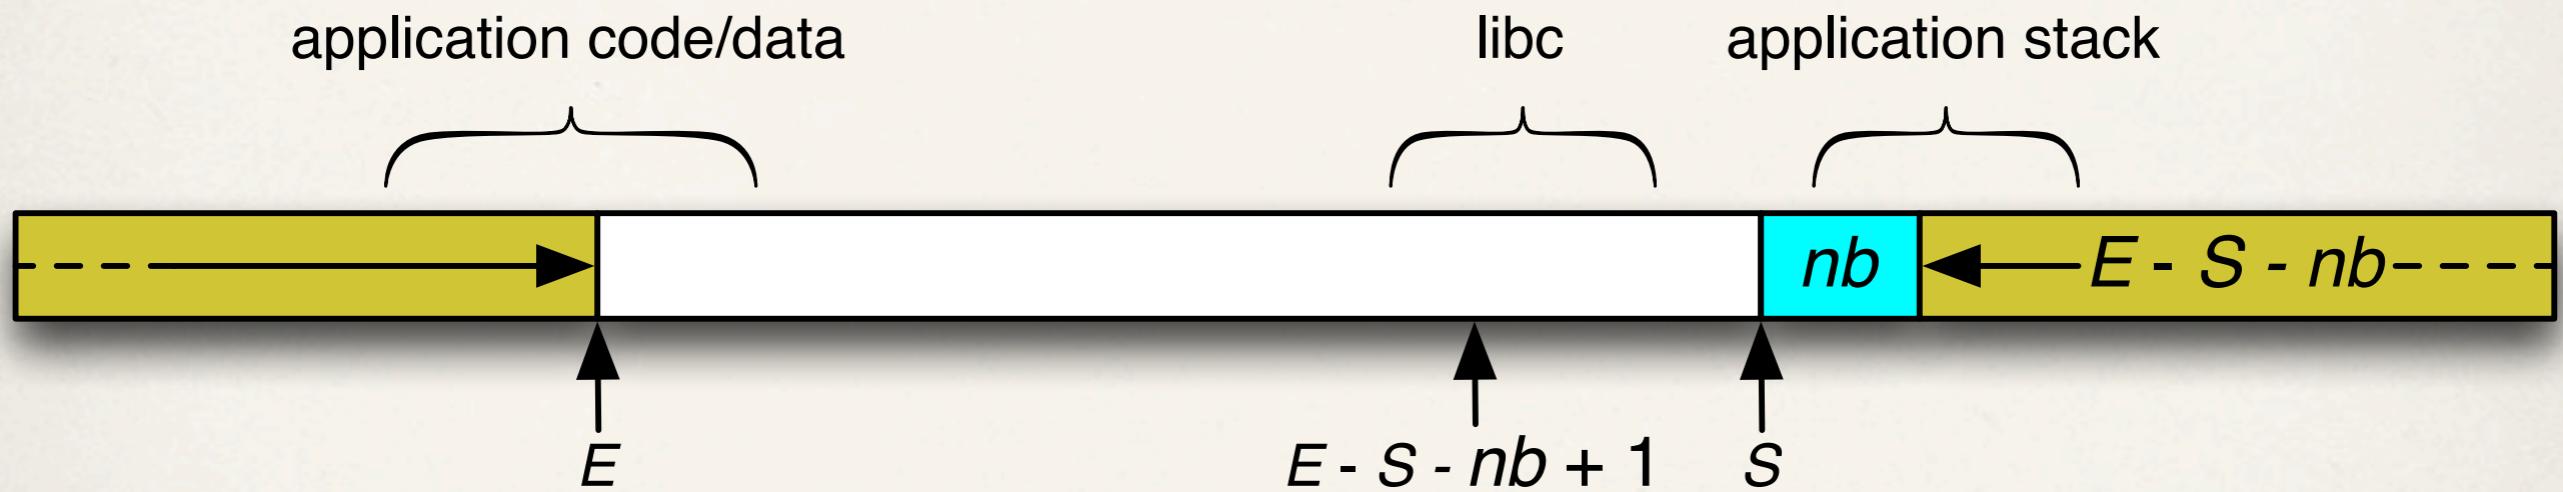
7. Write the size  $E - S - nb$  of the remaining top chunk at  $S + nb + 4$



# Attacking the lower half

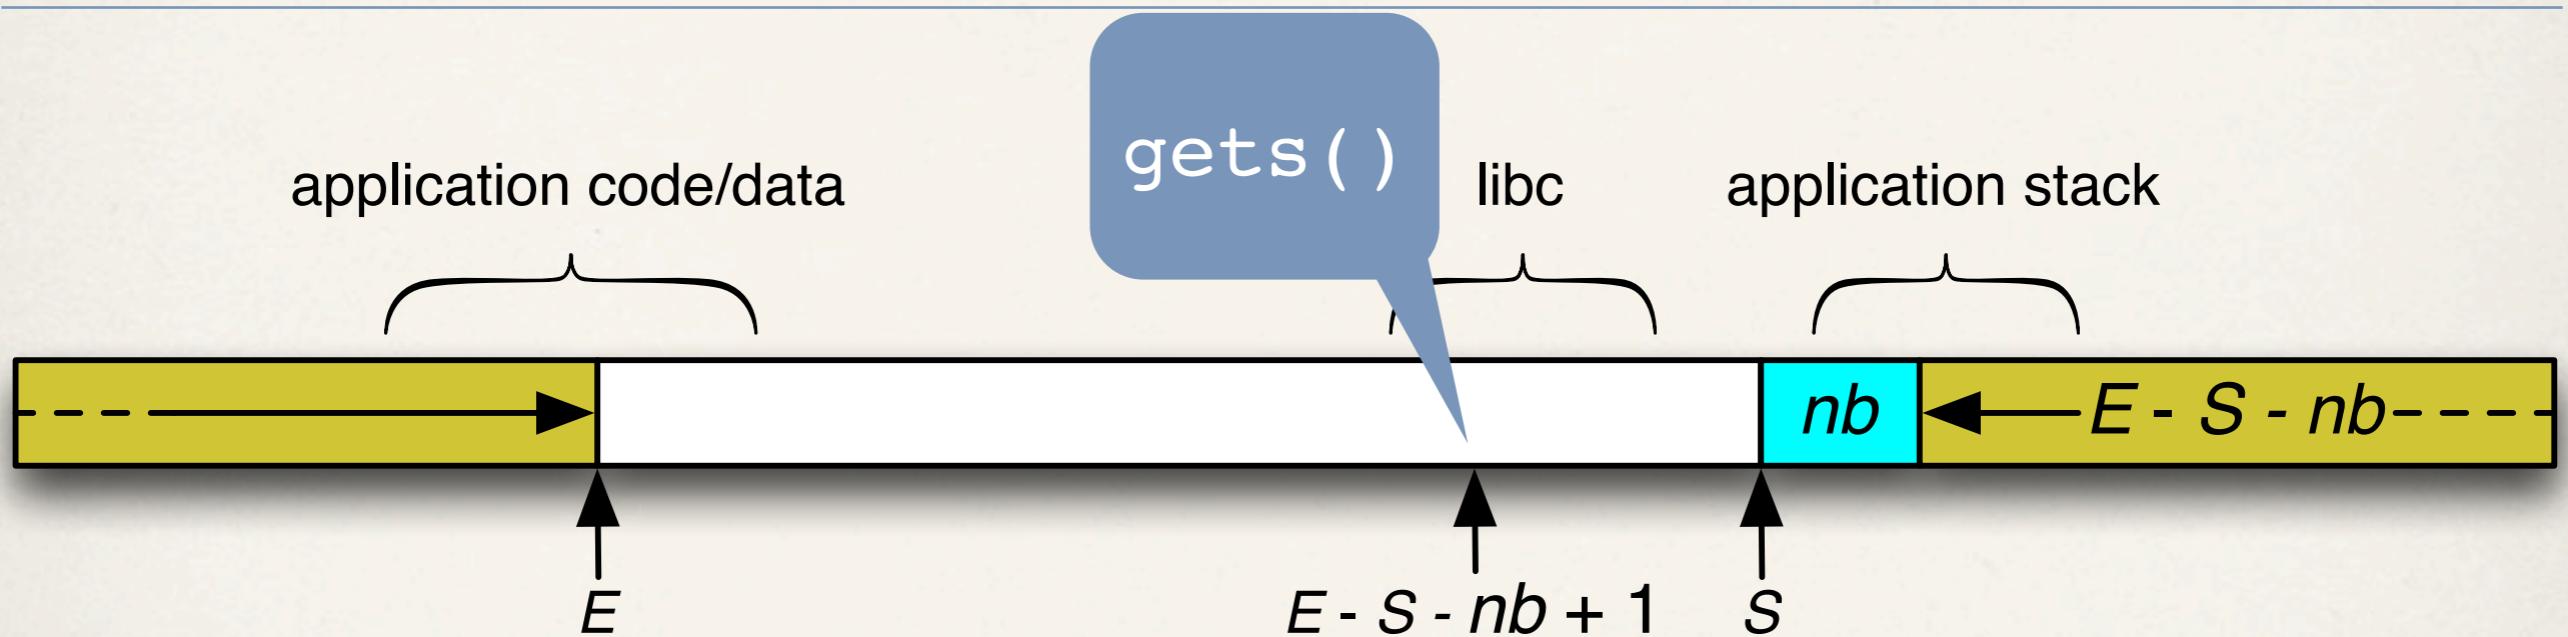


# Attacking the lower half



1. Choose  $S$  such that  $S + nb + 4$  is the address of a saved return address

# Attacking the lower half



1. Choose  $S$  such that  $S + nb + 4$  is the address of a saved return address
2. Choose  $E$  such that  $E - S - nb + 1$  is the address of `gets()`

# Step 3: Putting it all together; Iago attack

---



1. Malicious kernel responds to `brk`
2. `malloc()` writes address of `gets()` over saved return address
3. `gets()` allocates a buffer via `mmap()`
4. Kernel returns an address on the stack
5. `gets()` fills the buffer with `read()`
6. Kernel responds with a return-oriented program

# Conclusions

---



- ❖ The system call interface is a bad RPC mechanism
- ❖ Malicious kernels can take control of protected applications
- ❖ Options:
  1. Design a new system call interface
  2. Enable the hypervisor to check the validity of all system calls
  3. Paraverification (see the next talk!)

# Thank you

---



*Fin*  
33