

# CSCI 210: Computer Architecture

## Lecture 22: Floating Point

Stephen Checkoway

Oberlin College

Nov. 29, 2021

Slides from Cynthia Taylor

# Announcements


- Problem Set 7 due Friday
- Lab 6 due Sunday (it'll be up tonight)
- Office Hours tomorrow 13:30 – 14:30

# Review

- Unsigned 32-bit integers let us represent 0 to  $2^{32} - 1$
- Signed 32-bit integers let us represent  $-2^{31}$  to  $2^{31} - 1$
- 32-bit floating point numbers let us represent a wider range of values: larger, smaller, fractional

$$(-1)^s * 1.x * 2^e$$

- 1 bit for sign  $s$  (1 = negative, 0 = positive)
- 8 bits for exponent  $e$
- 0 bits for implicit leading 1 (called the “hidden bit”)
- 23 bits for significand (without hidden bit)/fraction/mantissa  $x$



# Want To Make Sorting Easy


- Can easily tell if number is positive or negative
  - Just check MSB bit
- Exponent is in higher magnitude bits than the fraction
  - Numbers with higher values will look bigger
  - 0 00000111 100000000000000000000000000000 =  $1.1 * 2^7$
  - 0 00001000 100000000000000000000000000000 =  $1.1 * 2^8$

# Problem with Two's Complement

- 0 00000111 100000000000000000000000000000 =  $1.1 * 2^7$
- 0 00001000 100000000000000000000000000000 =  $1.1 * 2^8$
- 0 11111000 100000000000000000000000000000 =  $1.1 * 2^{-8}$
- Solution: Get rid of negative exponents!
  - We can represent  $2^8 = 256$  numbers: normal exponents -126 to 127 and two special values for zero, infinity, (and NaN and subnormals)
  - Add 127 to value of exponent to encode it, subtract 127 to decode

$$(-1)^s * 1.x * 2^e$$

- 1 bit for sign  $s$  (1 = negative, 0 = positive)
- 8 bits for exponent  $e + 127$
- 0 bits for implicit leading 1 (called the “hidden bit”)
- 23 bits for significand (without hidden bit)/fraction/~~mantissa~~  $x$



$1.000000001 * 2^7$  in Floating Point

- A. 0 00000111 000000001000000000000000
- B. 0 00000111 100000000100000000000000
- C. 0 10000110 000000001000000000000000
- D. 0 10000110 100000000100000000000000
- E. None of the above



# How Can We Represent 0 in Floating Point (as described so far)?

- A. 0 00000000 000000000000000000000000
- B. 0 01111111 000000000000000000000000
- C. 1 00000000 000000000000000000000000
- D. More than one of the above
- E. We can't represent 0

# Special Cases



Object	Exponent	Significand
Zero	0	0
Subnormal	0	Nonzero
Infinity	255	0
NaN	255	Nonzero

- Subnormal number: Numbers with magnitude smaller than  $2^{-126}$ 
  - They have an implicit leading 0 bit
- NaN: Not a Number. Results from  $0/0$ ,  $0 * \infty$ ,  $(+\infty) + (-\infty)$ , etc.

# Overflow/underflow

- Overflow happens when a positive exponent becomes too large to fit in the exponent field
- Underflow happens when a negative exponent becomes too large (in magnitude) to fit in the exponent field
- One way to reduce the chance of underflow or overflow is to offer another format that has a larger exponent field
  - Double precision – takes two MIPS words

# Double precision in MIPS



# Adding

- Add together  $2.34 * 10^3$  and  $4.56 * 10^5$
- Normalize so both have the larger exponent
  - $0.0234 * 10^5 + 4.56 * 10^5$
- Add significands taking sign of numbers into account
  - $4.5834 * 10^5$
- Normalize to a single leading digit
  - $4.5834 * 10^5$

$$1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2}$$

A.  $0.001_2 \times 2^{-1}$

B.  $1.111_2 \times 2^{-1}$

C.  $1.011_2 \times 2^{-2}$

D.  $1.000_2 \times 2^{-4}$

E. None of the above

# What problems could we run into doing this in binary?

- A. Added fraction could be longer than 23 bits
- B. Normalized exponent could be greater than 127 or less than -126
- C. Shifting fraction to match largest exponent could take more than 23 bits
- D. More than one of the above

# Floats in higher-level languages

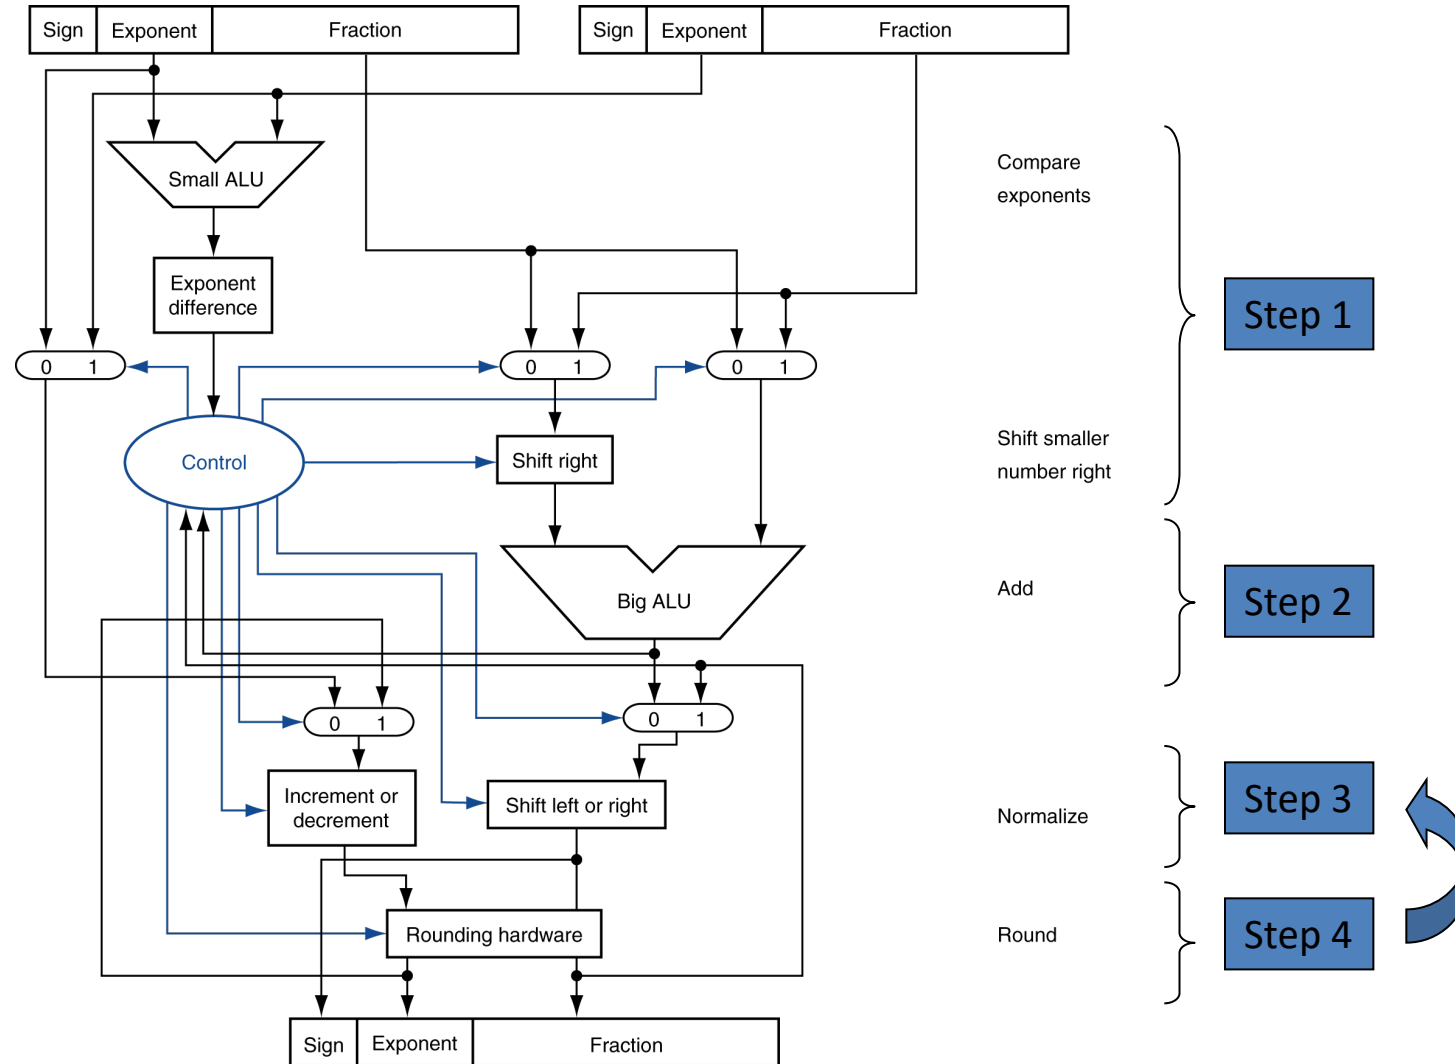
- C, Java: float, double
- JavaScript: numbers are always 64-bit double precision
- Rust: f32, f64
- Sometimes intermediate values (e.g.,  $x*y$  in  $x*y + z$ ) may be doubles (or larger types!) even when the inputs are all floats



# FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
  - Much longer than integer operations
  - Slower clock would penalize all instructions
- FP adder usually takes several cycles

# FP Adder Hardware



# Multiplication

- Multiply  $2.34 * 10^3$  and  $4.56 * 10^5$
- Add together exponents
  - $10^8$
- Multiply fractions (with appropriate signs)
  - $10.6704 * 10^8$
- Normalize
  - $1.06704 * 10^9$

$$1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2}$$

A.  $-1.110_2 \times 2^{-1}$

B.  $-1.110_2 \times 2^{-2}$

C.  $-1.110_2 \times 2^{-3}$

D.  $-1.110_2 \times 2^1$

What problems could we run into doing this in binary floating point?

- A. Adding bias in exponent in twice
- B. Shifted exponent could be greater than 127 or less than -126
- C. Multiplied fraction could be longer than 23 bit
- D. More than one of the above

# FP Instructions in MIPS

- FP hardware is coprocessor 1
  - Adjunct processor that extends the ISA
- Separate FP registers
  - 32 single-precision: \$f0, \$f1, ... \$f31
  - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
- FP instructions operate only on FP registers
  - Programs generally don't do integer ops on FP data, or vice versa
- FP load and store instructions
  - lwc1, ldc1, swc1, sdc1
    - e.g., ldc1 \$f8, 32(\$sp)
  - Psuedoinstructions are easier to read: l.s, l.d, s.s, s.d

# FP Instructions in MIPS

- Single-precision arithmetic
  - `add.s`, `sub.s`, `mul.s`, `div.s`
    - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic (operates on paired registers)
  - `add.d`, `sub.d`, `mul.d`, `div.d`
    - e.g., `mul.d $f4, $f4, $f6`

# Reading

- Next Lecture: Floating Point/Performance
- Problem Set 7