

CSCI 210: Computer Architecture

Lecture 37: Faster Caches

Stephen Checkoway

Oberlin College

Jan. 12, 2022

Slides from Cynthia Taylor

Announcements

- Problem Set 12 due Friday at midnight
- Cache Lab (final project) due on the final exam day
- Course Evals!
 - Extra credit for 90% response rate
 - Currently at 62% (as of 2022-01-12 at 09:00)
- Office Hours Friday 13:30 – 14:30
 - On Zoom

```
for(int i = 0; i<10,000,000;i++)  
    sum+=A[i];
```

Assume each element of A is 4 bytes and sum is kept in a register. Assume a direct-mapped 32 kB cache with 32 byte blocks. Which changes would help the hit rate of the above code?

Selection	Change
A	Increase to 2-way set associativity
B	Increase block size to 64 bytes
C	Increase cache size to 64 KB
D	A and C combined
E	A, B, and C combined

Three types of cache misses

- Compulsory (or cold-start) misses
 - first access to the data.
- Capacity misses
 - we missed only because the cache isn't big enough.
- Conflict misses
 - we missed because the data maps to the same index as other data that forced it out of the cache.

address:

```
4  00000100
8  00001000
12 00001100
4  00000100
8  00001000
20 00010100
4  00000100
8  00001000
20 00010100
24 00011000
12 00001100
8  00001000
4  00000100
```

tag	data

DM cache

Cache Miss Type

Suppose you experience a cache miss on a block (let's call it block A). You have accessed block A in the past. There have been precisely 1027 different blocks accessed between your last access to block A and your current miss. Your block size is 32-bytes and you have a 64 kB cache (recall a kB = 1024 bytes). What kind of miss was this?

Selection	Cache Miss
A	Compulsory
B	Capacity
C	Conflict
D	Both Capacity and Conflict
E	None of the above

Multilevel Caches

- Primary (or level-1) cache attached to CPU
 - Small, but fast
- Level-2 cache services misses from primary cache
 - Larger, slower, but still faster than main memory
- L-3 cache usually services multiple CPUs
- L-3 misses go to main memory

Multilevel Cache Example

- Given
 - CPU base CPI = 1, clock rate = 1 cycle/.25 ns
 - Miss rate/instruction = 2%
 - Main memory access time = 100 ns
- With just primary (L-1) cache
 - Miss penalty = $100 \text{ ns} / (0.25 \text{ ns/cycle}) = 400 \text{ cycles}$
 - Effective CPI = $1 + 0.02 \times 400 = 9$

- Now add L-2 cache
 - Access time = 5 ns
 - Global miss rate to main memory = 0.5%
- Primary miss with L-2 hit
 - Penalty = $5 \text{ ns} / (0.25 \text{ ns/cycle}) = 20 \text{ cycles}$
- Main memory still 400 cycle penalty, L1 miss rate of 2%
- The Total CPI will be
 - A. $1 + 2 \times 20 + 5 \times 400$
 - B. $1 + 0.02 \times 20 + 0.005 \times 400$
 - C. $1 + 0.02 \times 20 \times 0.005 \times 400$
 - D. $1 + 0.0195 \times 20 + 0.005 \times 400$

Multilevel Cache Considerations

- Primary cache
 - Focus on minimal hit time
- L-2 cache
 - Focus on low miss rate to avoid main memory access
 - Hit time has less overall impact
- Results
 - L-1 cache usually smaller than a single cache
 - L-1 block size smaller than L-2 block size
 - L-1 less associative than L-2

Interactions with Advanced CPUs

- Out-of-order CPUs can execute instructions during cache miss
 - Pending store stays in load/store unit
 - Dependent instructions wait in reservation stations
 - Independent instructions continue

Prefetching

- Hardware Prefetching
 - suppose you are walking through a single element in an array of large objects
 - hardware determines the “stride” and starts grabbing values early
- Software Prefetching
 - Compiler adds extra instructions to load data before it is needed

Which data structure will have better memory access times assuming you have a prefetcher?

- A. ArrayList
- B. Linked List
- C. There will not be any difference

Writing Cache-Aware Code

- Focus on your working set
- If your “working set” fits in L1 it will be vastly better than a “working set” that fits only on disk.
- If you have a large data set – do processing on it in chunks.
- Think about regularity in data structures (can a prefetcher guess where you are going – or are you pointer chasing)

Cache Simulator

- Take in a datatrace of load/stores from a real program
- Simulate running the program on a given cache
- Calculate how well a given cache would perform for that trace

Cache Parameters

- Always: Write-allocate, write-back, LRU replacement
- Change:
 - Cache size
 - Block size
 - Associativity
 - Miss penalty

Address Trace

Load/Store Address InstructionCount

```
# 0 7fffed80 1
# 0 10010000 10
# 0 10010060 3
# 1 10010030 4
# 0 10010004 6
# 0 10010064 3
# 1 10010034 4
```

L/S: 0 for load, 1 for store

Simulation Results

Simulation results:

execution time	52268708	cycles
instructions	5136716	
memory accesses	1957764	
overall miss rate	0.79	
load miss rate	0.88	
CPI	10.18	
average memory access time	24.07	cycles
dirty evictions	225876	
load_misses	1525974	
store_misses	30034	
load_hits	205909	
store_hits	195847	

What do you need to do?

- Create data structures that emulate a cache
- For each instruction, find where it would go in the cache, check if it's already there
- Calculate number of miss penalty cycles, load misses, store misses, instructions, etc

Reading

- Next lecture: Class Wrap Up