

CSCI 210: Computer Architecture

Lecture 4: Introduction to MIPS

Stephen Checkoway
Oberlin College
Oct. 11, 2021
Slides from Cynthia Taylor

Announcements

- Problem Set 0 due tonight at 11:59 pm
- Problem Set 1 due in one week (it'll be up soon)
- Office hours today 13:30 – 14:30

Why you should learn (a little) assembly

- Learn what your computer is fundamentally capable of
- By learning about how high level mechanisms are created in assembly, we learn what is fast, what is slow . . .
- Might use it for reverse engineering, embedded systems, compilers

The MIPS Instruction Set

- Used as the example throughout the book
- Stanford MIPS commercialized by MIPS Technologies (owned by John L. Hennessy, who wrote your book.)
- Used in Embedded Systems
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
- Typical of many modern ISAs

Three Types of Instruction

- Arithmetic (R)
- Immediate (I)
- Jump (J)

Arithmetic and Logical Operations

- Add and subtract, three operands

- Two sources and one destination

add a, b, c # a = b + c

sub a, b, c # a = b - c

and a, b, c # a = b & c (bit-wise AND)

- All arithmetic and logical operations have this form

Convert to MIPS: $f = (g + h) - (i + j);$

A.

```
add    f, g, h  
sub    f, i, j
```

B.

```
add    t0, g, h  
add    t1, i, j  
sub    f, t0, t1
```

C.

```
sub    f, (add g,h), (add i,j)
```

D. More than one of these is correct

Register Operands

- Arithmetic instructions use register operands
- MIPS has a 32×32 -bit register file
 - Numbered 0 to 31
 - 32-bit data called a “word”

Aside: MIPS Register Convention

Name	Register Number	Usage
\$zero	0	constant 0 (hardware)
\$at	1	reserved for assembler
\$v0–\$v1	2–3	returned values
\$a0–\$a3	4–7	arguments
\$t0–\$t7	8–15	temporaries
\$s0–\$s7	16–23	saved values
\$t8–\$t9	24–25	temporaries
\$gp	28	global pointer
\$sp	29	stack pointer
\$fp	30	frame pointer
\$ra	31	return addr (hardware)

Register Operand Example

- C code:

`f = (g + h) - (i + j);`

– f, g, h, and j in registers \$s0, \$s1, \$s2, \$s3, and \$s4

- Compiled MIPS code:

`add $t0, $s1, $s2`

`add $t1, $s3, $s4`

`sub $s0, $t0, $t1`

Some R-type instructions

- `add dest, src1, src2`
- `sub dest, src1, src2`
- `div dest, src1, src2` # Pseudoinstruction!
- `mul dest, src1, src2`
- `move dest, src` # `add dest, $zero, src`
- `and dest, src1, src2`
- `or dest, src1, src2`
- `nor dest, src1, src2`
- `xor dest, src1, src2`

Assume registers initially have the following values

\$a0	\$a1	\$t0	\$t1	\$v0
2	100	5	6	7

What values do they have after running this code?

```
move $t0, $a0
```

```
add  $t1, $a0, $a0
```

```
add  $t1, $t1, $t1
```

```
sub  $t0, $t1, $t0
```

```
add  $v0, $t0, $a1
```

	\$a0	\$a1	\$t0	\$t1	\$v0
A	2	100	5	6	7
B	2	100	6	8	106
C	5	-10	-17	22	7
D	5	100	15	20	115
E	None of the above				

Questions about Arithmetic Operations?

Memory Instructions

- `lw $t0, 0($t1)`
 - $\$t0 = \text{Mem}[\$t1+0]$
 - Loads 4 bytes from $\$t1$, $\$t1+1$, $\$t1+2$, and $\$t1+3$
- `sw $t0, 4($t1)`
 - $\text{Mem}[\$t1+4] = \$t0$
 - Stores 4 bytes at $\$t1+4$, $\$t1+5$, $\$t1+6$, and $\$t1+7$
- These instructions are the cornerstones of our being able to go to and from memory

Accessing the Operands

There are typically two locations for operands – **registers** (internal storage e.g., \$t0 or \$a0) and **memory**. In each column we have which—reg or mem—is better. Which row is correct?

	Faster access	Fewer bits to specify address	More locations
A	Mem	Mem	Reg
B	Mem	Reg	Mem
C	Reg	Mem	Reg
D	Reg	Reg	Mem
E	None of the above		

Load-store architectures

can do:

```
load r3, M(address)
```

```
add r1 = r2 + r3
```

can't do

```
add r1 = r2 + M(address)
```

⇒ forces heavy dependence
on registers, which is
exactly what you want in
today's CPUs

- more instructions
+ fast implementation

Memory

- Main memory used for composite data
 - Arrays, structures, dynamic data
- Memory is byte addressed
 - Each address identifies an 8-bit byte
- Words are aligned in memory
 - Address of a word must be a multiple of 4

Memory Organization

- Viewed as a large, single-dimension array, with an address.
- A memory address is an index into the array
- “Byte Addressing” means that the index points to a byte of memory.

0	8 bits of data
1	8 bits of data
2	8 bits of data
3	8 bits of data
4	8 bits of data
5	8 bits of data
6	8 bits of data

...

Memory Organization

- Bytes are nice, but most data items use larger "words"
- For MIPS, a word is 32 bits or 4 bytes.

0	32 bits of data
4	32 bits of data
8	32 bits of data
12	32 bits of data

Registers hold 32 bits of data

- 2^{32} bytes with byte addresses from 0 to $2^{32} - 1$
- 2^{30} words with byte addresses 0, 4, 8, ... $2^{32} - 4$

If you have a pointer to address 0x00001000 and you increment it by one to 0x00001001. What does the new pointer point to, relative to the original pointer?

- A) The next word in memory
- B) The next byte in memory
- C) Either the next word or byte – depends on if you use that address for a load byte or load word
- D) Pointers are a high level construct – they don't make sense pointing to raw memory addresses.
- E) None of the above.

If a 4-byte word is in memory at address 0x0420E074, what is the address of the next word in memory?

- A) 0x0420E078
- B) 0x0420E079
- C) 0x1420E078
- D) It depends on the value of the words in memory
- E) Since a word is 4 bytes, it's not possible to have one at address 0x0420E074

Reading

- Next lecture: Assembly
 - 2.3
- Problem Set 0: Due tonight at 11:59pm via Gradescope