

CSCI 210: Computer Architecture

Lecture 22: Floating Point

Stephen Checkoway

Oberlin College

Apr. 18, 2022

Slides from Cynthia Taylor

Announcements


- Problem Set 7 due Friday
- Lab 6 due Sunday
- Office Hours tomorrow 13:30 – 14:30

Review

- Unsigned 32-bit integers let us represent 0 to $2^{32} - 1$
- Signed 32-bit integers let us represent -2^{31} to $2^{31} - 1$
- 32-bit floating point numbers let us represent a wider range of values: larger, smaller, fractional

$$(-1)^s * 1.\textcolor{red}{x} * 2^{\textcolor{green}{e}}$$

- 1 bit for sign s (1 = negative, 0 = positive)
- 8 bits for exponent e
- 0 bits for implicit leading 1 (called the “hidden bit”)
- 23 bits for significand (without hidden bit)/fraction/~~mantissa~~ x



Want To Make Comparisons Easy


- Can easily tell if number is positive or negative
 - Just check MSB bit
- Exponent is in higher magnitude bits than the fraction
 - Numbers with higher values will look bigger
 - 0 00000111 10000000000000000000000000000000 = $1.1 * 2^7$
 - 0 00001000 10000000000000000000000000000000 = $1.1 * 2^8$

Problem with Two's Complement

- 0 00000111 100000000000000000000000000000 = $1.1 * 2^7$
- 0 00001000 100000000000000000000000000000 = $1.1 * 2^8$
- 0 11111000 100000000000000000000000000000 = $1.1 * 2^{-8}$
- Solution: Get rid of negative exponents!
 - We can represent $2^8 = 256$ values for the exponent:
 - normal exponents -126 to 127; and
 - two special values for zero, infinity, (and NaN and subnormals)
 - Add 127 to value of exponent to encode it, subtract 127 to decode

$$(-1)^s * 1.x * 2^e$$

- 1 bit for sign s (1 = negative, 0 = positive)
- 8 bits for exponent e + 127
- 0 bits for implicit leading 1 (called the “hidden bit”)
- 23 bits for significand (without hidden bit)/fraction/mantissa x



1.000000001 * 2⁷ in Floating Point

- A. 0 00000111 000000001000000000000000
- B. 0 00000111 100000000100000000000000
- C. 0 10000110 000000001000000000000000
- D. 0 10000110 100000000100000000000000
- E. None of the above

How Can We Represent 0 in Floating Point (as described so far)?

- A. 0 00000000 000000000000000000000000
- B. 0 01111111 000000000000000000000000
- C. 1 00000000 000000000000000000000000
- D. More than one of the above
- E. We can't represent 0

Special Cases



Object	Exponent	Significand
Zero	0	0
Subnormal	0	Nonzero
Infinity	255	0
NaN	255	Nonzero

- Subnormal number: Numbers with magnitude smaller than 2^{-126}
 - They have an implicit leading 0 bit and an exponent of 2^{-126}
- NaN: Not a Number. Results from $0/0$, $0 * \infty$, $(+\infty) + (-\infty)$, etc.

Overflow/underflow

- **Overflow** happens when a positive exponent becomes too large to fit in the exponent field
- **Underflow** happens when a negative exponent becomes too large (in magnitude) to fit in the exponent field
- One way to reduce the chance of underflow or overflow is to offer another format that has a larger exponent field
 - Double precision – takes two MIPS words

Double precision in MIPS



Floats in higher-level languages

- C, Java: float, double
- JavaScript: numbers are always 64-bit double precision
- Rust: f32, f64
- Sometimes intermediate values (e.g., $x*y$ in $x*y + z$) may be doubles (or larger types!) even when the inputs are all floats

Adding in base-10 scientific notation

- Add together $2.34 * 10^3$ and $4.56 * 10^5$
- Normalize so both have the larger exponent
 - $0.0234 * 10^5 + 4.56 * 10^5$
- Add significands taking sign of numbers into account
 - $4.5834 * 10^5$
- Normalize to a single leading digit
 - $4.5834 * 10^5$

Adding in floating point (assuming 4 fractional bits)

- Add together $1.1011 * 2^{-1}$ and $-1.0110 * 2^2$
- Normalize so both have the larger exponent
 - $0.0110 * 2^2 + 1.0110 * 2^2$
- Add significands taking sign of numbers into account
 - $0.0110 * 2^2 + 1.0110 * 2^2 = 1.1100$
- Normalize to a single leading digit
 - $1.1100 * 2^2$

What problems could we run into doing this in binary?

- A. Added fraction could be longer than 23 bits
- B. Normalized exponent could be greater than 127 or less than -126
- C. Shifting fraction to match largest exponent could take more than 23 bits
- D. The inputs could be zero or the result could be zero
- E. More than one of the above

Floating point addition algorithm

Input: two single-precision, floating point numbers x , and y

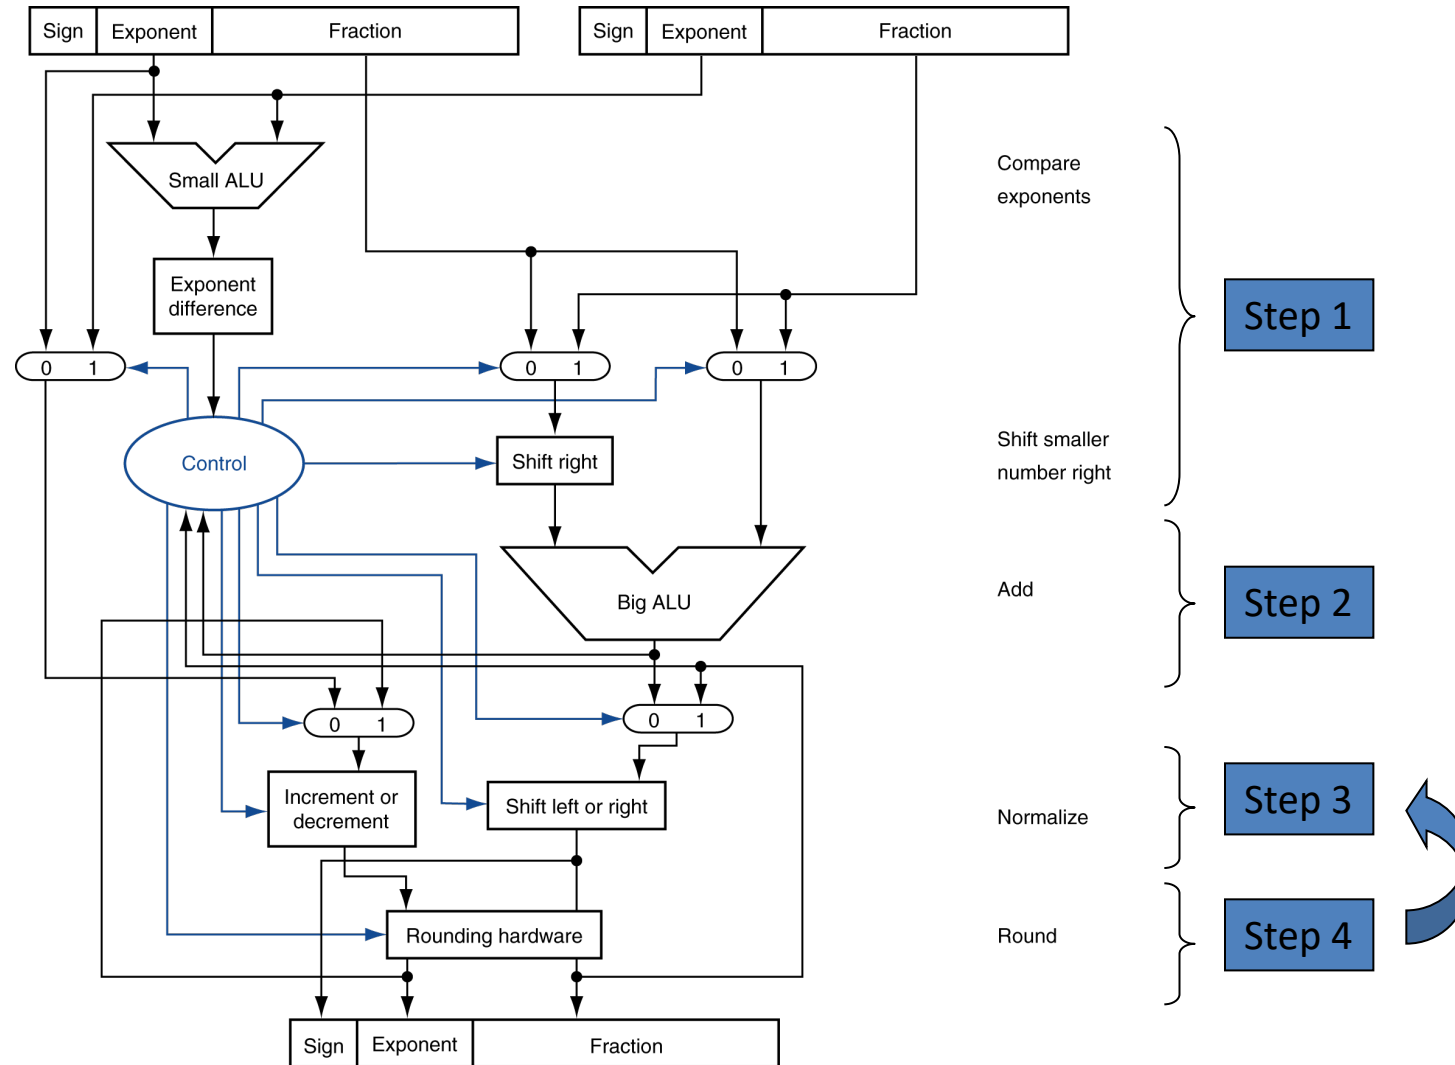
Output: $x + y$

1. If either x or y is 0, return the other one
2. Denormalize x or y to give them both the larger exponent
3. Add the significands, taking sign into account
4. If the result is 0, return 0
5. Normalize the result
6. Encode the result (bias the exponent, remove the hidden bit)

FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
 - Much longer than integer operations
 - Slower clock would penalize all instructions
- FP adder usually takes several cycles

FP Adder Hardware



Multiplication in base-10 scientific notation

- Multiply $2.34 * 10^3$ and $4.56 * 10^5$
- Add together exponents
 - 10^8
- Multiply fractions (with appropriate signs)
 - $10.6704 * 10^8$
- Normalize
 - $1.06704 * 10^9$

What problems could we run into doing this in binary floating point?

- A. Adding bias in exponent in twice
- B. Shifted exponent could be greater than 127 or less than -126
- C. Multiplied fraction could be longer than 23 bit
- D. More than one of the above

Floating point multiplication algorithm

Input: two single-precision, floating point numbers x , and y

Output: $x * y$

1. If either x or y is 0, return 0
2. Compute the sign of the result
3. Add the exponents
4. Multiply the significands as 64-bit integers and shift right by 23 bits
5. Normalize the result
6. Encode the result (bias the exponent, remove the hidden bit)

FP Instructions in MIPS

- FP hardware is coprocessor 1
 - Adjunct processor that extends the ISA
- Separate FP registers
 - 32 single-precision: \$f0, \$f1, ... \$f31
 - Paired for double-precision: \$f0/\$f1, \$f2/\$f3, ...
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
- FP load and store instructions
 - lwc1, ldc1, swc1, sdc1
 - e.g., ldc1 \$f8, 32(\$sp)
 - Psuedoinstructions are easier to read: l.s, l.d, s.s, s.d

FP Instructions in MIPS

- Single-precision arithmetic
 - `add.s`, `sub.s`, `mul.s`, `div.s`
 - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic (operates on paired registers)
 - `add.d`, `sub.d`, `mul.d`, `div.d`
 - e.g., `mul.d $f4, $f4, $f6`

Reading

- Next Lecture: Floating Point/Performance
- Problem Set 7