

CSCI 210: Computer Architecture Cache Simulator Final Project

Stephen Checkoway
Oberlin College

Cache Simulator

- Take in a trace of load/stores from a real program
- Simulate running the program on a given cache configuration
- Calculate how well that configuration would perform for that trace

Cache Parameters

- Always: Write-allocate, write-back, LRU replacement
- Change:
 - Cache size
 - Block size
 - Associativity
 - Miss penalty

Store-hit policy: write-back

- Only update the block in cache
 - Keep track of whether each block is “dirty” (i.e., it has a different value than in memory)
- When a dirty block is replaced
 - Write it back to memory
 - Can use a write buffer to allow replacing block to be read first
- Faster than write-through, but more complex

V	D	Tag	Data
1	0	000042	FE FF 3C ...
0			
1	1	001234	65 82 5C ...
0			
0			
1	0	000F3C	00 00 00 ...
0			
0			

Store-miss policy: write-allocate

- Read a block from memory into the cache (just like a load miss)
- Perform the write according to the store-hit policy (i.e., write in cache or write in both cache and memory)
- Good for when data is likely to be read shortly after being written (temporal locality)

Replacement Policy

- Least-recently used (LRU)
 - Choose the one unused for the longest time

Address Trace

Load/Store Address InstructionCount

```
# 0 7ffffed80 1
# 0 10010000 10
# 0 10010060 3
# 1 10010030 4
# 0 10010004 6
# 0 10010064 3
# 1 10010034 4
```

L/S: 0 for load, 1 for store

Simulation Results

Simulation results:

execution time	52268708 cycles
instructions	5136716
memory accesses	1957764
overall miss rate	0.79
load miss rate	0.88
CPI	10.18
average memory access time	24.07 cycles
dirty evictions	225876
load_misses	1525974
store_misses	30034
load_hits	205909
store_hits	195847

What do you need to do?

- Create data structures that emulates a cache
- For each instruction, find where it would go in the cache, check if it's already there
- Calculate number of miss penalty cycles, load misses, store misses, instructions, etc.
- Suggestion: Keep track of counts of events (e.g., load/store hits/misses, instruction count, dirty evictions) and use those counts to compute statistics like execution time, miss rates, etc.

Questions on Cache Final Project?

We have an 8 byte block size, direct-mapped, 16 kB cache. For a given byte address, to find the offset

- A. Mod by 8
- B. Mod by 16
- C. Divide by 8
- D. Divide by 16
- E. None of the above

If we are simulating a cache and not actually accessing data, do we need to use the offset?

A. Yes

B. No

We have an 8 byte block size, direct-mapped, 16 kB cache. How many rows will this cache have?

A. 16

*Recall a KB is 1024 bytes

B. $2 * 1024$

C. $8 * 1024$

D. $16 * 1024$

We have an 8 byte block size, direct-mapped, 16 KB cache. How can we find the index from an address?

- A. Divide by 8
- B. Divide by $(16 * 1024)$
- C. Mod by $16 * 1024$
- D. Divide by 8, then mod by $(2 * 1024)$
- E. None of the above

We have an 8 byte block size, **2-way set associative**, 16 kB cache. How can we find the index from an address?

- A. Divide by 8, then mod by 1024
- B. Divide by 8, then mod by $(2 * 1024)$
- C. Divide by 8, then mod by $(2 * 2 * 1024)$
- D. Divide by 16, then mod by $(2 * 1024)$
- E. None of the above

We have an 8 byte block size, direct mapped, 16 KB cache. How can we find the tag of an address?

- A. Divide by 8
- B. Divide by $8 * 1024$
- C. Divide by $16 * 1024$
- D. None of the above