

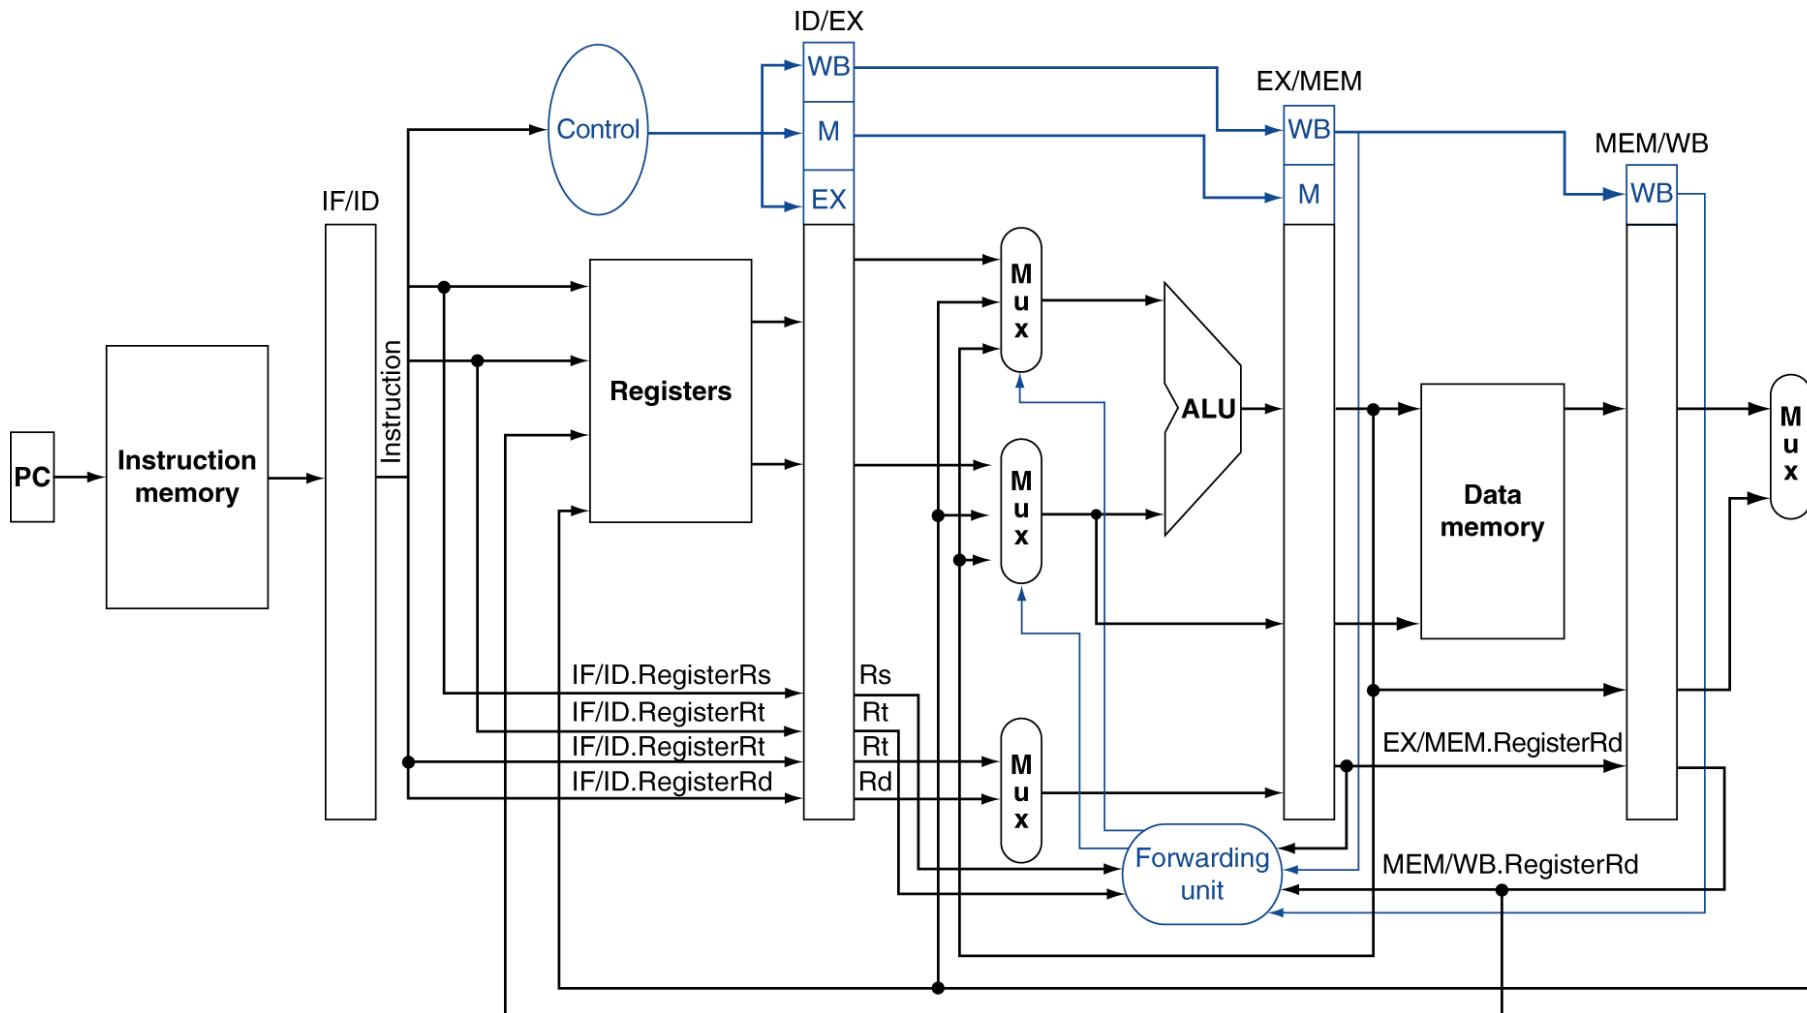
CSCI 210: Computer Architecture

Lecture 30: Data Hazards

Stephen Checkoway

Slides from Cynthia Taylor

Datapath with Forwarding



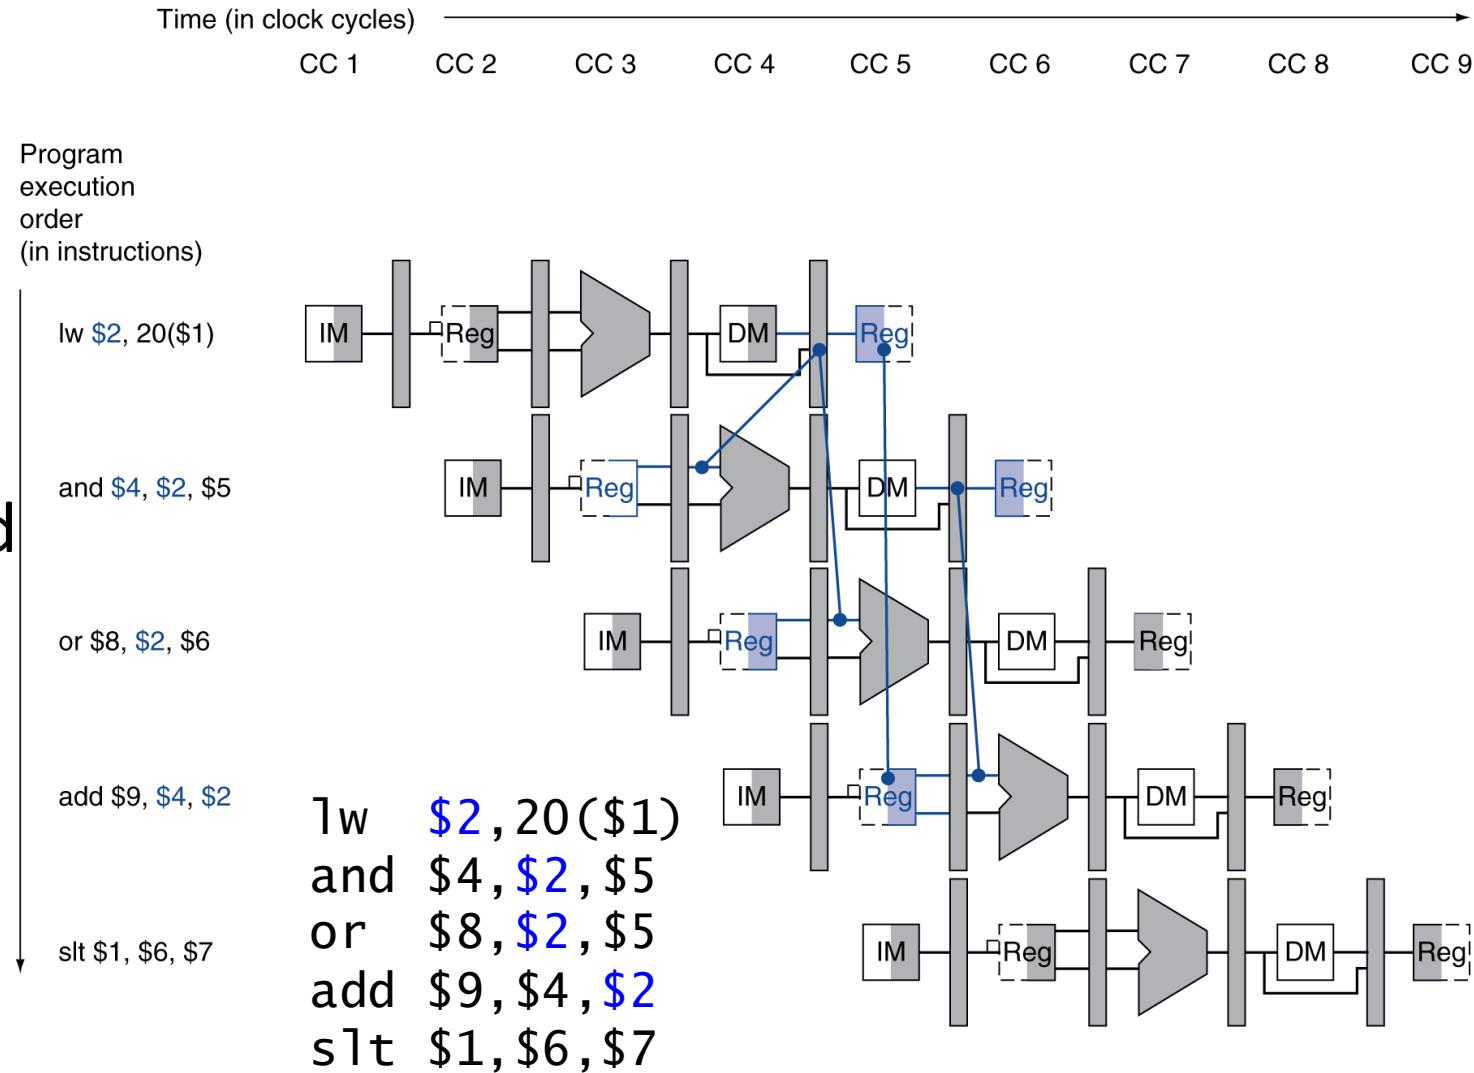
We can best solve these data hazards

A. By stalling.

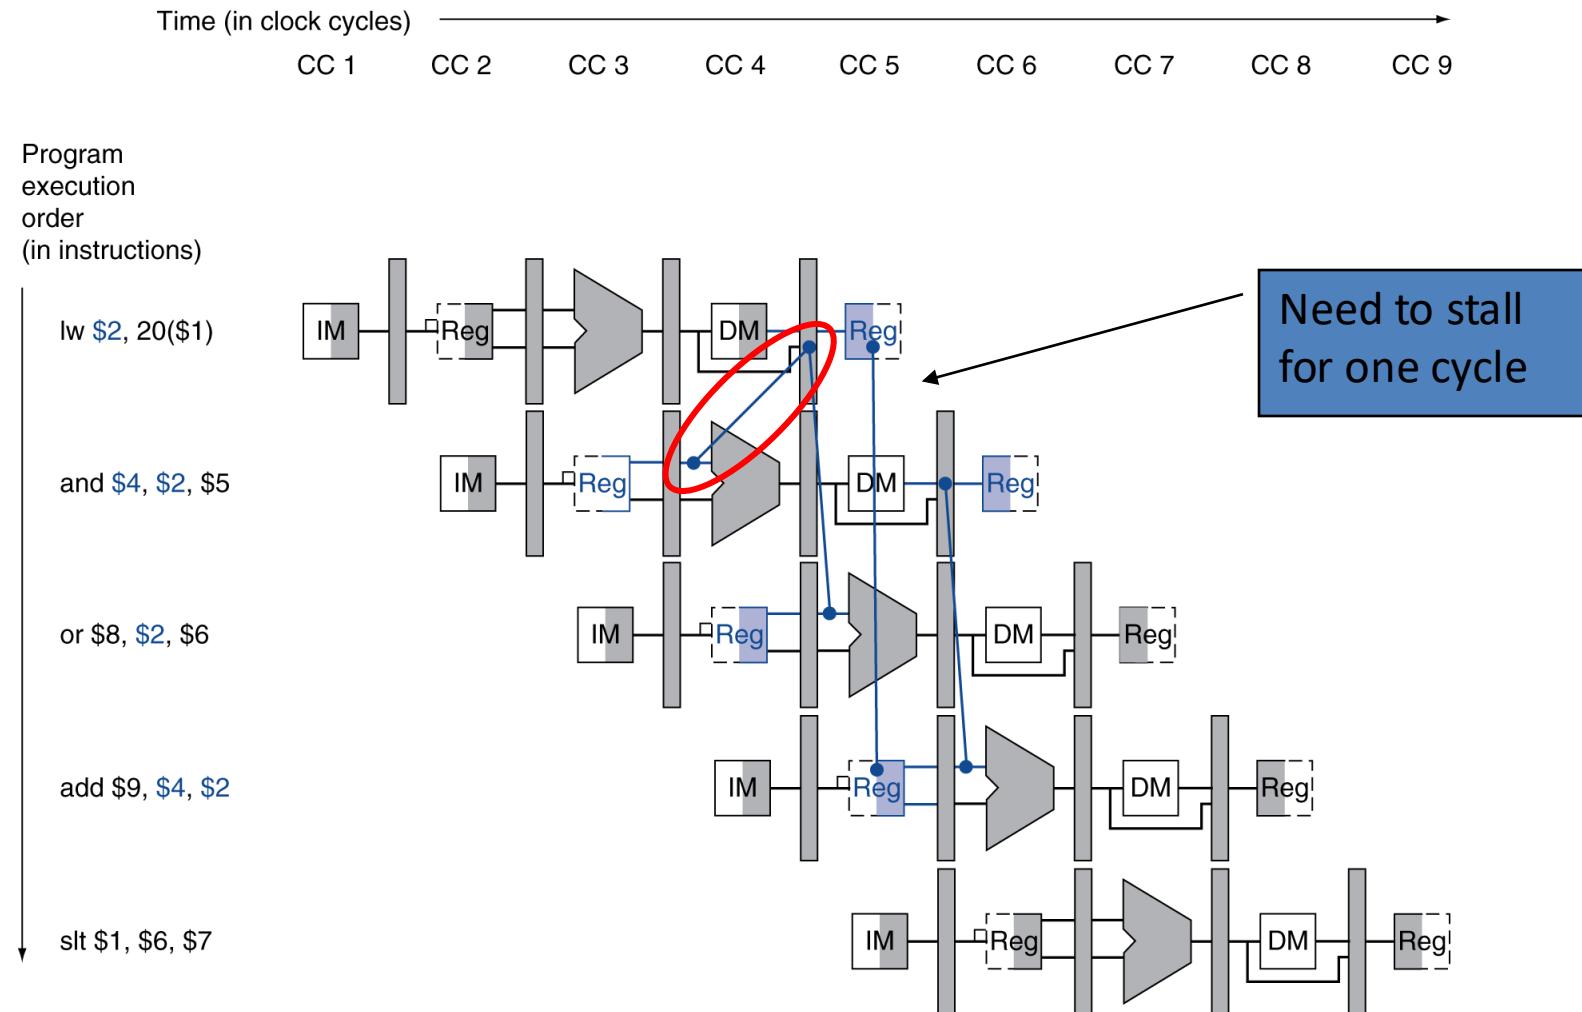
B. By forwarding.

C. By combining forwards and stalls.

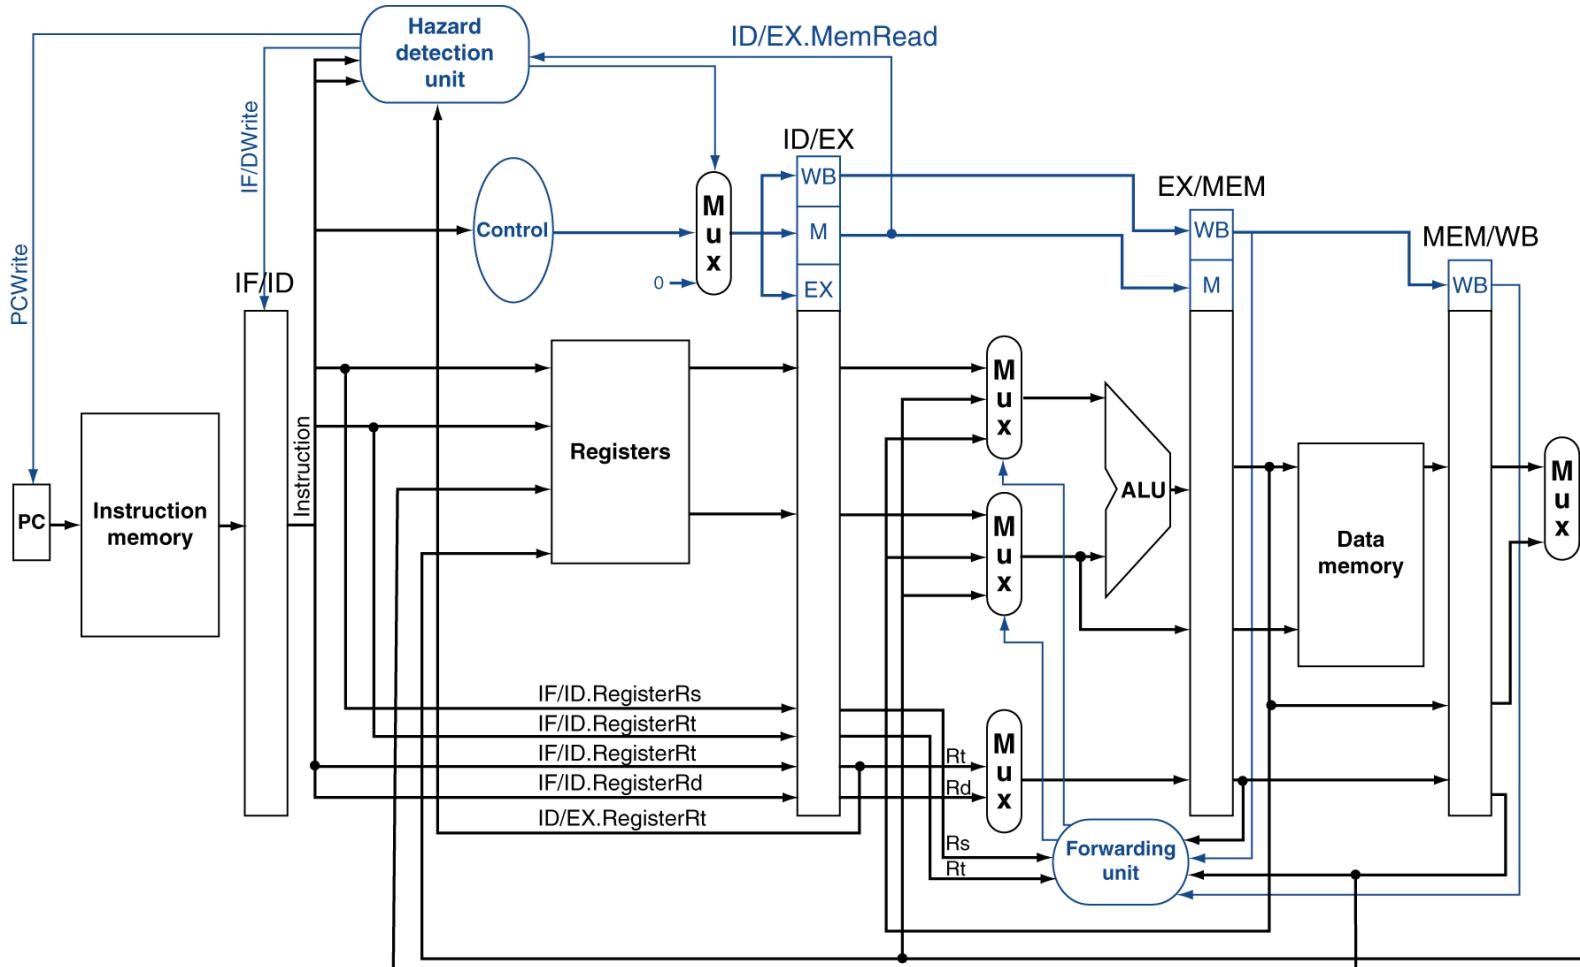
D. By doing something else.



Load-Use Data Hazard

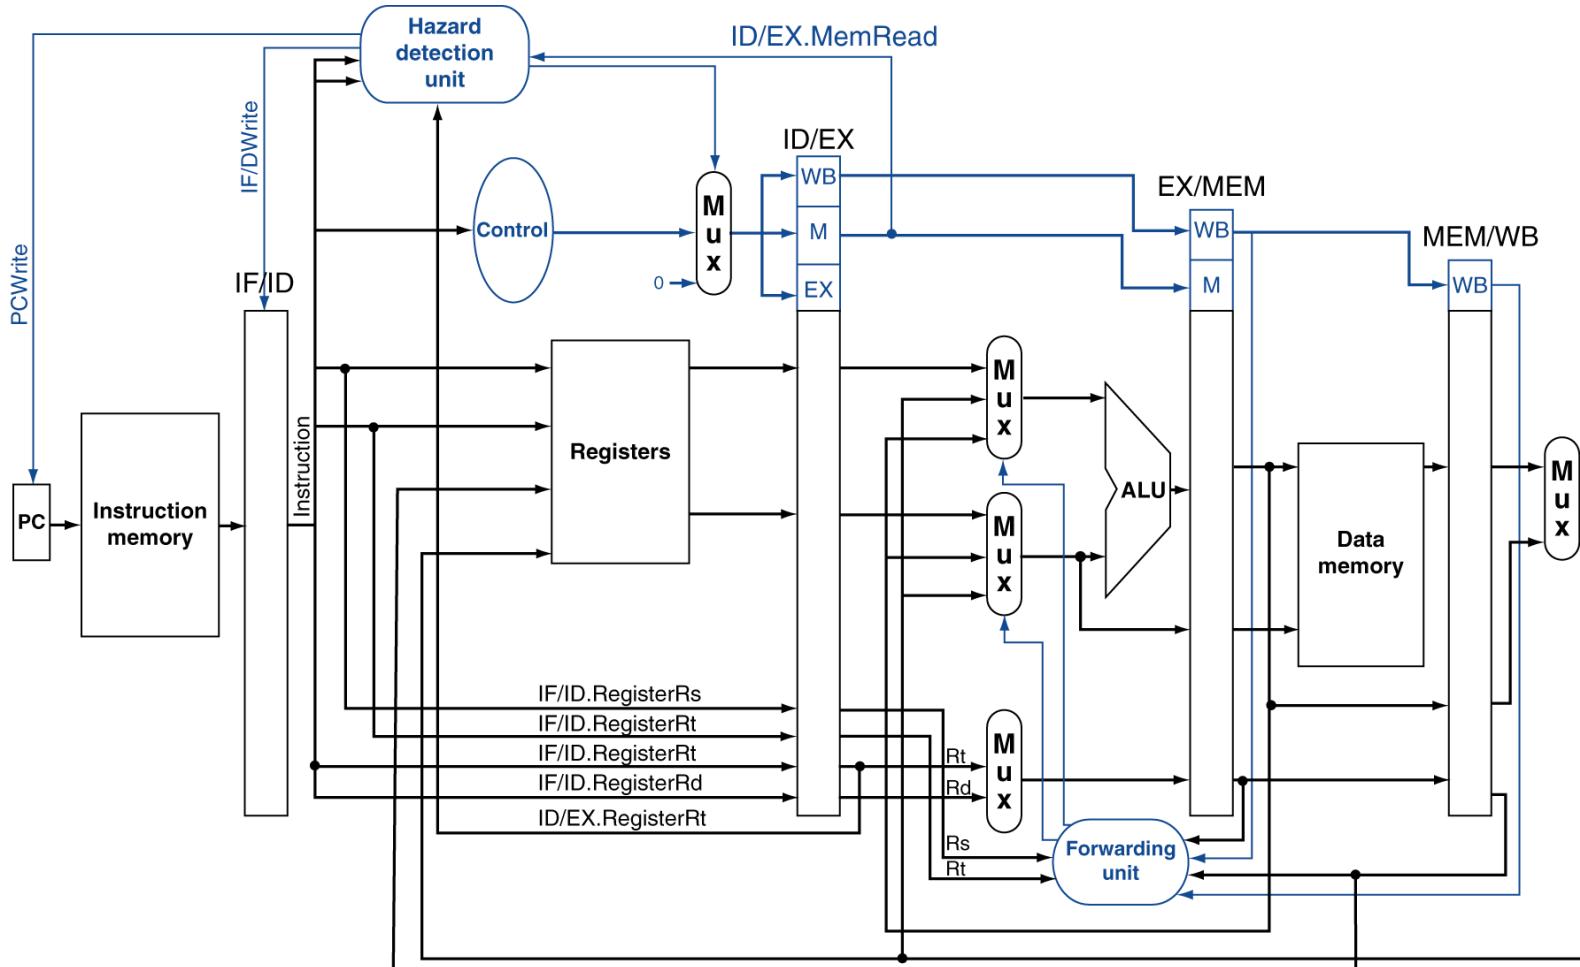


How to Stall the Pipeline



- Detect hazard in ID stage using Hazard detection unit
 - Check if instruction in EX stage is load with destination rs or rt

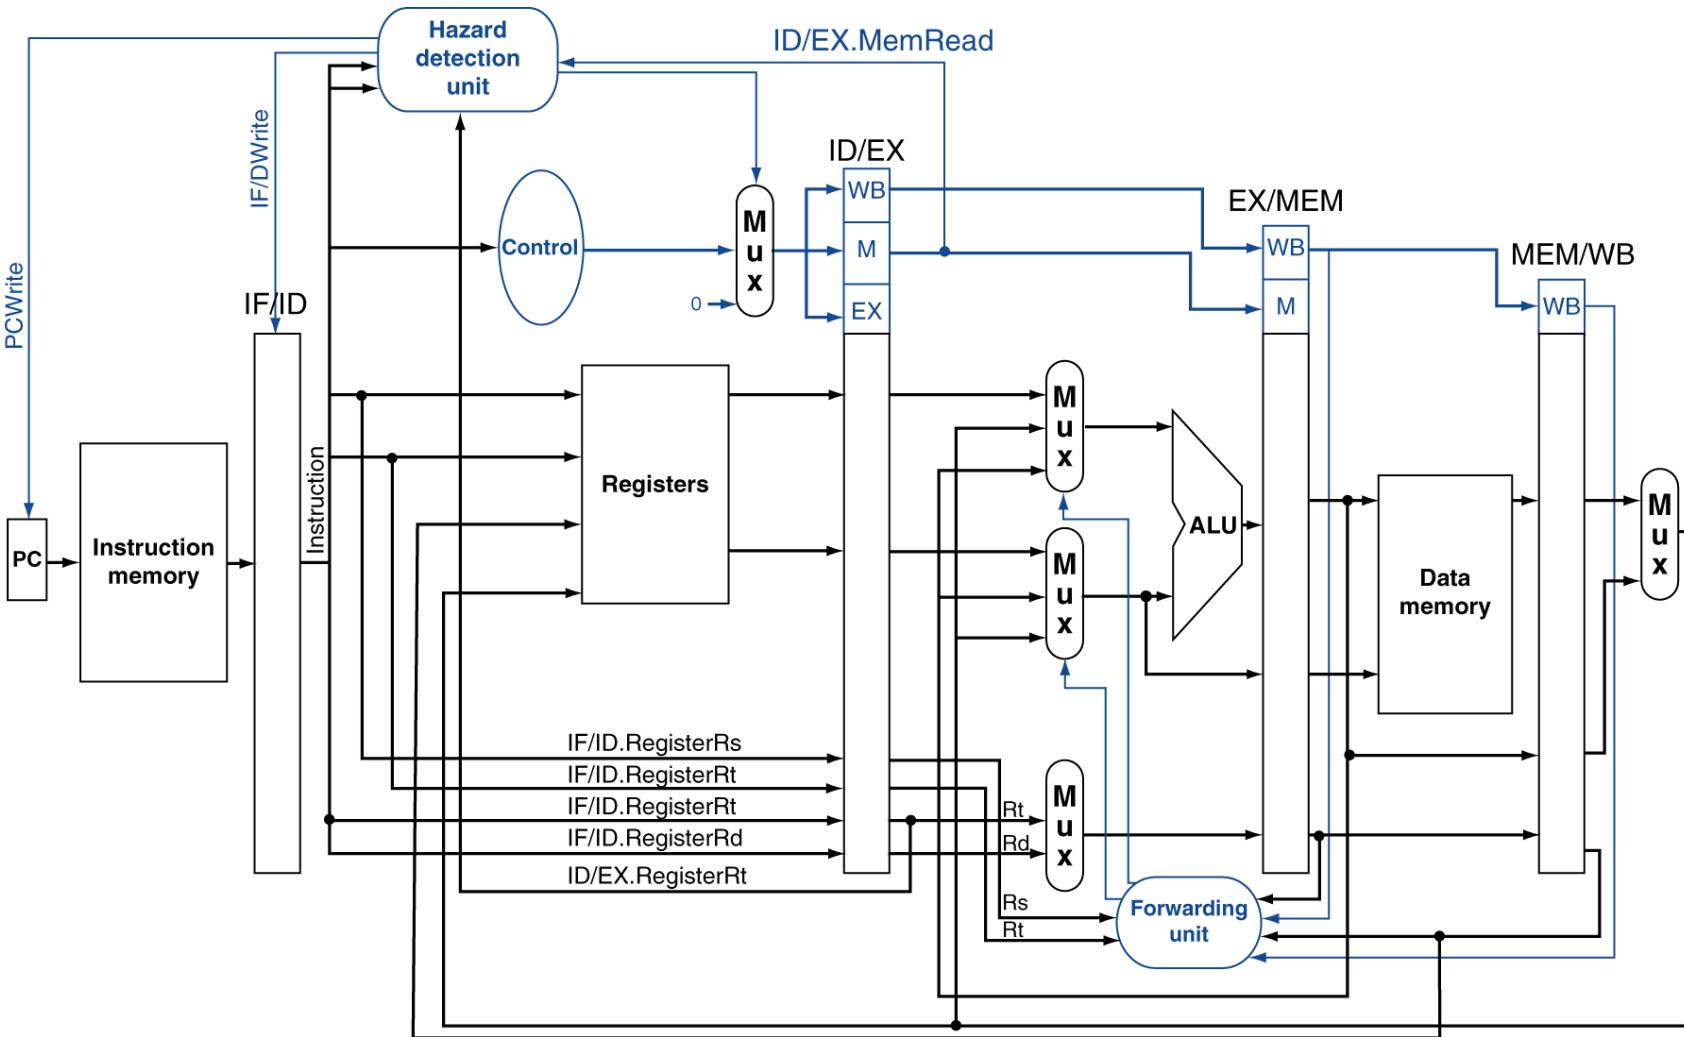
How to Stall the Pipeline



- Force control values in ID/EX register to 0
 - EX, MEM and WB do nop (no-operation)

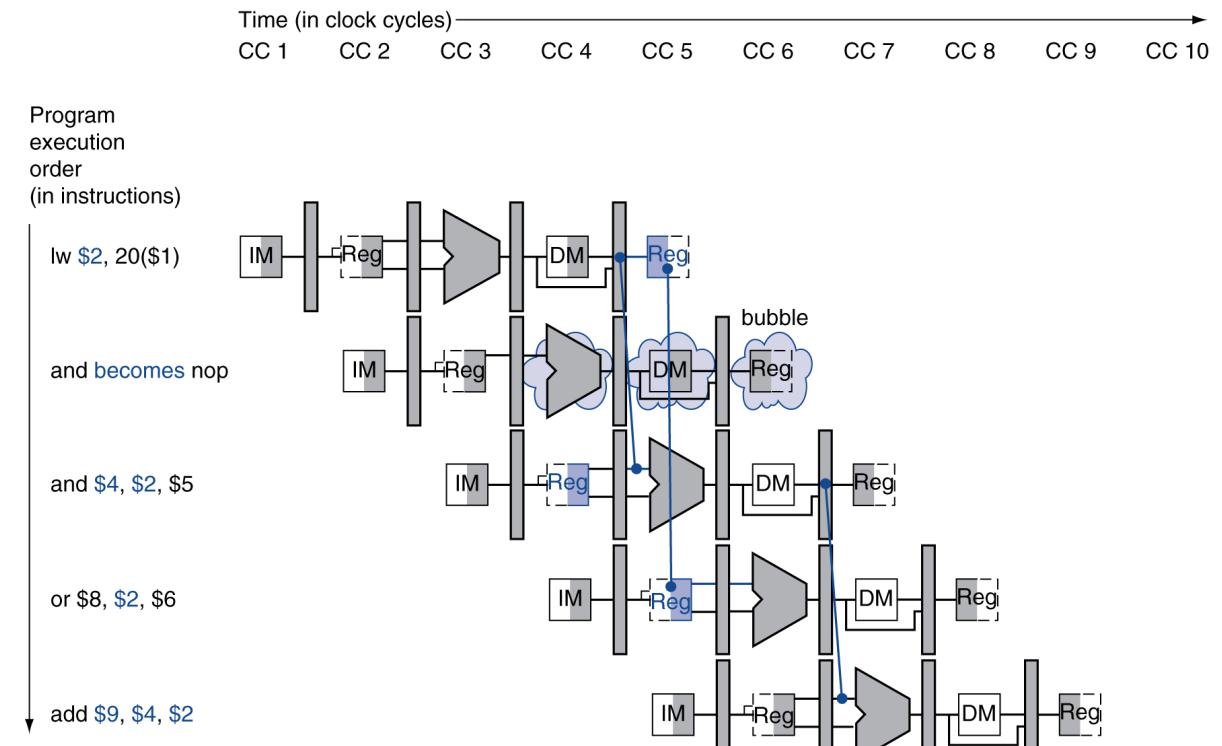
How to Stall the Pipeline

- Prevent update of PC and IF/ID register
 - Instruction with dependency is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1_w

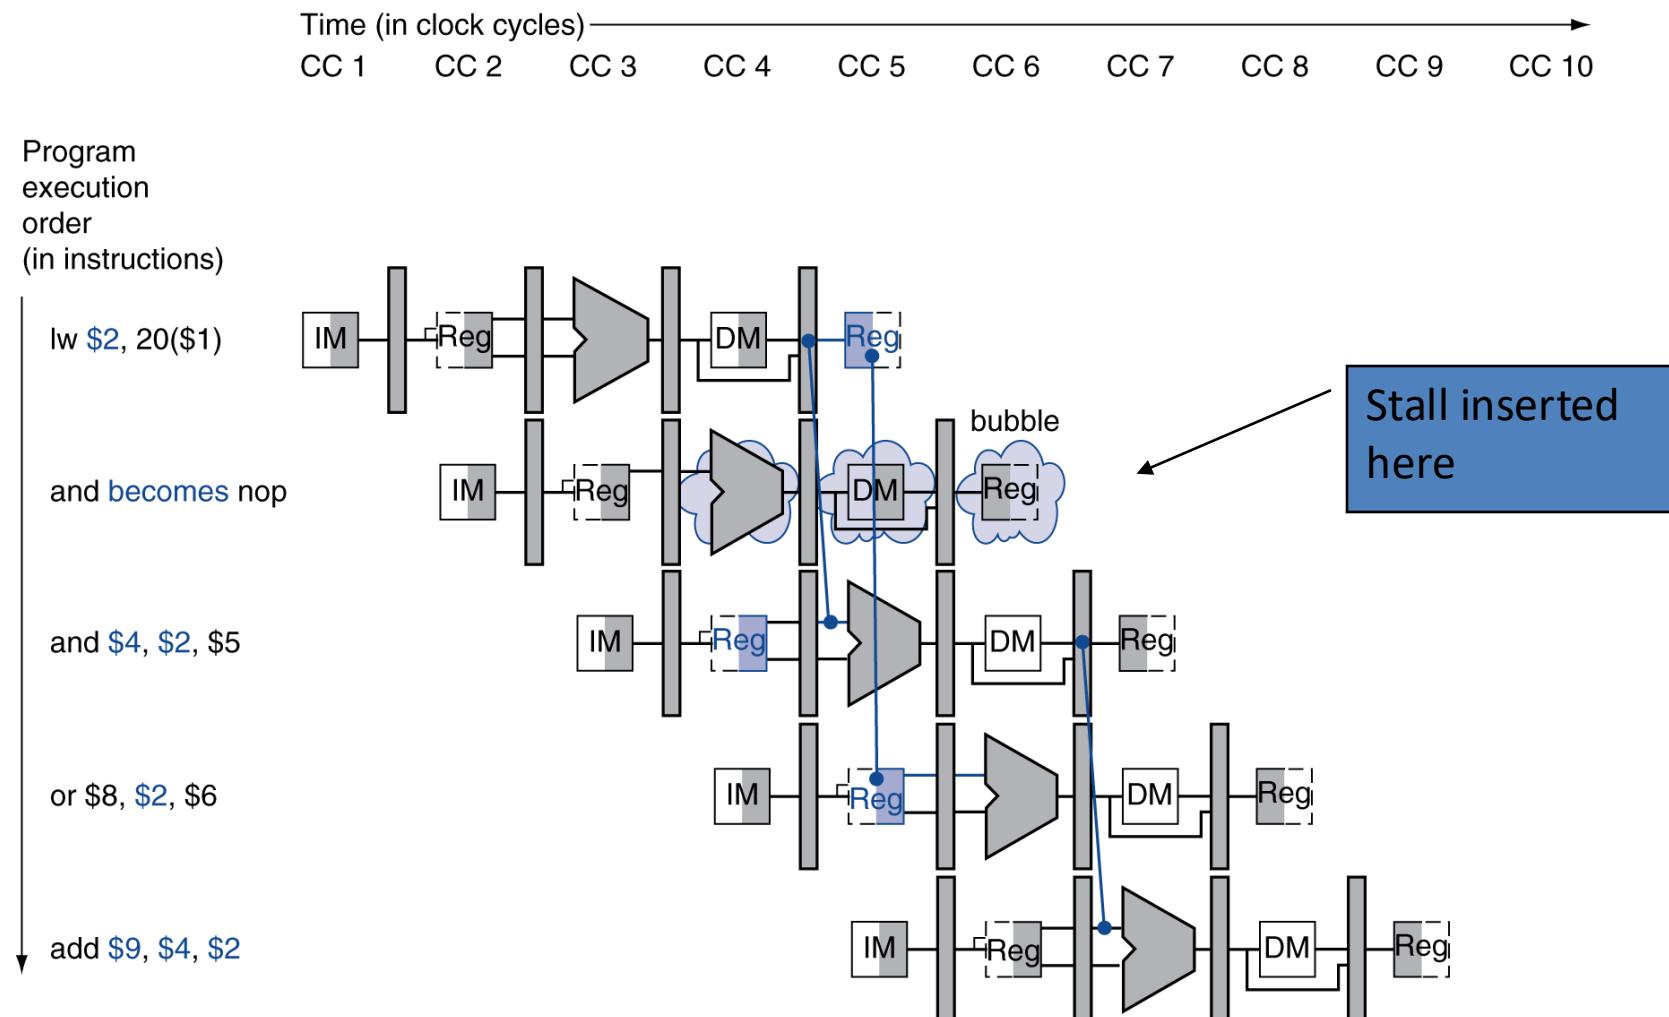


After we add the stall

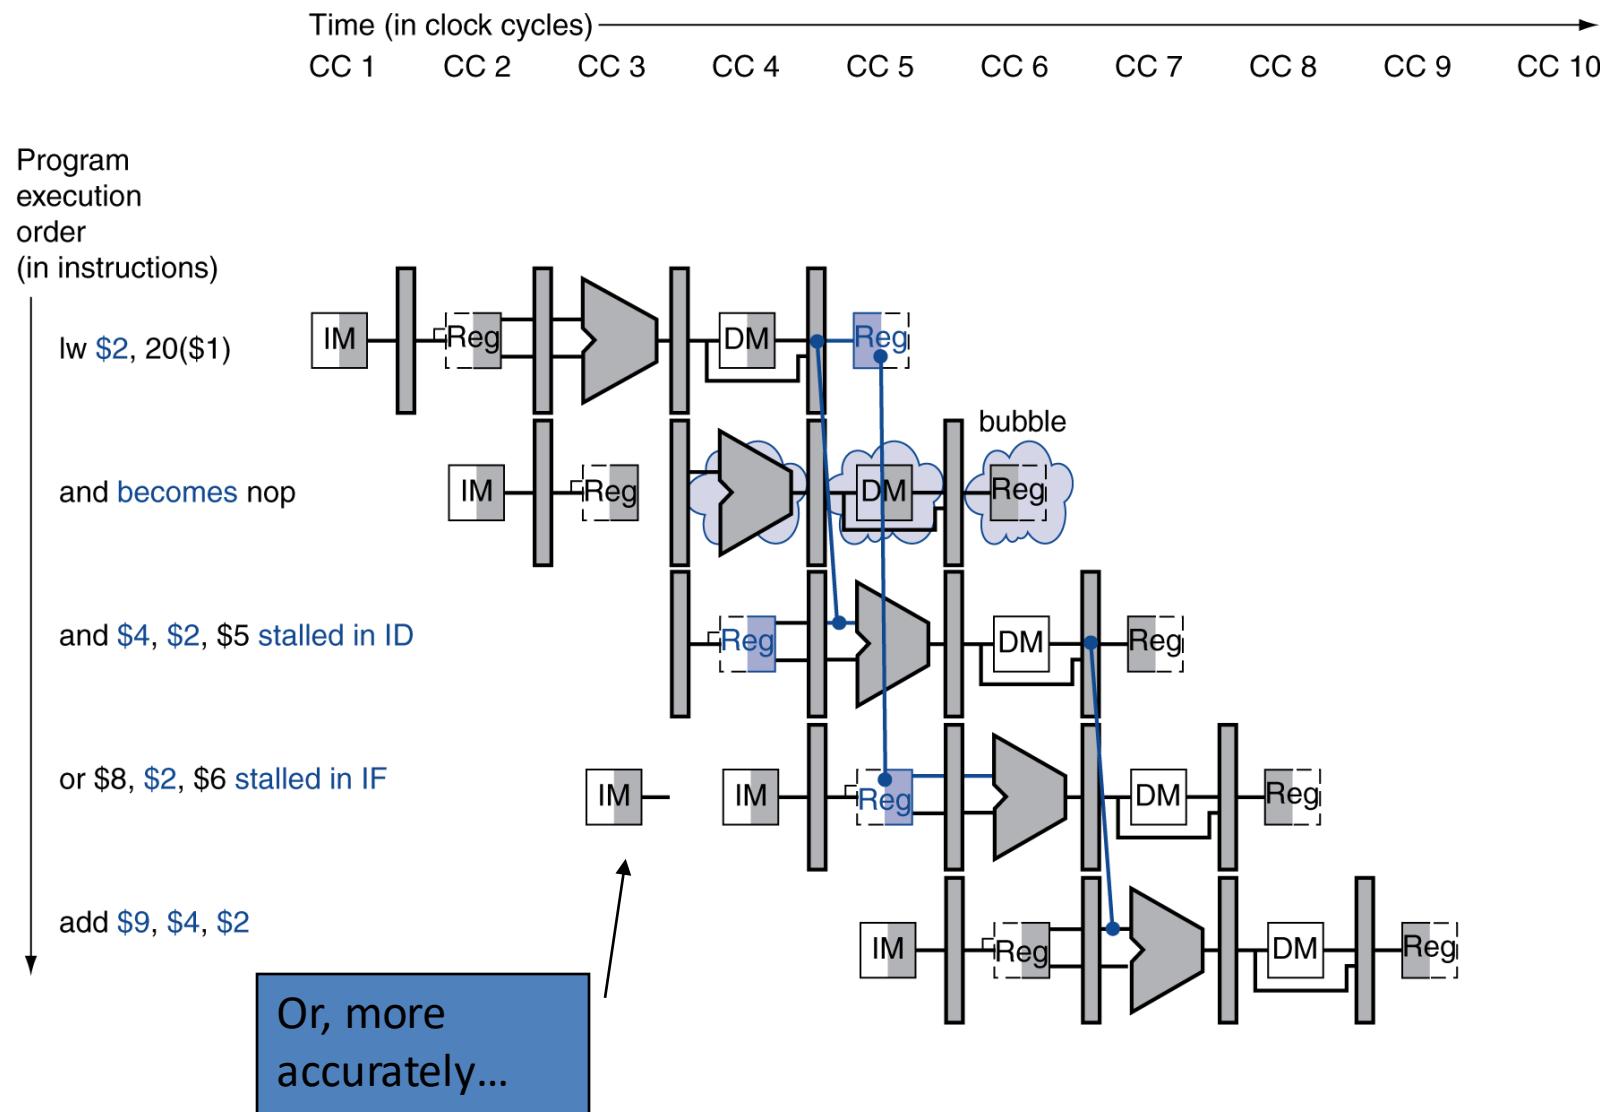
- A. Everything works with our existing forwarding
- B. We need to forward between the register files to solve the 2nd hazard
- C. We need to do something else



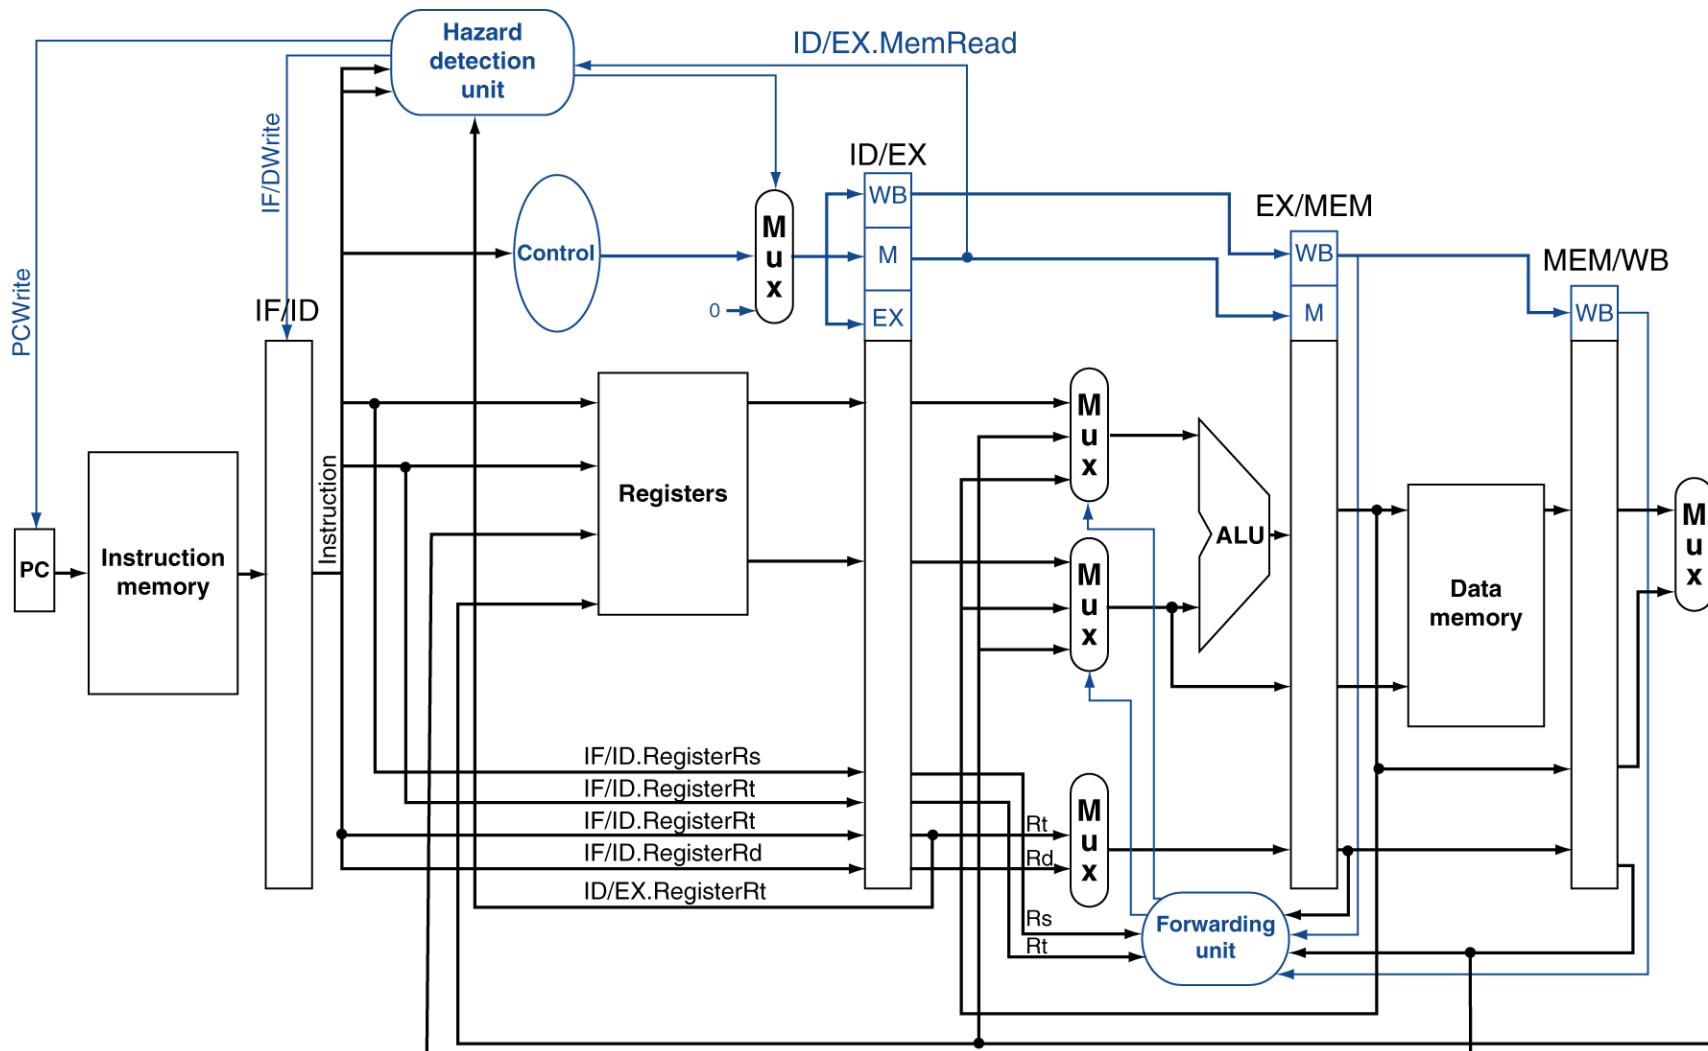
Stall/Bubble in the Pipeline



Stall/Bubble in the Pipeline



Questions about Data Hazards?



Consider the code

```
addi    $s0, $s0, 4  
lw      $t0, 0($s0)  
sub    $t1, $t2, $t2  
add    $t0, $t0, $t1
```

Does this code require a forward, a stall, both, or neither?

- A. Forward
- B. Stall
- C. Both
- D. Neither

Stalls and Performance

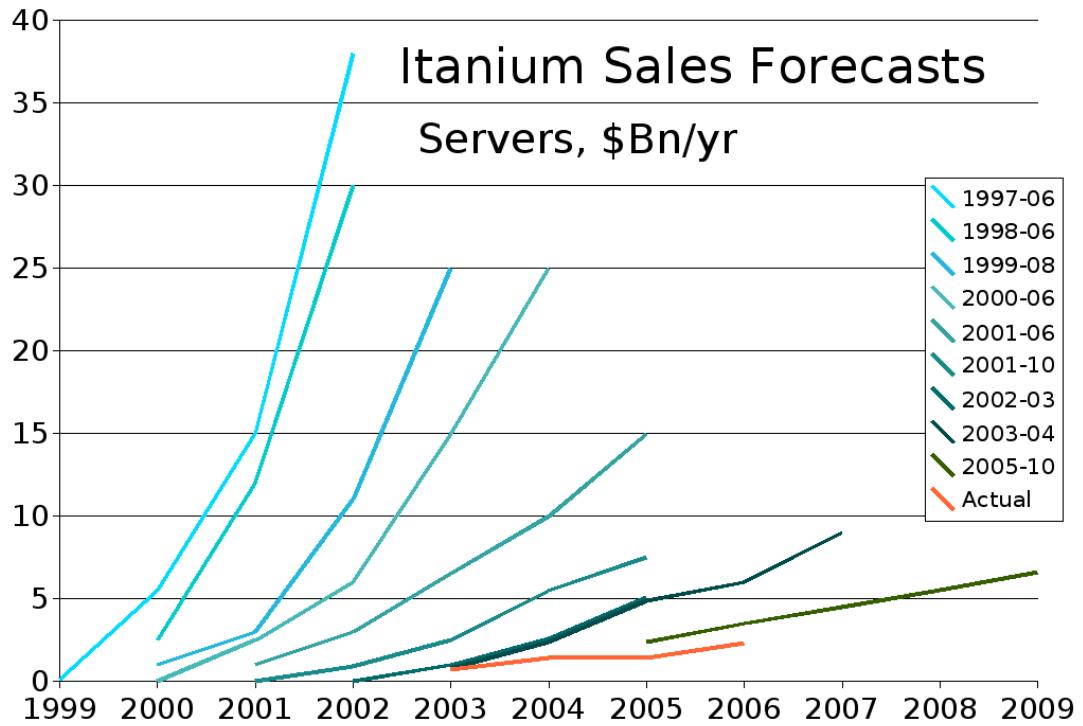
- Stalls reduce performance
 - But are required to get correct results
- Can rearrange code to avoid hazards and stalls

Dealing with Data Hazards

- As an ISA designer, you have a choice between reordering instructions in software or hardware. Which might you choose and why?

Selection	HW or SW	
A	Software	Compilers have a large window of instructions available to do reordering to eliminate hazards
B	Software	Detecting data hazards in hardware can be difficult and expensive
C	Hardware	Hardware knows at runtime the actual dependencies and can exploit that knowledge for better reordering
D	Hardware	Exposing the number of required stalls violates the abstraction between hardware and software

CS History: Intel Itanium Chip



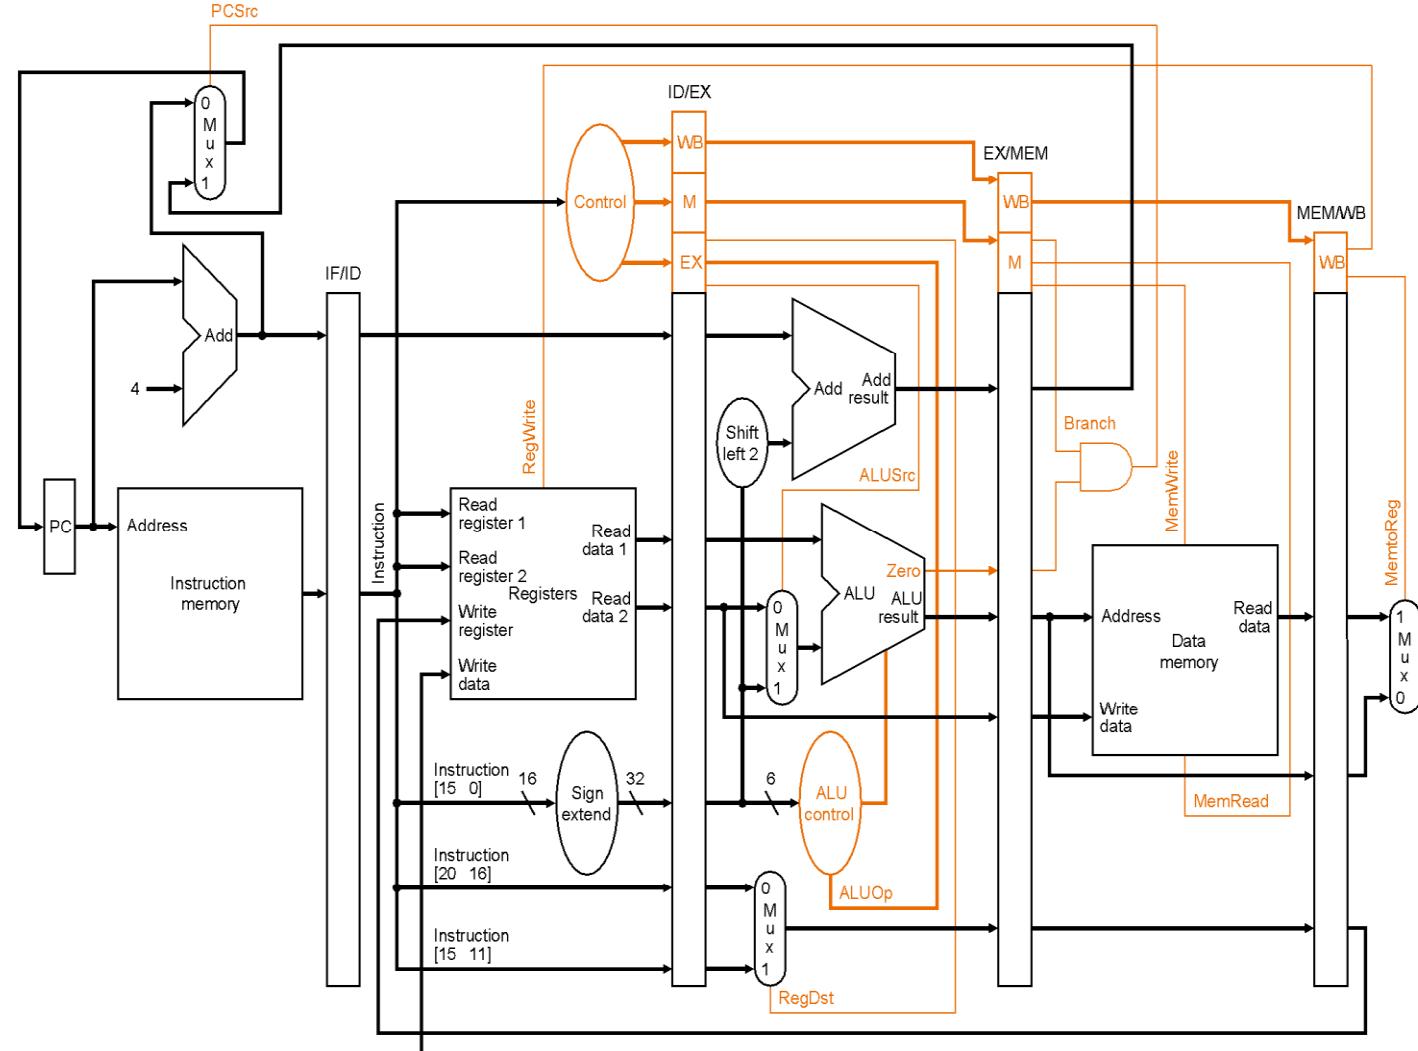
Arch dude, CC BY-SA, via Wikimedia Commons

- Intel Chip launched in 2001 that used the VLIW (Very Long Instruction Word) ISA
- This ISA was designed to do all code reordering at compile time, rather than at runtime
- Designed for servers/high-performance, goal eventually desktop market
- Performance was disappointing, especially when emulating x86
- “Itanium's promise ended up sunken by a lack of legacy 32-bit support and difficulties in working with the architecture for writing and maintaining software” - Techspot

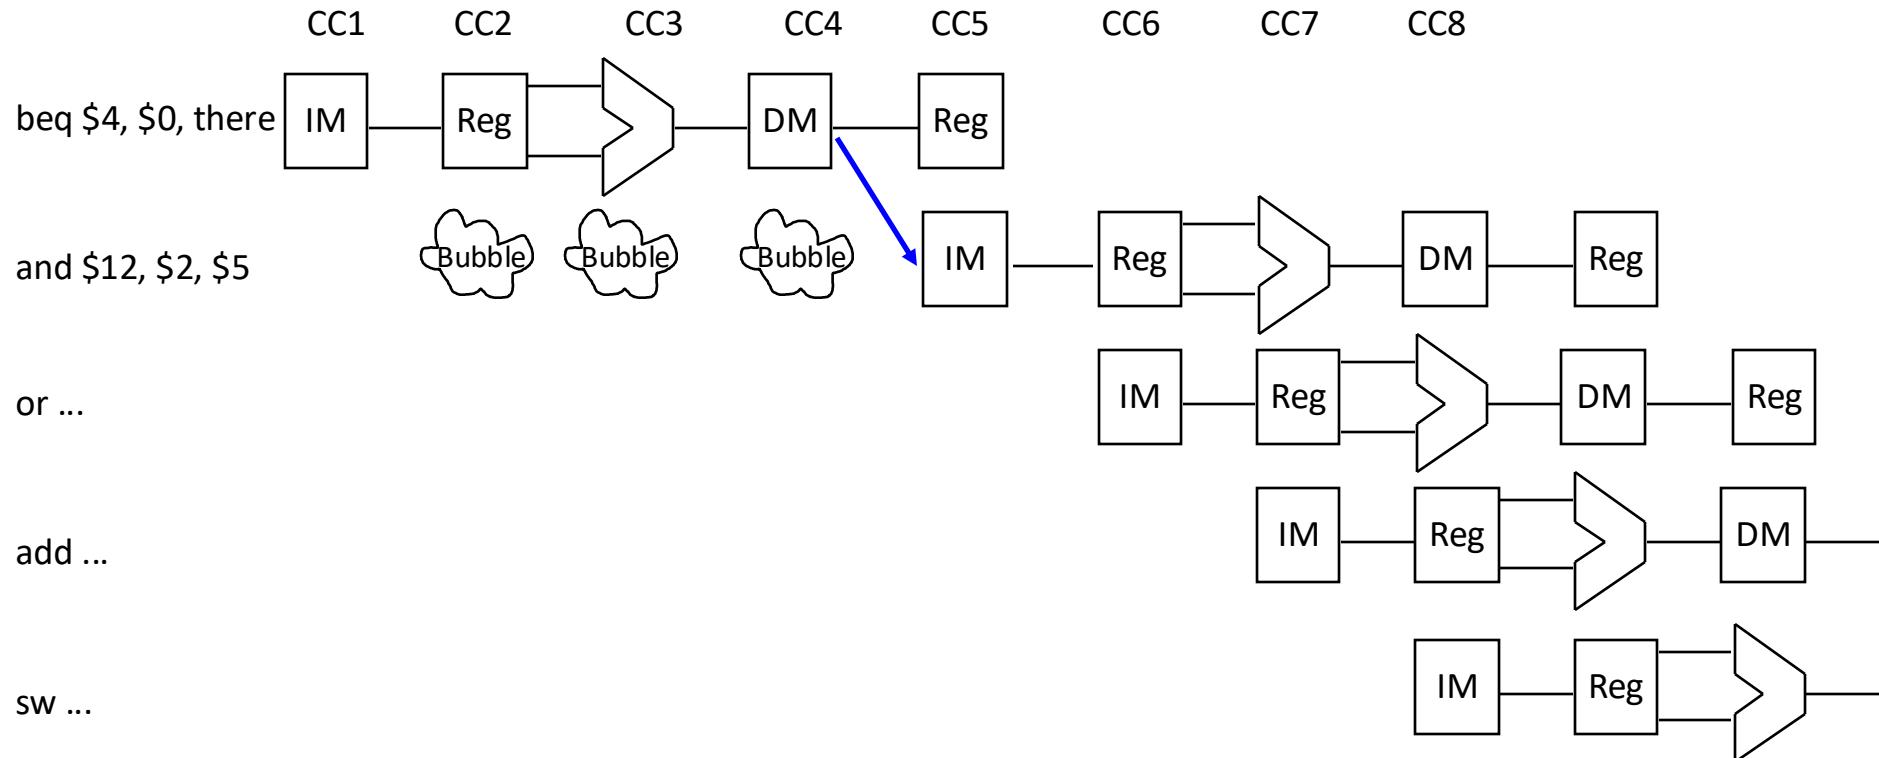
Stalling the pipeline

Given this pipeline where branches are resolved by the ALU – let's assume we stall until we know the branch outcome. How many stall cycles we need per branch?

Selection	cycles
A	0
B	1
C	2
D	3
E	4



Stalling for Branch Hazards

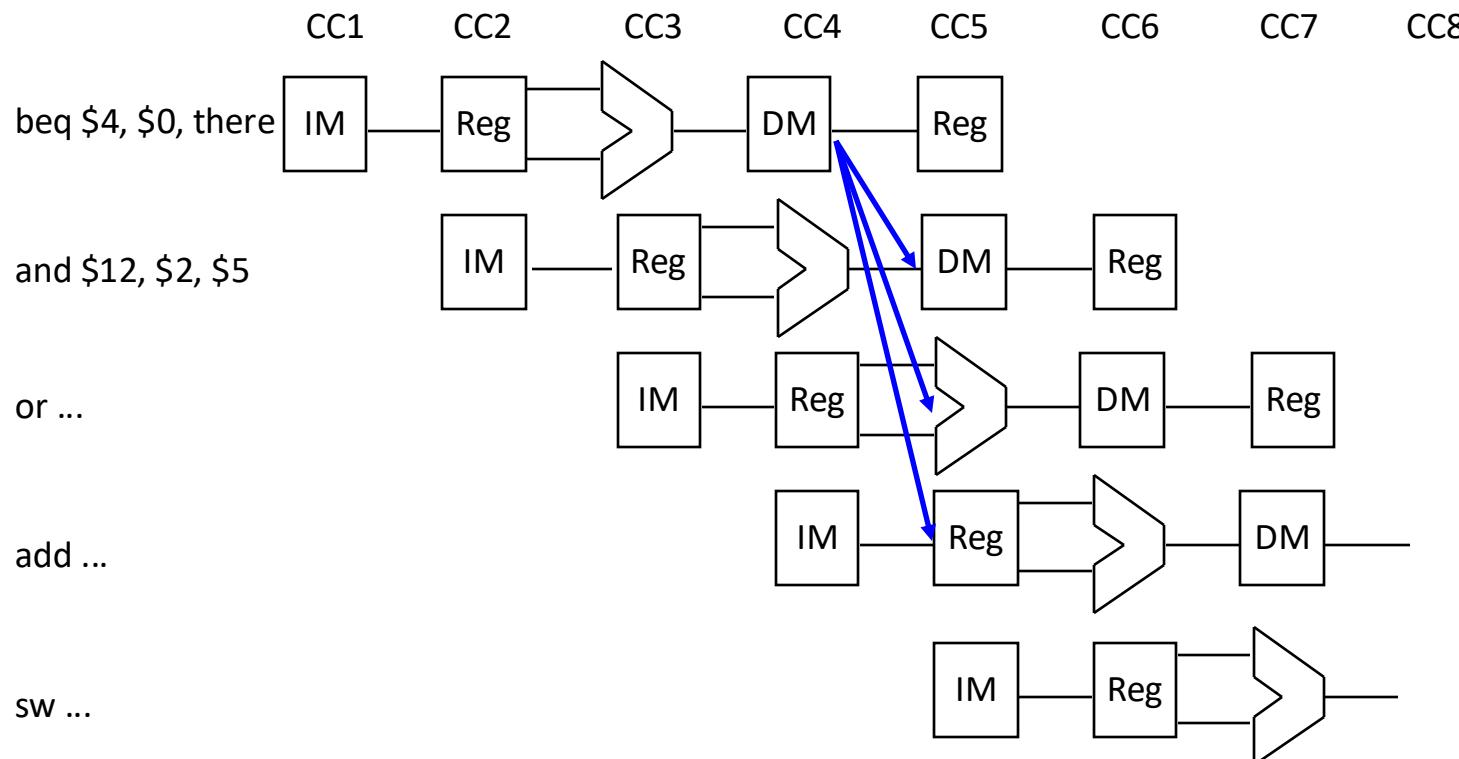


Stalling for Branch Hazards

- Seems wasteful, particularly when the branch isn't taken.
- Makes all branches cost 4 cycles.
- What if we just assume the branch isn't taken?

Assume Branch Not Taken

- works pretty well when you're right



Assume Branch Not Taken

- same performance as stalling when you're wrong

