

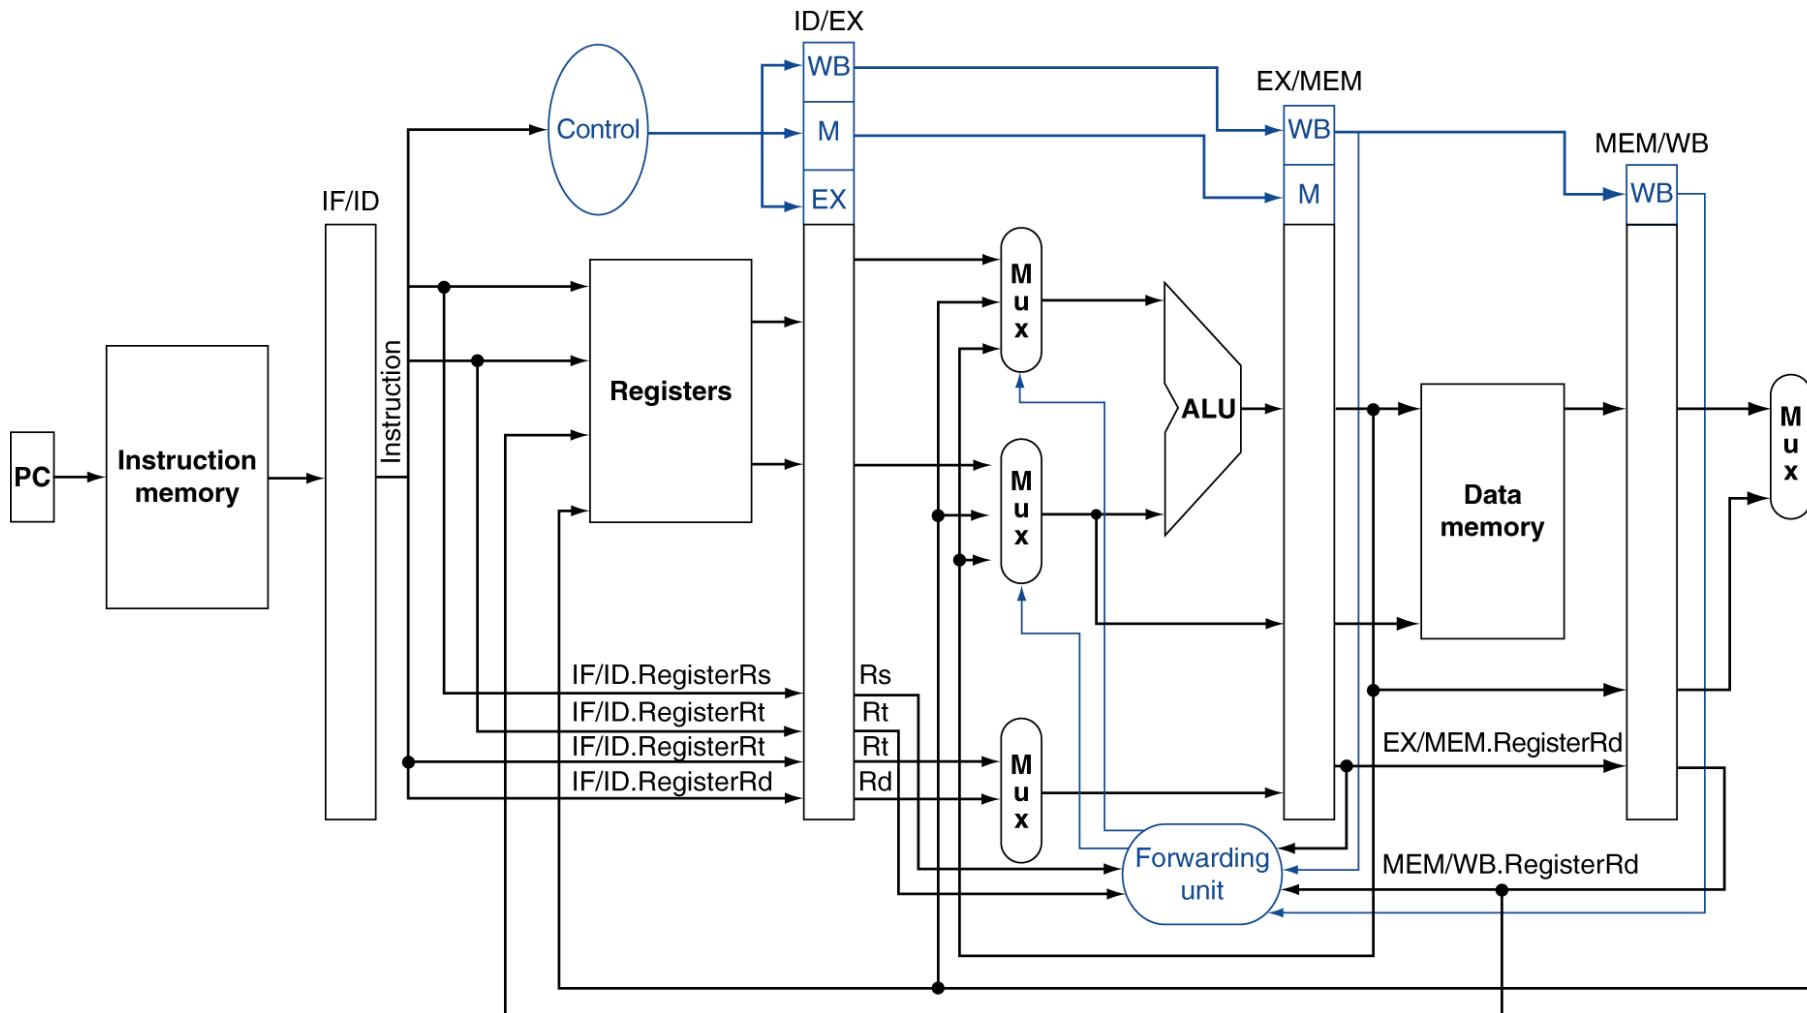
# CSCI 210: Computer Architecture

## Lecture 30: Data Hazards

Stephen Checkoway

Slides from Cynthia Taylor

# Datapath with Forwarding



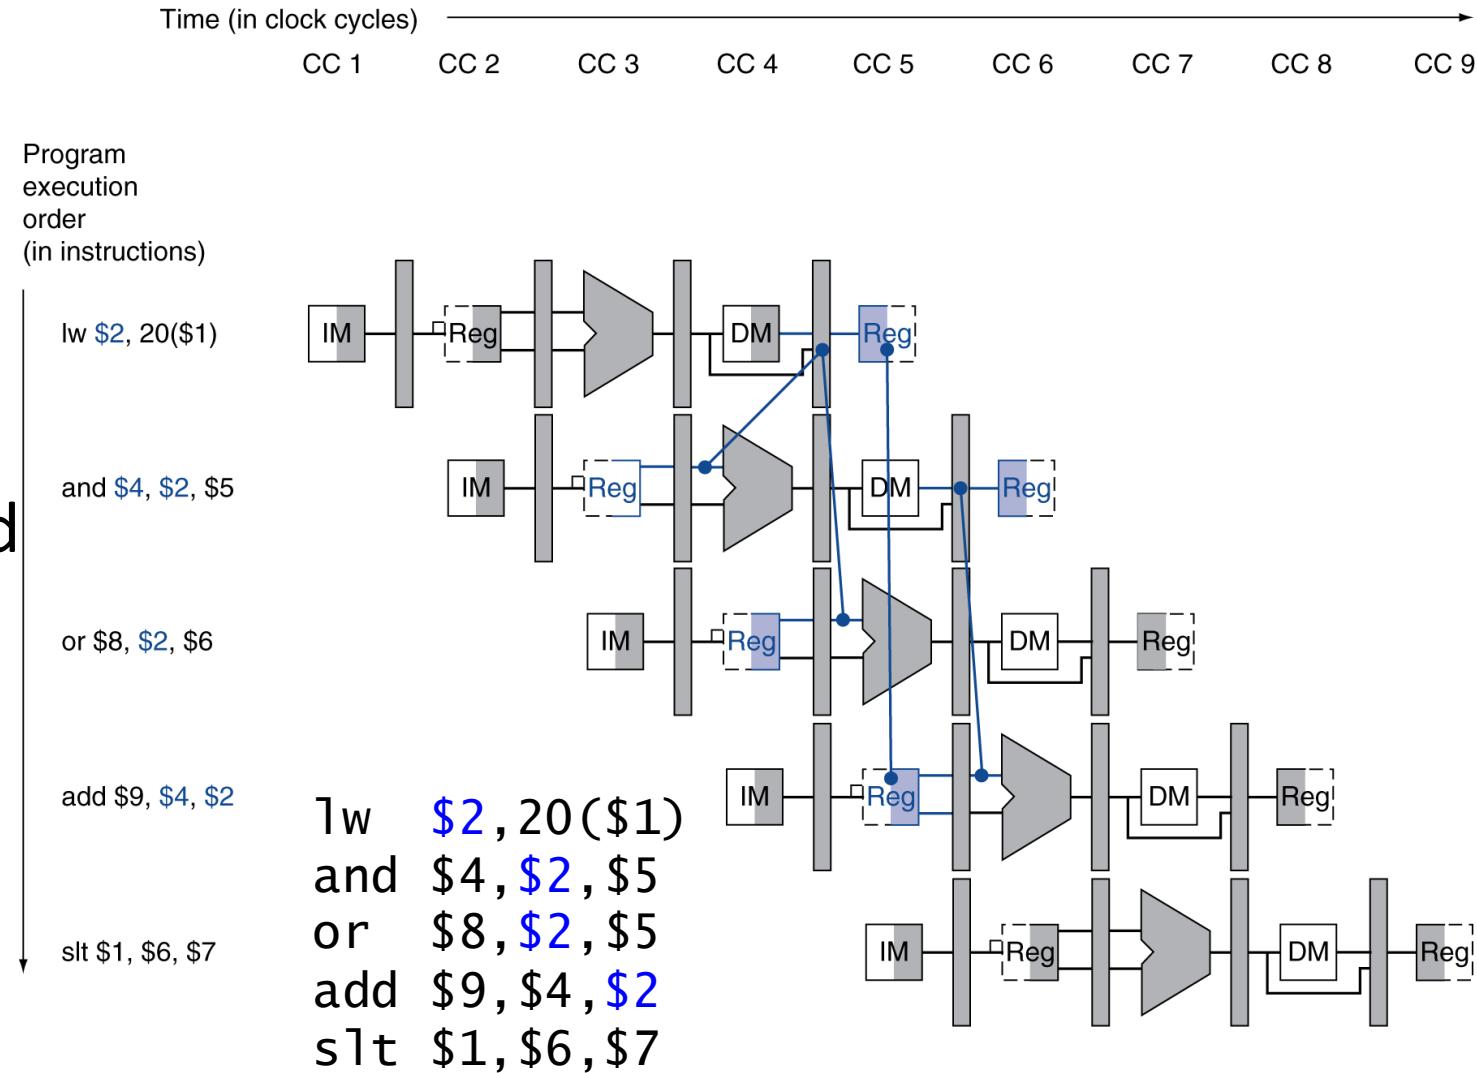
# We can best solve these data hazards

A. By stalling.

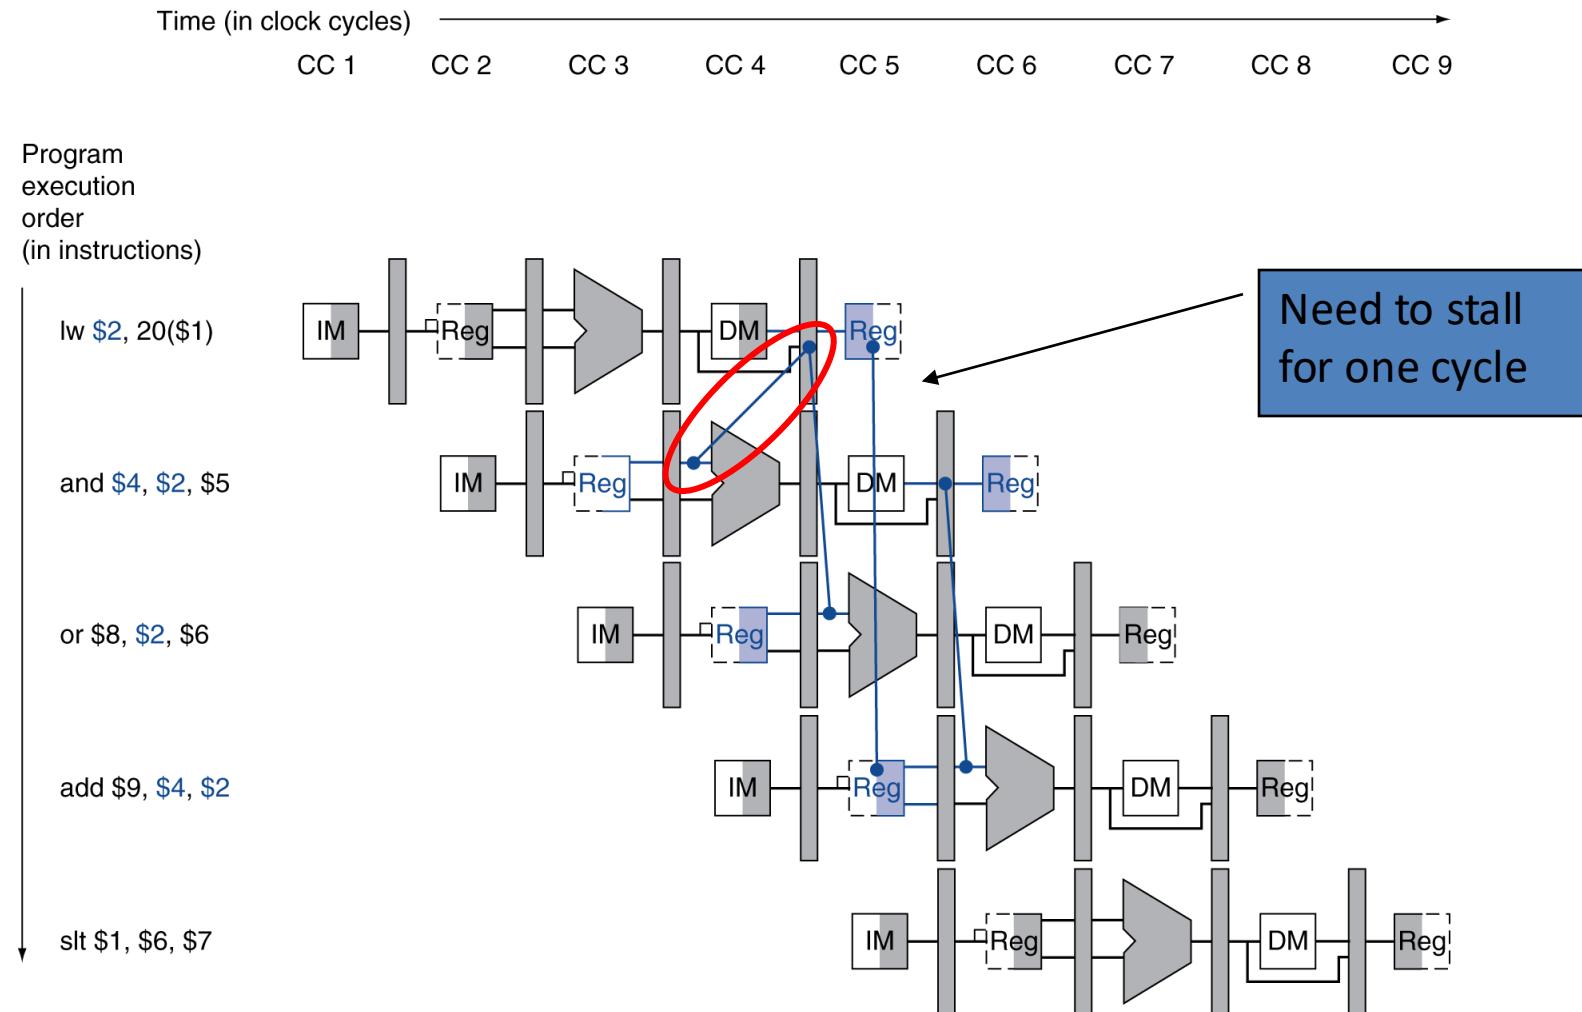
B. By forwarding.

C. By combining forwards and stalls.

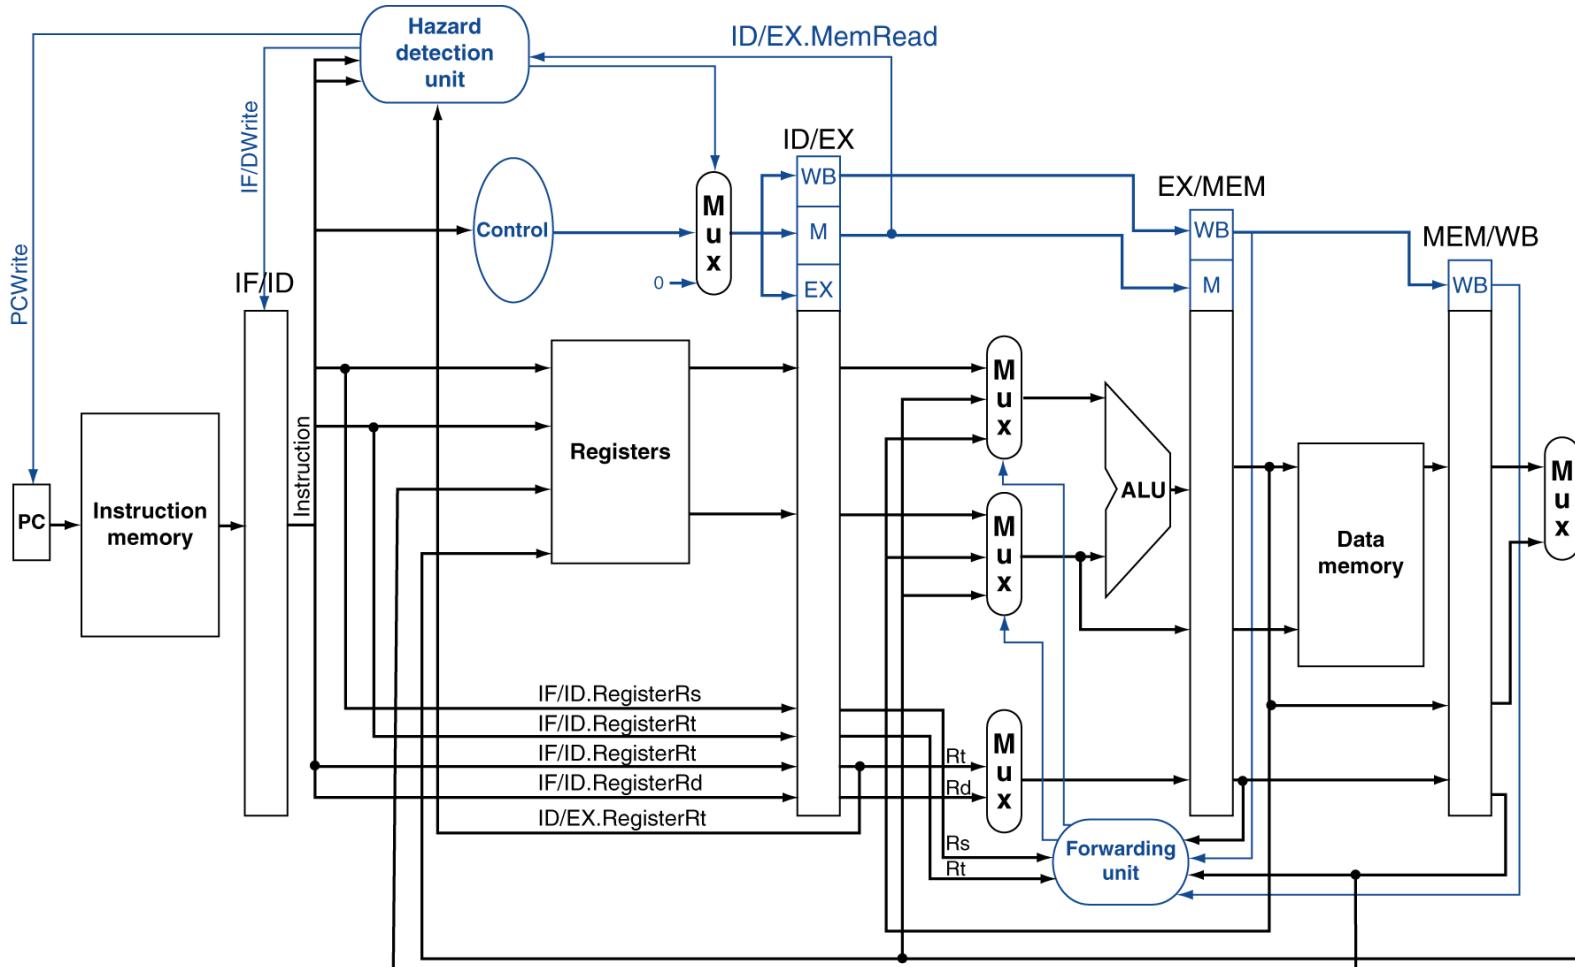
D. By doing something else.



# Load-Use Data Hazard

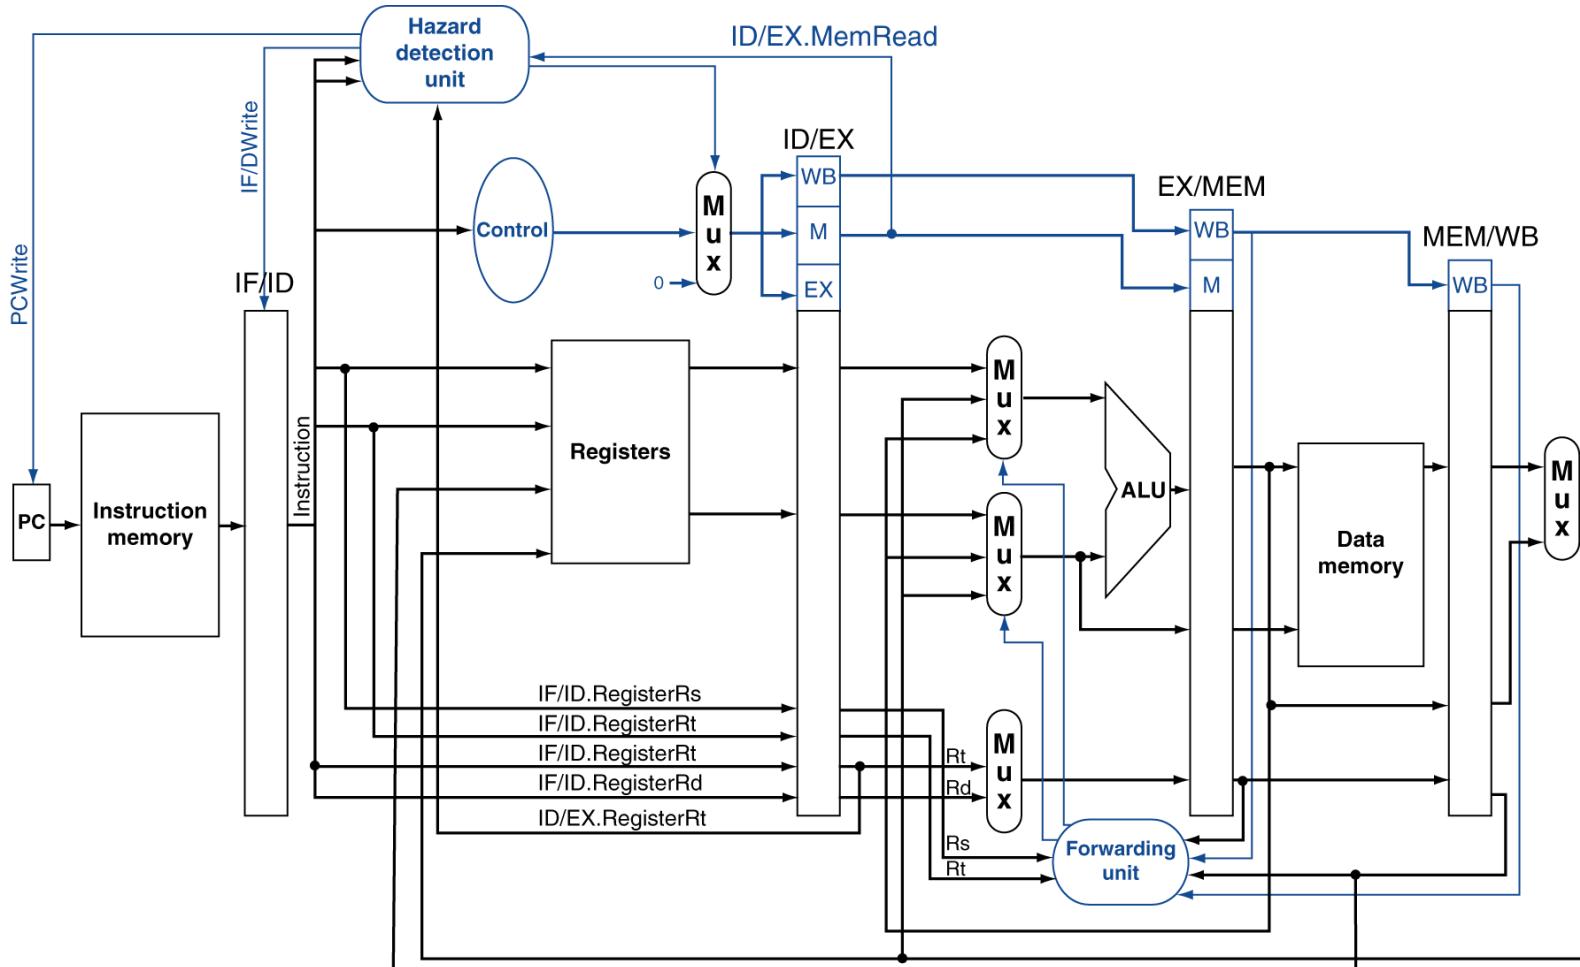


# How to Stall the Pipeline



- Detect hazard in ID stage using Hazard detection unit
  - Check if instruction in EX stage is a load with destination rs or rt

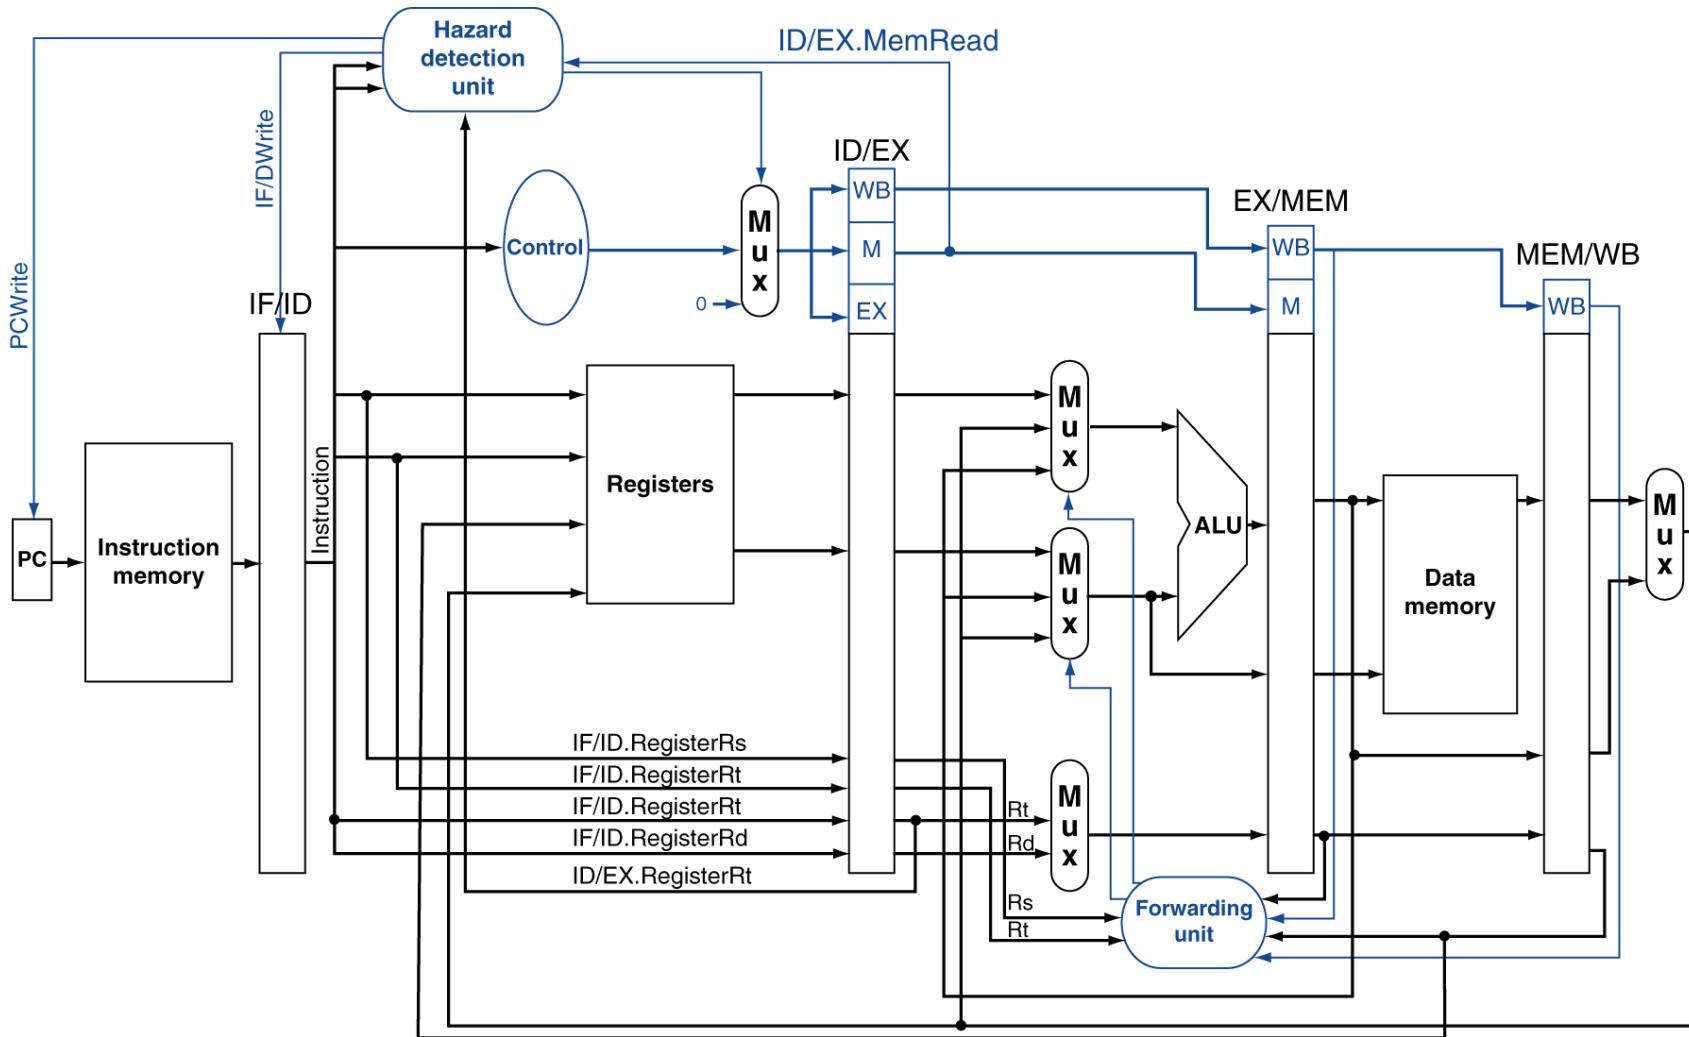
# How to Stall the Pipeline



- Force control values in ID/EX register to 0
  - EX, MEM and WB do nop (no-operation)

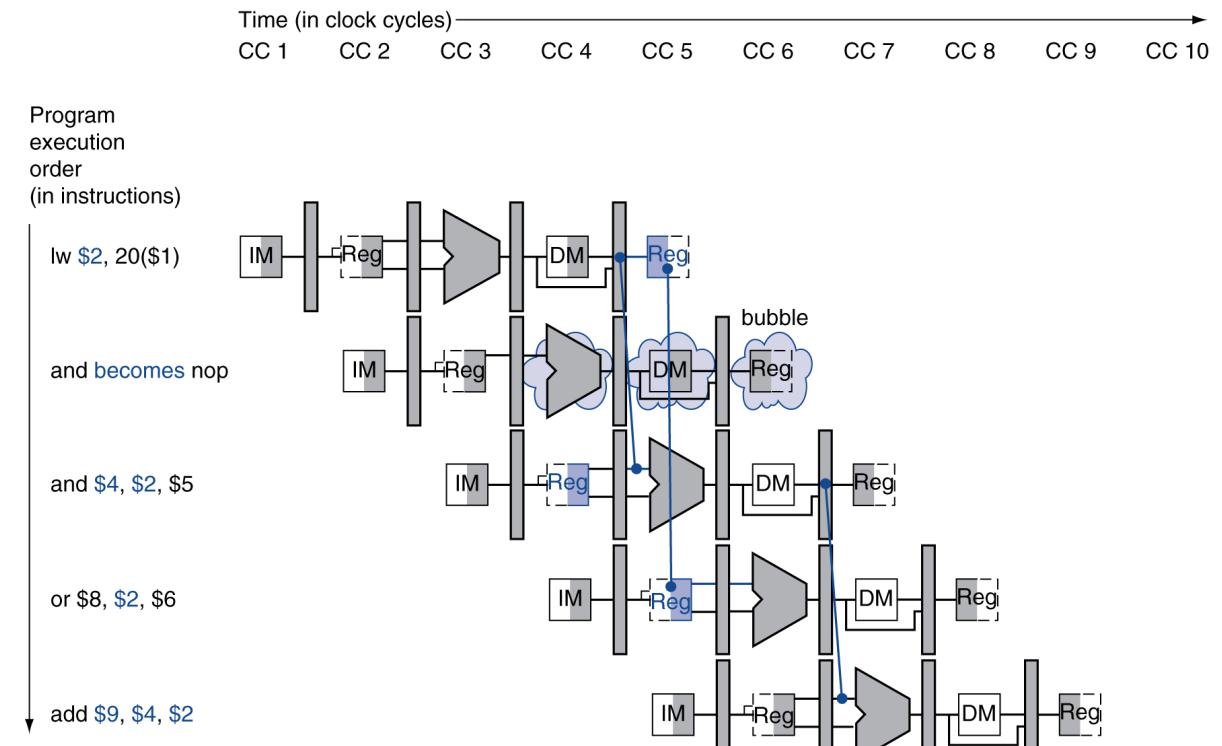
# How to Stall the Pipeline

- Prevent update of PC and IF/ID register
  - Instruction with dependency is decoded again
  - Following instruction is fetched again
  - 1-cycle stall allows MEM to read data for  $1_w$

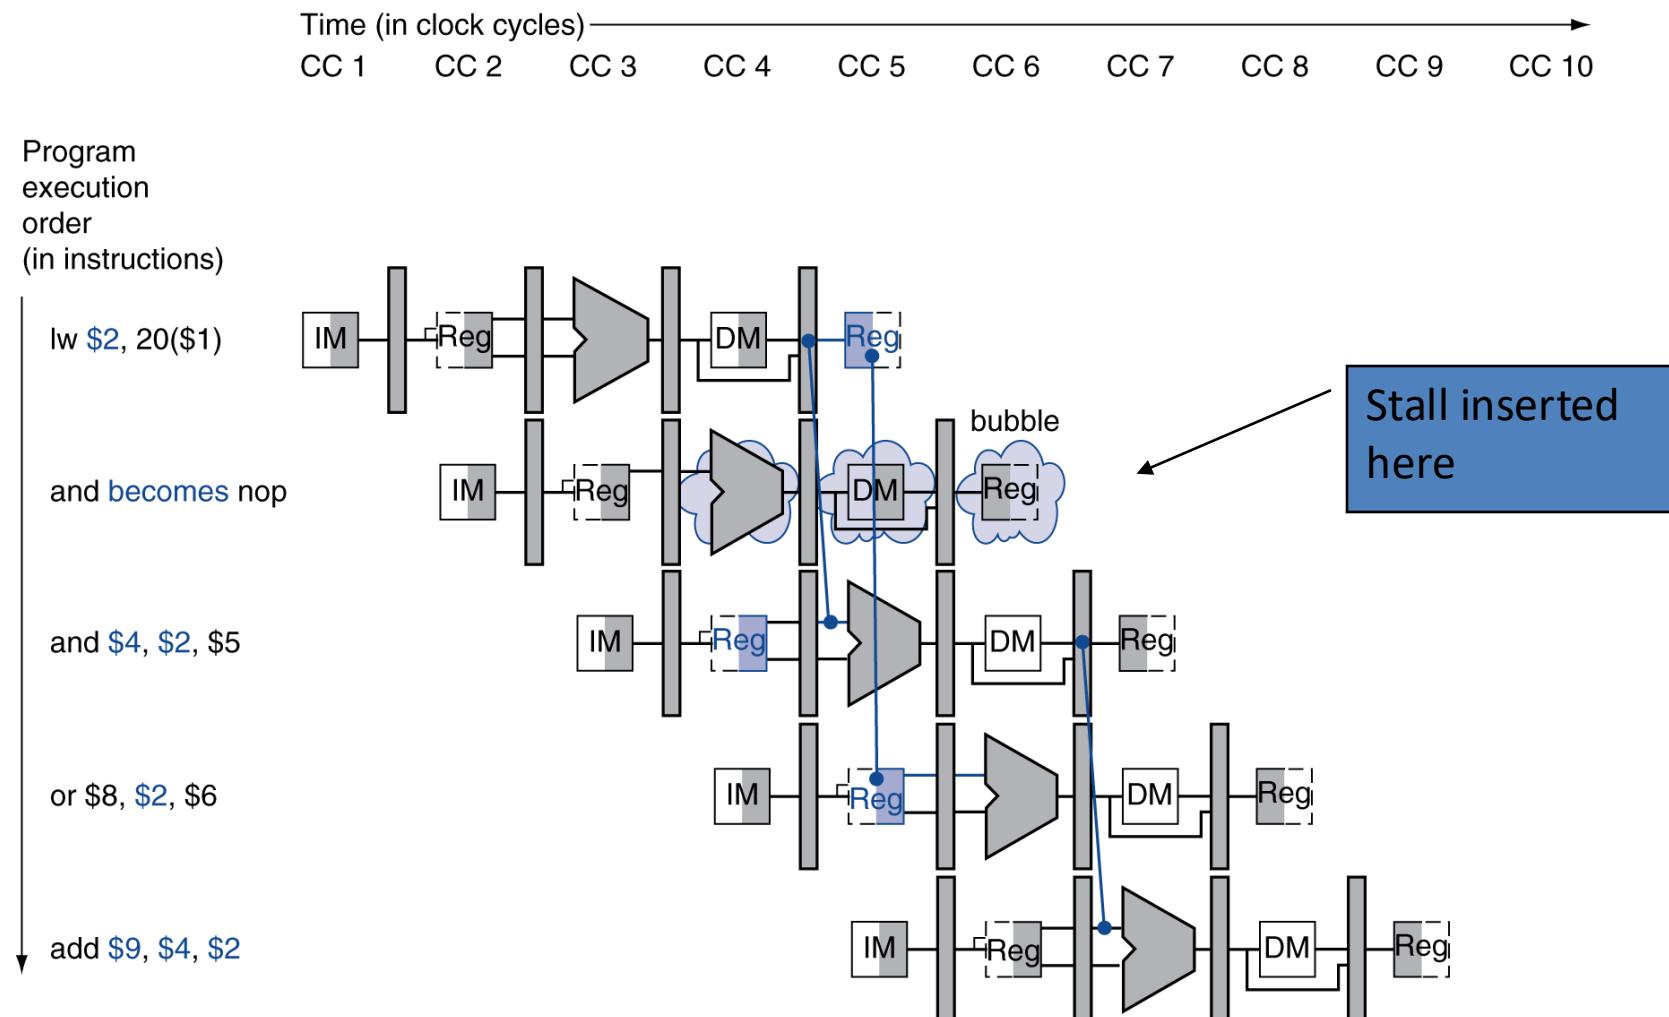


# After we add the stall

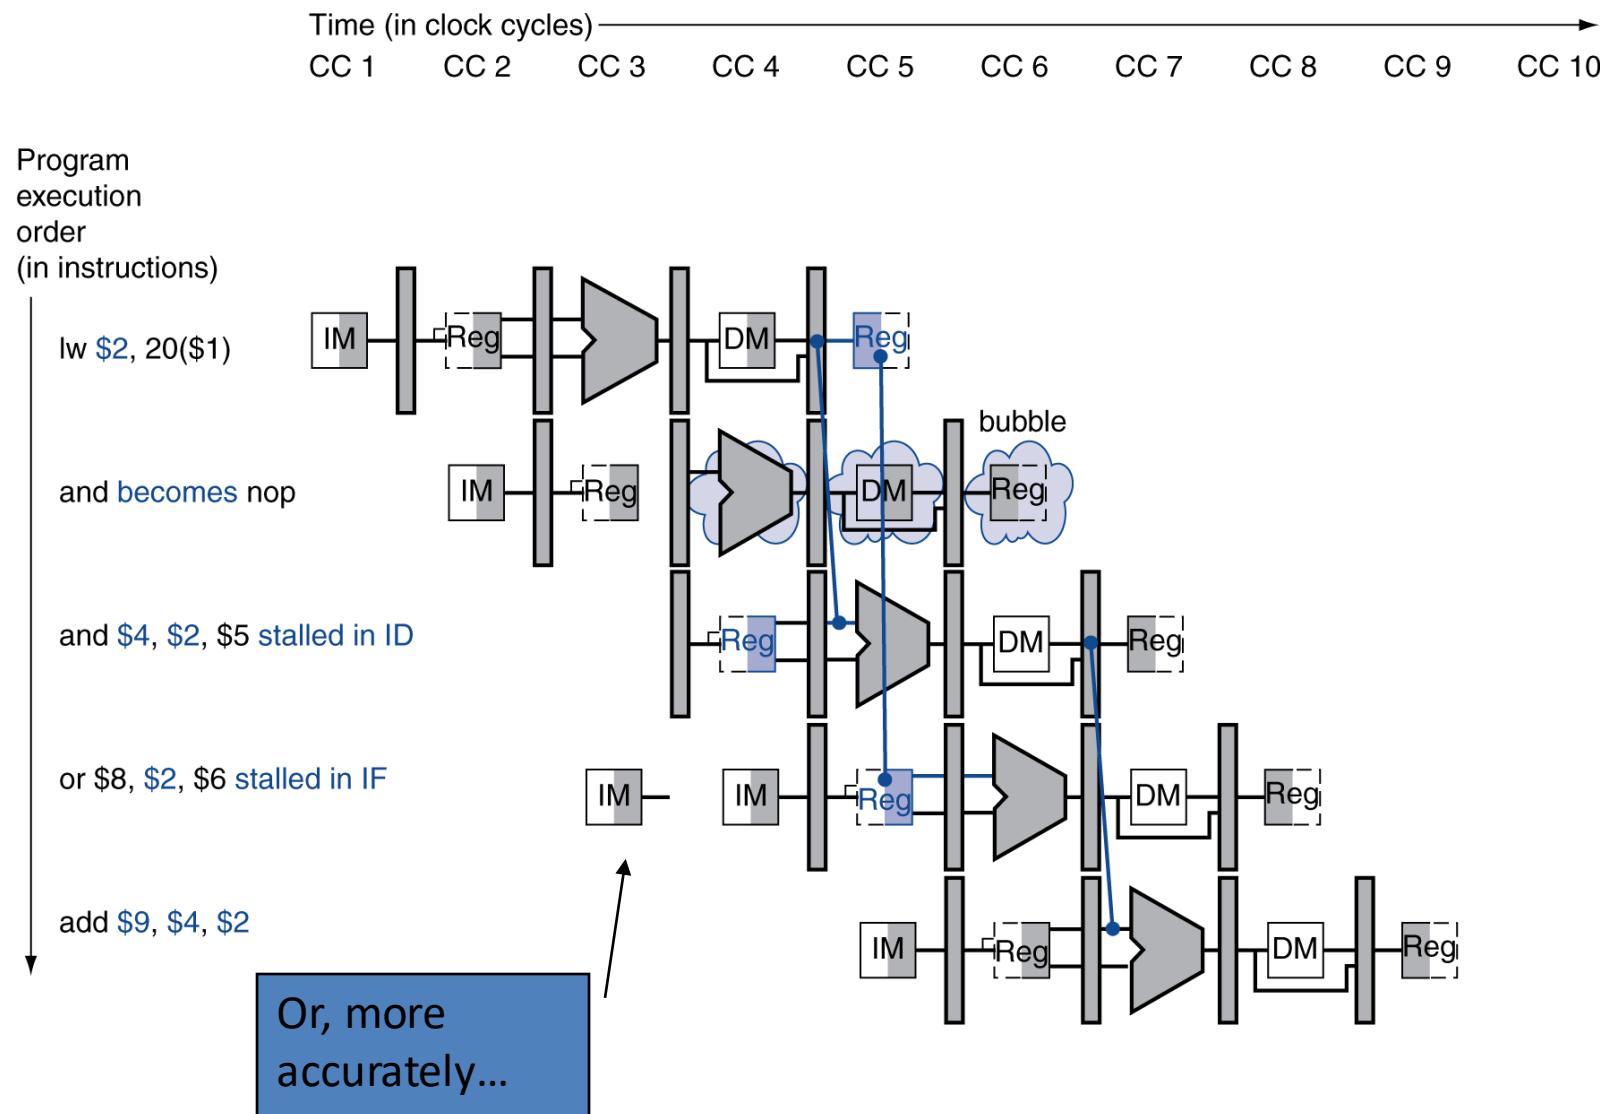
- A. Everything works with our existing forwarding
- B. We need to forward between the register files to solve the 2<sup>nd</sup> hazard
- C. We need to do something else



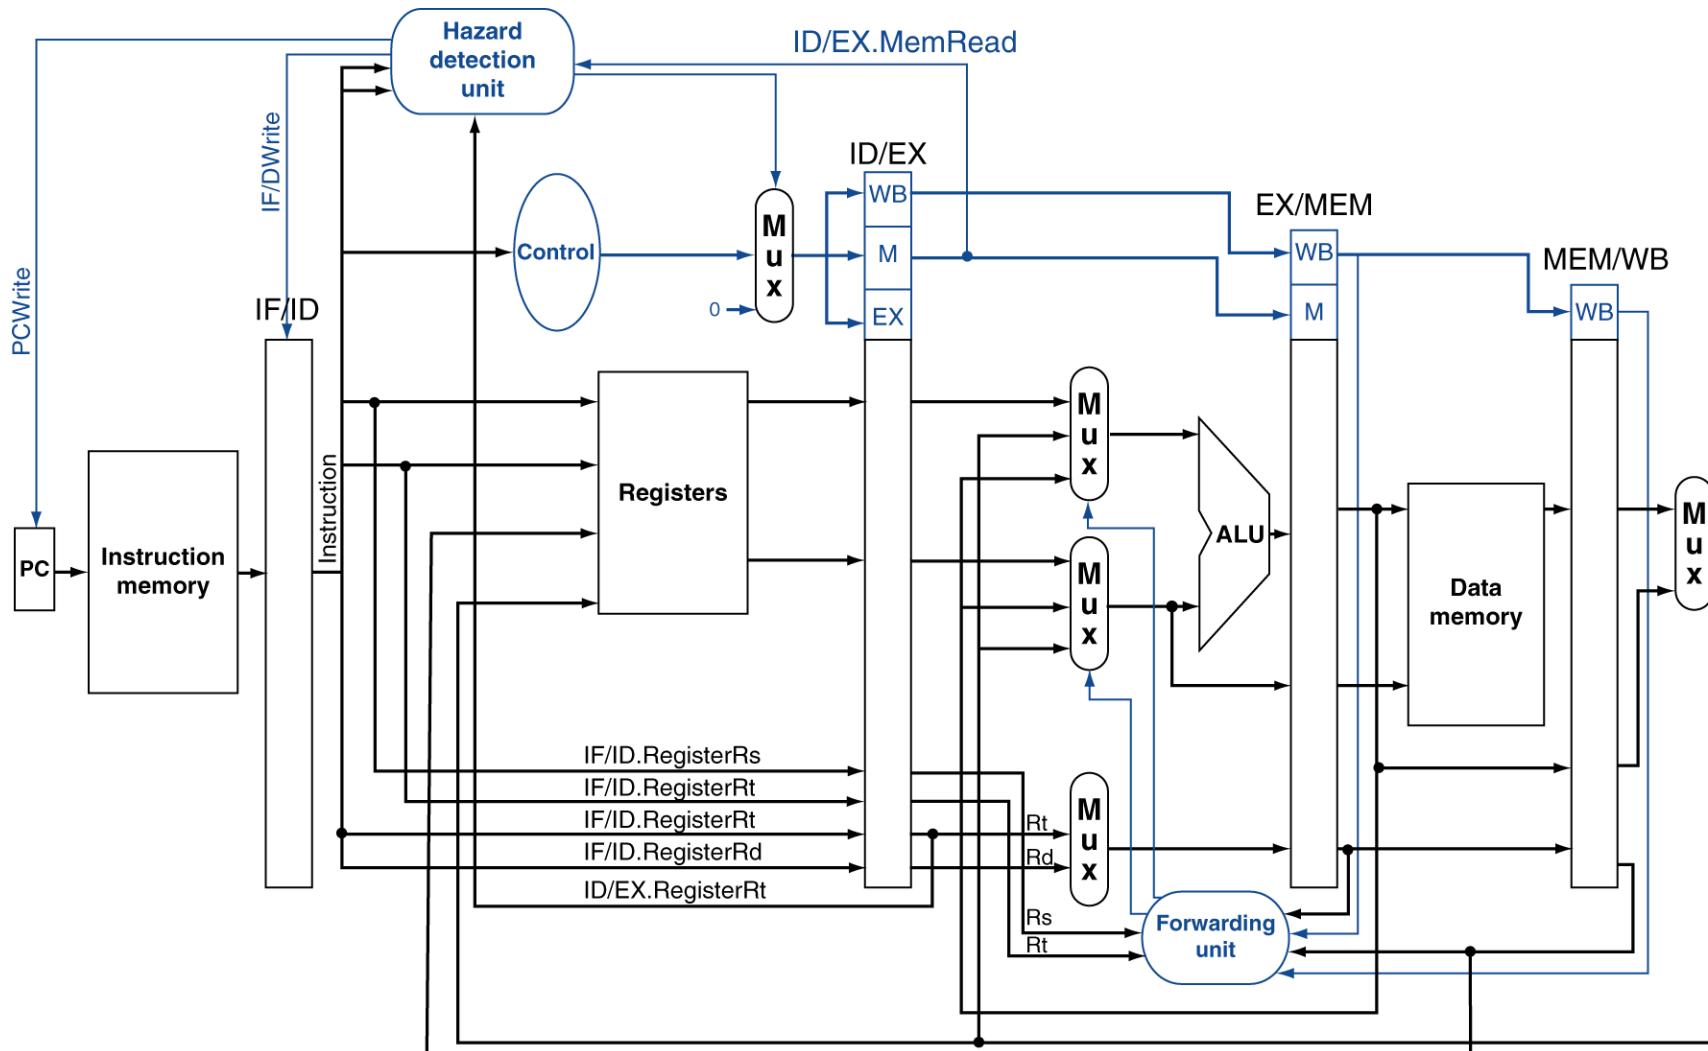
# Stall/Bubble in the Pipeline



# Stall/Bubble in the Pipeline



# Questions about Data Hazards?



Consider the code

```
addi    $s0, $s0, 4  
lw      $t0, 0($s0)  
sub    $t1, $t2, $t2  
add    $t0, $t0, $t1
```

Does this code require a forward, a stall, both, or neither?

- A. Forward
- B. Stall
- C. Both
- D. Neither

# Stalls and Performance

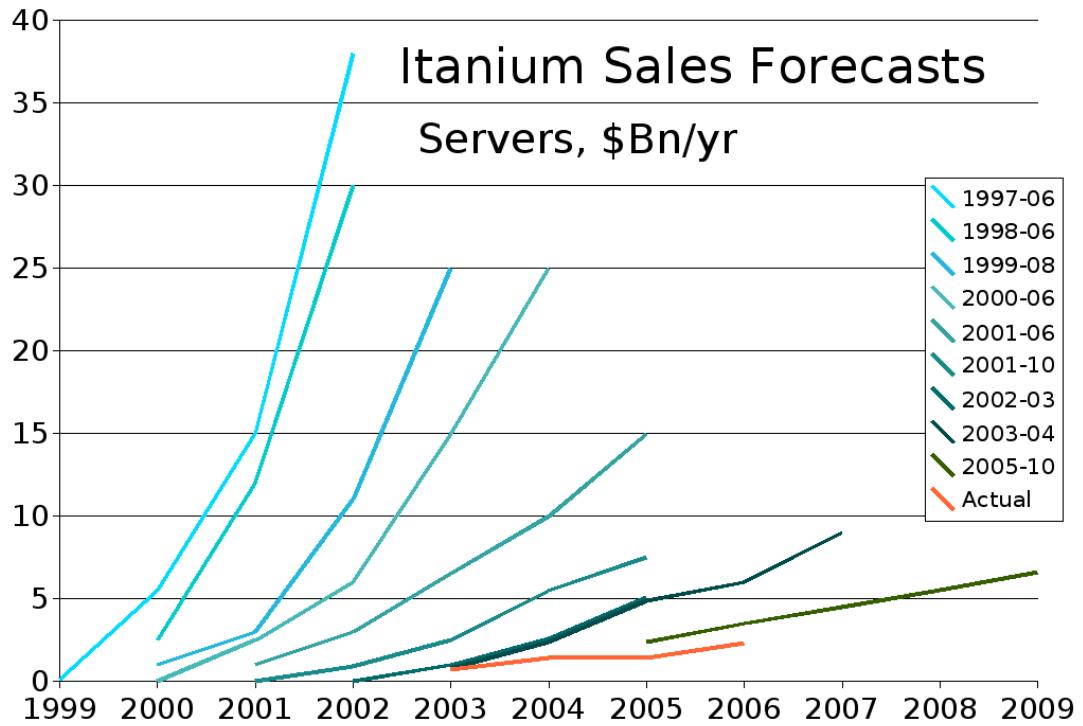
- Stalls reduce performance
  - But are required to get correct results
- Can rearrange code to avoid hazards and stalls

# Dealing with Data Hazards

- As an ISA designer, you have a choice between reordering instructions in software or hardware. Which might you choose and why?

Selection	HW or SW	
A	Software	Compilers have a large window of instructions available to do reordering to eliminate hazards
B	Software	Detecting data hazards in hardware can be difficult and expensive
C	Hardware	Hardware knows at runtime the actual dependencies and can exploit that knowledge for better reordering
D	Hardware	Exposing the number of required stalls violates the abstraction between hardware and software

# CS History: Intel Itanium Chip



Arch dude, CC BY-SA, via Wikimedia Commons

- Intel Chip launched in 2001 that used the VLIW (Very Long Instruction Word) ISA
- This ISA was designed to do all code reordering at compile time, rather than at runtime
- Designed for servers/high-performance, goal eventually desktop market
- Performance was disappointing, especially when emulating x86
- “Itanium's promise ended up sunken by a lack of legacy 32-bit support and difficulties in working with the architecture for writing and maintaining software” - Techspot