CSCI 210: Computer Architecture
Lecture 5: MIPS

Stephen Checkoway

Oberlin College

Slides from Cynthia Taylor

CS History: Nintendo 64

- Released in 1996
- Named after its 64-bit CPU
- Silicon Graphics Inc (SGI) adapted MIPS chips for supercomputers to use less power and be cheaper
- "At the heart of the system will be a version of the MIPS Multimedia Engine, a chip-set consisting of a 64-bit MIPS RISC microprocessor, a graphics co-processor chip and Application Specific Integrated Circuits" (SGI press release, 1993)
- "If it works at all, it could bring MIPS to levels of volume [SGI] never dreamed of." (Michael Slater, Microprocessor Report)
- Super Mario 64 features a rabbit named MIPS after the processor

Review: Memory Instructions

- `lw   $t0, 0($t1)`
  - `$t0 = Mem[$t1+0]`
  - Loads 4 bytes from `$t1, $t1+1, $t1+2, and $t1+3`
- `sw   $t0, 4($t1)`
  - `Mem[$t1+4] = $t0`
  - Stores 4 bytes at `$t1+4,  $t1+5,  $t1+6, and $t1+7`
- These instructions are the cornerstones of our being able to go to and from memory

Review: If you have a pointer to address 1000 and you increment it by one to 1001. What does the new pointer point to, relative to the original pointer?

A) The next word in memory

B) The next byte in memory

C) Either the next word or byte – depends on if you use that address for a load byte or load word

D) Pointers are a high level construct – they don't make sense pointing to raw memory addresses.

E) None of the above.

If a 4-byte word is in memory at address 4203084, what is the address of the next word in memory?

A) 4203085

B) 4203088

C) 14203084

D) It depends on the value of the words in memory

E) Since a word is 4 bytes, it's not possible to have one at address 4203084

Arrays

- Arrays are stored consecutively in memory
- The **base** address points to the first element in the array
- Accessing other elements in the array requires adding an **offset** to the base address
  - The offset to use is the index of the array element * the size of one element

Memory Operand Example 1

- C code:

  g = h + A[8];
  - g in $s1, h in $s2, base address of A in $s3, A is an array of 4 byte ints

- Compiled MIPS code:
  - Index 8 requires offset of 32

  ```
  lw  $t0, 32($s3)
  add $s1, $s2, $t0
  ```

Translate to MIPS

- C code: g = h + A[5];
  - g in $s1, h in $s2, base address of A in $s3.
  - A is an array of 4-byte ints

A.
```
lw   $t0, 5($s3)
add  $s1, $s2, $t0
```

B.
```
lw   $t0, 20($s3)
add  $s1, $s2, $t0
```

C.
```
lw   $t0, $s5
add  $s1, $s2, $t0
```

D.
```
lw   $t0, $s3
add  $s1, $s2, $t0
```

Memory Operand Example 2

- C code:

  A[12] = h + A[8];
  - h in $s2, base address of A in $s3

- Compiled MIPS code:

  - Index 8 requires offset of 32

```
lw  $t0, 32($s3)      # load word
add $t0, $s2, $t0
sw  $t0, 48($s3)      # store word
```

When a 2-byte integer is stored in byte-addressed memory (occupying two consecutive bytes), is the most significant byte (MSB) stored in the lower address or the higher address?

A. High

0000 1111	0000 0000	= 15
00241000	00241001	

B. Low

0000 0000	0000 1111	= 15
00241000	00241001	

C. It Depends

Byte ordering

- Big-endian: Most significant byte in lowest address
  - **MIPS**, Network byte order, Motorola 68000, PowerPC (usually), …
- Little-endian: Most significant byte in highest address
  - Intel x86, x86-64, ARM (usually), …
- Bi-endian: Switchable between big and little endian
  - ARM, PowerPC, Alpha, SPARC, …
- Middle-endian/mixed-endian
  - Bytes not stored in either order, at least in some cases
  - Words stored one endian, bytes within words stored another endian
  - PDP-11 stored 16-bit words in little-endian order, but 32-bit double words as pairs of 16-bit words with the most significant word of the double word stored first (i.e., big-endian order)

Big-endian means most significant byte/digit/piece comes first, little-endian means least significant byte/digit/piece comes first. Mixed-endian means not in order.

Which row of the table correctly identifies the endianness of date formats?

	US (MM-DD-YYYY)	Most of the world (DD-MM-YYYY)	ISO 8601 (YYYY-MM-DD)
A	Little	Mixed	Big
B	Big	Little	Mixed
C	Mixed	Little	Big
D	Mixed	Big	Little
E	Little	Big	Mixed

Questions about Memory?

Immediate Operands

- Constant data specified in an instruction
  - `addi    $s3, $s3, 4`
  - `li      $t0, -25      # Pseudoinstruction: addi $t0, $zero, -25`
  - `ori     $v0, $t8, 1`

Subtract 2 from $s0 and store in register $s1

A. `addi $s0, $s1, -2`

B. `addi $s1, $s0, -2`

C. `subi $s0, $s1, 2`

D. `subi $s1, $s0, 2`

E. More than one of the above

Pseudoinstructions

- `move dest, src` => `add dest, $zero, src`
- `subi dest, src, imm` => `addi dest, src, -imm`
- `li dest, imm` => `addi dest, $zero, imm`

- More complicated expansions are possible when the immediate value is too large, MARS simulator will show you how it expands pseudoinstructions

MIPS Design Principles

- Simplicity favors regularity
  - fixed size instructions
  - small number of instruction formats
- Smaller is faster
  - limited instruction set
  - limited number of registers in register file
- Make the common case fast
  - arithmetic operands from the register file (load-store machine)
  - **allow instructions to contain small immediate operands**

Loading a large number into a register

- Immediates are limited to 16 bits
  - -32768 to 32767 or 0 to 65535
- Numbers outside this range need to be loaded into registers before being used
- load upper immediate instruction sets the most-significant 16 bits of a register
  - ```
    lui   $t0, 0x1234
    ori   $t0, $t0, 0x5678
    ```
- When li is given a value that's too large, the assembler expands it to lui/ori

MIPS Questions?

Reading

- Next lecture: Number Representation
  - Section 2.4

- Problem Set 1 – due Friday