

CS 241: Systems Programming

Lecture 21. Binary and Formatted I/O

Spring 2020

Prof. Stephen Checkoway

Review from last time

```
// Open and close a file.  
FILE *fopen(char const *path, char const *mode);  
int fclose(FILE *stream);  
  
// Read or write a character.  
int fgetc(FILE *stream);  
int fputc(int ch, FILE *stream);  
  
// Read or write a string.  
char *fgets(char *str, int size, FILE *stream);  
int fputs(char const *str, FILE *stream);  
  
// Formatted output.  
int fprintf(FILE *stream, char const *format, ...);
```

re-reading a file

```
void rewind(FILE *stream);
```

- resets the file back to the start of the stream
- NOTE: no return value
 - zero out and check errno for problems

Actually is an alias for...

Changing location in a file

```
int fseek(FILE *stream, long offset, int whence);  
int fseeko(FILE *stream, off_t offset, int whence);
```

- ▶ Return 0 on success, -1 on failure (errno)

Reposition location in stream

- ▶ `offset` is number of bytes added to position specified by `whence`
 - `SEEK_SET` - start of the file
 - `SEEK_CUR` - current position
 - `SEEK_END` - end of the file

Getting location in file

```
long ftell(FILE *stream);
```

```
off_t ftello(FILE *stream);
```

- returns current offset on success
- returns -1 and sets errno on failure

fseeko and ftello are specified by POSIX but not C

`off_t` is an integral type representing file sizes (often 8 bytes)

How can we get the size of a file to which we have an open `FILE *stream`?

```
A. off_t size = ftello(stream);      E. off_t pos = ftello(stream);
                                       fseeko(stream, 0, SEEK_END);
B. fseeko(stream, 0, SEEK_SET);      off_t size = ftello(stream);
   off_t size = ftello(stream);      fseeko(stream, pos,
                                       SEEK_SET);
C. fseeko(stream, 0, SEEK_CUR);
   off_t size = ftello(stream);
D. fseeko(stream, 0, SEEK_END);
   off_t size = ftello(stream);
```

```
int get_file_size(char const *path, off_t *size) {  
    FILE *fp = fopen(path, "rb");  
    if (!fp)  
        return -1;  
  
    int ret = -1;  
    if (fseeko(fp, 0, SEEK_END) == 0) {  
        if ((*size = ftello(fp)) != -1)  
            ret = 0;  
    }  
    int err = errno;  
    fclose(fp);  
    errno = err;  
    return ret;  
}
```

```
int get_file_size(char const *path, off_t *size) {  
    FILE *fp = fopen(path, "rb");  
    if (!fp)  
        return -1;  
  
    int ret = -1;  
    if (fseeko(fp, 0, SEEK_END) == 0) {  
        if ((*size = ftello(fp)) != -1)  
            ret = 0;  
    }  
    int err = errno;  
    fclose(fp);  
    errno = err;  
    return ret;  
}
```

Returning the size via
a pointer parameters


```
int get_file_size(char const *path, off_t *size) {  
    FILE *fp = fopen(path, "rb");  
    if (!fp)  
        return -1;  
  
    int ret = -1;  
    if (fseeko(fp, 0, SEEK_END) == 0) {  
        if ((*size = ftello(fp)) != -1)  
            ret = 0;  
    }  
    int err = errno;  
    fclose(fp);  
    errno = err;  
    return ret;  
}
```

Returning the size via
a pointer parameters

fclose() might change errno

```
int main(int argc, char *argv[argc]) {
    for (int i = 1; i < argc; ++i) {
        off_t size;
        if (get_file_size(argv[i], &size) == -1) {
            perror(argv[i]);
        } else {
            intmax_t s = size; // Can't print off_t directly.
            printf("%s: %jd\n", argv[i], s);
        }
    }
    return 0;
}
```

Formatted input

```
int scanf(const char *format, ...);
```

- input analog to `printf()`
- reads input from `stdin`
- uses format string to determine types
- arguments must be **pointers**
 - common error
- Stops when
 - format string is done
 - input mismatch
- returns # of successfully matched items
- returns EOF on EOF (not 0)

Example

```
#include <stdio.h>
int main(void) {
    int pairs = 0;
    int x, y;
    while (scanf(" (%d , %d )", &x, &y) == 2)
        ++pairs;
    printf("Read in %d valid pairs.\n", pairs);
    return 0;
}
```

Spaces in the format match white space characters, the %d skips white space so (1 , 2) are (3 , 4) both valid, but (0 , 1 (2 , 3 (4 , 5 gives 3 valid pairs!

scanf format string interpretation

White space matches 0 or more white space characters in the input

Ordinary characters are matched against non-whitespace

Conversion specifications: e.g., %8lx

- ▶ % to indicate start (like printf)
- ▶ * indicates not to store the value
- ▶ number for field width
- ▶ hh, h, l: size of storage character
- ▶ conversion character (see printf)
- ▶ Most conversion specifiers skip white space (all but % [...])

Assume we have an integer variable x, how do I read in a decimal value from `stdin` and assign it to x?

```
int x;
```

A. `scanf("&x");`

B. `scanf("%d", x);`

C. `scanf("%d", &x);`

D. `scanf("%d", *x);`

E. `scanf("%x", &x);`

scanf family

```
int fscanf(FILE *stream, const char *format, ...);
```

```
int sscanf(const char *str, const char *format, ...);
```

Character ranges

`%s` matches a sequence of non-whitespace characters

`%[chars]` matches a range of characters, which can include whitespace

```
char html_tag[32];  
sscanf(line, "<%31[^>]>", html_tag);
```


Assume we have a char array word, how do I read in a word of text from `stdin`?

```
char word[16];
```

A. `scanf("&word");`

B. `scanf("%15s", word);`

C. `scanf("%16s", word);`

D. `scanf("%s", &word);`

E. `scanf("%s", *word);`

Useful input technique

`fgets()` / `sscanf()` pairing

- Read a line using `fgets()`
- Parses data using `sscanf()` from line

Always does bounds checking

Binary data

```
size_t fread(void *ptr, size_t size, size_t nitems,  
             FILE *stream);
```

```
size_t fwrite(void const *ptr, size_t size,  
             size_t nitems, FILE *stream);
```

- Read/write nitems number of size sized objects
- Returns the number of objects read/written which will be less than nitems for EOF or an error
- Must use feof() or ferror() to determine which occurred

```
int x = 42;
```

```
float y[8];
```

```
size_t num = fread(y, sizeof(float), 8, stream);
```

```
num = fwrite(&x, sizeof(int), 1, stream);
```

Aside: Printing to a string

```
int snprintf(char *str, size_t size,  
             char const *format, ...);
```

- Writes at most size-1 bytes into str and null terminates
- Returns number of bytes that are printed (or would be printed if the string were large enough), negative on error

```
char message[100];  
snprintf(message, sizeof message, "%s %d",  
         some_string, some_int);  
  
/* or */  
size_t size = snprintf(0, 0, "%s %d", str, x);  
char *message = malloc(size + 1);  
snprintf(message, size+1, "%s %d", str, x);
```

DANGER: Format String Attacks

Don't just print arbitrary users' strings

- `printf(str);`

If the attacker sets the value of line they can

- Cause it to reveal other program data by printing it from the program stack (e.g., "`%x%x%x`")
- Can cause it to change program data by using "`%n`" which stores # of chars printed so far

Use one of these instead

- `printf("%s", str);`
- `fputs(str, stdout);`

In-class exercise

<https://checkoway.net/teaching/cs241/2020-spring/exercises/Lecture-21.html>

Grab a laptop and a partner and try to get as much of that done as you can!