

CS 241: Systems Programming

Lecture 25. Function Pointers

Spring 2020

Prof. Stephen Checkoway

Function pointers

Function pointers are pointers that point to...functions

Syntax

- `return_type (*var)(parameters);`
- `int (*f1)(void);` // f1 is a pointer to a function returning an int
- `struct foo *(*f2)(double, size_t) = blah;`

Calling a function pointer (two options)

- Pretend it's a function: `int x = f1();`
- Dereference it first: `struct foo *p = (*f2)(2.3, 82);`

Aside: C is *super* weird

Function call operator (...) only applies to function pointers

Functions **decay** to pointers to the function

When calling `foo(5)`, `foo` decays to a pointer and then the call happens

Assuming we have a function `void foo(int x)`, these are identical

- ▶ `foo(3)` // decay -> call
- ▶ `(&foo)(3)` // address of -> call
- ▶ `(*foo)(3)` // decay -> dereference -> decay -> call
- ▶ `(*&foo)(3)` // address of -> dereference -> decay -> call
- ▶ `(&*foo)(3)` // decay -> dereference -> address of -> call

Example

```
#include <stdio.h>
```

```
void foo(void) { puts("foo"); }
```

```
void bar(void) { puts("bar"); }
```

```
void qux(void) { puts("qux"); }
```

```
// An array of function pointers
```

```
void (*table[])(void) = { foo, bar, qux };
```

```
int main(int argc, char *argv[argc]) {
```

```
    void (*ptr)(void) = table[argc % 3];
```

```
    ptr();
```

```
    return 0;
```

```
}
```

An actual use case

```
int atexit(void (*handler) (void));
```

- Call `atexit` and pass it a function (pointer)
- When the program exits normally (via `exit(3)` or returning from `main`), the function is called
- `_exit(2)` [defined by POSIX] or `_Exit(3)` [defined by C] don't call the `atexit` handlers
- `Atexit` handlers are called in reverse order
- `Atexit` handlers must not call `exit(3)`

What does this code print?

```
#include <stdio.h>
#include <stdlib.h>

void foo(void) { puts("1"); }
void bar(void) { puts("2"); }

int main(void) {
    atexit(foo);
    puts("3");
    atexit(bar);
    exit(0);
    puts("4");
    return 0;
}
```

A. 1
2
3
4

B. 1
3
2

C. 3
1
2

D. 3
4
2
1

E. 3
2
1

Generic sorting

```
void qsort(void *base, size_t nel, size_t width,  
          int (*compare)(void const *, void const *));
```

Takes an array, `base`, of `nel` elements, each of size `width` and a comparison function, `compare` and sorts the array

`compare` gets a pointer to two elements `x` and `y` and returns `<0`, `0`, or `>0` depending on the `x < y`, `x = y`, or `x > y`

Void pointers (void *)

Void pointers are allowed to point to any object

- `int i = 5;`
`float f = 8.2f;`
`void *p = NULL; // Valid`
`p = &i; // Valid`
`p = &f; // Valid`

Void pointers can be assigned to any other pointer type

- `void *p = /* ... */;`
`double *q = p; // Valid`
`struct foo *r = p; // Valid`


```

#include <stdlib.h>
#include <stdio.h>
#include <string.h>

enum Rank { ASSISTANT, ASSOCIATE, FULL };
char const *const ranks[] = { "Assistant", "Associate", "Full" };

struct Professor {
    enum Rank rank;
    char const *name;
};

struct Professor profs[] = {
    { .rank = ASSISTANT, .name = "Roberto Hoyle" },
    { .rank = ASSISTANT, .name = "Adam Eck" },
    { .rank = FULL, .name = "John Donaldson" },
    { .rank = ASSISTANT, .name = "Sam Taggart" },
    { .rank = FULL, .name = "Bob Geitz" },
    { .rank = ASSISTANT, .name = "Cynthia Taylor" },
    { .rank = ASSISTANT, .name = "Stephen Checkoway" },
    { .rank = ASSISTANT, .name = "Sanchari Das" }, // New faculty, yay!
};

```

```

// Compare by descending rank and then ascending names.
int compare_profs(void const *x, void const *y) {
    struct Professor const *p1 = x;
    struct Professor const *p2 = y;
    if (p1->rank > p2->rank)
        return -1;
    if (p1->rank < p2->rank)
        return 1;
    return strcmp(p1->name, p2->name);
}

int main(void) {
    size_t num_profs = sizeof profs / sizeof profs[0];

    qsort(profs, num_profs, sizeof profs[0], compare_profs);

    for (size_t i = 0; i < num_profs; ++i)
        printf("%s, %s Professor\n", profs[i].name, ranks[profs[i].rank]);

    return EXIT_SUCCESS;
}

```

\$./profs

Bob Geitz, Full Professor

John Donaldson, Full Professor

Adam Eck, Assistant Professor

Cynthia Taylor, Assistant Professor

Roberto Hoyle, Assistant Professor

Sam Taggart, Assistant Professor

Sanchari Das, Assistant Professor

Stephen Checkoway, Assistant Professor

```
// Compare by names only.  
int compare_by_names(void const *x, void const *y) {  
    struct Professor const *p1 = x;  
    struct Professor const *p2 = y;  
    return strcmp(p1->name, p2->name);  
}
```

```
$ ./profs
```

```
Adam Eck, Assistant Professor
```

```
Bob Geitz, Full Professor
```

```
Cynthia Taylor, Assistant Professor
```

```
John Donaldson, Full Professor
```

```
Roberto Hoyle, Assistant Professor
```

```
Sam Taggart, Assistant Professor
```

```
Sanchari Das, Assistant Professor
```

```
Stephen Checkoway, Assistant Professor
```

Generic binary search

```
void *bsearch(void const *key, void const *base,  
             size_t nel, size_t width,  
             int (*compare)(void const *, void const *));
```

Takes a key; a sorted array, base, of nel elements each of size width; and a comparison function and returns a pointer to the element matching the key or **NULL** if none do

```
int compare(void const *key, void const *elem);
```

- Compares the key with the element, returning <0, 0, or >0
- key and elem need not point to the same type

```
int find_by_name(void const *key, void const *elem) {
    char const *name = key;
    struct Professor const *p = elem;
    return strcmp(name, p->name);
}

// Assuming profs is sorted according to name.
struct Professor *steve;
steve = bsearch("Stephen Checkoway", profs, num_profs,
               sizeof profs[0], find_by_name);
if (steve)
    puts(ranks[steve->rank]); // Prints "Assistant".
```

What happens if we call `bsearch ()` on an array that isn't sorted? Assume that the array contains an element that matches the given key.

- A. A pointer to the matching element is returned.
- B. **NULL** is returned.
- C. Either a pointer to the matching element or **NULL** is returned, but it's impossible to say which
- D. `bsearch ()` raises an exception

Signals (brief intro)

Signals are the mechanism the OS uses to communicate with UNIX processes

There are a whole bunch of signals (see `signal(7)` or run `$ kill -l`)

`SIGINT` is the signal that is sent when the user presses control-c

A signal handler can be installed for many (but not all) signals

- Signal handlers are ***extremely*** limited
- They can't call most library functions (including `malloc(3)` and `printf(3)`)
- They should essentially set a variable of type **`volatile sig_atomic_t`** and return

C is ridiculous again

The signal function takes an int and a function pointer as arguments and returns a function pointer:

```
void (*signal(int signum, void (*handler)(int)))(int);
```

This is totally unreadable.

Use a typedef!

```
▸ typedef void (*sighandler_t)(int);  
   sighandler_t signal(int signum, sighandler_t handler);
```

```

#include <signal.h>
#include <stdio.h>
#include <time.h>
#include <unistd.h>

static volatile sig_atomic_t done;
static void handler(int signum) { done = 1; }

int main(void) {
    signal(SIGINT, handler);

    time_t start_time = time(0);
    time_t now = start_time;
    while (!done) {
        printf("The current time is %s", ctime(&now));
        sleep(10);
        now = time(0);
    }
    long diff = now - start_time;
    printf("\e[G\e[K%ld seconds elapsed\n", diff);
    return 0;
}

```

```
#include <signal.h>
#include <stdio.h>
#include <time.h>
#include <unistd.h>
```

```
static volatile sig_atomic_t done;
static void handler(int signum) { done = 1; }
```

```
int main(void) {
    signal(SIGINT, handler);

    time_t start_time = time(0);
    time_t now = start_time;
    while (!done) {
        printf("The current time is %s", ctime(&now));
        sleep(10);
        now = time(0);
    }
    long diff = now - start_time;
    printf("\e[G\e[K%ld seconds elapsed\n", diff);
    return 0;
}
```

```
$ ./a.out
The current time is Sun Nov  3 18:36:43 2019
The current time is Sun Nov  3 18:36:53 2019
The current time is Sun Nov  3 18:37:03 2019
26 seconds elapsed
```

In the previous example, after the signal handler runs, the code essentially performs

```
long diff = time(0) - start_time;  
printf("seconds elapsed\n", diff);  
exit(0);
```

Could this code be placed into the signal handler instead and would that be a better approach? (Assume start_time were changed to be global.)

- A. Yes, that would be better
- B. Yes, but it's not any better
- C. Yes, but it would be worse
- D. No, this code cannot be placed into the signal handler

In-class exercise

<https://checkoway.net/teaching/cs241/2020-spring/exercises/Lecture-25.html>

Grab a laptop and a partner and try to get as much of that done as you can!