

# CS 241: Systems Programming

## Lecture 8. Introduction to Rust

Fall 2025

Prof. Stephen Checkoway

# Why Rust?

Systems programming usually taught in C, but C is (nearly) impossible to write correctly!

Bad, unsafe C code leads to very bad vulnerabilities

Rust focuses on safety

Rust isn't just for academics! US government issued report urging use of memory-safe languages

Rust is starting to appear in operating system kernels and is replacing critical software like sudo and Python's cryptography module

# Hello, World!

```
fn main() {  
    println!("Hello world!");  
}
```

Every program needs a main function

println!() prints a string and a newline to stdout

All of the executable code lives in a function (unlike Python)

# Compiling and running

Use rustc to compile (will perform both compiling and linking by default)

- ▶ `$ rustc helloworld.rs`

rustc produces the executable helloworld

To run a program from the current directory, use `./` as usual:

- ▶ `$ ./helloworld`  
Hello world!

# Jobs of a Compiler

## Inputs

- Rust program files and options
- Libraries


## Compilation phases

- Compilation — Turns source files into object files
- Linking — Combines object files into executables


## Outputs

- Executable
- Warnings and errors

# Compilation



Java Model



Rust Model

# Basic types

## Integer types

- ▶ Signed integer types (can be negative): `i8`, `i16`, **`i32`**, `i64`, `i128`
  - Equivalent to Java's `byte`, `short`, `int`, and `long`
  - `i32` is the default when not specified
- ▶ Unsigned integer types (only nonnegative): `u8`, `u16`, `u32`, `u64`, `u128`

## Floating point types

- ▶ `f32` and `f64`
- ▶ Equivalent to Java's `float` and `double`

## String types



- ▶ `String` and `&str`

# More basic types

Boolean type: `bool`

- Values are `true` and `false`

Character type: `char`

- 4-bytes in size, holds one Unicode code point which represents one simple character like `B` or `한` or  but not complex characters like 

Platform-dependent integer types

- `usize`: 32-bit or 64-bit unsigned integer
  - Used as an index or as a count of items in a collection
- `isize`: signed version of `usize`



# Unit type: ()

The unit type () has one value: ()

```
let unit: () = ();
```

There isn't much you can do with it, but we'll actually be seeing it quite a bit

# Introduce variables with let

```
let variable_name: type = value;
```

```
fn compute_area() {  
    let width: u64 = 100;  
    let height: u64 = 24;  
    let area = width * height;  
  
    println!("{width} x {height} = {area}");  
}
```

# Function arguments/return value

```
fn function_name(arg1: type1, arg2: type2) -> return_type {}
```

```
fn compute_area(width: u64, height: u64) -> u64 {  
    let area = width * height;  
    return area;  
}
```

```
fn main() {  
    let area = compute_area(20, 40);  
    println!("The area is {area}");  
}
```

You're designing a function, `neg()`, that takes an argument of type `i32` and returns an `i32` with the opposite sign (i.e., positive values become negative and negative values become positive). Which of the options is the correct way to specify this?

A. `i32 neg(i32 val) {  
 return -val;  
}`

B. `fn i32 neg(val: i32) {  
 return -val;  
}`

C. `fn neg(val: i32) -> i32 {  
 return -val;  
}`

D. `fn neg(i32 val) -> i32 {  
 return -val;  
}`

# Returning a String

```
fn rectangle_description(width: u64, height: u64) -> String {  
    let desc: String;  
  
    if width == height {  
        desc = format!("{width} x {width} square");  
    } else {  
        desc = format!("{width} x {height} rectangle");  
    }  
    return desc;  
}
```

# Blocks have values

```
let val = {  
  let x = 10;  
  let y = 20;  
  x + y  
};
```

The value of a block of code in braces is the value of the last expression in the block

Notice the lack of ; at the end of the block and the ; after the block

The value of an **if** expression is the value of the last expression of its branches

# Variables' scope ends at the end of their containing block

The **scope** of a variable is the region of code where the variable is accessible

```
fn main() {  
    let val = {  
        let x = 10;  
        let y = 20;  
        x + y  
    };  
    println!("{val}"); // OK  
    println!("{x} {y}"); // Not OK  
}
```

error[E0425]: cannot find value `x` in this scope

--> foo.rs:8:16

```
143 |     println!("{x} {y}"); // Not OK  
    |                   ^ not found in this scope
```

# `if` is an expression, it has a value

```
fn rectangle_description(width: u64, height: u64) -> String {  
    let desc = if width == height {  
        format!("{width} x {width} square")  
    } else {  
        format!("{width} x {height} rectangle")  
    };  
    return desc;  
}
```

The value of an `if` expression is the value of the last expression of its taken branch

Notice the lack of `;` at the end of both blocks of the `if` and the `;` after the `if`



# Last expression in a function is returned

```
fn rectangle_description(width: u64, height: u64) -> String {  
    let desc = if width == height {  
        format!("{width} x {width} square")  
    } else {  
        format!("{width} x {height} rectangle")  
    };  
    desc  
}
```

The `return` is gone as is the semicolon

# Idiomatic Rust

```
fn rectangle_description(width: u64, height: u64) -> String {  
    if width == height {  
        format!("{width} x {width} square")  
    } else {  
        format!("{width} x {height} rectangle")  
    }  
}
```

The value returned from the function is the value of the last expression: the `if`

The value of the `if` is the value of the last expression of the taken branch of the `if`

What is the “Rusty” way to write the `neg()` function? Meaning, which of these is the best practice?

A. `fn neg(val: i32) -> i32 {  
 -val  
}`

B. `fn neg(val: i32) -> i32 {  
 -val;  
}`

C. `fn neg(val: i32) -> i32 {  
 return -val  
}`

D. `fn neg(val: i32) -> i32 {  
 return -val;  
}`

# Mutability

Variables are immutable by default (they cannot be changed)

Let's experiment with the Rust Playground

<https://play.rust-lang.org>

# Cannot assign twice to immutable variable

Error indicates we tried to modify an immutable variable

Error message indicates a solution

- help: consider making this binding mutable: `mut x`

```
let mut x = 10;  
println!("{x}");  
x = 20;  
println!("{x}");
```

Group discussion: Why do you think variables are immutable by default in Rust when most languages make them mutable by default?

A. Select this answer

# Strings

A `String` holds an *owned* collection of characters

- Owned means the collection of characters belongs to the `String` value

A `&str` is an *immutable reference* to a string

- References are a way to share values

Text in double quotes is a `&str`, a reference to an immutable string

We can create a `String` from a `&str` using `String::from()`

```
let s1: &str = "Điếc không sợ súng.";
let s2: String = String::from("Ignorance is bliss.");
```

# Omitting the type

```
let s1 = "Điếc không sợ súng.";
let s2 = String::from("Ignorance is bliss.");
```

The type of a variable is often omitted when it is clear from context

- ▶ Strings in double quotes are always `&str` so the type is omitted
- ▶ When the type name appears on the right-hand side of the `=`, the type is omitted



# Converting between &str and String

`String::from(s)` creates a `String` from a `&str` by making a copy of the string




`"foo".to_string()` also creates a `String` from a `&str` by making a copy of the string

A `String`'s `as_str()` method returns a `&str` reference to itself, no copy is made

```
let s1 = String::from("blah");  
let s2 = s1.as_str();
```

# Example

```
fn main() {  
    let s = String::from("Hello world!"); // L2  
    let r = s.as_str(); // L3  
    let s2 = s.to_string(); // L4  
}
```



# Passing strings to functions

```
fn foo(arg: String) {}
fn bar(arg: &str) {}

fn main() {
    let s = String::from("abc");
    foo(s);          // Valid, moves s into foo
    foo("abc");      // Invalid, foo() expects a String

    let t = String::from("xyz");
    bar(&t);          // Valid, passes a reference to t to bar
    bar("xyz");      // Valid
}
```

# Returning &str is hard

There are two problems with this function:

```
fn foo(num: i32) -> &str {  
    let s = format!("num = {num}");  
    return &s;  
}
```

1. Rustc gives an error, “expected named lifetime parameter” (we’ll talk about lifetimes later)
2. More importantly, s goes away when the function ends so the reference to it would be invalid; Rust prevents this.

# Variables' scope ends at the end of their containing block

The **scope** of a variable is the region of code where the variable is accessible

```
fn main() {  
    let val = {  
        let x = 10;  
        let y = 20;  
        x + y  
    };  
    println!("{val}"); // OK  
    println!("{x} {y}"); // Not OK  
}
```

error[E0425]: cannot find value `x` in this scope

--> foo.rs:8:16

```
143 |     println!("{x} {y}"); // Not OK  
    |                   ^ not found in this scope
```

# Aside, C doesn't prevent this!

```
#include <stdio.h>

char *foo(int num) {
    char s[100];
    snprintf(s, sizeof(s), "num = %d", num);
    return s;
}

int main() {
    char *s = foo(123);
    puts(s);
    return 0;
}
```

What happens  
when we run  
this?

```
$ ./example  
\M
```

```
$ ./example  
?
```

```
$ ./example  
??
```

# General rule of strings

When passing a string to a function, use a `&str` reference

When returning a string from a function, return a `String`

These rules don't always hold, later we'll see how to return a `&str` in some cases