

# CS 241: Systems Programming

## Lecture 11. Ownership in Rust

Fall 2025

Prof. Stephen Checkoway

# Memory safety

Data in programs are stored in memory (RAM)

One reasonable way to think about RAM is as a giant array of bytes

All of the data (and the program code!) is stored somewhere in that array

When you create an i64 variable, 8 bytes of memory are allocated somewhere in the array for that variable

When you create a String, some bytes of memory hold the contents of the string, some bytes of memory hold a **pointer** to the contents, some other bytes hold the length of the string

# Memory safety

It's critical that it's not possible to confuse which bytes are which

E.g., if our program can become confused about whether some memory is an i64 or is a pointer to our string contents, *anything could happen!*

- E.g., changing the i64 could cause the pointer to change and point at something that's not a string or some region of memory that isn't allocated at all

Memory safety is all about ensuring that it's impossible for these sorts of errors to occur

# Memory safety and Rust

Rust ensures that program are memory safe, e.g.,

- It's impossible to confuse a pointer with an integer
- It's impossible to access out-of-bounds data in an array/Vec

Most modern languages (Python, Java, Go, Haskell, Ruby, etc.) are memory-safe

Most systems languages (C and C++) are not!

- Memory safety errors are common and lead to real harm

# Ownership

Rust ensures memory safety through a concept of ownership

These are rules that the rust compiler enforces to prevent **undefined behavior**

# Stack frames

Variables live in a region of memory called the stack

The stack is organized into **frames**


Local variables in functions live in a stack frame

Each function that is called pushes a new frame onto the stack to hold its local variables

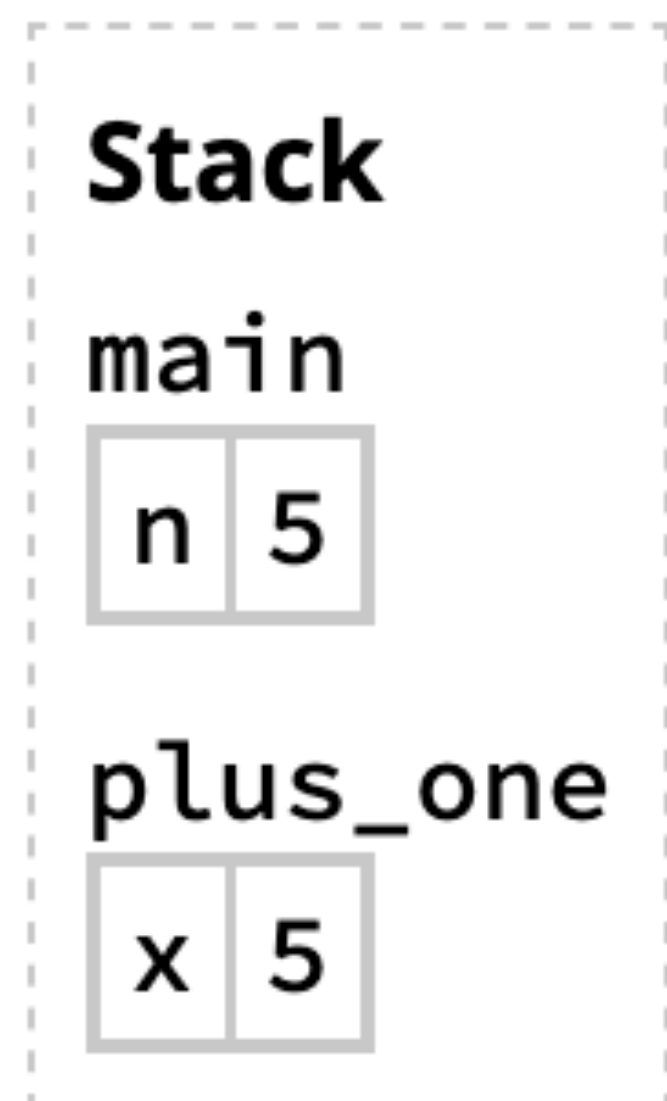
Each function that returns pops its stack frame off the stack

```
fn main() {  
    let n = 5; L1  
    let y = plus_one(n); L3  
    println!("The value of y is: {y}");  
}  
  
fn plus_one(x: i32) -> i32 {  
    L2 x + 1  
}
```


L1



L2




L3




# Every variable gets its own slot

```
let a = 5; L1  
let mut b = a; L2  
b += 1; L3
```


L1



L2



L3





```

fn foo() {
    // What does the stack look
    // like in this function...
}

fn bar() {
    foo();
}

fn main() {
    foo();
    bar();
    foo(); // ...when called here?
}

```

A. main

B. foo

C. main  
foo

D. main  
bar  
foo

E. main  
foo  
bar  
foo  
foo

# Local variables

Local variables in functions live on the stack (in a stack frame)


When the function returns, variables in the stack frame for the function are **dropped**

Once a variable is dropped, it can no longer be accessed

Returning a value from a function copies it into the stack frame of the function being returned to

Other than returning a variable, there's no way for the *variable* to live longer than the function (data can live longer as we're about to see)

# Memory layout (simplified)



# Heap

Data in the heap lives longer than an individual function

Strings and Vecs store their contents on the heap

A String or Vec variable holds a **pointer** to the contents

Any data type that needs to hold a variable amount of data works the same way:

- Contents in the heap
- Pointers to the contents

# Pointers

A pointer says where data can be located in memory

At a hardware level, a pointer is nothing more than an index into memory where the data can be found

In Java, every Object lives in the heap and is accessed via a pointer

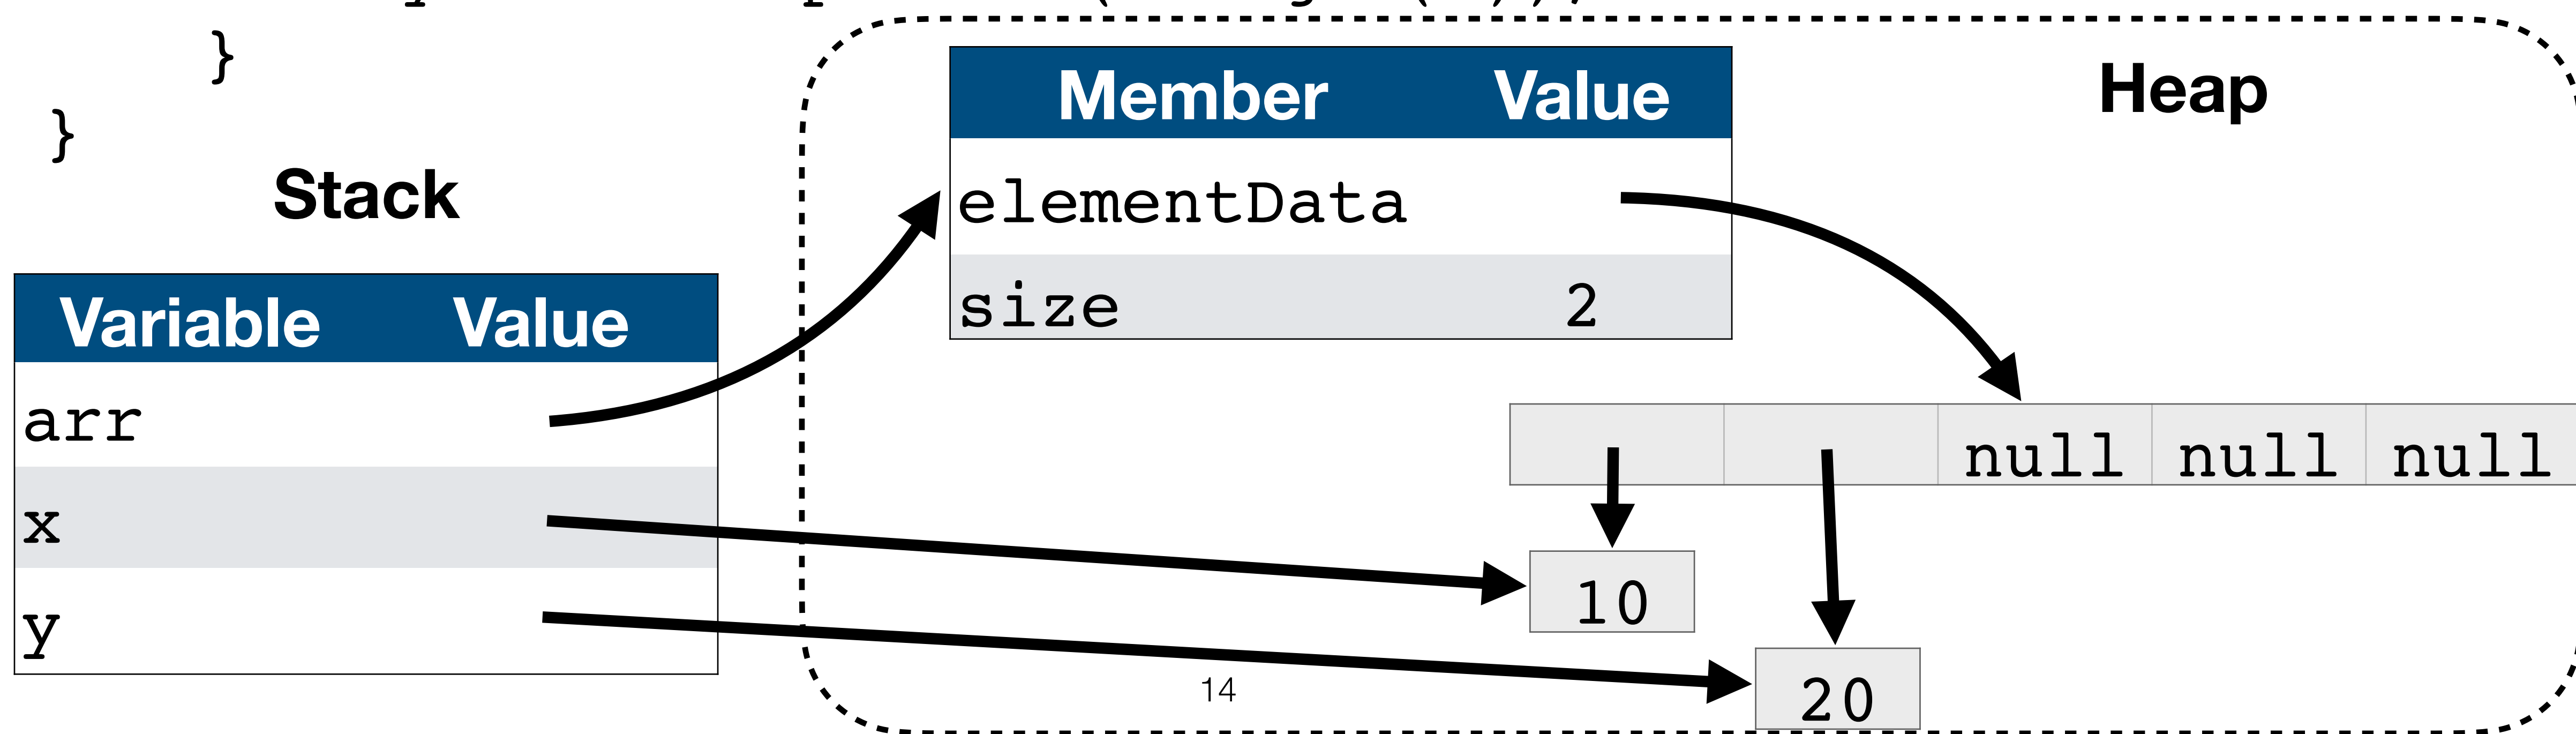
- The variables are pointers

In Rust, objects can live on the stack or in the heap


- Many objects (like String and Vec) contain pointers to heap memory

```
import java.util.ArrayList;

public class A {
    public static void main(String[] args) {
        ArrayList<Integer> arr = new ArrayList<Integer>();
        Integer x = new Integer(10);
        Integer y = new Integer(20);
        arr.add(x);
        arr.add(y);
        System.out.println(arr.get(1));
    }
}
```

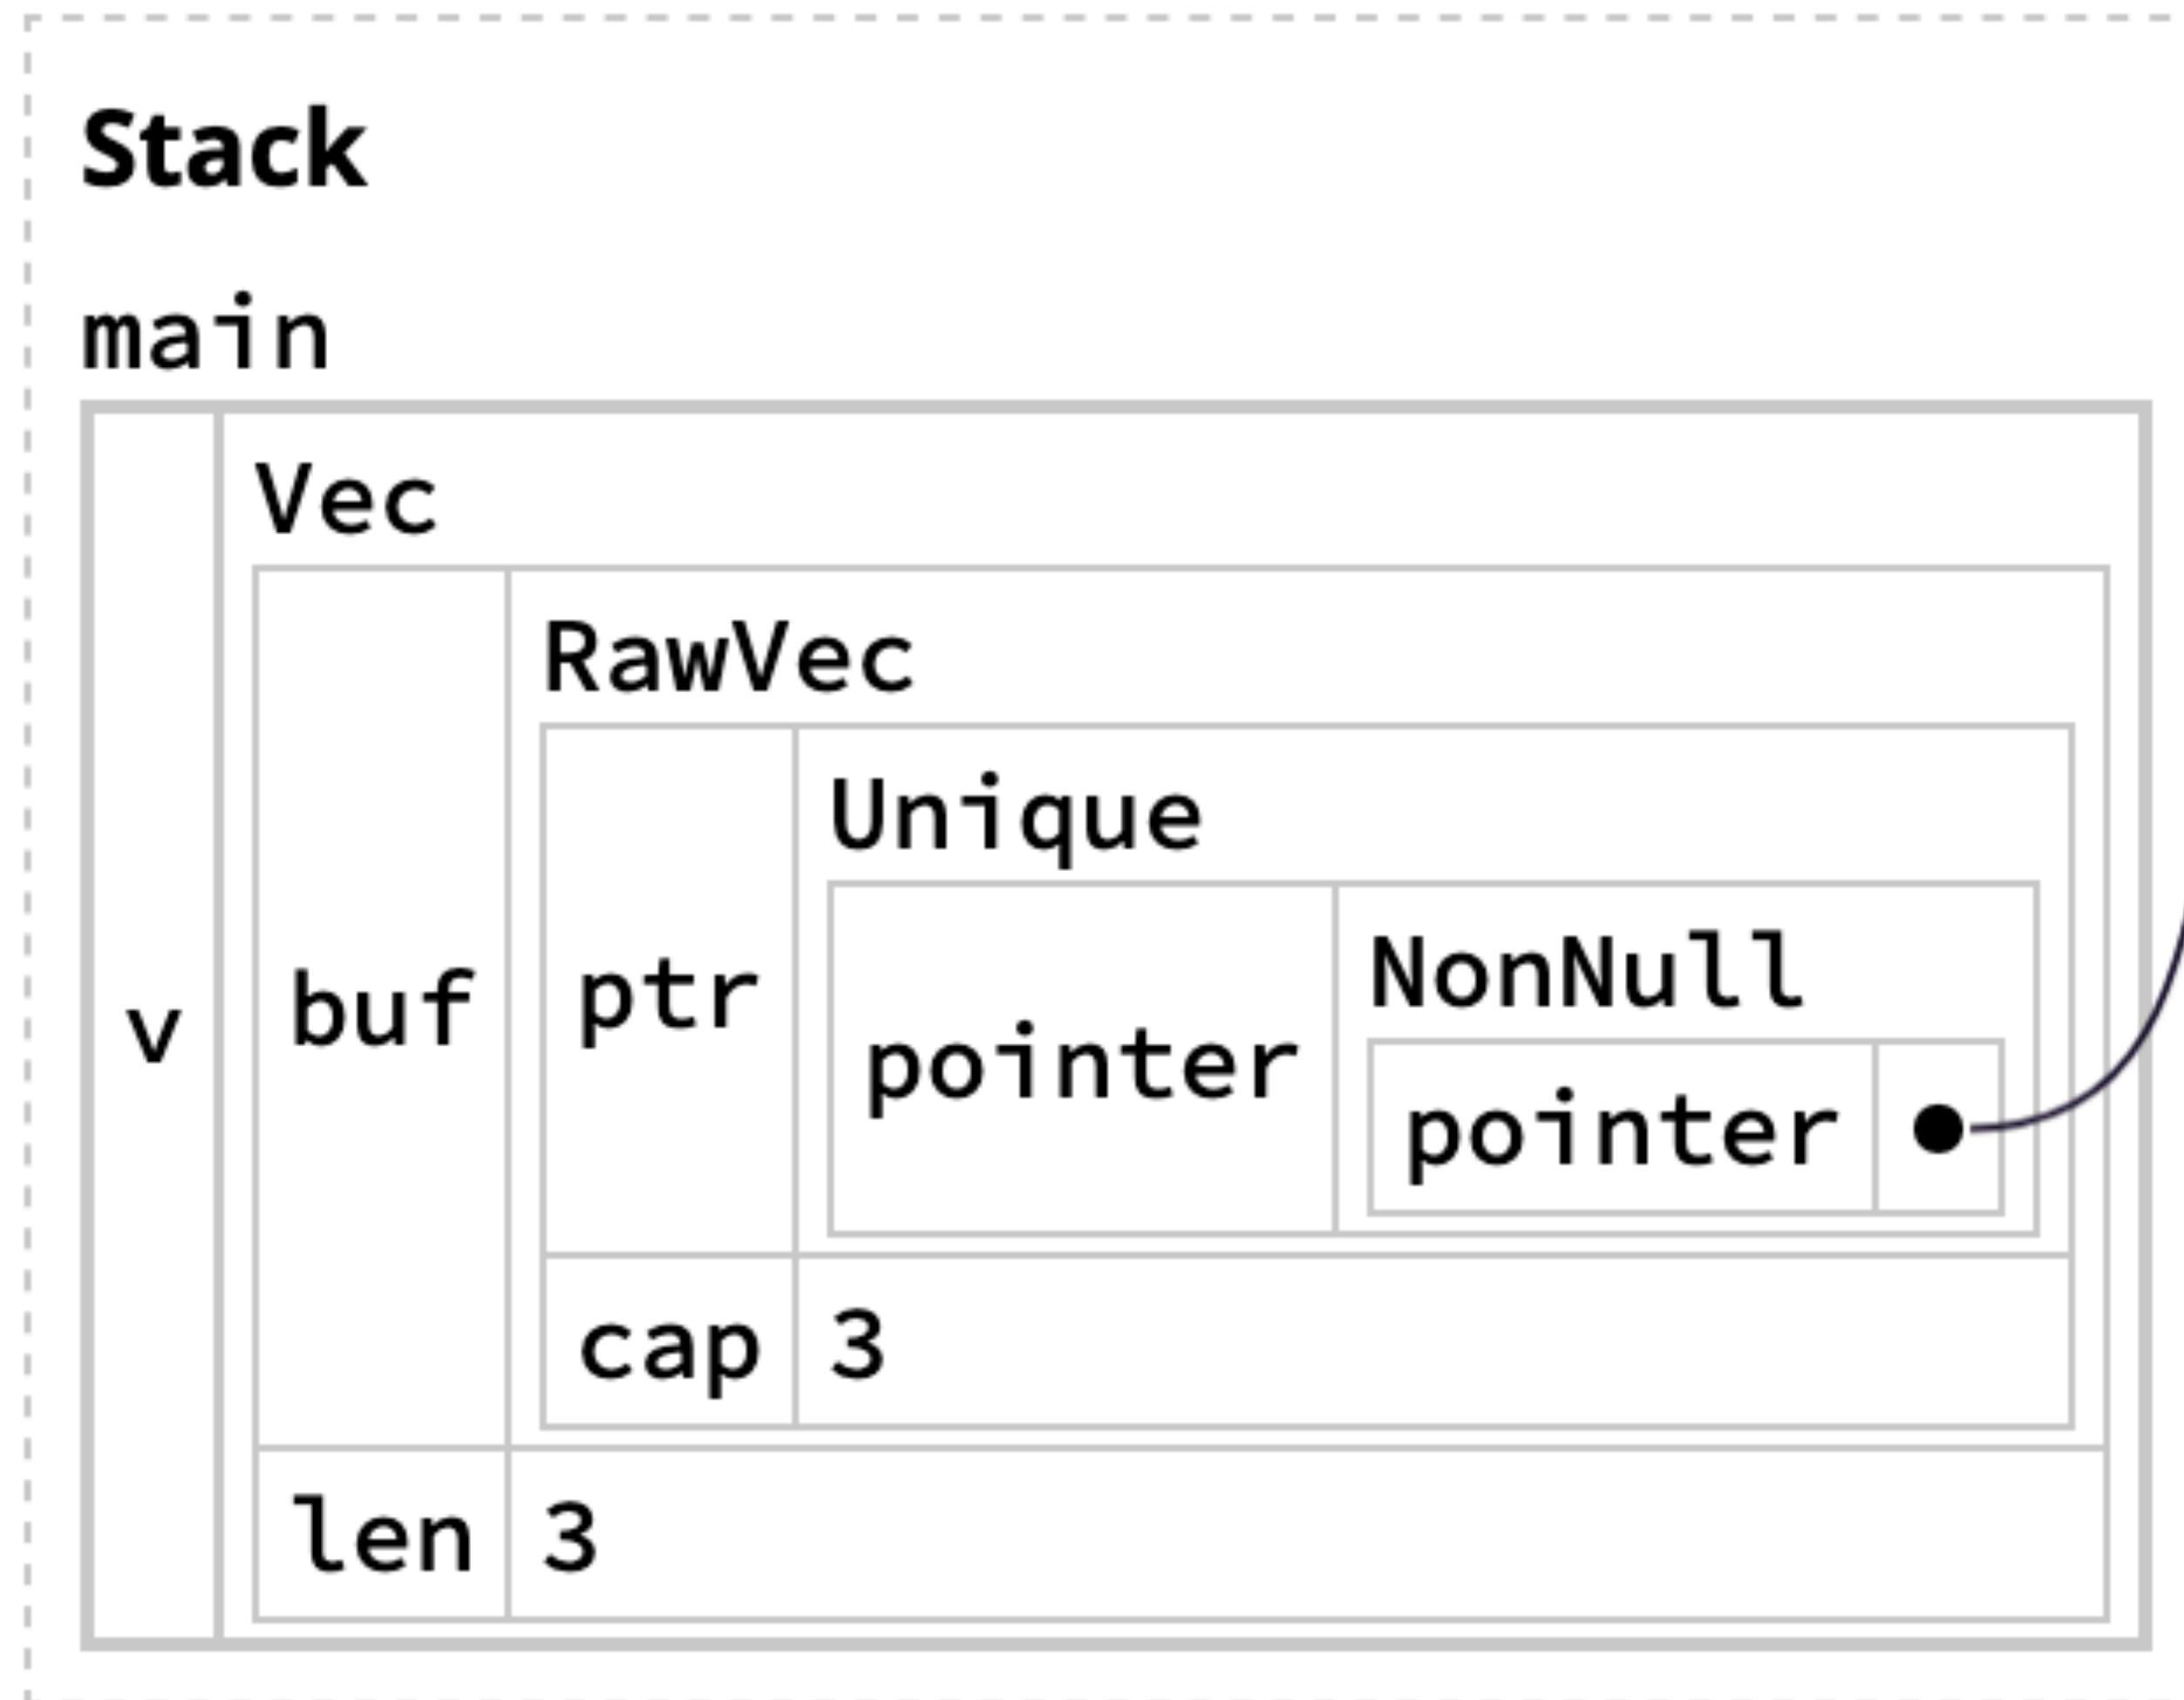


```
fn main() {  
    let mut arr: Vec<i32> = Vec::new();  
    let x = 10;  
    let y = 20;  
    arr.push(x);  
    arr.push(y);  
    println!("{arr:?}");  
}
```



```
let mut v: Vec<i32> = vec![1, 2, 3]; L1
```

L1



**Heap**

1 2 3

Vec's three members are all on the stack

- len
- cap
- pointer

(implementation detail: cap and pointer are nested inside other structures, but still on the stack)



Which of the following statements are true?

1. A local variable in a function can outlive the function.
2. Data in the heap can outlive the function that created it.
3. Variable-length data (usually) live in the heap.
4. Data on the heap is accessed using pointers

A. 1 and 2

D. 2, 3, and 4

B. 1, 2, and 3,

E. 1, 2, 3, 4, and 5

C. 3, and 4,

# Boxes—an owning pointer

We can store data in the heap by putting it in a Box

```
let b: Box<i64; 1000> = Box::new([42; 1000]);  
let b2: Box<usize> = Box::new(37);
```

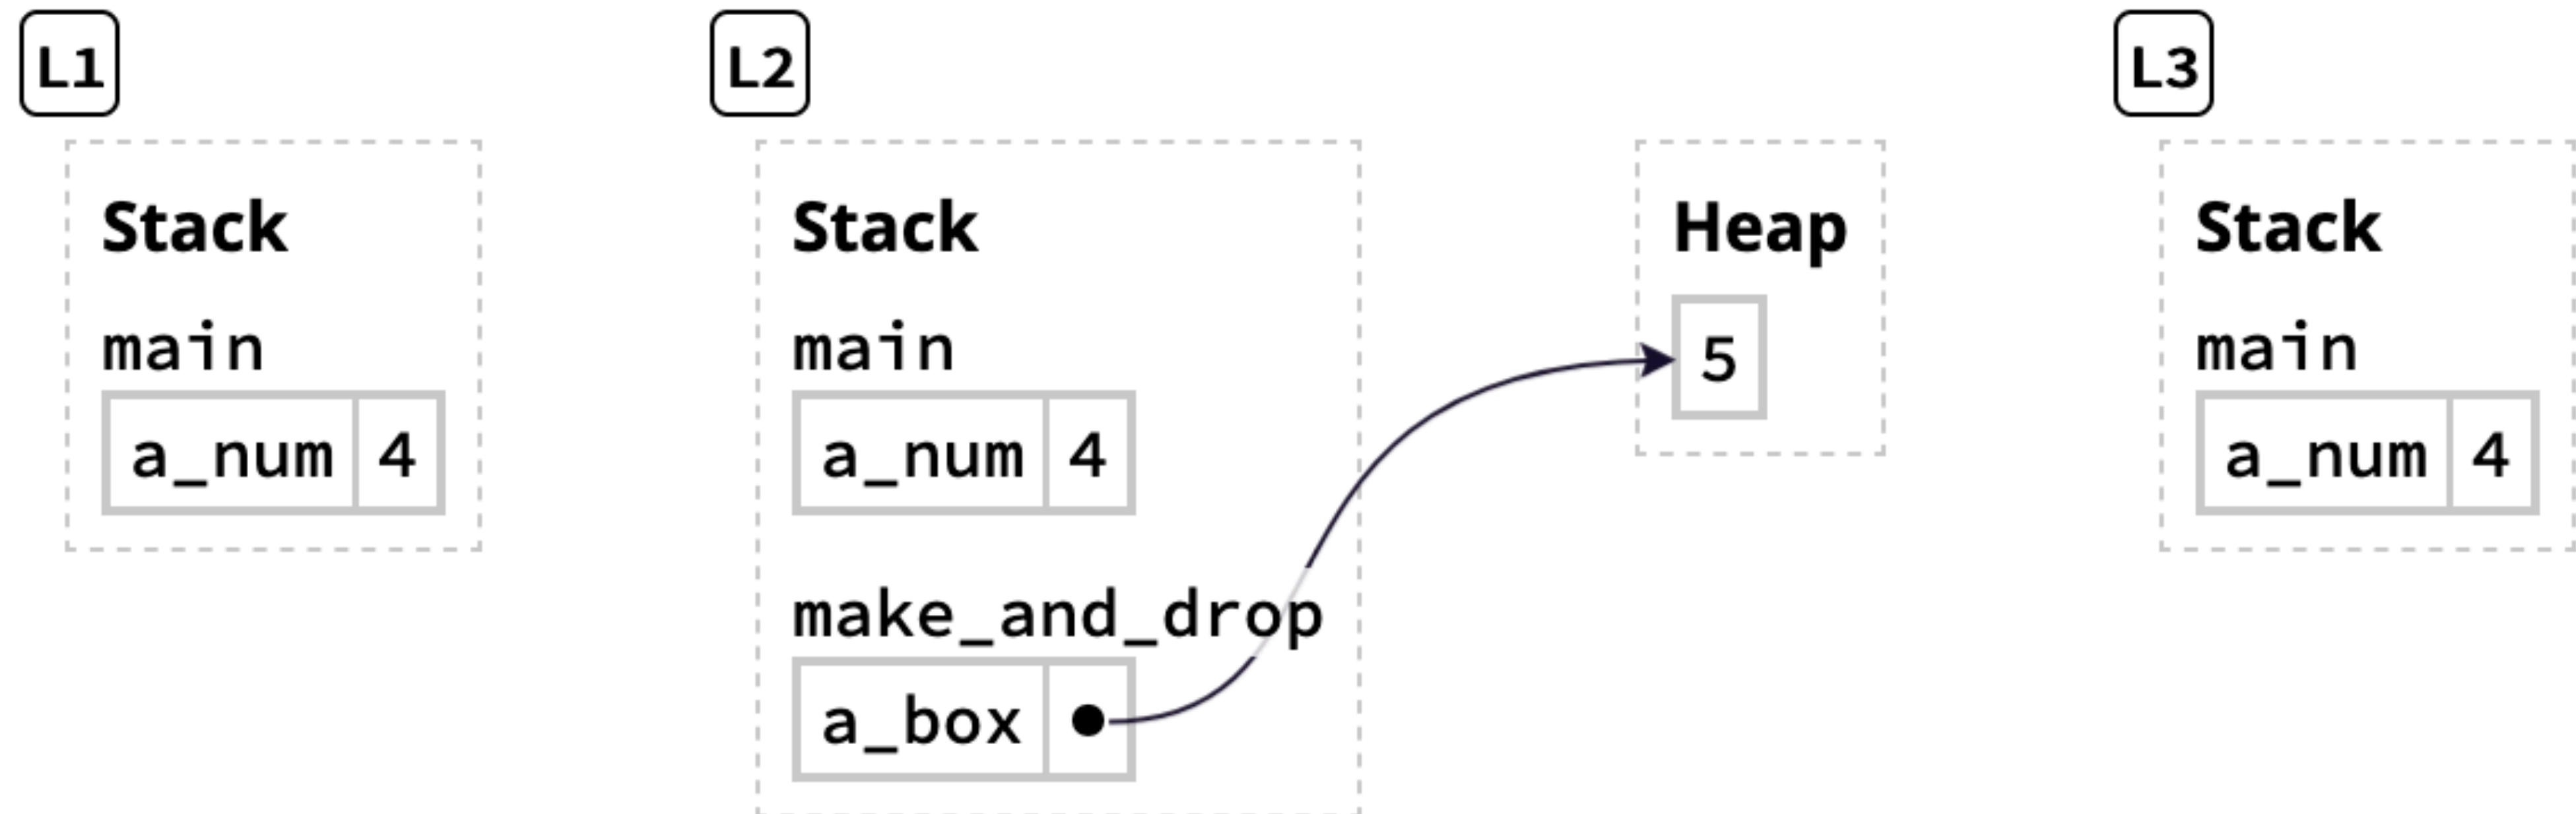
A Box is a type of pointer that always points to valid data in the heap

**The Box owns the data it points to**

When a Box variable is dropped (e.g., because the function whose frame contains the variable returns), the data in the heap is **freed**

Once data has been freed, it is no longer accessible

```
fn main() {  
    let a_num = 4; L1  
    make_and_drop(); L3  
}  
  
fn make_and_drop() {  
    let a_box = Box::new(5); L2  
}
```



# Data on the stack vs. heap

```
let a = [0; 1_000_000]; L1
let b = a; L2
```

L1



```
main
```



**L2**



```
main
```




```
let a = Box::new([0; 1_000_000]); L1
let b = a; L2
```

L1





```
main
```



**L2**



```
main
```



# No manual memory management

Languages like C and C++ let programmers allocate and free heap memory

- malloc(n) allocates n bytes of heap memory and returns a pointer to it
- free(p) frees the memory pointed to by the pointer p
- This is a **massive** source of security vulnerabilities

Rust doesn't permit manual memory management

- Once you allocate a Box, the data remains valid and accessible until the Box is dropped
- Once the Box is dropped, the data is freed and no longer accessible

# No double frees


We need to reconcile two facts

- ▶ When a Box is dropped (e.g., because the function returns), the heap memory is freed
- ▶ When we assign a Box to a new variable, the new variable points to the same heap memory

```
let a = Box::new([0; 1_000_000]);  
let b = a;
```

When main returns, it seems like both a and b will be dropped and the heap memory will be freed twice!

**Undefined behavior!**




# No double frees due to ownership!

```
let a = Box::new([0; 1_000_000]);  
let b = a;
```

Double frees don't happen because the box was **moved**, not copied

After moving data, it can no longer be accessed by the old name

We say that b **owns** the Box






# Box deallocation principle

If a variable owns a Box, when Rust deallocates the variable's frame, then Rust deallocates the Box's heap memory

In the example,  
`let b = a;`  
moved the ownership of the Box from a to b

Therefore the heap memory is only freed once





# No use-after-free

A common vulnerability in C and C++ code is

- Allocate some heap memory
- Free the allocated memory
- Use the freed memory; this is **undefined behavior!**

In Rust, that might look something like


```
let b = Box::new(10);  
drop(b); // Frees the allocated memory  
println!("{b}");
```

Rust gives a compile time error

# Cannot use a variable after moving it

```
fn main() {  
    let first = String::from("Ferris");  
    let full = add_suffix(first);  
    println!("{full}, originally {first}"); L1 // first is now used here  
}  
  
fn add_suffix(mut name: String) -> String {  
    name.push_str(" Jr.");  
    name  
}
```

**L1** undefined behavior: pointer used  
after its pointee is freed



Appending the string “ Jr.” causes  
the string to be reallocated

If we could continue to access first,  
it would point to freed memory!  
**Undefined behavior!**

# Cloning


Primitive types like numeric types (`i32`, `u64`, `usize`, etc.) can be copied

Types that involve pointers (`Box`, `String`, `Vec`, etc.) cannot; they are moved

If we want to make a deep copy of a type, we can use the `clone()` method

# Cloning

```
fn main() {  
    let arr = vec![1, 2, 3, 4, 5];  
    let arr2 = arr.clone();  
    let arr3 = arr;  
}
```



```
fn foo(s: String) { /* ... */ }

fn main() {
    let clickers = String::from("Clickers!");
    foo(XXX); // <-- Here
    println!("{clickers}");
}
```

What should we replace XXX with to pass the clickers string to foo( ) ?

- A. clickers
- B. &clickers
- C. clickers.clone()
- D. clone(clickers)
- E. More than one of the above

# Collections

Collections like String, Vec, and HashMap use a Box internally\*

When the String or Vec variable is dropped, the contents is freed

When a collection is passed as an argument to a function or returned from a function, only the pointer needs to be copied, not the contents

```
fn make_evens(num_evens: u32) -> Vec<u32> {  
    let mut result = Vec::new();  
    for num in 0..num_evens {  
        result.push(num * 2);  
    }  
    result  
}
```

\* It's not actually a Box, but it behaves similarly