

# CS 241: Systems Programming

## Lecture 18. System Calls I

Fall 2025

Prof. Stephen Checkoway

# What is an operating system?

An **Operating System (OS)** is software that manages all of the resources of a computer. It acts an interface between software and hardware.

# Operating system tasks

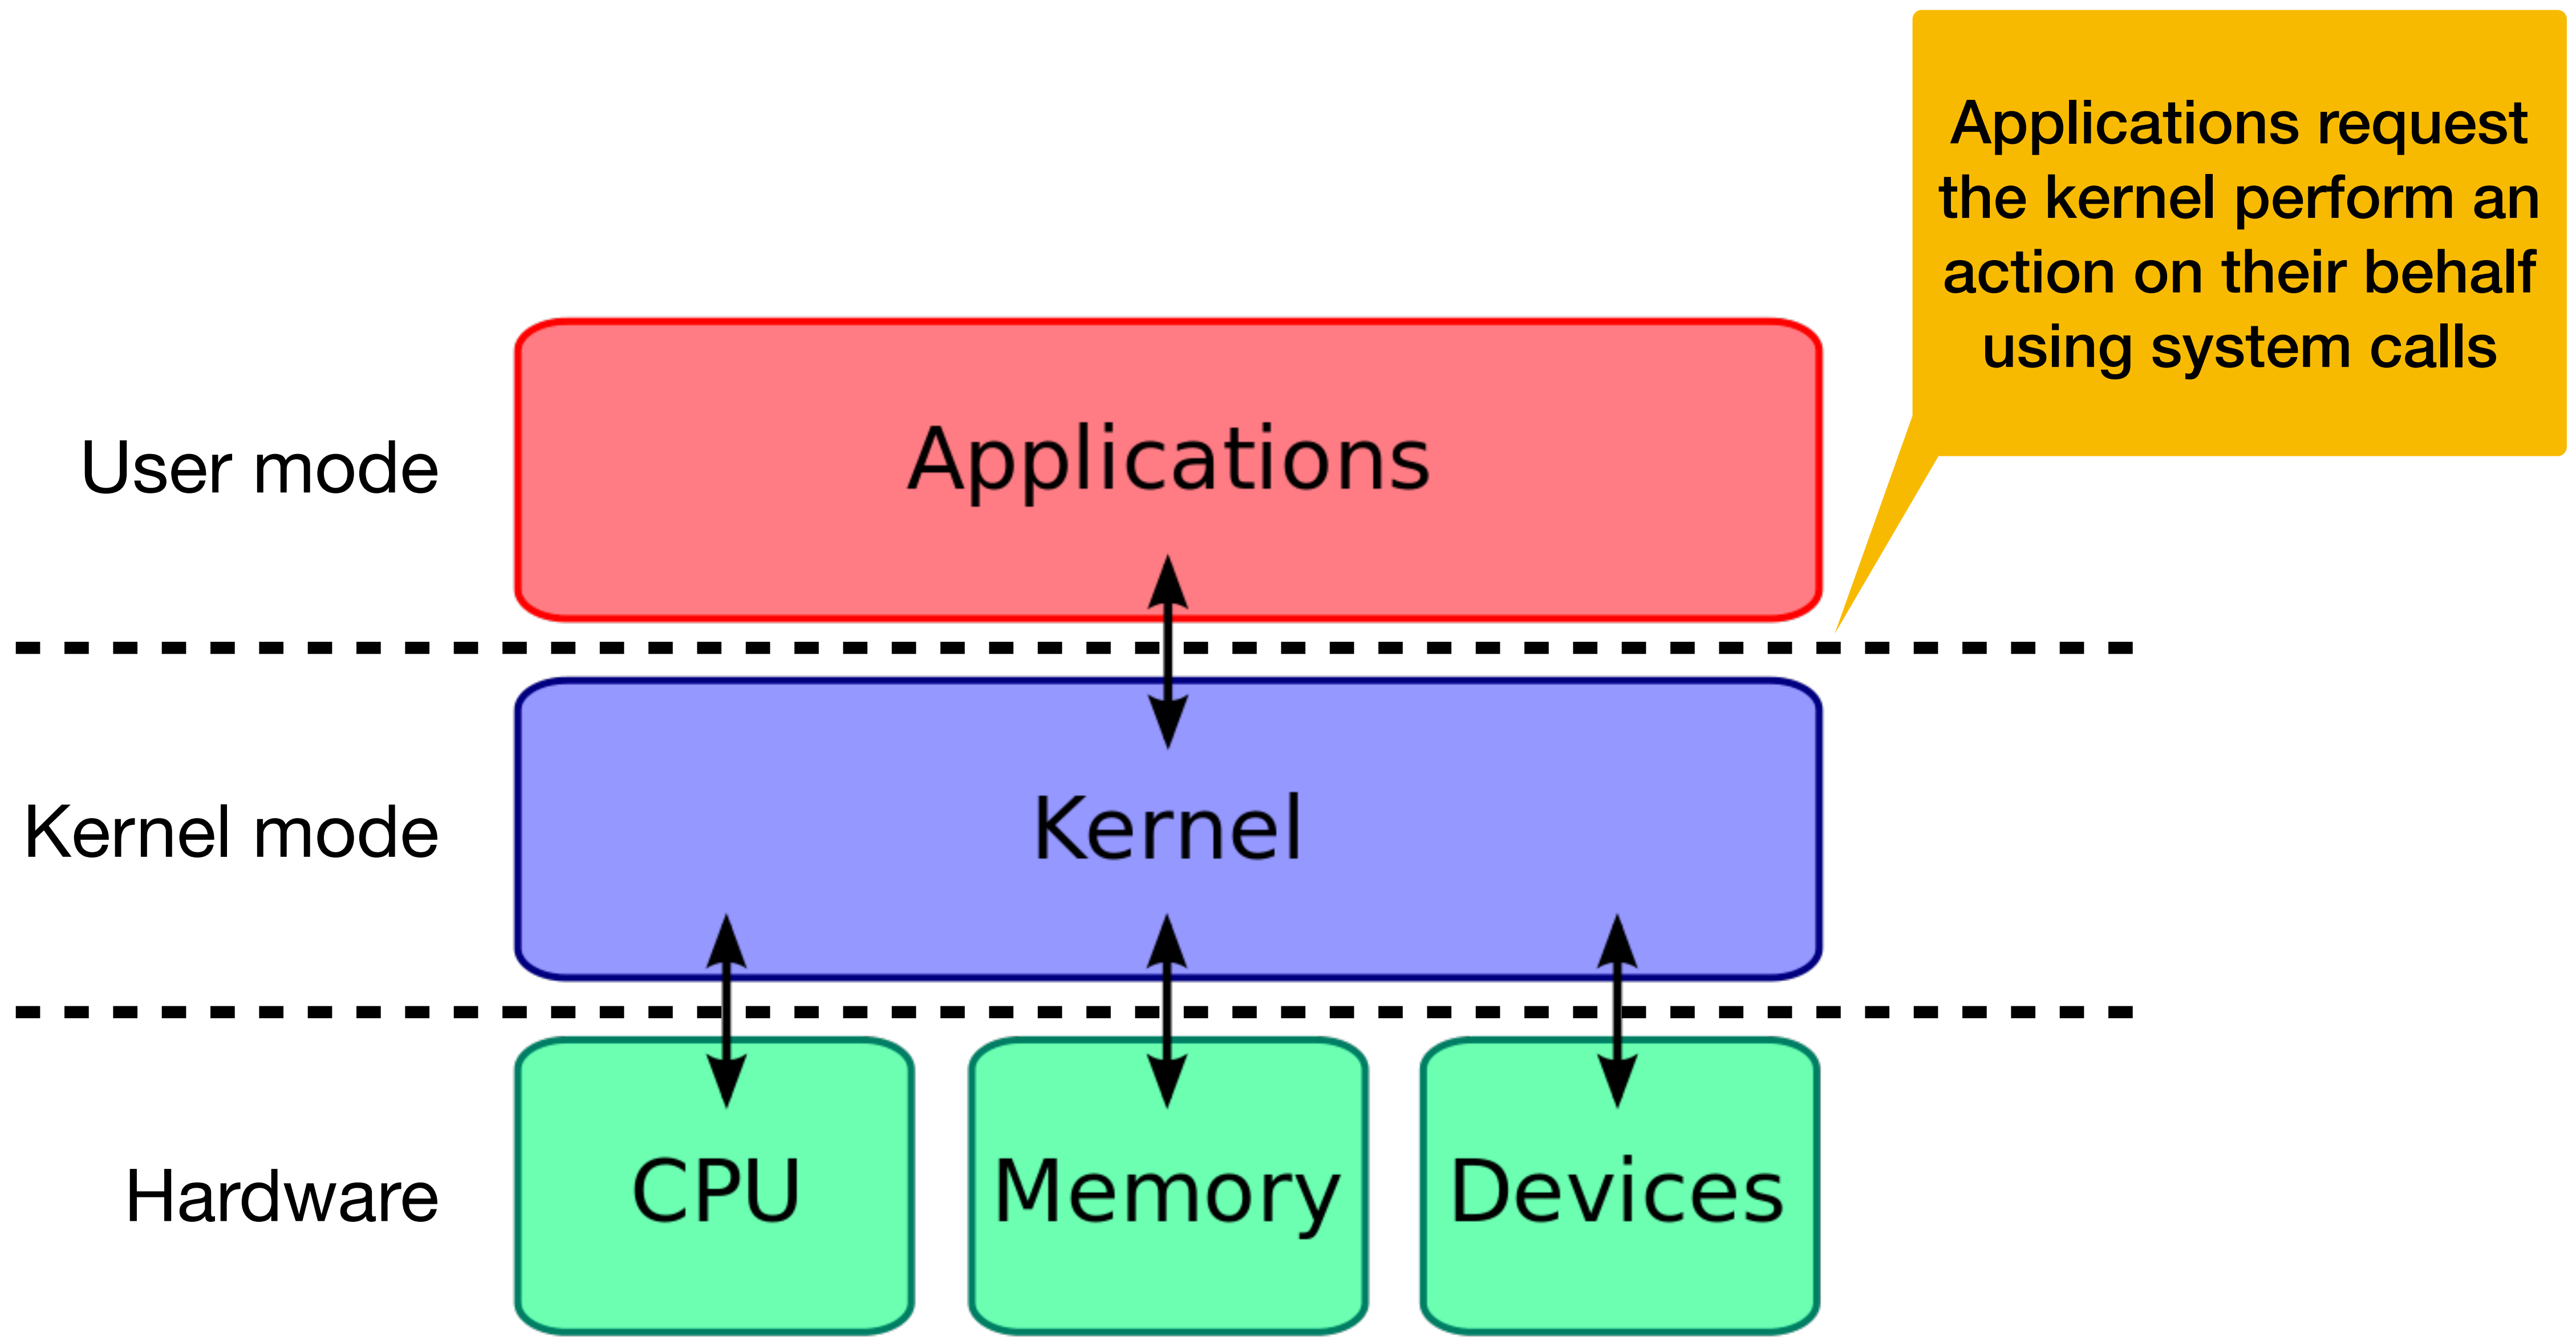
Managing the resources of a computer

- CPU, memory, network, etc.

Coordinate the running of all other programs

OS can be considered as a set of programs

- **kernel** – name given to the core OS “program”



Does one need an operating system?

A. Yes

B. No

C. I don't know/I'm not sure

# System calls

Programs talk to the OS via system calls

- Set of functions to request access to resources of the machine
- System calls vary by operating system and computer architecture

Types of system calls

- Input/output (may be terminal, network, or file I/O)
- File system manipulation (e.g., creating/deleting files/directories)
- Process control (e.g., process creation/termination)
- Resource allocation (e.g., memory)
- Device management (e.g., talking to USB devices)
- Inter-process communication (e.g., pipes and sockets)
- ...

# Most basic UNIX system call: `exit`

Programs (normally) end by calling `exit ( )` or returning from `main ( )` ...  
which calls `exit ( )`

The `exit` system call takes an exit status as its only parameter

When the kernel receives an `exit` system call from a program, it

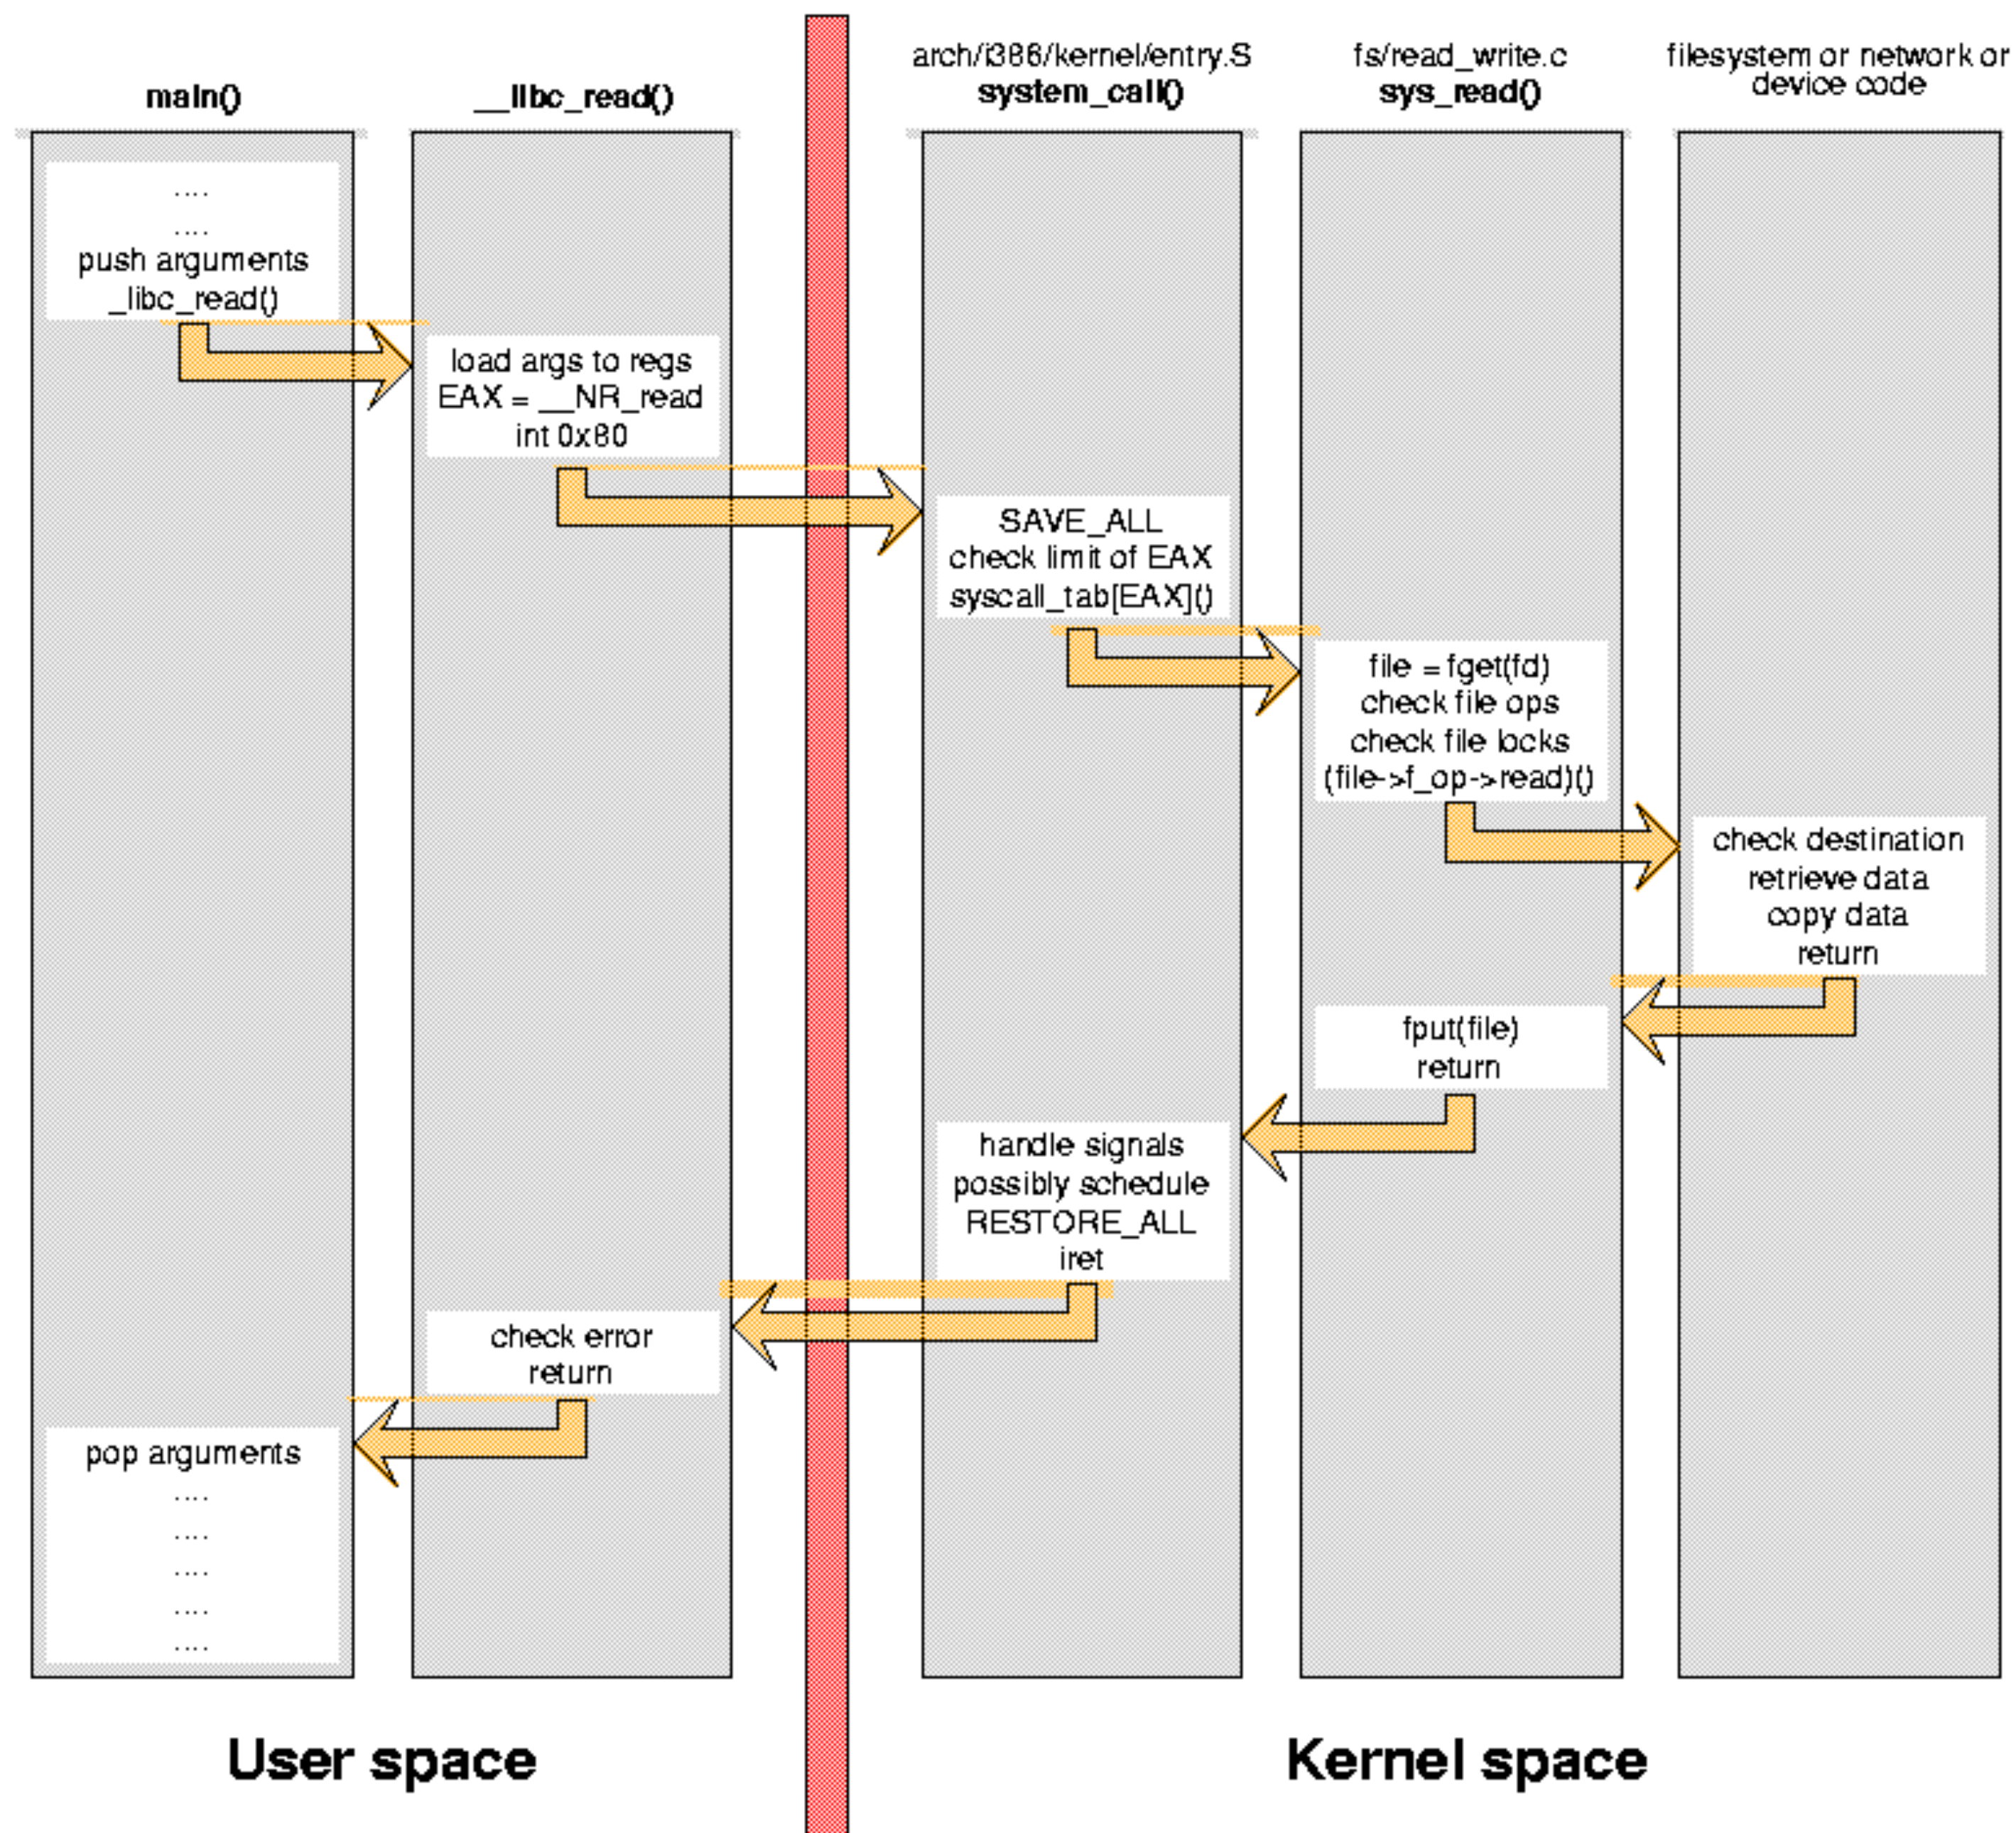
- cleans up all of the resources associated with the program
- notifies the program that created the exiting program (the parent) that a child has exited

# System calls as API

System calls are an example of an **application programming interface** (API)

- Each system call is assigned a small integer (the system call number)
- System calls are performed by setting up the arguments (often in registers) and using a dedicated "system call" or "interrupt" instruction
- The kernel's system call handler calls an appropriate function based on the system call number
- Data (and success/failure) is returned to the application





# System calls and libc

## C standard library

- Some functions make no system calls (e.g., `strcpy(3)`)
- Some functions "wrap" a single system call (e.g., `open(2)`)
- Some functions have complex behavior and might make a variable number of system calls (e.g., `malloc(3)`)

We're going to focus on the libc wrappers for the system calls

- These live in section 2 of the manual: `open(2)`, `_exit(2)`, `fork(2)`

# System calls and Rust

OS vendors make changes to their system calls over time

Different computer architectures use different system call numbers

To deal with this, the system call interface lives in libc:

- To make a system call, applications call functions in libc
- libc places the system call number and arguments in the correct registers and traps into the kernel

In Rust, we have two options

1. Use higher-level functionality provided by the standard library
2. Call functions in libc

# Unsafe Rust

# Rust lets you be unsafe

Rust has an `unsafe` keyword that lets you perform unsafe operations

- **Call functions marked as unsafe** (including everything in the libc)
- **Dereference raw pointers** (we'll talk about these shortly)
- Modify a mutable global variable
- Implement an unsafe trait (we'll talk about traits in a few lectures)
- Access fields of a C-style union

To make system calls, we'll need `unsafe` for the first two

# The purpose of unsafe

The compiler (and language) is *conservatively correct*

It ensures that the programs (that don't use `unsafe`) are memory safe

It rejects programs that are safe due to its inability to prove them safe

`unsafe` provides a way to bypass those limitations

`unsafe` limits the scope of where memory errors can occur to precisely those regions of the code marked `unsafe`



# Unsafe functions/methods

Functions and methods can be marked as unsafe by using the `unsafe` keyword

Unsafe functions can only be called from within an `unsafe` block (or `unsafe` function)

```
unsafe fn does_unsafe_things() -> i32 { 0 }
```

```
fn main() {  
    let x = unsafe {  
        does_unsafe_things()  
    };  
    println!("{x}");  
}
```

# Functions in other languages

Rust can call functions in other languages (usually C functions)

All such external functions are unsafe and can only be called from `unsafe` blocks



Why did the Rust designers require that functions written in other languages be called from an unsafe block?

A. Select A when you have an answer

# Raw pointers

We've seen pointers in Rust

- Shared references (e.g., `&i32`)
- Mutable references (e.g., `&mut i32`)
- Boxes
- Slices

Rust has two additional pointer types

- Constant pointer (e.g., `*const i32`)
- Mutable pointer (e.g., `*mut i32`)

# Raw pointers (often just called pointers)

Raw pointers are like their reference counterparts but without some restrictions

References must always point to valid, aligned objects of the appropriate type

Additionally, mutable references may not be aliased

Pointers may be invalid (including null) or point to a misaligned object

Mutable pointers may alias

# Alignment

Alignment of a value refers to its address in memory

An **aligned** value is one whose address is a multiple of its size in bytes (at least for primitive types like `i32` or `usize`, structs are aligned at the alignment of their largest member)

A **misaligned** value is one whose address is not a multiple of its size (or its largest member)

Rust (and most programming languages) require values be aligned

This restriction comes from hardware which often doesn't support misaligned memory reads/writes or performs them more slowly

# Creating a pointer from a reference

```
fn pointer_stuff(ptr: *const i32) { }

fn main() {
    let x = 10;

    // Cast the reference to a pointer
    let p = &x as *const i32;
    pointer_stuff(p);

    // Implicit conversion
    pointer_stuff(&x);
}
```

# Creating a mutable pointer

```
fn pointer_stuff(ptr: *const i32) { }

fn main() {
    let mut x = 10;

    // Cast the mutable reference to a mutable pointer
    let p = &mut x as *mut i32;
    // Implicit conversion from *mut i32 to *const i32
    pointer_stuff(p);

    // Implicit conversion from &mut i32 to *const i32
    pointer_stuff(&mut x);
}
```

The reason to create a pointer is to read or write the memory it points to

```
fn main() {  
    let mut x = 10;  
    let mut y = 20;  
    let x_ptr = &mut x as *mut _;  
    let y_ptr = &mut y as *mut _;
```

The `_` causes the compiler to use type inference to determine the type, in this case: `*mut i32`

```
unsafe {  
    let tmp = *x_ptr; // Read  
    *x_ptr = *y_ptr; // Read + write  
    *y_ptr = tmp; // Write  
}  
// after  
}
```

What will x and y be at the “after” line in this code?

	X	Y
a	10	20
b	10	10
c	20	20
d	20	10

# Reading or writing values

The reason one wants to create a pointer is to read or write the memory it points to

```
fn main() {  
    let mut x = 10;  
    let mut y = 20;  
    let x_ptr = &mut x as *mut _;  
    let y_ptr = &mut y as *mut _;  
  
    println!("Before: x={x} y={y}");  
    unsafe {  
        let tmp = *x_ptr; // Read  
        *x_ptr = *y_ptr; // Read + write  
        *y_ptr = tmp; // Write  
    }  
    println!("After:  x={x} y={y}");  
}
```

The \_ causes the compiler to use type inference to determine the type, in this case: \*mut i32

Output:
Before:  x=10  y=20
After:   x=20  y=10



# Pointer from a slice (or Vec or String)

```
let v: Vec<char> = vec!['🤪'; 1000];  
let p: *const char = v.as_ptr();
```

```
let s = String::from("Pointers!");  
let p: *const u8 = s.as_ptr();
```

Strings hold their characters  
UTF-8 encoded in a Vec<u8>

Gives a pointer to the first element of the slice

Use `.as_mut_ptr()` to get a `*mut _` rather than `*const _`

Is the `.as_ptr()` method necessary or can we just cast the reference?

```
let s = String::from("Pointers!");  
let p = &s as *const u8;
```

- A. `.as_ptr()` is necessary
- B. Casting the reference also works
- C. `.as_ptr()` is necessary for a `String` but casting would work for a `Vec`

# Creating pointers from other pointers

4	28	-47	36	0	1	-8	234	72	9	74	-9	4	2	5	8	10
---	----	-----	----	---	---	----	-----	----	---	----	----	---	---	---	---	----

↑  
p

↑  
q

```
fn main() {  
    let v = vec![4, 28, -47, 36, 0, 1, -8, 234,  
                 72, 9, 74, -9, 4, 2, 5, 8, 10];  
    let p = v.as_ptr();  
  
    unsafe {  
        let q = p.offset(10);  
        println!("*p = {}; *q = {}", *p, *q);  
    }  
}
```

# Null pointers

Use `std::ptr::null()` and `std::ptr::null_mut()` to create `*const _` or `*mut _`

Use `.is_null()` to test if a pointer is null

```
let ptr: *mut i32 = std::ptr::null_mut();  
println!("{}", ptr.is_null());
```

# libc crate

The libc crate exposes libc functions/types/constants in Rust

System call wrapper in libc	Rust declaration of the function	Notes
<code>void _exit(int status)</code>	<code>fn _exit(status: c_int) -&gt; !</code>	Exit (doesn't return)
<code>pid_t getpid(void)</code>	<code>fn getpid() -&gt; pid_t</code>	Get the process ID
<code>ssize_t read(int fildes, void *buf, size_t nbyte);</code>	<code>fn read(fd: c_int, buf: *mut c_void, count: size_t) -&gt; ssize_t</code>	Read data from a file
<code>int rename(const char *old, const char *new)</code>	<code>fn rename(oldname: *const c_char, newname: *const c_char) -&gt; c_int</code>	Renames files

# Types of arguments

Arguments to syscalls fall into a few basic types

C type	Libc crate's equivalent	Normal Rust equivalent (on many platforms)	Notes
int	c_int	i32	Normal integer
size_t/ssize_t	size_t/ssize_t	usize/usize	Represents the size of things ssize_t is used to return -1 as an error
char *	*const c_char *mut c_char	&CStr CString *const i8/*mut i8	0-byte terminated string
void *	*const c_void *mut c_void		A pointer to <i>anything</i>
Pointers to structs	Pointers to structs with the same name		References can be converted to pointers

# C strings

C considers a string to be a sequence of nonzero (often signed) bytes followed by a byte with value 0

Here's "Hello 🌀"

72	101	108	108	111	32	-16	-97	-104	-75	-30	-128	-115	-16	-97	-110	-85	0
----	-----	-----	-----	-----	----	-----	-----	------	-----	-----	------	------	-----	-----	------	-----	---

Rust considers a string to be a sequence of u8 of UTF-8 encoded characters and an associated length

72	101	108	108	111	32	240	159	152	181	226	128	141	240	159	146	171
----	-----	-----	-----	-----	----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----

len = 17

There's no difference between -16 and 240 other than interpretation. Both have binary value 11110000

# The kernel uses C strings

The kernel, being written in C, uses C strings

More importantly, the system call interface uses C strings

Converting between a C string and a Rust string isn't difficult but can be subtle

- Who owns the data?
- Is the string from the kernel valid UTF-8?



# &CStr and CString

To pass a Rust string to the kernel, use `std::ffi::CString`

```
let cstr = CString::new(some_str)?;  
let ptr = cstr.as_ptr();
```

To convert a C string into a Rust string, use `std::ffi::CStr`

```
let normal_str = CStr::from_ptr(ptr).to_str()?.to_string();
```

`CString::new()` will return an `Err(err)` if `some_str` contains a 0 byte

`.to_str()` will return an `Err(err)` if the `CStr` points to non UTF-8 data

# Example: Home directories

```
struct passwd *getpwnam(const char *name);

struct passwd {
    char    *pw_name;           /* username */
    char    *pw_passwd;        /* user password */
    uid_t   pw_uid;            /* user ID */
    gid_t   pw_gid;            /* group ID */
    char    *pw_gecos;         /* user information */
    char    *pw_dir;           /* home directory */
    char    *pw_shell;         /* shell program */
};
```

# Declarations in Rust's libc crate

```
pub unsafe extern "C" fn getpwnam(name: *const c_char)
-> *mut passwd;
```

```
#[repr(C)]
pub struct passwd {
    pub pw_name: *mut c_char,
    pub pw_passwd: *mut c_char,
    pub pw_uid: uid_t,
    pub pw_gid: gid_t,
    pub pw_gecos: *mut c_char,
    pub pw_dir: *mut c_char,
    pub pw_shell: *mut c_char,
}
```

# Handling errors

When a system call fails

- the C wrapper returns -1 (or NULL, in the case of returning a pointer)
- the per-thread global variable `errno` is set to an integer specifying the reason

Rust's `std::io::last_os_error()` reads `errno` and constructs a `std::io::Error` which we can use with a `Result`

```

use std::ffi::{CStr, CString};

type Result<T> = std::result::Result<T, Box<dyn std::error::Error>>;

fn get_home_dir_for_user(user: &str) -> Result<String> {
    let user = CString::new(user)?;
    unsafe {
        let pwd: *const libc::passwd = libc::getpwnam(user.as_ptr());

        if pwd.is_null() {
            return Err(std::io::Error::last_os_error().into());
        }

        let dir: *const libc::c_char = (*pwd).pw_dir;
        if dir.is_null() {
            return Err("No home directory found".into());
        }
        let home_dir = CStr::from_ptr(dir).to_str()?.to_string();
        Ok(home_dir)
    }
}

```

Why do we use system calls instead of making a function call directly to the function in the kernel that will handle our system call request?

Discuss with your group and select A on your clickers when you have a reason (or multiple reasons)

# Input/output system calls

# Open a file: open(2)

... indicates 0 or more additional arguments. In this case, open() takes exactly 2 or 3 arguments

```
#include <fcntl.h>
```

```
int open(char const *path, int oflag, ...);
```

- ▶ O\_RDONLY open for reading only
- ▶ O\_WRONLY open for writing only
- ▶ O\_RDWR open for reading and writing
- ▶ O\_APPEND append on each write
- ▶ O\_TRUNC truncate size to 0
- ▶ O\_CREAT create file if it does not exist
- ▶ O\_EXCL error if O\_CREAT and the file exists
- ▶ O\_NONBLOCK do not block on open or for data to become available

Bitwise OR the flags together,  
e.g.,  
O\_WRONLY | O\_CREAT

Last arg is the "int mode" -- see chmod(2) and umask(2)

Returns file descriptor on success, -1 on error



# File descriptors

Integer index into OS file table for this process

3 are automatically created for you

- **STDIN\_FILENO** 0 standard input
- **STDOUT\_FILENO** 1 standard output
- **STDERR\_FILENO** 2 standard error

These are what are used in shell redirection

- `$ ./a.out 2> errors.txt`

# Read data: read(2)

```
#include <unistd.h>
```

```
ssize_t read(int fildes, void *buf, size_t nbyte);
```

- Attempts to read nbytes from fildes storing data in buf
- Returns the number of bytes read
- Upon **EOF**, returns 0
- Upon error, returns -1 and sets **errno**

# Write data: write(2)

```
#include <unistd.h>
```

```
ssize_t write(int fildes, void const *buf, size_t nbyte);
```

- Attempts to write nbyte of data to the object referred to by fildes from the buffer buf
- Upon success, returns number of bytes are written
- On error, returns -1 and sets errno

In what scenario would these functions return less than the number of bytes specified?

- A. Reading until the end of the file
- B. Writing to the end of the file
- C. Reading from a network
- D. Writing to a network
- E. More than one of the above (which ones?)

# Seek in file: lseek(2)

```
#include <sys/types.h>
#include <unistd.h>
```

```
off_t lseek(int fd, off_t offset, int whence);
```

- whence is one of **SEEK\_SET**, **SEEK\_CUR**, **SEEK\_END**
- On success, returns the resultant offset in terms of bytes from the beginning of the file
- On error, returns (**off\_t**) - 1 and sets **errno**

# Close files: close(2)

```
#include <unistd.h>
```

```
int close(int fildes);
```

- Closes `fildes`, returns 0 on success
- Returns -1 and sets `errno` on error

# Reading a file with system calls

1. Open the file with `libc::open()` and handle errors
2. Reserve space in a `Vec<u8>`
3. Read some data with `libc::read()` and handle errors
4. If all of the data was not read, go back to step 2
5. Close the file with `libc::close()`

Why do we have an `open()` call, as opposed to just `read()`ing or `write()`ing a file path and checking permissions each time?

- A. To improve security
- B. To improve performance
- C. It makes interacting with the hard drive easier
- D. Two of the above
- E. All of the above



# Opening the file

```
use std::ffi::CString;
use std::io;

fn read_file(path: &str) -> io::Result<Vec<u8>> {
    let path = CString::new(path)?;
    let mut data: Vec<u8> = Vec::new();

    unsafe {
        let fd = libc::open(path.as_ptr(), libc::O_RDONLY);
        if fd == -1 {
            return Err(io::Error::last_os_error());
        }
        // Read the data here
        libc::close(fd);
    }
    Ok(data)
}
```

Construct a 0-terminated  
C string

# Reserve space

```
fn read_file(path: &str) -> io::Result<Vec<u8>> {  
    // ...  
    loop {  
        if data.capacity() - data.len() < 4096 {  
            data.reserve(4096);  
        }  
        // ...  
    }  
    // ...  
    Ok(data)  
}
```

# Read some data

```
fn read_file(path: &str) -> io::Result<Vec<u8>> {  
    // ...  
    loop {  
        // ...  
        let ptr: *mut libc::c_void = data.as_mut_ptr()  
            .offset(data.len() as isize)  
            .cast();  
        let amount = libc::read(fd, ptr, data.capacity() - data.len());  
        if amount < 0 {  
            let err = io::Error::last_os_error();  
            libc::close(fd);  
            return Err(err);  
        }  
        if amount == 0 {  
            break;  
        }  
        data.set_len(data.len() + amount as usize);  
    }  
    // ...  
    Ok(data)  
}
```

Easy to forget to close  
the file!

```

fn read_file(path: &str) -> io::Result<Vec<u8>> {
    let path = CString::new(path)?;
    let mut data: Vec<u8> = Vec::new();

    unsafe {
        let fd = libc::open(path.as_ptr(), libc::O_RDONLY);
        if fd == -1 {
            return Err(io::Error::last_os_error());
        }
        loop {
            if data.capacity() - data.len() < 4096 {
                data.reserve(4096);
            }
            let ptr: *mut libc::c_void = data.as_mut_ptr().offset(data.len() as isize).cast();
            let amount = libc::read(fd, ptr, data.capacity() - data.len());
            if amount < 0 {
                let err = io::Error::last_os_error();
                libc::close(fd);
                return Err(err);
            }
            if amount == 0 {
                break;
            }
            data.set_len(data.len() + amount as usize);
        }
        libc::close(fd);
    }
    Ok(data)
}

```

# Contrast with normal Rust

```
fn read_file(path: &str) -> io::Result<Vec<u8>> {  
    use std::io::Read;  
    let mut file = File::open(path)?;  
    let mut data = Vec::new();  
    file.read_to_end(&mut data)?;  
    Ok(data)  
}
```

open system call

1 or more read system  
calls

close system call when  
file is dropped

# Or even easier

```
fn main() {  
    let data1 = read_file("example.txt").unwrap();  
    let data2 = std::fs::read("example.txt").unwrap();  
    assert_eq!(data1, data2);  
}
```



One function to call.  
It'll call open(), read(), and close()

```

#include <errno.h>
#include <fcntl.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void *read_file(char const *path,
               size_t *len_ptr)
{
    int fd = open(path, O_RDONLY);
    if (fd == -1) {
        return NULL;
    }
    *len_ptr = 0;
    char *data = NULL;
    size_t len = 0;
    size_t cap = 0;
    while (1) {
        if (cap - len < 4096) {
            cap += 4096;
            char *new_data = realloc(data,
                                     cap);

            if (new_data == NULL) {
                int old_errno = errno;
                free(data);
                close(fd);
                errno = old_errno;
                return NULL;
            }
            data = new_data;
        }
        ssize_t amount = read(fd,
                              &data[len],
                              cap - len);

        if (amount < 0) {
            int old_errno = errno;
            free(data);
            close(fd);
            errno = old_errno;
            return NULL;
        }
        if (amount == 0) {
            break;
        }
        len += amount;
    }
    close(fd);
    *len_ptr = len;
    return data;
}

```

# **File system manipulation system calls**



# Delete files: unlink(2)

```
#include <unistd.h>
```

```
int unlink(char const *path);
```

- Removes path, returns 0 on success
- Returns -1 and sets **errno** on error

# Rename files: rename(2)

```
#include <stdio.h>
```

```
int rename(char const *oldpath, char const *newpath);
```

- Renames oldpath to newpath, returns 0 on success
- Returns -1 and sets **errno** on error
- This can change directories, but not file systems!

# Get current directory: getcwd(3)

```
#include <unistd.h>
```

```
char *getcwd(char *buf, size_t size);
```

- Copies absolute path of current working directory to buf
  - length of array is "size"
  - if path is too long (including null byte), NULL/ERANGE
- Linux allows NULL for buf for dynamic allocation, see man page

Basically just a wrapper around the getcwd system call plus some memory allocations

# Change directories: chdir(2)

```
#include <unistd.h>
```

```
int chdir(const char *path);
```

```
int fchdir(int fildes);
```

Change working directory of calling process

- How "cd" is implemented
- fchdir( ) is only in certain standards, but widely available
- fchdir( ) lets you return to a directory referenced by a file descriptor from open(2)ing a directory

0 on success, -1/**errno** on error

# Create/delete a directory

```
#include <sys/stat.h>
```

```
#include <sys/types.h>
```

```
int mkdir(char const *path, mode_t mode);
```

- Create a directory called path
- Don't forget execute bits in mode!

```
#include <unistd.h>
```

```
int rmdir(char const *path);
```

- Delete the directory specified by path

0 for success, -1/**errno** on error

# Reading directories

`opendir(3)`, `readdir(3)`, `closedir(3)`

- Enables the application to read the contents of directories

These are actually just higher-level wrappers around `open(2)`, `getdents(2)`, and `close(2)` which are themselves wrappers around the corresponding system calls

# Libc crate and normal Rust

The libc crate declares all of these functions

The std::fs module has Rust-versions

- remove\_file() for unlink(2)
- rename()
- create\_dir() for mkdir(2)
- remove\_dir() for rmdir(2)
- read\_dir() for opening, reading, and closing directories

The std::env module has some other related functions

- current\_dir() for getcwd(2)
- set\_current\_dir() for chdir(2)

What are some reasons you would programmatically create, delete, and change directories?

A. Discuss and select any letter on your clicker when you are done