

CS 241: Systems Programming

Lecture 19. System Calls II

Fall 2025

Prof. Stephen Checkoway

Creating a new process

Two schools of thought

- Windows way: single system call
 - `CreateProcess("calc.exe", /* other params */)`
- Unix way: two (or more) system calls
 - Create a copy of the currently running process: `fork()`
 - The copy transforms itself into a new process:
`execve("/usr/bin/bc", args, env)`

Process IDs

Every Unix process has a unique identifier

- Integer, used to index into a kernel process table
- `$ ps ax` # Print a list of all running processes and their PIDs

```
pid_t getpid(void);  
std::process::id() -> u32;
```

Every process has a parent process

- processes are "reparented" to the `init` process if their parent already exited

```
pid_t getppid(void);  
std::os::unix::process::parent_id() -> u32;
```

Creating a new process

```
#include <unistd.h>
#include <sys/types.h>
```

```
pid_t fork(void);
```

Creates an (almost) identical copy of the running program with one big exception

- Returns 0 to the child but PID of child to the parent
- -1 on error and sets **errno**

This includes a copy of memory, code, file descriptors and most other bit of process state (but not all)

What will print out after running this code if the child's PID is 5? Include output from all processes.

```
fn main() {  
    let pid = unsafe { libc::fork() };  
    if pid == 0 {  
        println!("I'm the child");  
    } else {  
        println!("Child is {pid}");  
    }  
}
```

A. "I'm the child"

B. "Child is 5"

C. "Child is 0"

D. More than 1 of the above

Parent

```
fn main() {  
→ let pid = unsafe { libc::fork() };  
  if pid == 0 {  
    println!("I'm the child");  
  } else {  
    println!("Child is {pid}");  
  }  
}
```

Stack

pid

Parent

```
fn main() {  
    let pid = unsafe { libc::fork() };  
    → if pid == 0 {  
        println!("I'm the child");  
    } else {  
        println!("Child is {pid}");  
    }  
}
```

Stack

pid

5

Child

```
fn main() {  
    let pid = unsafe { libc::fork() };  
    → if pid == 0 {  
        println!("I'm the child");  
    } else {  
        println!("Child is {pid}");  
    }  
}
```

Stack

pid

0

In which order will the two statements be printed?

```
fn main() {  
    let pid = unsafe { libc::fork() };  
    if pid == 0 {  
        println!("I'm the child");  
    } else {  
        println!("Child is {pid}");  
    }  
}
```

- A. First the parent and then the child
- B. First the child and then the parent
- C. In either order


```
fn whoami(s: &str) {  
    let pid = std::process::id();  
    let ppid = std::os::unix::process::parent_id();  
    println!("{s:<8} pid={pid:<8} ppid={ppid}");  
}
```

```
fn main() -> io::Result<()> {
```

```
    Ok(())
```

```
}
```

```
fn whoami(s: &str) {  
    let pid = std::process::id();  
    let ppid = std::os::unix::process::parent_id();  
    println!("{s:<8} pid={pid:<8} ppid={ppid}");  
}
```

```
fn main() -> io::Result<()> {  
    whoami("Prefork:");  
    Ok(())  
}
```

```
Prefork: pid=88361      ppid=86581
```

```
fn whoami(s: &str) {  
    let pid = std::process::id();  
    let ppid = std::os::unix::process::parent_id();  
    println!("{s:<8} pid={pid:<8} ppid={ppid}");  
}
```

```
fn main() -> io::Result<()> {  
    whoami("Prefork:");  
    let pid = unsafe { libc::fork() };  
  
    Ok(())  
}
```

Prefork: pid=88361 ppid=86581

```
fn whoami(s: &str) {  
    let pid = std::process::id();  
    let ppid = std::os::unix::process::parent_id();  
    println!("{s:<8} pid={pid:<8} ppid={ppid}");  
}
```

```
fn main() -> io::Result<()> {  
    whoami("Prefork:");  
    let pid = unsafe { libc::fork() };  
    if pid < 0 {  
        return Err(io::Error::last_os_error());  
    }  
}
```

```
Ok(())
```

```
}
```

Prefork: pid=88361 ppid=86581

```
fn whoami(s: &str) {
    let pid = std::process::id();
    let ppid = std::os::unix::process::parent_id();
    println!("{s:<8} pid={pid:<8} ppid={ppid}");
}
```

```
fn main() -> io::Result<()> {
    whoami("Prefork:");
    let pid = unsafe { libc::fork() };
    if pid < 0 {
        return Err(io::Error::last_os_error());
    }
    if pid == 0 {
        whoami("Child:");
    } else {
        whoami("Parent:");
    }
    Ok(())
}
```

Prefork:	pid=88361	ppid=86581
Parent:	pid=88361	ppid=86581
Child:	pid=88362	ppid=88361

fork/exec

Usually used together

`fork` to create a duplicate process

`exec` (one of the `exec` family that is) to run a new program

`fork` and `exec` both preserve file descriptors

- This is how `bash` operates: it forks, sets file descriptors, and execs

Running another program

```
int execve(char const *path, char *const argv[],  
           char *const envp[]);
```

- ▶ Last element of `argv[]` and `envp[]` must be 0 (**NULL**)
- ▶ If successful, `execve` won't return, instead, the OS will remove all of the process's code and data and load the program from `path` in its place and start running that
- ▶ The PID of the process doesn't change
- ▶ The open file descriptors remain open (unless marked close on exec)
- ▶ Returns `-1` and sets **errno** on error

exec(3) family

```
int execl(const char *path, const char *arg0, ...,
          (char *)0);
```

```
int execlp(const char *path, const char *arg0, ...,
           (char *)0, char *const envp[]);
```

```
int execlp(const char *program, const char *arg0, ...,
           (char *)0);
```

```
int execv(const char *path, char *const argv[]);
```

```
int execvp(const char *program, char *const argv[]);
```

- execl, execlp, execlp take 0-terminated variable number of arguments
- The argv and envp arrays must be 0-terminated
- execlp and execvp search PATH for the program
- glibc has an execlpe which is like execlp but searches the PATH

Fork + exec

```
let pid = unsafe { libc::fork() };
if pid == 0 {
    let path = CString::new("/bin/ls")?;
    let arg = CString::new("-l")?;
    unsafe {
        libc::execl(path.as_ptr(), // path
                    path.as_ptr(), // arg0
                    arg.as_ptr(),   // arg1
                    std::ptr::null::<i8>());
    };
}
```

Child process will be /bin/ls -l

After the child executes `exec1()`, which of the `pid`, `path`, and `arg` variables does the stack for the child process contain?

```
let pid = unsafe { libc::fork() };
if pid == 0 {
    let path = CString::new("/bin/ls")?;
    let arg = CString::new("-l")?;
    unsafe {
        libc::exec1(path.as_ptr(), // path
                    path.as_ptr(), // arg0
                    arg.as_ptr(),  // arg1
                    std::ptr::null::<i8>());
    };
}
```

A. `pid` only

B. `path` and `arg` only

C. `pid`, `path`, and `arg`

D. None of the above

Which of the following statements about `execve ()` is false?

- A. If `execve()` is successful, the new program replaces the calling program.
- B. The file descriptors that were open before `execve()` are open in the new program (except for those marked as close on exec).
- C. If `execve ()` has an error, it returns -1 and sets **errno**.
- D. If `execve ()` is successful, it returns 0.

After a `fork`, you have two copies of a program, the parent and the child, and...

- A. Either the parent or the child must call `exec ()` immediately
- B. The parent gets a PID and the child gets a 0 as return values from `fork`
- C. The child gets a PID and the parent gets a 0 as return values from `fork`
- D. Both parent and child get PIDs as the return values from `fork`
- E. Both parent and child must call `exec` to proceed

Process exit status

Can wait for a child process to exit (or be stopped, e.g., by a debugger)

```
#include <sys/wait.h>
```

```
int status;
```

```
pid_t pid = wait(&status);
```

Suspends execution until child exits, returns the PID of the child

Checking exit status

Use macros to examine exit status

WIFEXITED(status)

- True if the process exited normally

WEXITSTATUS(status)

- Returns actual return/exit value if **WIFEXITED**(status) is true

WIFSIGNALED(status)

- True if the process was terminated by a signal (e.g., **SIGINT** from ctrl-C)

WTERMSIG(status)

- Returns the signal that terminated the process if **WIFSIGNALED**(status)

Wait gets exit status from the process table. What if a process has called exit, but its parent process has not called wait?

- A. The process will not be allowed to exit.
- B. The entry will remain in the process table after the process exits.
- C. The process will exit, and when the parent calls wait, it will receive an error.

Zombies and Orphans

If a process exits but its parent has not called wait, it remains in the process table

- “Kill” command has no effect

If a process' parent exits before it does, it is adopted by the init process, which will call wait

Creating a new process, the Rust way

```
use std::os::unix::process::ExitStatusExt;
use std::process::Command;
```

```
fn main() -> io::Result<()> {
    let mut child = Command::new("/bin/ls")
        .args(["-l", "/etc/hosts"])
        .spawn()?;
```

Command uses the
“builder pattern” to
configure which
process to spawn.

.spawn() returns a Result<Child>

```
println!("Spawned process with id {}", child.id());
let status = child.wait()?;
if let Some(code) = status.code() {
    println!("Process exited with code {code}");
} else if let Some(sig) = status.signal() {
    println!("Process exited with signal {sig}");
}
Ok(())
}
```

“Builder” pattern in Rust

Create a builder object which will (eventually) construct the actual object

- Most methods take `&mut self` and return a `&mut Self` (they return `self`)
- One method will return the actual object you want

This lets you chain together method calls

```
let mut child = Command::new("/bin/ls")  
    .args(["-l", "/etc/hosts"])  
    .spawn()?;
```

is equivalent to

```
let mut cmd = Command::new("/bin/ls");  
cmd.args(["-l", "/etc/hosts"]);  
let mut child = cmd.spawn()?;
```

Another builder example

The open system call takes a bunch of different options (look at the man page for `open(2)`)

The basic `File::open()` and `File::create()` handle the two most common cases: opening a file for reading and creating a file to write

`std::fs::OpenOptions` is another builder pattern

- You call methods to configure reading, writing, appending, truncating, etc.
- Then you call `.open()` to actually perform the open system call and return a new `File` object

OpenOptions example

To open a file for reading and writing, creating the file if it doesn't exist, use

```
let file = OpenOptions::new()  
    .read(true)  
    .write(true)  
    .create(true)  
    .open("foo.txt")?;
```

`OpenOptions::new()` returns an `OpenOptions`

`.read()`, `.write()`, `.create()` all return `self`

`.open()` returns an `io::Result<File>`

strace(1)

strace(1)

strace is a Linux program that prints out the system calls a program uses

- `-e trace=open,openat,close,read,write` will trace those system calls
- `-f` will trace children too
- `-s size` will show up to `size` bytes of strings

strace(1)

strace is a Linux program that prints out the system calls a program uses

- `-e trace=open,openat,close,read,write` will trace those system calls
- `-f` will trace children too
- `-s size` will show up to `size` bytes of strings

```
$ strace -e trace=open,openat,close,read,write cat Makefile
```

```
...
```

```
openat(AT_FDCWD, "Makefile", O_RDONLY) = 3
```

```
read(3, "CC := clang\nCFLAGS := -Wall -std"..., 1048576) = 176
```

```
write(1, "CC := clang\nCFLAGS := -Wall -std"..., 176) = 176
```

```
read(3, "", 1048576) = 0
```

```
close(3) = 0
```

```
...
```