

CS 241: Systems Programming

Lecture 31. Regular Expressions I

Fall 2025

Prof. Stephen Checkoway

Announcement

- SETs are available

Today - Regular Expressions


- A way to specify patterns used in many programming languages and tools
- Used in Rust, Python, Java, C, etc. and command-line tools like grep, sed, and awk

Theory of regular languages

Mathematical theory of sets of strings

- ▶ You'll see this in CS 383

Connection to finite state machines




Theory of regular languages

Mathematical theory of sets of strings

- ▶ You'll see this in CS 383

Connection to finite state machines



We're going to skip all of this for this course!

Problem we want to solve

Identify and/or extract text that matches a given **pattern**

Examples

- ▶ Determine if a text string matches the pattern
- ▶ Find all lines of text in a file containing a given word
- ▶ Extract all phone numbers from a file
- ▶ Extract fields from structured text
- ▶ Classify types of text (e.g., compilers need to determine if some text is a number like 0x7D2 or symbols like == or keywords like fn)
- ▶ ~~Find all of the tags in an HTML file (Don't use regex for this please)~~

Approach: Use a **regular expression** to specify the **pattern**

What is a regular expression?

Text that describes a **search pattern**

Comes in a variety of “flavors”

- ▶ Basic Regular Expression (**BRE**)
- ▶ Extended Regular Expression (**ERE**)
- ▶ Perl-Compatible Regular Expression (**PCRE**)

Be careful not to confuse with file globbing which uses similar special characters like * and ? but with slightly different meanings (we talked about this with bash)

Baseline regex characters

Baseline regex characters

- (period) any single character except newline

Baseline regex characters

- (period) any single character except newline
- * 0 or more of the preceding item (greedy)

Baseline regex characters

- (period) any single character except newline
- * 0 or more of the preceding item (greedy)
- ^ start of a line

Baseline regex characters

- (period) any single character except newline
- * 0 or more of the preceding item (greedy)
- ^ start of a line
- \$ end of the line

Baseline regex characters

- (period) any single character except newline
- * 0 or more of the preceding item (greedy)
- ^ start of a line
- \$ end of the line
- [] match one of the enclosed characters
 - ▶ [**a-z**] matches a range
 - ▶ [^] reverses the sense of match
 - ▶ put] or – at start to be a member of the list

Baseline regex characters

- (period) any single character except newline
- * 0 or more of the preceding item (greedy)
- ^ start of a line
- \$ end of the line
- [] match one of the enclosed characters
 - ▶ [a-z] matches a range
 - ▶ [^] reverses the sense of match
 - ▶ put] or – at start to be a member of the list

Every other character just matches itself; precede any of the above with \ to treat as a normal character that must literally match

Examples

Examples

a

Anything with the letter 'a'

Examples

a

Anything with the letter 'a'

abc

Anything with the string 'abc'

Examples

a

Anything with the letter 'a'

abc

Anything with the string 'abc'

a . c

'a' followed by any char then 'c'

Examples

a

Anything with the letter 'a'

abc

Anything with the string 'abc'

a . c

'a' followed by any char then 'c'

^a

Line starting with 'a'

Examples

a

Anything with the letter 'a'

abc

Anything with the string 'abc'

a . c

'a' followed by any char then 'c'

^a

Line starting with 'a'

a\$

Line ending with 'a'

Examples

a

Anything with the letter 'a'

abc

Anything with the string 'abc'

a . c

'a' followed by any char then 'c'

^a

Line starting with 'a'

a\$

Line ending with 'a'

^a\$

Line consisting of a single 'a' on it

Examples

a	Anything with the letter 'a'
abc	Anything with the string 'abc'
a . c	'a' followed by any char then 'c'
^a	Line starting with 'a'
a\$	Line ending with 'a'
^a\$	Line consisting of a single 'a' on it
a . *b	'a' then anything else, then 'b' (includes 'ab')

Examples

a	Anything with the letter 'a'
abc	Anything with the string 'abc'
a . c	'a' followed by any char then 'c'
^a	Line starting with 'a'
a\$	Line ending with 'a'
^a\$	Line consisting of a single 'a' on it
a . *b	'a' then anything else, then 'b' (includes 'ab')
[abc]	One of 'a', 'b', or 'c'

Examples

a	Anything with the letter 'a'
abc	Anything with the string 'abc'
a . c	'a' followed by any char then 'c'
^a	Line starting with 'a'
a\$	Line ending with 'a'
^a\$	Line consisting of a single 'a' on it
a . *b	'a' then anything else, then 'b' (includes 'ab')
[abc]	One of 'a', 'b', or 'c'
[a-zA-Z0-9]	Anything containing a letter or number

Valid identifiers in Rust* (things like variable or function names)

1. start with either a letter or an underscore; and
2. consist of letters, numbers, or underscores.

E.g., `main`, `foo_bar`, `_Okay123XY` are valid identifiers;
but `32x`, `foo-bar`, and `&blah` are not

Which regular expression describes valid Rust identifiers?

- A. `[a-zA-Z0-9_]*`
- B. `[a-zA-Z0-9_][a-zA-Z0-9_]*`
- C. `[a-zA-Z_][a-zA-Z0-9_]*`
- D. `[^0-9][a-zA-Z0-9_]*`

*Not totally true. Rust has “raw” identifiers as well, ignore those

Extended regex (modern)

{m, n} match previous item at least **m** times, but at most **n** times

() group and save enclosed pattern match

+ match 1 or more of the previous {1, }

? match previous 0 or 1 time {0, 1}

| match the regex either before or after the pipe

▶ apple|banana

(ab|c+) {2} 'abab', 'abc', 'abcccc', 'cab', 'cccab' 'cccccccccc'

What are some expressions that match this extended regex?

(**a****b** | **c**) {2}

A. Select any option on your clicker

{m, n} match previous item at least **m** times, but at most **n** times

() group and save enclosed pattern match

+ match 1 or more of the previous {1,}

? match previous 0 or 1 time {0,1}

| match regex either before or after

Basic regex (obsolete)

`\{m, n\}` match previous item at least **m** times, but at most **n** times

`\{m\}` match previous item exactly **m** times

`\{m, \}` match previous item at least **m** times

`\(\)` group and save enclosed pattern match

▶ `\1` the first saved match

▶ `\5` the fifth saved match

▶ Using such "back references" makes it not a real regular expression and should be avoided

POSIX character classes

POSIX character classes

Within brackets [], we can use character classes corresponding to those in ctype.h by surrounding the name with [: and :]

- ▶ alnum, digit, punct, alpha, graph, space, blank, lower, upper, cntrl, print, xdigit
- ▶ E.g., [[:digit:] [:space:]]

POSIX character classes

Within brackets [], we can use character classes corresponding to those in ctype.h by surrounding the name with [: and :]

- ▶ alnum, digit, punct, alpha, graph, space, blank, lower, upper, cntrl, print, xdigit
- ▶ E.g., [:digit:] [:space:]

Shortcuts (needs "enhanced" basic or extended regular expressions):

- ▶ \d is [:digit:] \D is ^[:digit:]
- ▶ \s is [:space:] \S is ^[:space:]
- ▶ \w is [:alnum:] \W is ^[:alnum:]

Which string does the ERE

`\([:digit:] \{3\} \) [:digit:] \{3\}-[:digit:] \{4\}`

match?

A. ([1] {3}) [2] {3} - [3] {4}

B. 123 456-7890

C. (123) 456-7890

D. \(123\) 456-7890

grep(1)

grep matches lines of input against a given regular expression (regex), printing each line that matches (or does not match)

```
$ grep 'Computer Science' file  
  ▷ prints each line of file that contains the string "Computer Science"
```

More generally,

```
$ grep regex file
```

will print each line of file that matches the regular expression regex

grep(1)

Name comes from ed(1) program command g/re/p

grep -E re files	use extended regex (or use egrep)
egrep -l re files	just list file names
egrep -c re files	just list count of matches
egrep -n re files	just list line numbers
egrep -i re files	ignore case
egrep -v re files	show non-matching lines

Example

```
mhogan@mcnulty:~/cs241/F24/lab6$ egrep -r process src/*
src/bin/ps.rs:use process::{Process, Result, user, device_map};
src/bin/ps.rs:#[command(author, version, about = "ps - report process status", long_about = None)]
src/bin/ps.rs:    ///Write information for all processes with controlling terminals. Omit session leaders.
src/bin/ps.rs:    /// Write information for all processes.
src/bin/ps.rs:    /// Write information for all processes, except session leaders
src/bin/ps.rs:    let vec = Process::all_processes()?;
src/bin/ps.rs:        std::process::exit(1);
src/bin/runnable.rs:use process::Result;
src/bin/runnable.rs:// (processes and threads) and the total number of kernel scheduling entities.
src/bin/runnable.rs:        std::process::exit(1);
src/bin/whoami.rs:use process::user;
src/proc.rs:// Models a Linux process.
src/proc.rs:    /// Look up information about a running process with the given PID.
src/proc.rs:    /// Look up information for the current process.
src/proc.rs:    let pid: i32 = std::process::id() as i32;
src/proc.rs:    /// Returns a list of all running processes.
src/proc.rs:    pub fn all_processes() -> Result<Vec<Self>> {
src/proc.rs:        /// Returns `true` if the process is a session leader.
src/proc.rs:        /// Returns `true` if the process has a controlling terminal.
```

Example

```
mhogan@mcnulty:~/cs241/F24/lab6$ egrep -r "(P|p)rocess" src/*
src/bin/ps.rs:use process::{Process, Result, user, device_map};
src/bin/ps.rs:#[command(author, version, about = "ps - report process status", long_about = None)]
src/bin/ps.rs:    ///Write information for all processes with controlling terminals. Omit session leaders.
src/bin/ps.rs:    /// Write information for all processes.
src/bin/ps.rs:    /// Write information for all processes, except session leaders
src/bin/ps.rs:fn print_full(p: &Process, full: bool, long: bool, ticks:i64, tty:&str) {
src/bin/ps.rs:    let me = Process::for_self()?;
src/bin/ps.rs:    let vec = Process::all_processes()?;
src/bin/ps.rs:    std::process::exit(1);
src/bin/runnable.rs:use process::Result;
src/bin/runnable.rs:// (processes and threads) and the total number of kernel scheduling entities.
src/bin/runnable.rs:        std::process::exit(1);
src/bin/whoami.rs:use process::user;
src/lib.rs:pub use proc::Process;
src/proc.rs:// Models a Linux process.
src/proc.rs:pub struct Process {
src/proc.rs:    /// Process ID.
src/proc.rs:impl Process {
src/proc.rs:    /// Look up information about a running process with the given PID.
src/proc.rs:    /// Look up information for the current process.
src/proc.rs:    let pid: i32 = std::process::id() as i32;
src/proc.rs:    /// Returns a list of all running processes.
src/proc.rs:    pub fn all_processes() -> Result<Vec<Self>> {
src/proc.rs:        /// Returns `true` if the process is a session leader.
src/proc.rs:        /// Returns `true` if the process has a controlling terminal.
```

Which command will return only “oberlin” and “snow”

- A. egrep o 241.txt
- B. egrep “[a-z]o[a-z]*” 241.txt
- C. egrep “[a-z]*[o[a-z]* ” 241.txt
- D. egrep o* 241.txt
- E. More than one of the above

oberlin
winter
snow
Rust
Ferris
turtle
241.txt

awk(1)

Named after the developers

- ▶ Alfred Aho
- ▶ Peter Weinberger
- ▶ Brian Kernighan

Programming language for working on files - used for text processing and data extraction

Consists of a sequence of pattern-action statements of the form

- ▶ **pattern { action }**
- ▶ Each line of the input is matched compared to each **pattern** in order; each matching **pattern** has its associated **action** run

Running AWK

Running

- ▶ \$ awk -f foo.awk files # foo.awk contains the program
- ▶ \$ awk prog files # prog is patterns-actions separated by ;

Understands whitespace separated fields (can change this via -F option)

Awk programs can manipulate the fields in a line with

- ▶ **\$1, \$2, \$3** are the first three fields
- ▶ **\$0** is the whole line

Other variables, just use their names

Simple AWK program

Simple AWK program

Prints the lines of a file with START and END

Simple AWK program

Prints the lines of a file with START and END

```
BEGIN { print "START" }
      { print }
END   { print "END" }
```

Simple AWK program

Prints the lines of a file with START and END

```
BEGIN { print "START" }
      { print }
END   { print "END" }
```

The first line has the special pattern **BEGIN** whose action runs before looking at any lines

Simple AWK program

Prints the lines of a file with START and END

```
BEGIN { print "START" }
        { print }
END   { print "END" }
```

The first line has the special pattern **BEGIN** whose action runs before looking at any lines

The second line does not have a pattern so its action (print the line) runs for each line

Simple AWK program

Prints the lines of a file with START and END

```
BEGIN { print "START" }
        { print }
END   { print "END" }
```

The first line has the special pattern **BEGIN** whose action runs before looking at any lines

The second line does not have a pattern so its action (print the line) runs for each line

The final line has the special pattern **END** whose action runs after all lines

Sum up a list of numbers

```
BEGIN { SUM = 0 }
      { SUM += $1 }
END   { print "Total is", SUM }
```

Sum up a list of numbers

```
BEGIN { SUM = 0 }
       { SUM += $1 }
END   { print "Total is", SUM }
```

```
$ cat nums
10
39
48
22
51
$ awk -f sum.awk nums
Total is 170
```

Patterns

/re/ matches the regular expression **re**

BEGIN matches before any input is used (can be used to set variables)

END matches after all input is used (e.g., can print things)

expr matches if the expression is nonzero

p1, p2 matches all lines between the line matching p1 and the line matching p2 (including those lines)

(empty pattern) matches every line

Expressions in patterns

Examples:

- ▶ `$3 == "foo" { ... }` Matches when field 3 is the string foo
- ▶ `$2 ~ /re/ { ... }` Matches when field 2 matches the regex re

You can use relational operators: `<`, `<=`, `==`, `!=`, `>`, and `>=`

You can use match operators: `expr ~ /re/` and `expr !~ /re/`

A bunch of builtin functions including `substr`, `length`, and `sub` (substitute)

The action(s) are performed when the pattern expression evaluates to true

Actions

An action is a sequence of statements inside { } separated by ;

- ▶ assignment statements var = value
- ▶ conditionals/loops: **if**, **while**, **for**, **do-while**, **break**, **continue**,
- ▶ **for** (var **in** array) stmt
- ▶ **print** expr-list
- ▶ **printf** format, expr-list

A missing action means to print the line

AWK example

Prints lines longer than 72 characters

```
length($0) > 72 { print }
```

Missing action block means print

```
length($0) > 72
```

Print size and owner from ls -l

```
$ ls -l | awk '{ print $5, "\t", $3 }'
```

```
mhogan@mcnulty:~$ ls -l
total 48
drwx----- 17 mhogan mhogan 4096 Nov 21 22:39 cs241
drwxr-xr-x  2 mhogan mhogan 4096 Sep 10 15:32 Desktop
drwxr-xr-x  2 mhogan mhogan 4096 Sep 10 15:32 Documents
drwxr-xr-x  2 mhogan mhogan 4096 Sep 10 15:32 Downloads
...
mhogan@mcnulty:~$
```

```
ls -l | awk '{ print $5, "\t", $3 }'
```

```
4096 mhogan
4096 mhogan
4096 mhogan
4096 mhogan
```

Given pop.txt with lines containing zip code, county, population, e.g.,

44001 Lorain 20769

44011 Lorain 21193

what is the awk command to print out the population of Oberlin (zip code 44074)?

A. \$ awk '/44074/ { print \$3 }'

B. \$ awk '\$0 == 44074 { print \$2 }'

C. \$ awk '\$1 == 44074 { print \$3 }'

D. \$ awk '44074 { print \$2 }'