

Lecture 03 – Control Flow

Stephen Checkoway

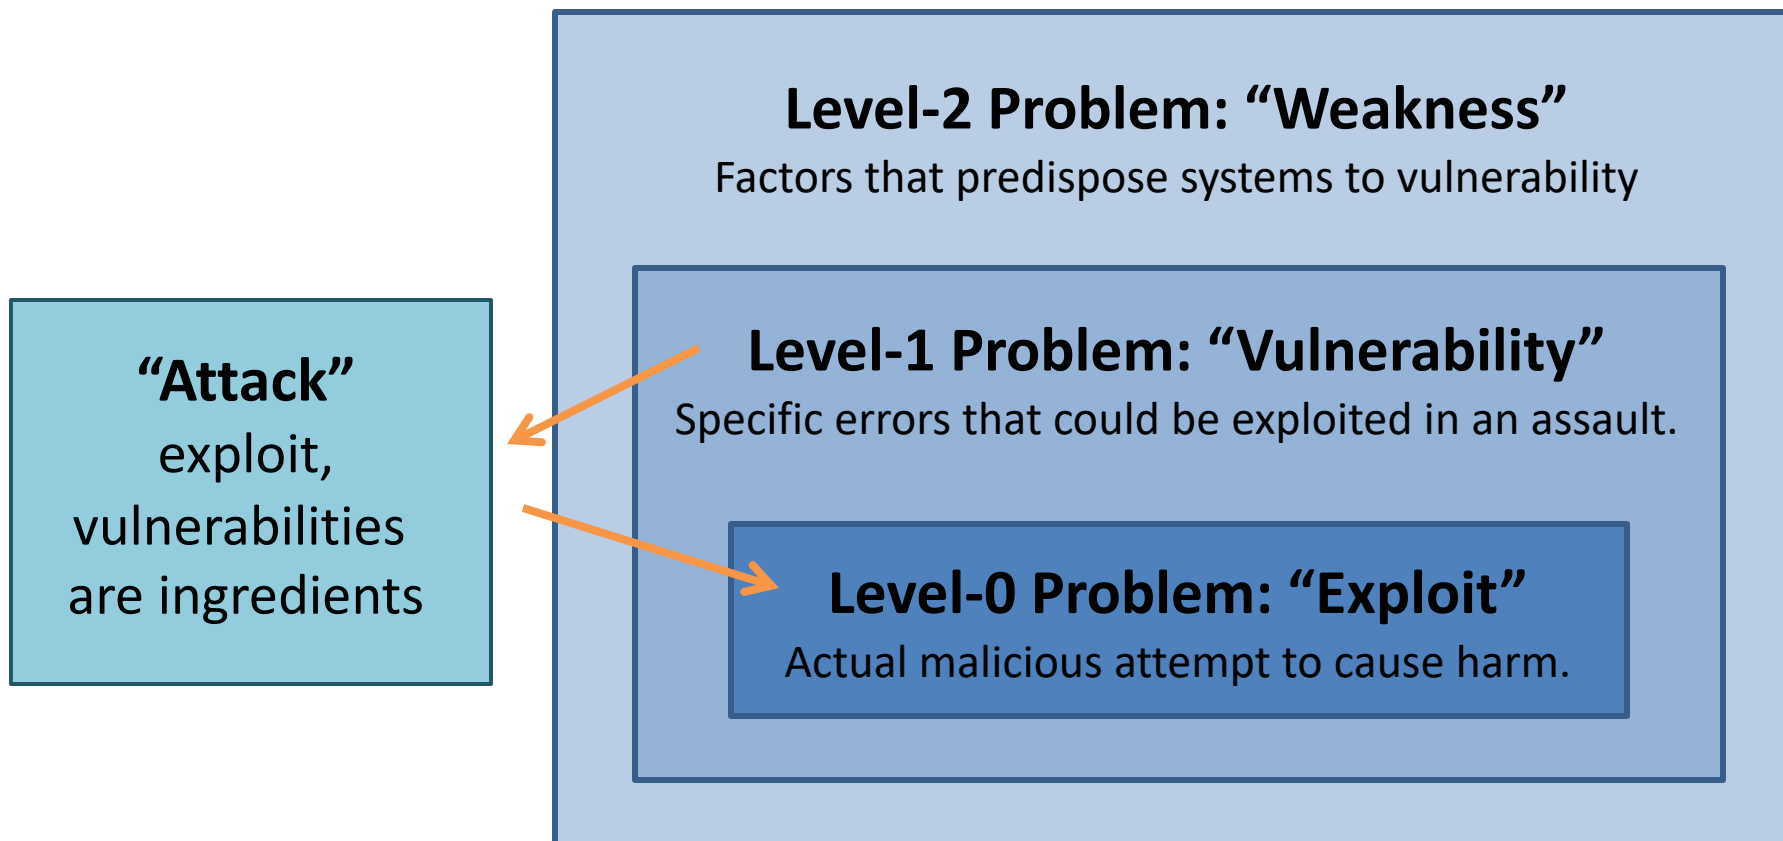
CS 343 – Fall 2020

Adapted from Michael Bailey's ECE 422

Outline

- Computer
 - CPU
 - Instructions
- The Stack (x86)
 - What is a stack
 - How it is used by programs
 - Technical details
- Attacks
- Buffer overflows
- Adapted from Aleph One's "Smashing the Stack for Fun and Profit"

“Insecurity”?



Why Study Attacks?

- Identify vulnerabilities so they can be fixed.
- Create incentives for vendors to be careful.
- Learn about new classes of threats.
 - Determine what we need to defend against.
 - Help designers build stronger systems.
 - Help users more accurately evaluate risk.

```

static OSStatus
SSLVerifySignedServerKeyExchange(SSLContext *ctx, bool isRsa, SSLBuffer signedParams,
                                uint8_t *signature, UInt16 signatureLen)
{
    OSStatus    err;
    ...

    if ((err = SSLHashSHA1.update(&hashCtx, &serverRandom)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.update(&hashCtx, &signedParams)) != 0)
        goto fail;
    if ((err = SSLHashSHA1.final(&hashCtx, &hashOut)) != 0)
        goto fail;
    ...

fail:
    SSLFreeBuffer(&signedHashes);
    SSLFreeBuffer(&hashCtx);
    return err;
}

```


Virtual memory

- Each running process has its own virtual memory space
 - Your computer has a bunch of RAM
 - RAM is an array of bytes indexed from 0
 - It would be bad if any process could read/write any byte of memory
 - The OS and hardware carve up memory and hand it out to processes




Virtual address space

- OS presents each process with the fiction that it has access to the entire valid range of memory from index 0 to the maximum index ($2^{32} - 1$ or $2^{64} - 1$)
- It does this by mapping virtual addresses used by processes to physical addresses used by the hardware



Virtual address space layout

- Each function called in a program is allocated a *stack frame* on the call stack; it stores
 - The return address
 - Local variables
 - Arguments to functions it calls
- The software maintains two pointers
 - Stack pointer: points to the top (lowest address) of the stack
 - Frame pointer: points to the call frame (optional)




example.c


```
void foo(int a, int b) {  
    char buf1[10];  
}
```

```
void main() {  
    foo(3, 6);  
}
```


C stack frames




C stack frames




C stack frames




C stack frames



C stack frames



C stack frames



32-bit x86 architecture overview

- 8 general purpose registers `eax`, `ebx`, `ecx`, `edx`, `esi`, `edi`, `ebp`, `esp`
 - `esp` is the stack pointer
 - `ebp` is the frame pointer (optional)
 - Others are used for integer and pointer operations
 - 16- and 8-bit parts of the registers can be named (`ax` is least significant 16 bits of `eax`, `al` is least sig. 8 bits of `eax`, etc.)
- Instruction pointer `eip` holds the address of the next instruction to execute

Some x86 instructions (AT&T notation)

- `add src, dest` ; Computes $\text{dest} + \text{src}$, stores in `dest`
- `sub src, dest` ; Computes $\text{dest} - \text{src}$, stores in `dest`
- `mov src, dest` ; Copies `src` to `dest`
- `push src` ; Decrements `esp` by 4, writes `src` to stack
- `pop dest` ; Reads top of stack into `dest`, increments `esp` by 4
- `call foo` ; Calls the function `foo`, pushes the address of the next instruction onto the stack
- `ret` ; Pops the top of the stack into `eip` (returns from a function)

Instruction suffixes

- l — (long) 32 bits
- w — (word) 16 bits
- b — (byte) 8 bits
- Examples
 - `movw %ax, %dx` ; Copies least sig. 16 bits of `eax` to least sig. 16 bits of `edx`
 - `pushl %edi`
 - `subl $16, %esp` ; Decrements `esp` by 16

x86 operands

- Constants are prefixed with \$
- Registers are prefixed with %
 - `movb $8, %bl`
- Read/writing to memory has several forms
 - `(%eax)` ; Refers to the 1, 2, or 4 bytes at address stored in `eax`
 - `-8(%esp)` ; Address is `%esp - 8`
 - `4(%esi, %eax)` ; Address is `esi + eax - 4`
 - `16(%eax, %edx, 4)` ; Address is `eax + 4*edx + 16`

Using memory operands

- Load 4 bytes from `ebp + 4` into `eax`
 - `movl 4(%ebp), %eax`
- Store 1 byte from `dl` (least sig. 8-bits of `edx`) to address `edi`
 - `movl %dl, (%edi)`
- Add 4 bytes from address `edx` to `eax` and store in `eax`
 - `addl (%edx), %eax`
- Xor the constant `0x5555AAAA` with 4 bytes at address `8+ebp`
 - `xorl $0x5555AAAA, 8(%ebp)`

C stack frames (x86 specific)


Grows toward lower address

Starts ~end of VA space

Two related registers

`%esp` - Stack Pointer

`%ebp` - Frame Pointer



example.c

```
void foo(int a, int b) {  
    char buf1[16];  
}
```

```
int main() {  
    foo(3, 6);  
}
```

example.s (x86)

main:

pushl %ebp

movl %esp, %ebp

subl \$8, %esp

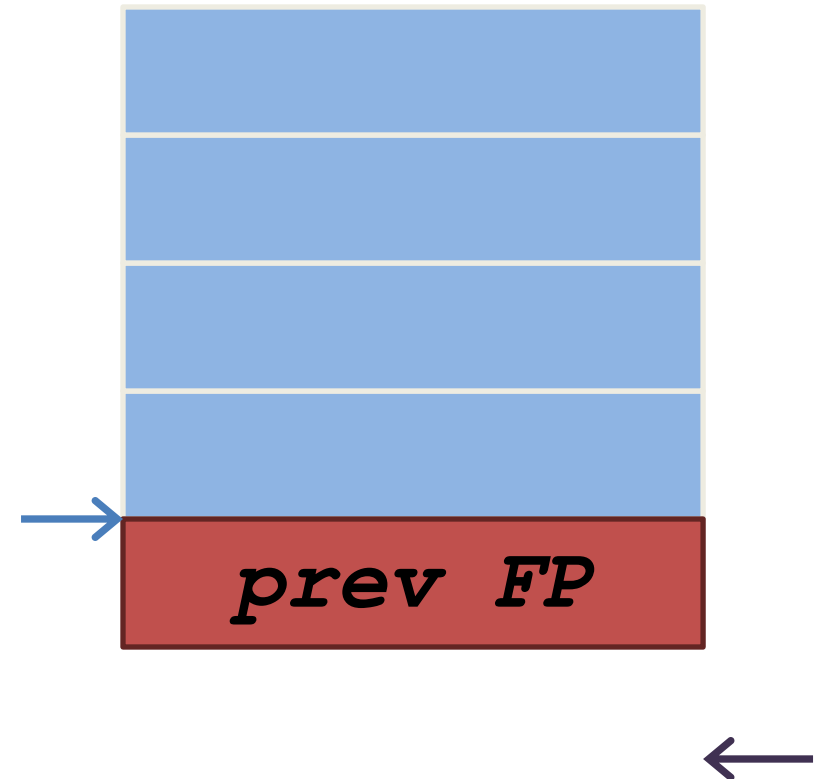
movl \$6, 4(%esp)

movl \$3, (%esp)

call foo

leave

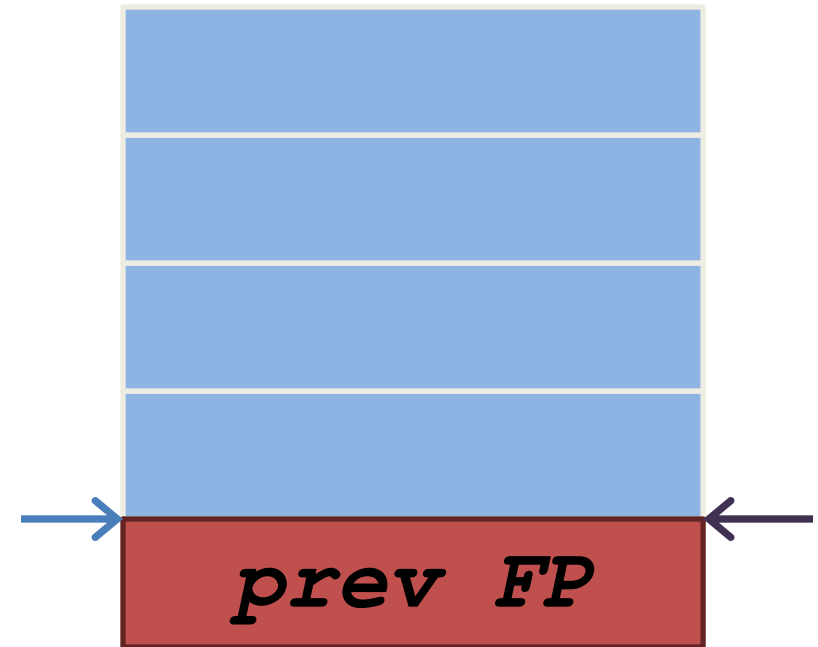
ret



example.s (x86)

main:

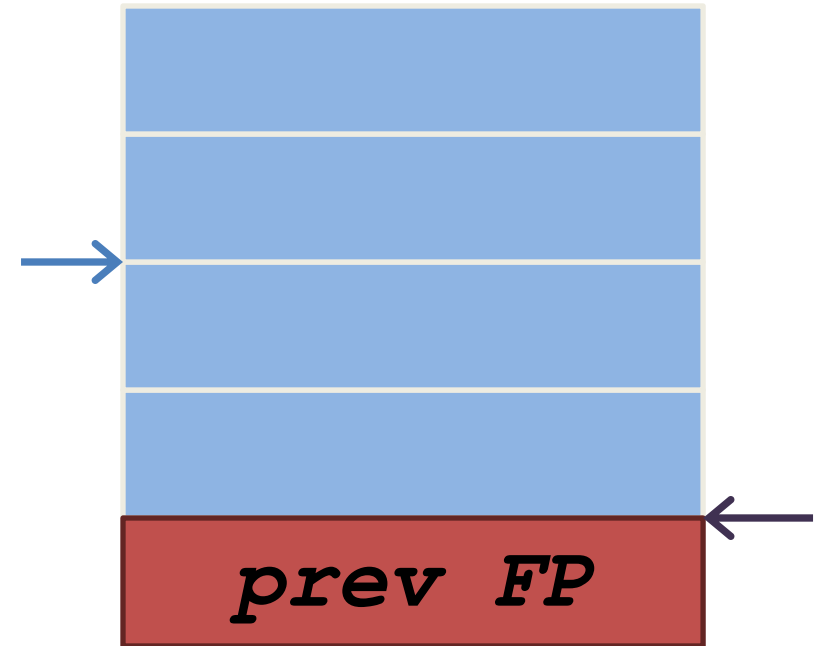
```
    pushl    %ebp
    movl     %esp, %ebp
    subl     $8, %esp
    movl     $6, 4(%esp)
    movl     $3, (%esp)
    call     foo
    leave
    ret
```



example.s (x86)

main:

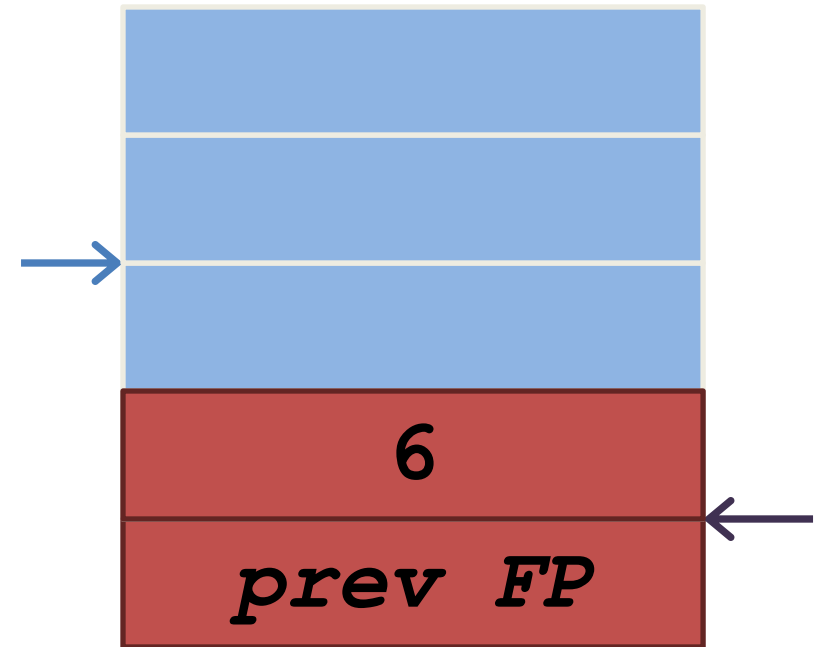
```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call     foo
leave
ret
```



example.s (x86)

main:

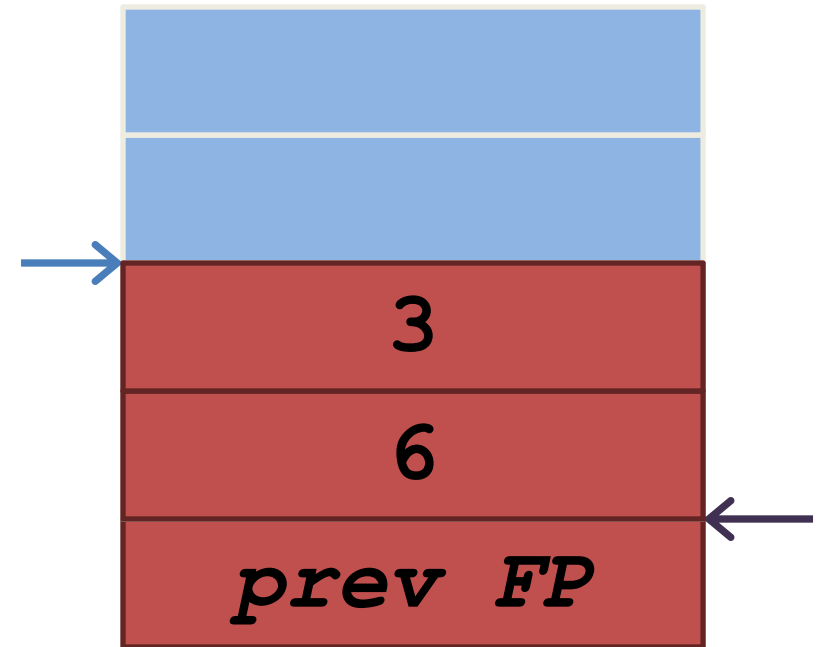
```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call     foo
leave
ret
```



example.s (x86)

main:

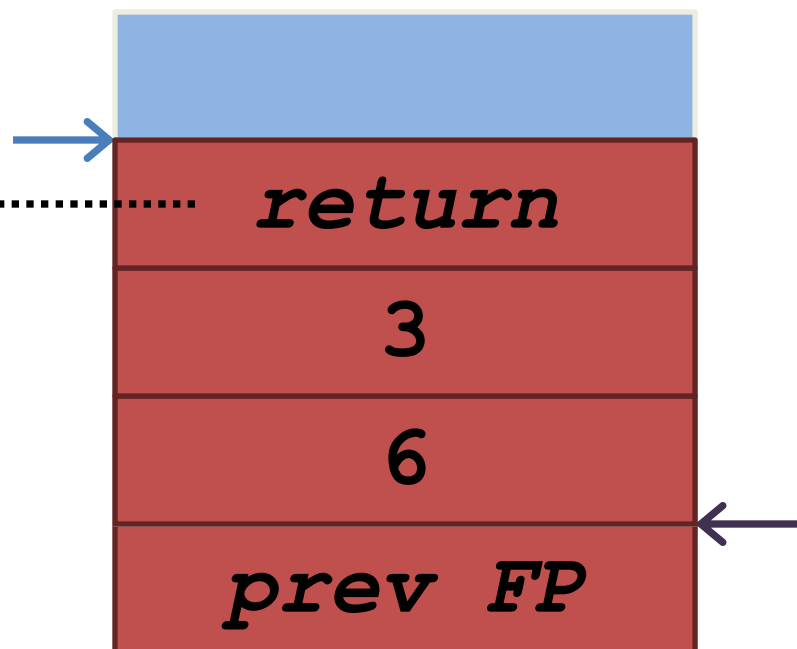
```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call     foo
leave
ret
```



example.s (x86)

main:

```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call    foo
leave    ←
ret
```



example.s (x86)

foo:


pushl **%ebp**

movl **%esp, %ebp**

subl **\$16, %esp**

leave


ret



example.s (x86)

foo:

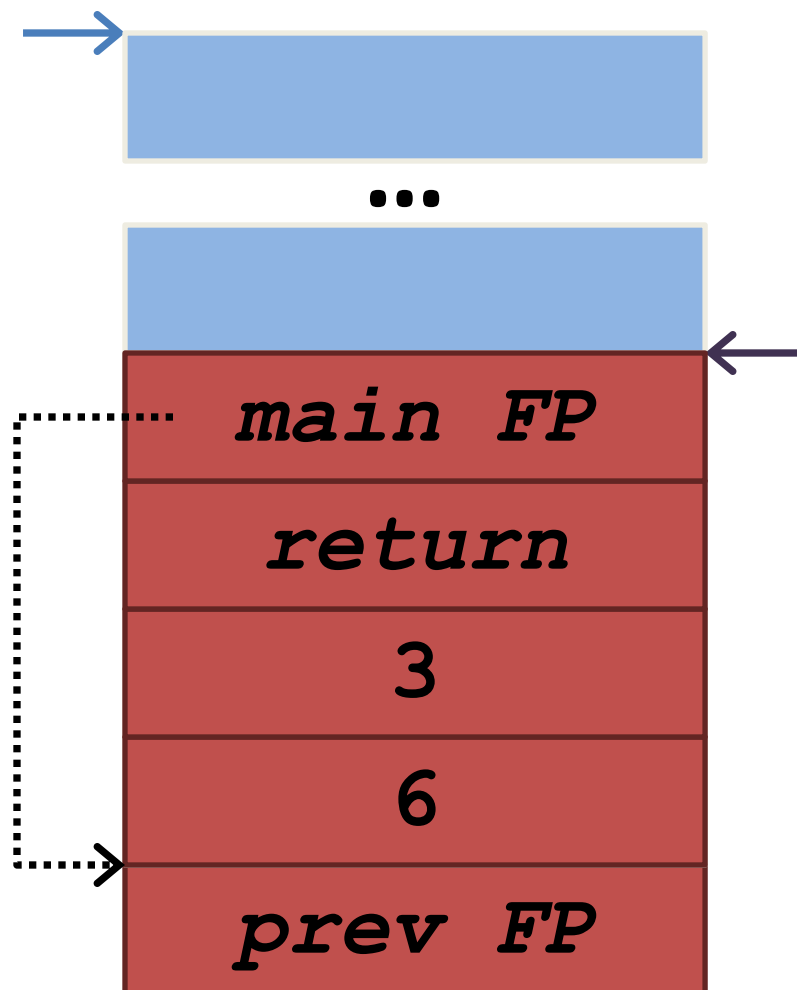
```
pushl    %ebp  
movl     %esp, %ebp  
subl     $16, %esp  
leave  
ret
```



example.s (x86)

foo:

```
pushl    %ebp
movl     %esp, %ebp
subl     $16, %esp
leave
ret
```

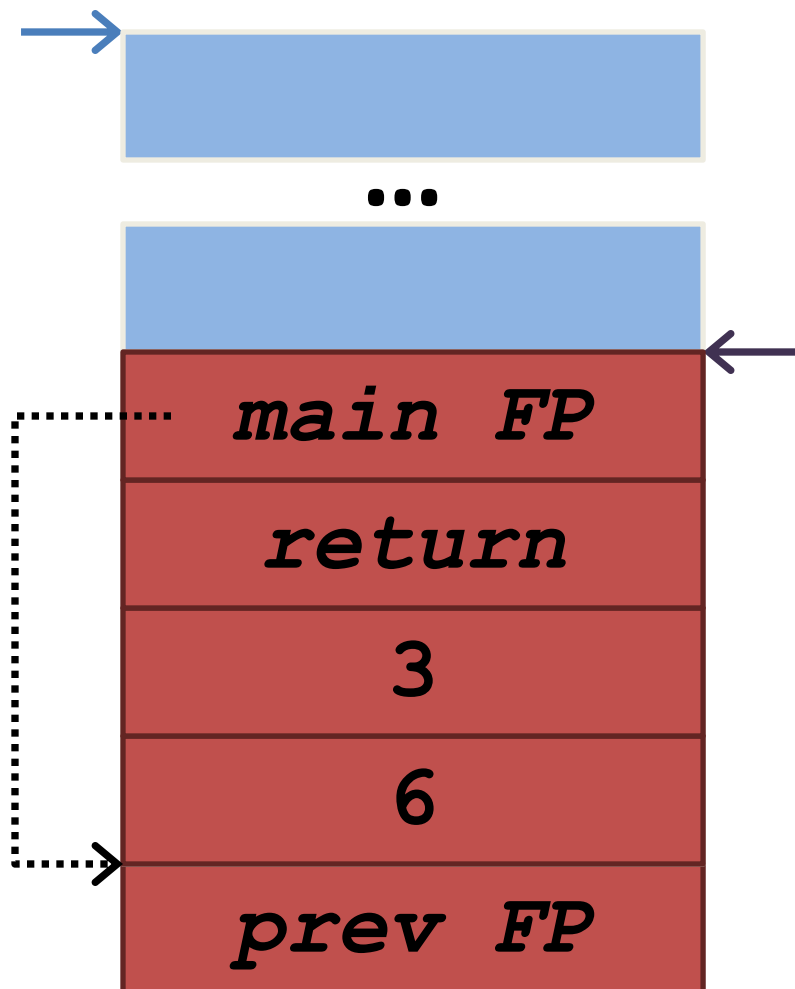


example.s (x86)

foo:

```
pushl    %ebp
movl     %esp, %ebp
subl     $16, %esp
leave
ret
```

```
mov %ebp, %esp
pop %ebp
```




example.s (x86)

foo:

```
pushl    %ebp
movl     %esp, %ebp
subl     $16, %esp
leave
ret
```

```
mov %ebp, %esp
pop %ebp
```




example.s (x86)

foo:

```
pushl    %ebp
movl     %esp, %ebp
subl     $16, %esp
leave
ret
```

```
mov %ebp, %esp
pop %ebp
```




example.s (x86)

foo:

```
pushl    %ebp
movl     %esp, %ebp
subl     $16, %esp
leave
```

ret

```
mov %ebp, %esp
pop %ebp
```




example.s (x86)

main:

```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call     foo
leave
ret
```

```
mov %ebp, %esp
pop %ebp
```



example.s (x86)


main:

```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call     foo
```

leave

```
ret
```

```
mov %ebp, %esp
pop %ebp
```



example.s (x86)

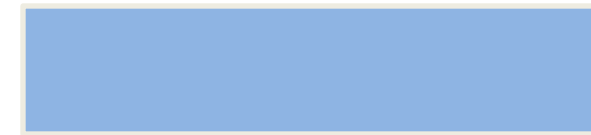
main:

```
pushl    %ebp
movl     %esp, %ebp
subl     $8, %esp
movl     $6, 4(%esp)
movl     $3, (%esp)
call     foo
```

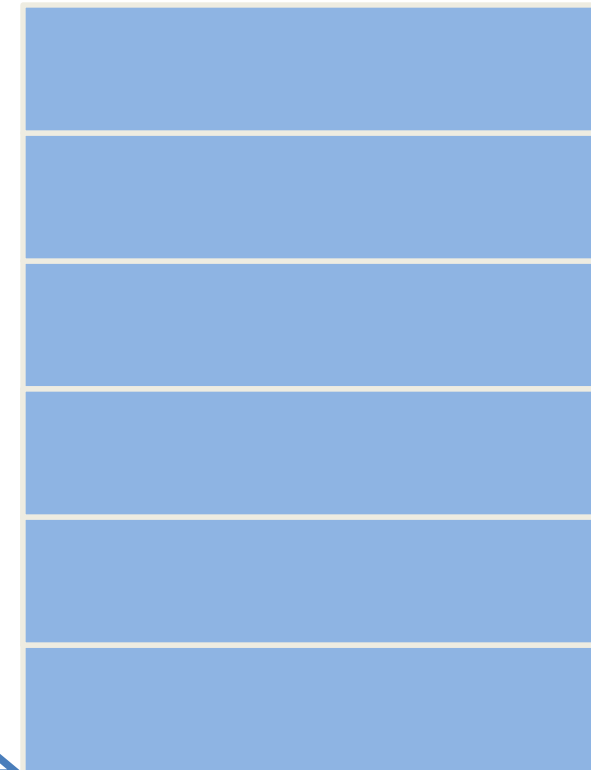
leave

```
ret
```

```
mov %ebp, %esp
pop %ebp
```



...

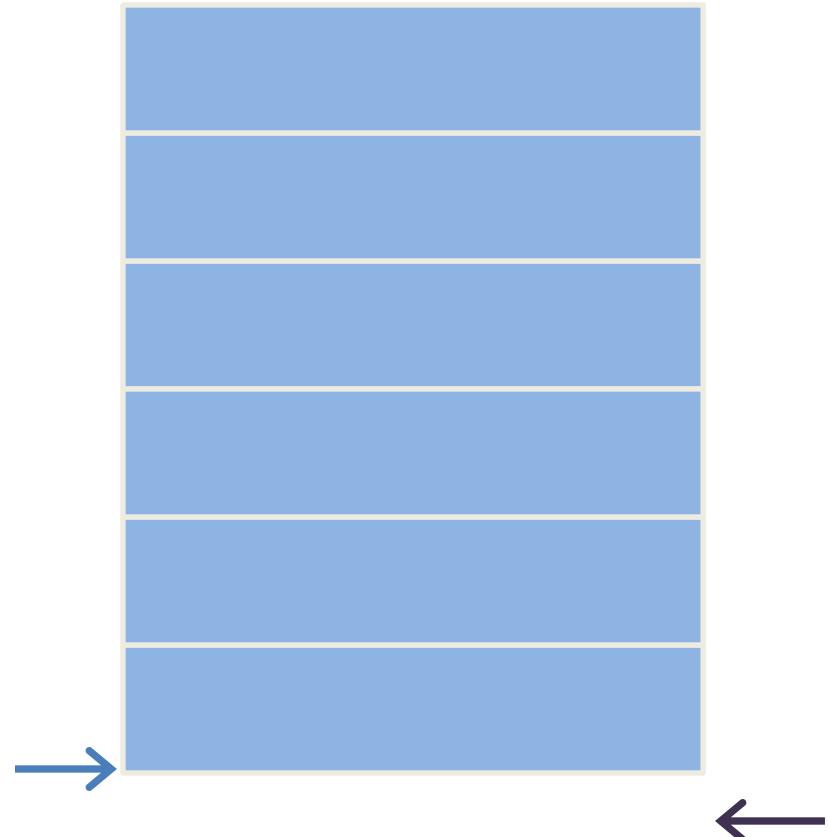


Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```

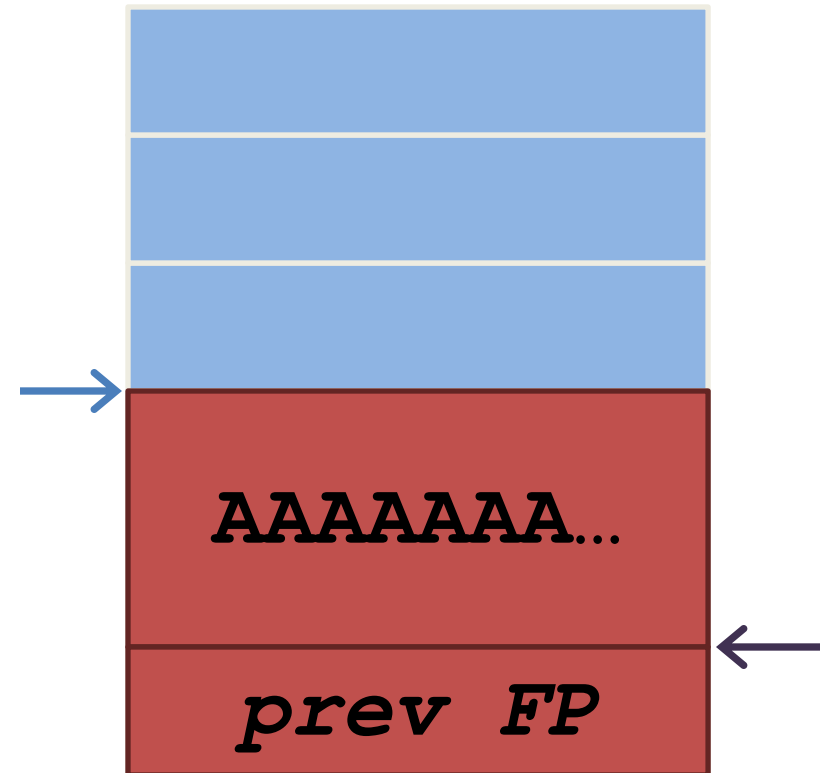
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



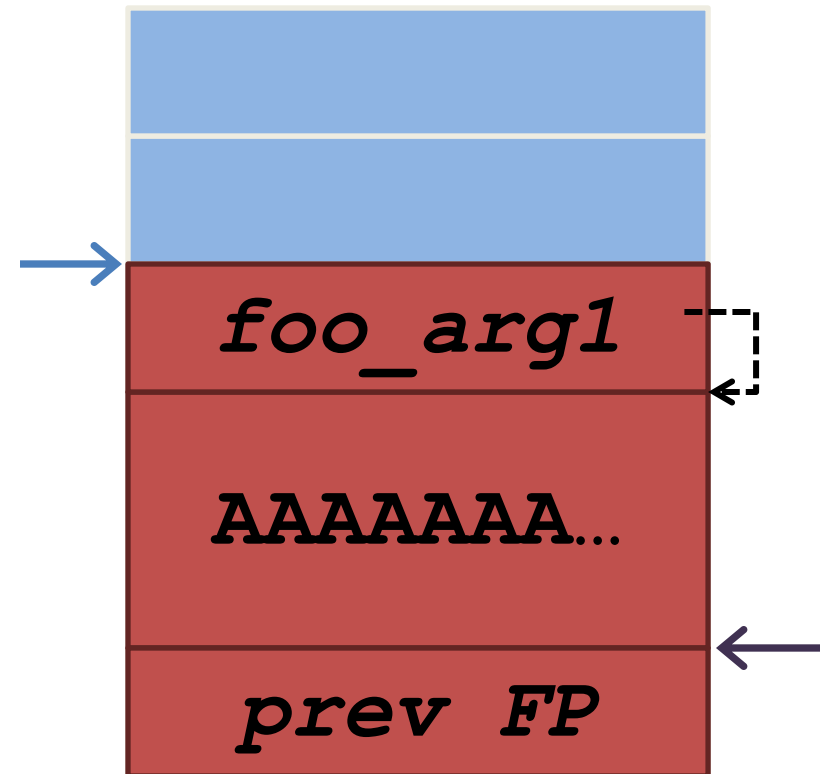
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



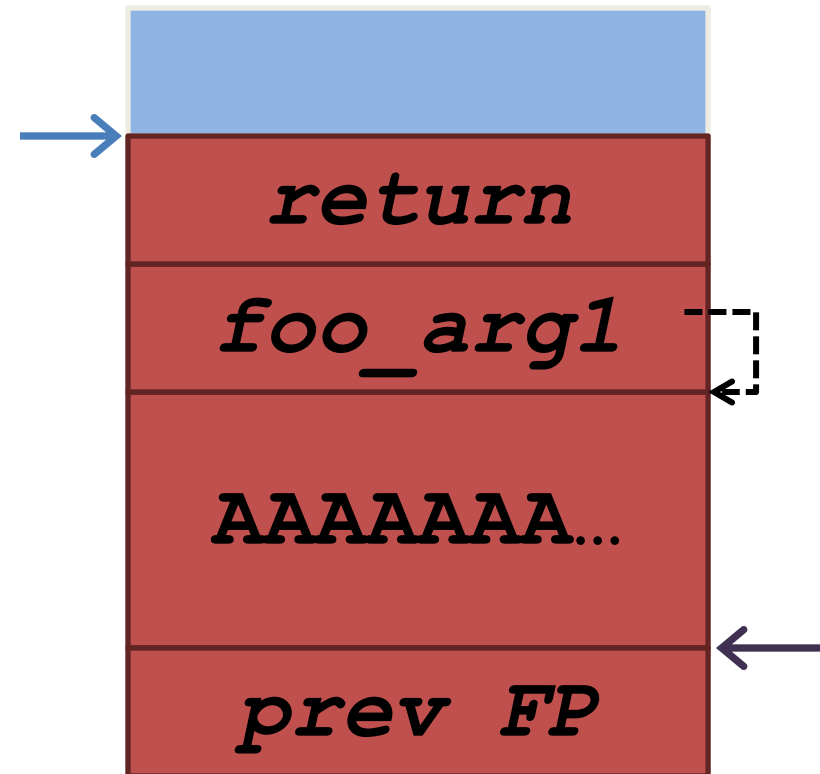
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```



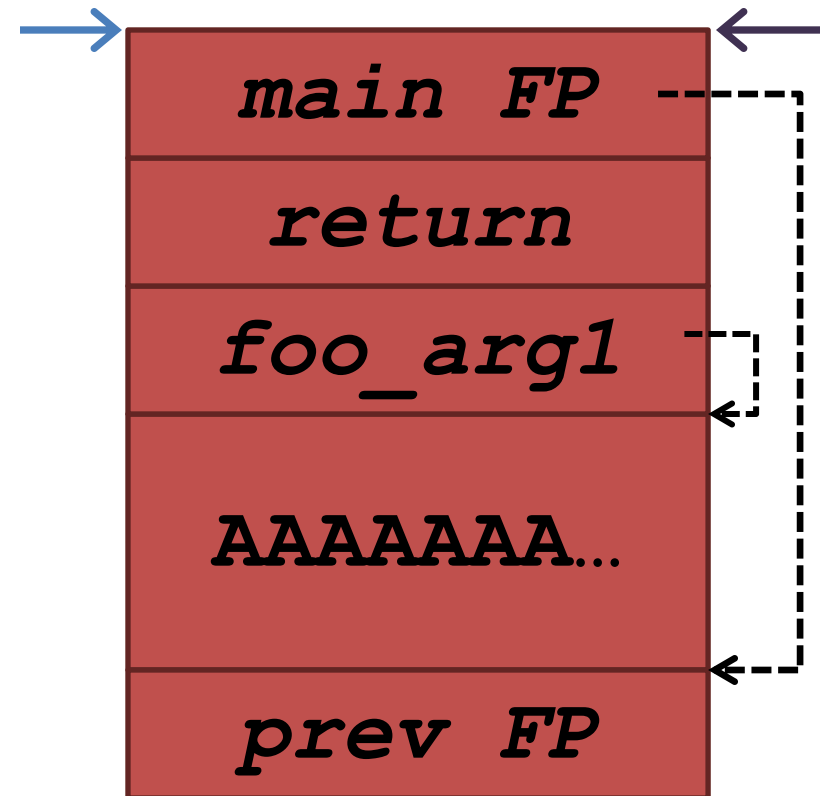
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



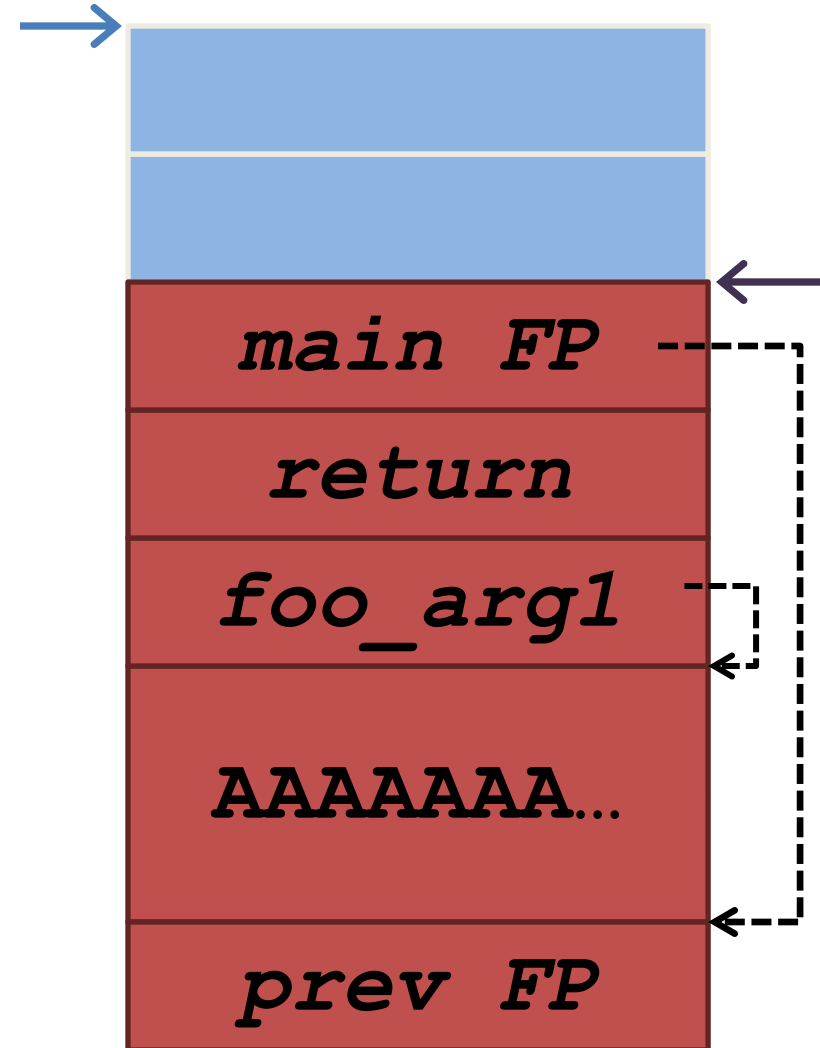
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



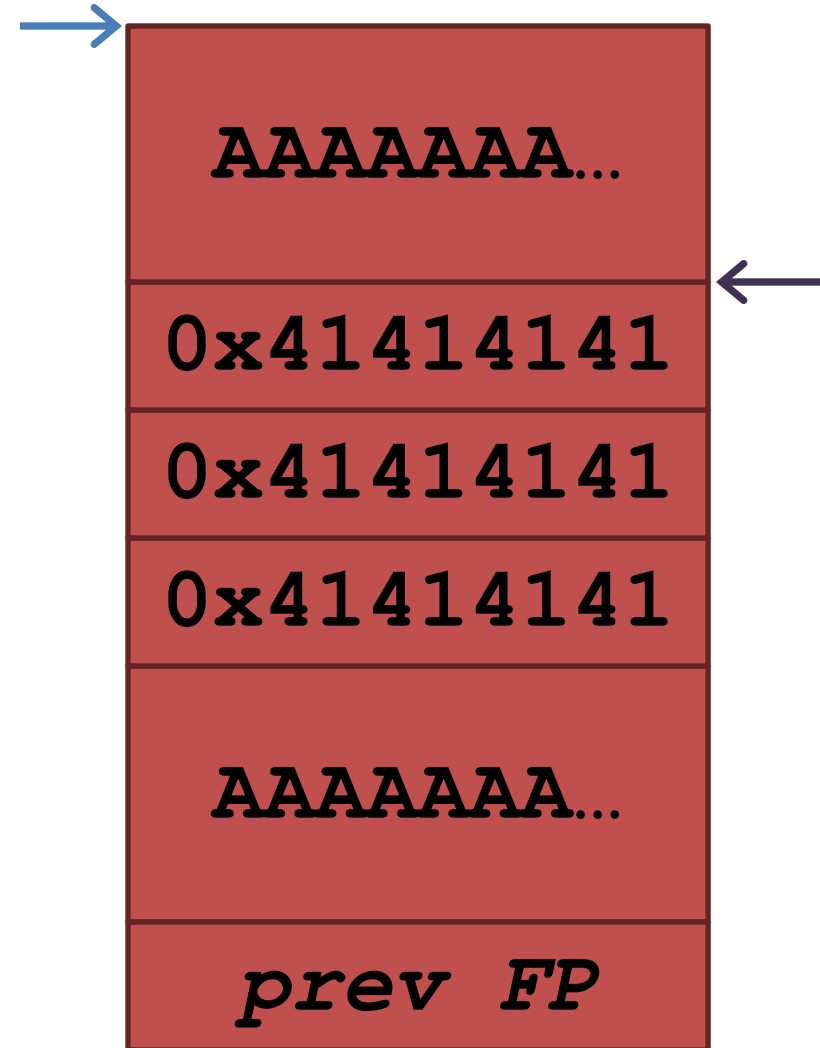
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



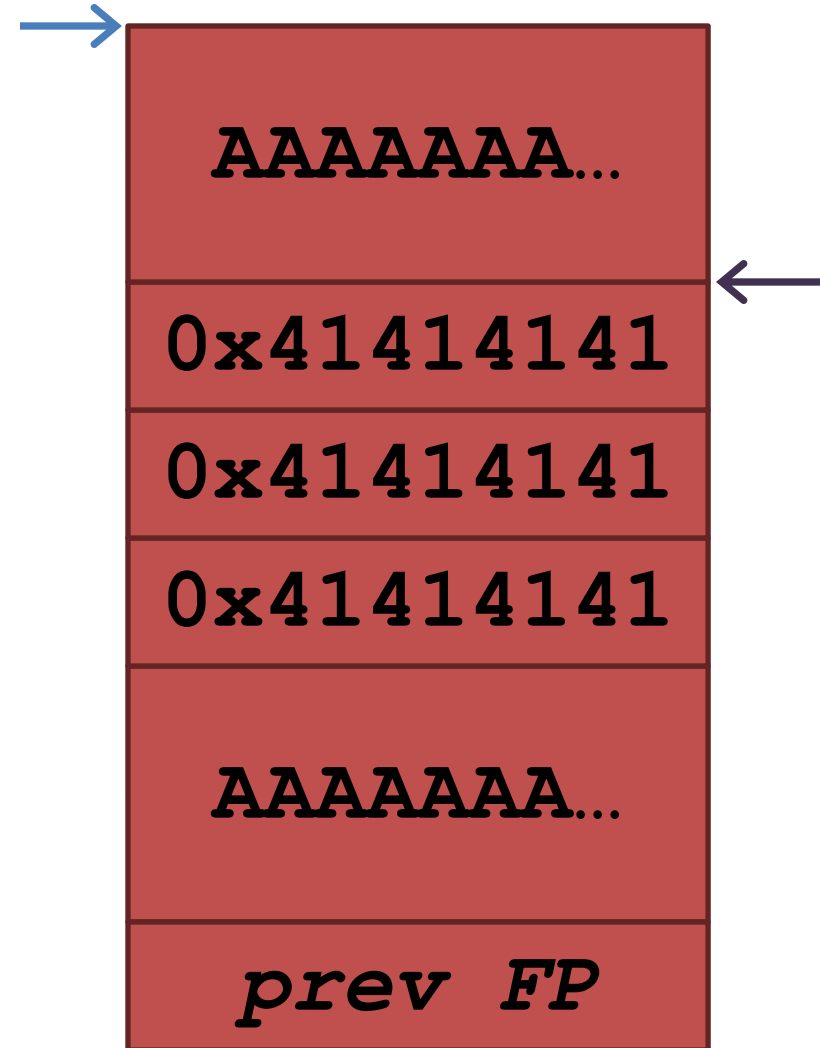
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\x00';  
    foo(buf);  
}
```



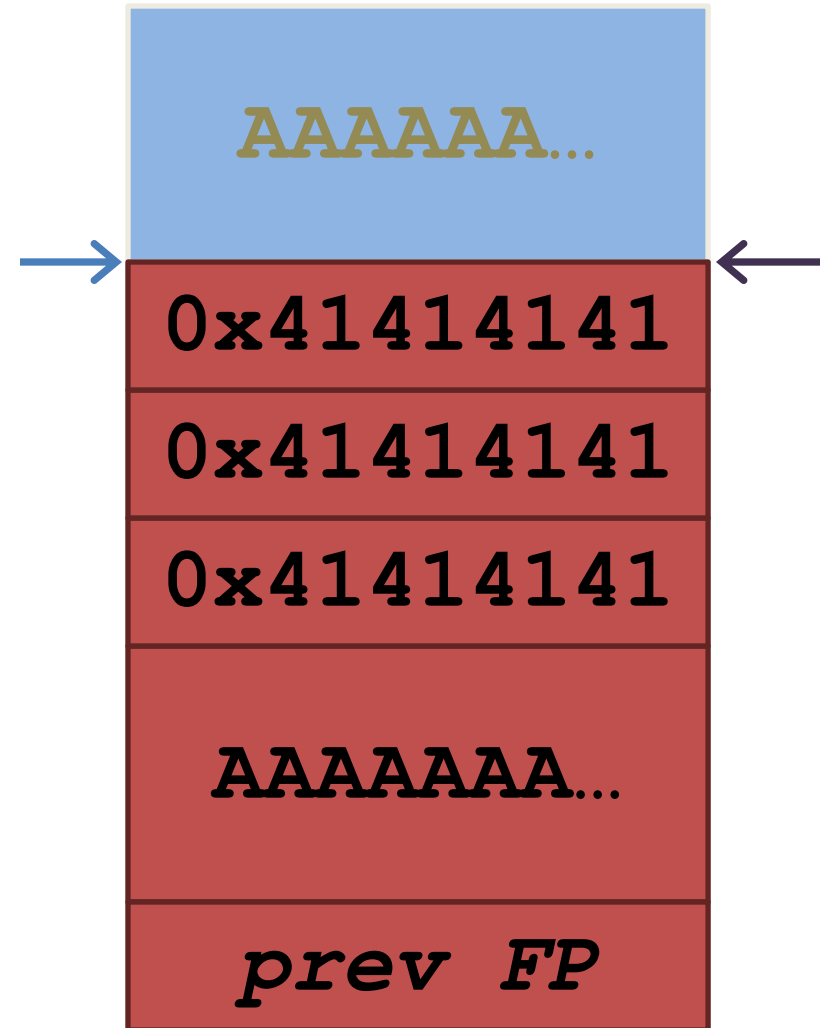
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
    mov %ebp, %esp  
    pop %ebp  
    ret  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



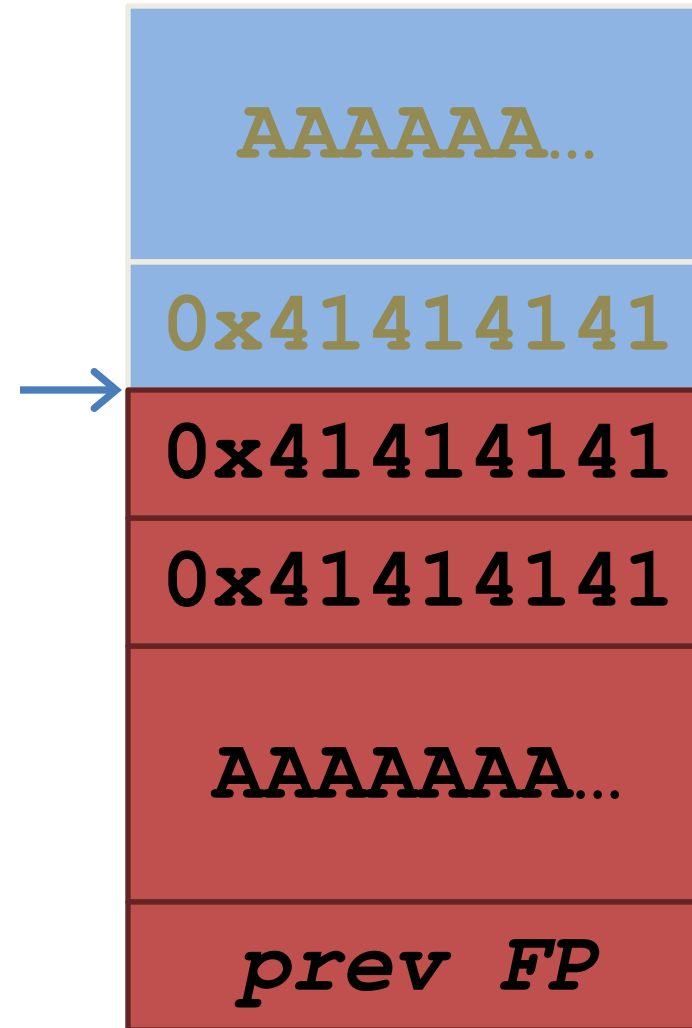
Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
    mov %ebp, %esp  
    pop %ebp  
    ret  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
    mov %ebp, %esp  
    pop %ebp  
    ret  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



Buffer overflow example

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
    mov %ebp, %esp  
    pop %ebp  
    ret  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    foo(buf);  
}
```



Buffer overflow example

`eip = 0x41414141`

`???`

`? ←`



Buffer overflow FTW

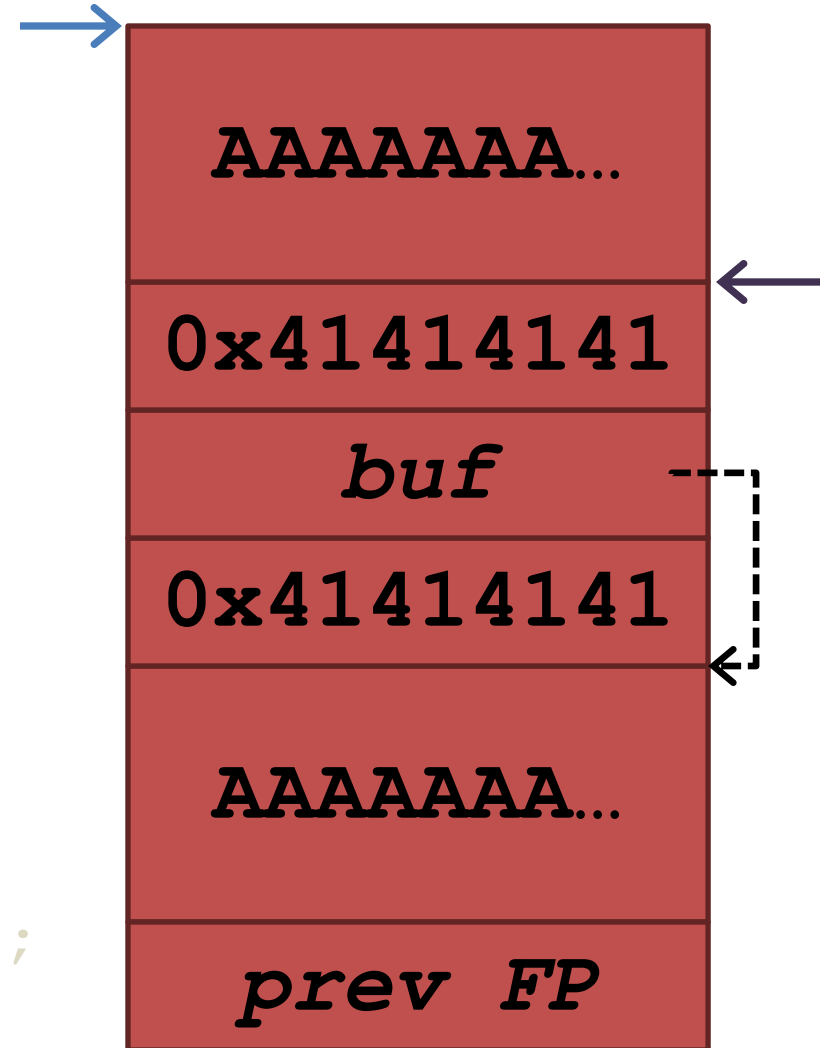
- Success! Program crashed!
- Can we do better?
 - Yes
 - How?

Exploiting buffer overflows

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    ((long*)buf)[5] = (long)buf;  
    foo(buf);  
}
```

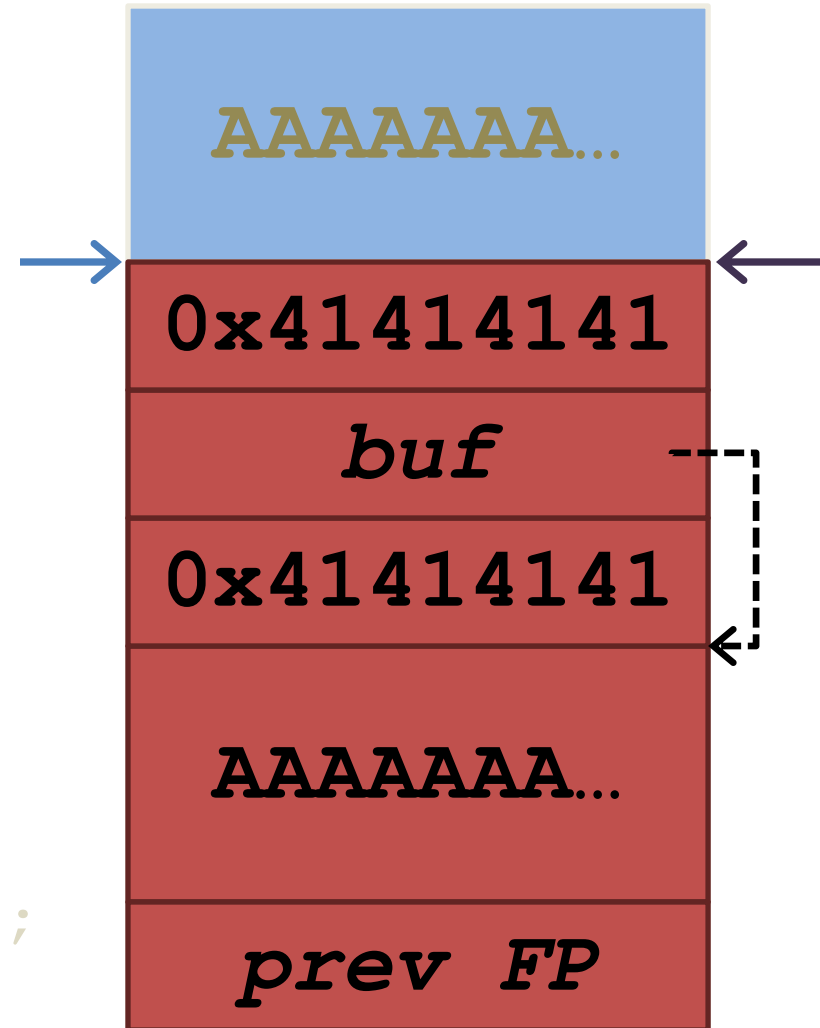
Exploiting buffer overflows

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    ((int*)buf)[5] = (int)buf;  
    foo(buf);  
}
```



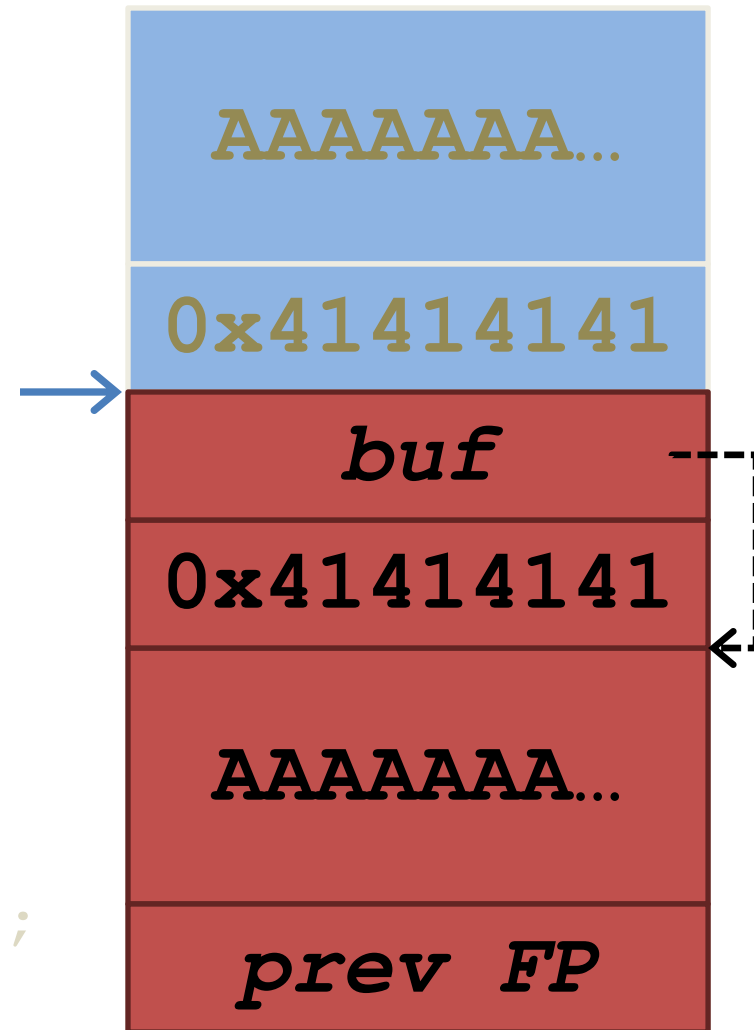
Exploiting buffer overflows

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
    mov %ebp, %esp  
    pop %ebp  
    ret  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    ((int*)buf)[5] = (int)buf;  
    foo(buf);  
}
```



Exploiting buffer overflows


```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
    mov %ebp, %esp  
    pop %ebp  
    ret  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    ((int*)buf)[5] = (int)buf;  
    foo(buf);  
}
```



Exploiting buffer overflows

```
void foo(char *str) {  
    char buffer[16];  
    strcpy(buffer, str);  
}  
  
int main() {  
    char buf[256];  
    memset(buf, 'A', 255);  
    buf[255] = '\\x00';  
    ((int*)buf)[5] = (int)buf;  
    foo(buf);  
}
```

```
mov %ebp, %esp  
pop %ebp  
ret
```



What's the Use?

- If you control the source?
- If you run the program?
- If you control the inputs?

(slightly) more realistic vulnerability

```
int main()  
{  
    char buffer[100];  
    printf("Enter name: ");  
    gets(buffer);  
    printf("Hello, %s!\n", buffer);  
}
```

(slightly) more realistic vulnerability

```
int main()  
{  
    char buffer[100];  
    printf("Enter name: ");  
    gets(buffer);  
    printf("Hello, %s!\n", buffer);  
}
```

```
python -c "print '\x90'*110 + \  
'\xeb\xfe' + '\x00\x00\xff\xff'" | \  
./a.out
```